

HANDOFF — ABRA, 2026-07-28 (overnight session)

ARCHIVED 2026-08-05 — PROVENANCE RECORD, NOT CURRENT STATE. Kept because the trail is the evidence: what was believed, when, and what broke it. Do not take a number out of this file. `node engine/status.js` is the state. - **Claimed:** the overnight corrections session — that three things the previous handoff asserted were wrong (Armor Tail / Queenly Majesty / Dazzling were already modelled, Contrary does not invert `myOffenseStage`, `pory.py` is not a raw-store reader), plus the mechanics work of that night. - **Written:** 2026-07-28, CHANGELOG at 3.27.x. - **Replaced by:** nothing wholesale — this is a dated record. Its mechanics claims are superseded by `docs/MECHANICS-COVERAGE.md` and `data/tags.json`, both generated. - **Retracted inside:** None registered. It corrects `HANDOFF-2026-07-27.md` rather than being corrected by anything, and the corrections held.

Read this **after** `HANDOFF-2026-07-27.md`, which it corrects in several places. Repo:

`C:\Users\willj\Projects\Pokemon\ABRA Suite: SHOWDOWN_PATH=C:/Users/willj/Projects/Pokemon/pokemon-showdown node tests/run-all.js` → **19 passed, 3 failed, 1 skipped**. The 3 failures are the same deliberate ones as before.

o. THE PREVIOUS HANDOFF WAS WRONG ABOUT THREE THINGS

Recorded first, because they were acted on as fact.

Armor Tail / Queenly Majesty / Dazzling were NOT absent. They were already modelled at 99% and 98% through the measured `ability-blocks.json` table, which `abilityBlockProb` applies. The handoff said "grep returns zero", and it does — `board.js` deliberately names almost nothing. **Grep is not a coverage test for this file.** Only the *ally-side* half was genuinely missing, and that is now fixed.

Contrary does NOT invert `myOffenseStage`. That feature reads stat stages already on the board, recorded from real protocol events, and is correct whatever produced them. Contrary changes the *predicted* effect of a move under consideration — `movesBoostMe` and `movesLowerFoe`, and only those.

`pory.py` **is not one of the raw-store readers.** It filters to clean ids via `store.load_games(clean=True)` and refuses to run without them. `pory_baseline.py` and `train_value.py` are on that list; PORY itself is not.

1. WHAT LANDED

The corpus grew 35%, and every earlier number is stale

The forfeit rule now excludes only forfeits **before any action**, on Will's ruling. The old rationale ("a forfeit records who quit, not who was winning") was never tested and is false: of 1,528 open-sheet forfeits with at least one action, the player who quit was **behind on mons 86.8%** of the time and **ahead 0.9%**. It was discarding 1,528 decided games to exclude 4 undecided ones.

clean open-sheet 2,114 -> 2,860 clean ladder 2,653 -> 3,571

Mean game length moves 8.46 → 8.09 turns; median (8) and p90 (12) do not move. **Anything computed on a clean corpus before 2026-07-28 was computed on a different population.**

The stores are tracked compressed

`games.ladder.jsonl` was 84.6 MB against GitHub's **hard** 100 MB limit — about 38 hours of collection from the point where every push fails, including the ingest Action's. All three stores now track as `.gz` (137.6 MB → 15.4 MB). `quality.js / quality.py` read either form, preferring the plain file when present. `build/compress-stores.js --check` reports staleness.

board.js batch — one refit, seven changes

- **Speed is investment-aware.** It was base stats; `spreadLines()` already existed and the *damage* calc already used it. The error is large and non-uniform: Garchomp 102→167, Whimsicott 116→179, but Kingambit 50→73. Base stats compress the range and flip orders.
- **Choice Scarf, paralysis, and the four weather-speed abilities** now reach move order, all probed from `onModifySpe` rather than listed.
- **Side-wide blockers protect their partner.** Classified by handler name (`onFoeTryMove`).
- **Megas are resolved from the sheet's stone.** 27.5% of sheet entries hold one; their abilities were being read off the base forme.
- **Contrary/Simple** via `onChangeBoost` .
- `stallIntoEncore` — Will's feature, and it is real (see §2).

PORYGON2 — a new value function, deliberately not a new PORY

k-NN over self-play positions, evaluated on held-out **human** games. **62.8%** against **60.5%** for the material baseline, and better on Brier. PORY is untouched so the two stay separately quotable.

The risk lever (`engine/variance.js`)

Seek variance as an underdog, decline it as a favourite. Credit: Nate Silver, *On the Edge*. Defaults to an exact no-op and must be switched on with a stated belief.

2. WHAT IS VERIFIED, AND WHAT IS NOT

VERIFIED

- `stallIntoEncore` **-1.171** [-1.686, -0.655] — 12th of 48 weights, largest mover under reweighting (-1.993). Will's domain claim is a measured regularity.
- Ally-side blocking: Fake Out at a Sinistcha reads 0.000 beside Garchomp, **0.992** beside Farigiraf, 0.984 beside Tsareena.
- Contrary: Swords Dance on Staraptor reads 1.000 bare, **0.000** holding Staraptite.
- Held-out fit: 30.4% top-1, -1.7694 logL, against 22.8% for the behaviour clone.
- Simulator speed: **60 ms** per complete game with a random policy, **318 ms** with MAG.

MEASURED NEGATIVES — these are results, not gaps

- **PORYGON2 plateaus.** 8x the training data moved accuracy **-0.2 points**. Saturation at ~9,000 positions. More self-play is *not* the lever; the ceiling is the feature set.
- **Weighting the k-NN distance by learned importance made it worse** (62.5% vs 62.8%). It amplifies material and crushes the features the k-NN was using non-linearly.
- **Fixing four real mechanics barely moved opponent prediction** — joint recall at k=5 went 32.8% → 33.3%. The opponent-prediction ceiling is not missing mechanics.
- **MAG cannot narrow the opponent's turn.** Joint recall 33.3% at k=5, 74.5% at k=20 against ~56 joint actions. There is no k that is both affordable and honest.

INCONCLUSIVE — the falsifier ran, and it settled nothing

- 1,200 games with one side asserting a `skillGap` of 0.10 it does not have, then the same 1,200 seeds with the lever off so the runs pair game-for-game.

unpaired risk side won 48.6% of 1,176 decisive games, 95% CI [45.7, 51.4] paired 1,158 of 1,176 ended with the SAME winner (98.5%) 18 discordant: 7 losses became wins, 11 wins became losses exact two-sided $p = 0.481$

The direction is what the falsifier predicts — asserting an edge you do not have cost 4 net games — and it is indistinguishable from chance. **Reported as inconclusive, not as weak support**; reading a direction off $p = 0.48$ is the error this project keeps making.

- The useful part: **the lever is nearly inert at this setting**. It changes 2.1% of picks and flips 1.5% of games. Power over that needs ~200 discordant pairs — about 13,000 games, ~2 hours — or a larger `strength`, which nobody has calibrated. `variance.js` is a mechanism with a verified shape and no verified effect. It stays off.

NOT VERIFIED — do not quote

- MAG against a human. Still never measured (task 24).
- Everything in `HANDOFF-2026-07-27.md` §4 marked NOT VERIFIED.

3. TRAPS — THE OLD ONES STILL APPLY, PLUS THREE

Never edit `board.js` **while a fit or self-play run is in flight**. Unchanged and still true.

Grep is not a coverage test. `board.js` names almost nothing on purpose. Use `docs/MECHANICS-COVERAGE.md` (generated by `engine/mechanics_coverage.js`), which reports per *channel*. 192 abilities in this format: 92 reachable (71.5% of usage), **98 with dex handlers and no channel (22.4%)**.

An ignore rule does not apply to a file git already tracks. `data/games.ots.jsonl` was in `.gitignore` and tracked, because it was committed before the rule was written. The rule silently did nothing for 30.4 MB. `git rm --cached` is what makes it take effect.

Duplicate games are corrosive to a k-NN specifically. `mew.js` enumerates matchups deterministically from its seed, so two runs to the same `--out` produce the same games twice. For a logistic that is a doubled sample; for a nearest-neighbour model an exact duplicate is its own neighbour at distance zero carrying its own outcome, so the model *appears* to predict well by retrieving copies of the answer. `porygon2.py` now dedupes by id.

4. THE COLLINEARITY, WHICH IS NOW THE BIGGEST KNOWN DEFECT

`koTarget` is **-0.206** in the fit. Read literally that says people avoid killing. Measured model-free on held-out decisions, humans pick a killing move **1.41x** the base rate, a move that kills and moves first **1.82x**.

Will's diagnostic — fit each feature alone — settles it:

`koTarget` +0.764 alone -> -0.206 in fit SIGN FLIPS `koFirst` +0.902 -> +0.089 absorbed killsThreat +0.584 -> -0.098 SIGN FLIPS `dmgFrac` +0.648 -> -0.004 SIGN FLIPS `movesFirst` +0.412 -> +0.127 absorbed <- NEW this session

The family is inter-correlated at 0.5–0.67 and `lambda` selected on held-out data is **0**, so the coefficients are unconstrained. `eff4` and `immune`, which correlate with nothing, do not move at all.

`movesFirst` joining the block is the instructive part: it was *improved* this session, and improving it is what pulled it in. **Feature quality does not rescue a collinear design.**

The model predicts fine. The individual weights inside that block are not statements about the game and must not be quoted. Fix is ridge, or collapsing the block — neither attempted.

5. SLOWKING'S INTERVALS ARE UNUSABLE — DIAGNOSED, NOT FIXED

The intervals do not contain their own point estimates:

exploitability nash 0.0003 CI [0.0006, 0.0033] greedy - nash gap 0.3887 CI [0.0244, 0.3777]

Cause: `slowing_preview.py:145` solves with `iters=15000` ; line 157 solves every bootstrap replicate with `iters=1000` . Regret matching converges with iterations, so the replicates are under-converged and their residual exploitability is solver error, not matrix uncertainty.

Not fixed because `tests/test-slowking.py` asserts on these numbers and the published non-transitive-cycle claim rests on them. **That is Will's call.** Until then treat both intervals as unusable.

5b. THE SECOND OVERNIGHT PASS — FOUR RESULTS, THREE OF THEM NEGATIVE

The mechanics batch did not make MAG play better. Measured properly this time, both builds on opposite sides of the same battles, each with its own `board.js` and weights:

new MAG (48 features) vs pre-batch (commit 2a60c1d, 47 features) 1,368 decisive games — new MAG won 654 = 47.8%, 95% CI [45.2, 50.5] swapped half 46.8%, straight half 48.8% — both agree, so not a slot artefact
Inconclusive, point estimate **below** 50%. Three independent measurements agree: head-to-head 47.8%, top-1 31.13% → 30.37%, opponent joint recall +0.5 points. **Correctness is not strength.** Every mechanic added is individually verified and none of it reached the play.

My collinearity diagnosis was wrong. I blamed rivalry inside the kill block. Deleting `koFirst` , `killsThreat` and `dmgFrac` and refitting leaves `koTarget` at **-0.173** — it does not recover. The siblings are not the absorber. What `koTarget` actually correlates with once its family is set aside is `tgtHurt` at **+0.524**: a move kills largely *because* the target is nearly dead, and the model already knows that. **Redundancy, not rivalry** — which means every "fix the collinear block" plan, including the one at the top of my own list, was aimed at the wrong thing.

Three repairs, none beats the baseline:

baseline (lambda 0) logL -1.7694 top-1 30.38% ridge lambda 0.03 logL -1.8268 top-1 28.87% signs fixed, costs 1.5 points collapse to one logL -1.7698 top-1 30.38% coefficient still -0.085 drop the other three logL -1.7694 top-1 30.38% `koTarget` still -0.173

The signs can be bought with ridge; **prediction is the price.**

The coverage document was rebuilt because it was wrong twice over. Grepping source credited `speedboost` and `prankster` to the damage engine (they appear in a species lookup table) and missed twelve type-immunity abilities (bare object keys). A "fix" then did nothing, because `␣` in a JS template literal is a backspace character — and the regenerated numbers were identical and reported as corrected. **Both earlier figures (22.37% and 59.41%) were artefacts.** Rebuilt as an experiment: swap the ability, rerun the real code, see if anything moves.

192 abilities: 35 responsive (34.47%) 155 blind (59.36%)

The blind list is concentrated in `onStart` / `onSwitchInPriority` — Intimidate 5.65%, weather setters 9.4% combined. Those are the **conditional, one-step-ahead** mechanics no static feature vector can hold. "Blind" means cannot *anticipate*, not cannot observe.

One positive. PORYGON2 gained the conversion terms (`matchup_edge` , `speed_edge` , `type_threat`). `matchup_edge` is immediately the 6th most predictive feature of 17, and weighting the k-NN distance **flipped from hurting to helping** — best is now weighted k=50 at **63.6%** [62.7, 64.5] against 62.8% before. Intervals overlap, so suggestive rather than settled; the qualitative flip is the part that is not noise.

6. WHAT I WOULD DO NEXT, IN ORDER

1. **Read the falsifier result** (`engine/mew.js --risk-a`). It decides whether `engine/variance.js` stays or goes.
2. ~~Fix the collinearity~~ — **tested and dead**. Deleting the block does not recover the sign; the absorber is `tgtHurt` , not the siblings. Ridge fixes the signs and costs 1.5 points of top-1. Nothing to do here.
3. **Give PORYGON2 the conversion terms** — `koTarget` , `movesFirst` , type matchup. The learning curve says data is not the lever and the feature set is. This is the single highest-value item.
4. **Decide the SLOWKING CI** (§5).
5. **Measure MAG against a human** (task 24). Still the only thing that would tell you whether any of this plays.

Do not spend compute generating more self-play for PORYGON2. That was the plan and it is measured wrong.