

The whole-game differential — design, and the research it is taken from

Version 3.59.0 · Last updated 2026-08-06

ROADMAP #68. Written 2026-08-06 from Will's specification: *"we would want to play n games with like a thousand different teams to really test every single mechanic in the game"*, under his bar: *"i dont care if it takes a while, i just want it done correctly so we can trust it before using it to make all these decisions."*

1. Why the instruments we already have cannot answer the question

Every part is proven and nothing proves the whole.

instrument	reads	what it establishes	what it cannot see
mechanics census	235 probed, 234 live, 235 armed, 0 unarmed, 0 directCall	each mechanic fires, and each probe was shown RED on a broken engine first	anything about sequencing
damage differential	149 / 150	damage matches Showdown on staged pairs	any error under its 12% tolerance , uncounted
interaction matrix	1,624 / 1,643	pairs of mechanics compose	still staged, still pairwise
data/game-diff.json	5 games	five hand-written scenarios, 6 turns each	its <code>not_compared</code> list excludes hp amounts, accuracy misses, chance secondaries, crits, any pair where one engine KOs and the other does not , the <code>protect</code> volatile, and Showdown-only volatiles

Two of those deserve their own note.

The damage differential's tolerance is not absorbing roll variance. `showdownDamage()` is called at `roll=0` and `roll=15` against MEDICHAM's min and max, so both engines emit a pinned 16-roll range and the comparison is endpoint-to-endpoint. **A correct engine scores 0.000 on every row.** The 12% is free slack, and the artifact publishes only `worst`, so whether the 149 passing rows sat at 0.000 or at 0.119 is recorded by nothing.

The 31.1-point win-probability disagreement quoted in `docs/SUMMARY.md` **predates WIRES 123–132.** `engine/rerun_list.js` classes it unquotable. We do not currently know what that gap is.

2. The research this design is taken from

Each citation drives one decision below. They are not decoration.

2.1 Differential testing itself — McKeeman (1998)

Differential Testing for Software, Bill McKeeman, Digital Technical Journal. [PDF](#)

The founding statement: when two or more implementations of the same specification exist, feed both the same mechanically generated input and compare. **The oracle is free.** This is a materially stronger position than fuzzing, which can only see crashes and hangs — a differential harness sees *wrong answers*, which is the entire class of bug this project keeps finding.

McKeeman's other contribution is the **input tier**: he tested C compilers across seven tiers, from random ASCII up to model-conforming C programs. Deeper tiers find deeper bugs. Our tiers are already built and we are standing on the third of four: staged pair → staged interaction → **real game** → real game with real decisions.

Decision it drives: the harness compares two engines rather than asserting against expectations. We already own the authority (`engine/champions_sim.js` , the official simulator) and ADR-002 already says Showdown is it.

2.2 The input must make a difference *definitionally* a bug — Csmith (PLDI 2011)

Finding and Understanding Bugs in C Compilers, Yang, Chen, Eide, Regehr. [preprint](#) · [ACM](#)

325+ previously unknown bugs over three years. **Every compiler they tested both crashed and silently generated wrong code on valid input.** The methodological core is not the random generator — it is that Csmith generates programs *avoiding undefined and unspecified behaviour*, because where the standard permits two answers, a difference is not evidence of a bug and the oracle collapses.

Our undefined behaviour is **the dice**. Where either engine may legally roll, a divergence proves nothing. This project has already paid for that lesson once: CHANGELOG 3.45.0 records that the interaction matrix's *"two pinned dice were not the same die, so every sub-100-accuracy move had been MISSING in the reference engine while medicham2 hit"* — a whole class of false disagreement from one unpinned roll.

Decision it drives: two modes, and never one blurred mode. See §4.

2.3 Diversity comes from OMITTING features, not adding them — Swarm Testing (ISSTA 2012)

Swarm Testing, Groce, Zhang, Eide, Chen, Regehr. [PDF](#) · [ACM](#)

This is the paper that turns Will's *"a thousand different teams"* from a good instinct into a technique, and it says something counter-intuitive enough to be worth stating carefully.

Swarm testing generates a large population of configurations **each of which deliberately omits some features**, rather than one configuration that enables everything. In a week of testing it found **42% more distinct ways to crash a collection of C compilers than the heavily hand-tuned default configuration** of the same random tester.

The mechanism is the part that matters to us, and the paper gives it in two halves:

1. **Some features actively suppress the behaviour you want to reach.** Their example is a stack: including `pop` calls prevents the tester from ever driving the stack into its overflow-detection logic. **Our `pop` is Protect.** It is on almost every competitive set, it is the single most common click in the format, and a turn spent Protecting is a turn in which no damage, no secondary, no ability trigger and no field interaction occurs. A swarm that samples real teams uniformly will spend an enormous fraction of its turns testing nothing.
2. **Features compete for space.** Every slot spent on one mechanic is a slot not spent on another, so a uniform sample explores each mechanic shallowly. Omission frees the depth.

Decision it drives: team sampling is a *swarm over feature omission*, not a uniform draw from the meta. See §3.

2.4 A divergence is not actionable until it is minimal — Delta Debugging

Simplifying and Isolating Failure-Inducing Input, Zeller & Hildebrandt (the `ddmin` algorithm). [retrospective](#) · [The Fuzzing Book's treatment](#)

`ddmin` systematically removes parts of a failing input and re-runs, isolating a **1-minimal** subset — one where removing any single remaining element makes the failure disappear. Worst case is $O(n^2)$; in practice far better, and later work (HDD, Perses) exploits input structure to do better still.

Why we need it: a divergence found in turn 5 of a 4-vs-4 game with eight Pokémon, thirty-two moves and a field state is a *report*, not a *bug*. Nobody can act on it. Reduced to "these two Pokémon, this one move, turn 1", it is a WIRE.

Decision it drives: the reducer is part of the harness, not a follow-up. See §5.

3. Team selection — a swarm, not a sample

Teams come from `data/games.ladder.jsonl`, 46,612 real games with real items, spreads and movesets. Real teams matter because a generated team is a team nobody would bring, and the mechanics that matter are the ones people actually click.

But drawing 1,000 teams uniformly from the meta is the *heavily hand-tuned default configuration* that Swarm Testing beats by 42%. Every top team shares the same features — Protect on three of four, Fake Out, the same two weather setters — so a uniform draw tests those deeply and everything else never.

So: **sample a swarm of configurations, each omitting a feature class**, and draw teams that satisfy each. Candidate omission axes, each of which unblocks something it currently suppresses:

omit	what it lets the run finally reach
Protect	turns that actually resolve — damage, secondaries, contact abilities, item triggers
priority moves	the ordinary speed-order path, and dynamic re-sorting (WIRE 118)
weather setters	terrain, and the no-weather damage path
Intimidate	attack stages at their declared values, and the crit interaction in §6
KO-capable attackers	long games, which is where residual damage, sleep counters and field expiry live
single-target moves	spread damage, Wide Guard, redirection

OMISSION IS NOT THE ONLY AXIS, AND WILL NAMED THE CASE THAT PROVES IT — "*some moves hit thru protect, like feint or unseen fist.*" A mechanic that exists only as an *interaction with* a feature cannot be reached by omitting that feature. Drop Protect from every team and you have made Feint (397 uses), Phantom Force (424) and Unseen Fist untestable, because there is nothing left for them to punch through.

So the swarm needs **paired configurations** alongside the omitting ones: Protect present *and* concentrated with its breakers, redirection present *and* concentrated with the moves that ignore it, weather present *and* concentrated with the abilities that overwrite it. Swarm Testing's own framing supports this — the paper's claim is about *variance* in which features are present, not a monotonic preference for fewer. Omission is what buys **depth** in the features that remain; pairing is what buys the **interaction** surface, which is precisely where this project's bugs have actually lived (the interaction matrix found twelve that no single-mechanic probe could reach).

The check on whether the swarm is working is §5's coverage report, not intuition: if a configuration never produced a turn where a Protect was broken, it did not test breaking Protect, and it must say so rather than being counted as a run.

Report the swarm's composition in the artifact. A configuration that produced no games is a configuration that tested nothing, and must say so.

3.1 A blind spot the swarm must not inherit: we only read the BASE form's ability

Will, 2026-08-06: "you seem to miss a lot of mega abilities because you only read the base form ability." That is the root cause of two separate findings from the same evening, and it is one broken assumption rather than a list of missing tags.

uses in data/tags.json counts sheet-declared abilities. A team sheet declares what a Pokemon starts with. A mega's ability exists only after it evolves. So every ability that is mega-only reads uses: 0 , never clears the 99%-of-usage coverage bar, and is therefore never scheduled to be wired — by construction, silently, forever.

Measured: of the 63 entries in MEGA_ABIL , twelve read uses = 0 or are absent from tags.json entirely, and six have no mechanic derived at all — trace on Mega Alakazam and Mega Meowstic (207 declared uses, none of them the mega's), berserk on Mega Drampa, stalwart on Mega Skarmory, unseenfist on Mega Golurk, piercingdrill on Mega Excadrill. Fairy Aura on Mega Floette is the same bug and reaches 10.5% of ladder sides (ROADMAP #64); Unseen Fist is ROADMAP #69.

This is CLAUDE.md's prefer OBSERVED over DECLARED failing inside the instrument that decides what gets wired, which is the worst place for it: the bar cannot report a gap it is structurally unable to see.

Consequence for this harness. Two, and the second is the one that generalises:

- 1. Team sampling must draw from observed post-mega state, which the store already records directly — floettemega ab=Fairy Aura , golurkmega ab=Unseen Fist — rather than from sheet declarations. A swarm built on declared abilities will faithfully reproduce the blind spot.
- 2. The coverage report in §5 must count the ability a body actually had when it acted, not the one its sheet declared. Otherwise the harness reports "Trace exercised: 0" when it megaed into Trace on turn one and used it, and reads as a clean sweep of something it never touched — the exact failure §5 exists to prevent, arriving through the door §3 left open.

3.2 The swarm alone will never reach the fringe — directed swarm testing

Will, 2026-08-06: "we want swarms of all the different modes, but we also need to test the fringes, like upper hand, trick room (especially!) feint, ilusiion, disguise, i could go on and on."

He is right, and the arithmetic is unforgiving. Usage below is the uses field of data/tags.json (sheet-declared counts); the per-team and per-game figures later in this section come from the open-sheet corpora data/games.bo3.jsonl and data/games.ots.jsonl , and the swarm's own composition is data/diff-swarm.json . The two are different denominators and must not be mixed — the same mistake was made and corrected on 2026-08-06 when tags.json counts were quoted against tests/regulation_usage.js counts, which disagree by up to 13× (ROADMAP #70).

Re-derived 2026-08-06 (3.59.0) from the current data/tags.json (sheet_entries 119,616). The column below used to hold a stale snapshot and tests/test-docs-current.js could only see two cells of it — Trick Room and Wide Guard, whose old values 7,423 and 3,353 appear in no artifact. The other thirteen passed by coincidence: 74,245, 13,292, 424 and 76 all happen to occur in unrelated files under data/ . That is the honest reason this table is being rewritten whole rather than patched in two places.

tier	mechanic	uses	how rare, relative to Protect
dominant	Protect	80,328	the reference
	Fake Out	14,380	saturated
global	Trick Room	8,077	plenty of games, few of the interactions that matter
	Prankster	8,061	fine
	Sucker Punch	7,704	fine

tier	mechanic	uses	how rare, relative to Protect
	Wide Guard	3,568	fine
fringe	Phantom Force	465	~170x rarer
	Feint	437	~180x rarer
	Illusion	242	~330x rarer
	Trace	221	~360x rarer — and untagged , so 0
	Ally Switch	190	~420x rarer
	Quash	175	~460x rarer
	Instruct	173	~460x rarer
	Disguise	133	~600x rarer
	Upper Hand	80	~1,000x rarer

The last column changed meaning, and that is a correction rather than a cosmetic edit. It used to read "*what a uniform 1,000-game sample gets*" and gave counts — ~9, ~4, ~3, ~1.6. **The conversion factor from a declared use to an expected exercise is recorded nowhere in this document, in any artifact, or in any script**, and it is not `uses / games`, so those counts could not be reproduced by anyone including their author. They are replaced by a ratio that is arithmetic on the column beside it. The argument is unaffected and is in fact stated more plainly: **Upper Hand is three orders of magnitude rarer than Protect**, so a uniform swarm sized to exercise Protect will reach Upper Hand about once.

A handful of exercises is not coverage, it is an accident. And the fringe is where the bugs are: Upper Hand and Sucker Punch share a tag and have *different* conditions (ROADMAP #60); Illusion mis-attributes every move it disguises (#67); Instruct is absent by construction; Ally Switch fails on consecutive uses. Every one is a mechanic that decides a game outright when it fires.

TRICK ROOM IS A DIFFERENT CASE AND EARNS ITS EMPHASIS. At 8,077 uses it is not rare — a uniform swarm sees it constantly. It is dangerous because it is the only mechanic that **inverts the meaning of every speed comparison on the field for five turns**, and 3.49.0 made speed order *dynamic*, re-sorted before every action (WIRE 118). So Trick Room does not add one branch; it doubles the meaning of every existing ordering branch, and it multiplies with priority brackets (which it does **not** invert), with Prankster, with Quash and After You, with dynamic re-sorting, and with speed control that expires mid-turn. Sampling teams that *run* Trick Room is easy; reaching the ordering interactions *inside* it is the hard part, and no amount of uniform sampling does it on purpose.

The technique is directed swarm testing — Groce, Alipour, Zhang, Chen, Regehr, ISSTA 2016, *Generating Focused Random Tests Using Directed Swarm Testing* (PDF). Its result is the natural sequel to the 2012 paper: you can use the statistics a swarm run already produced to learn **which configurations correlate with hitting a rarely-covered target**, then bias generation toward those configurations to hit it far more often — without hand-writing a test for it.

So the harness runs a loop, not a pass:

1. **Swarm** for breadth — §3.
2. **Read the coverage report** — §5, which already counts exercises per census mechanic. This is not a new instrument; it is the one already required, used as the feedback signal.
3. **Direct** at every mechanic under its floor. Bias team sampling toward the configurations the swarm's own statistics show correlate with reaching it.
4. **Repeat until every mechanic clears its floor**, and report the floor each one cleared.

The floor is not uniform, and the rule for it already exists in this project. The coverage bar in `tests/test-medicham-coverage.js` carries Will's carve-out — *anything that can turn a CERTAINTY into a FAILURE, regardless of usage*. Upper Hand at 76 uses belongs to that class: a bot that thinks Upper Hand beats an Earthquake makes a confident, wrong, game-losing click. Usage sets the floor for ordinary mechanics; the carve-out sets it for the ones that flip outcomes.

This is also what makes Will's "*i could go on and on*" a non-problem. **The list must not be hand-written**, or it inherits every weakness of the hand-maintained ban list this project already replaced with a mechanism. The 235-row census IS the list; the coverage report says which rows are starved; the directed phase feeds them. Nobody has to remember Upper Hand.

3.3 DECIDED: the driver is scripted, and it is coverage-seeking rather than skilful

Will, 2026-08-06: *"yes i agree lets do scripted to make sure we test every move, item, ability, and mon."*

The two candidates were **replay** (feed both engines the decisions real humans made, reconstructed from the store's turn events) and **scripted** (drive both from a policy this harness controls). Scripted wins for Mode A, for a reason that is about isolation rather than convenience: Mode A's claim is *"the engines are deterministic functions of the same input, so any difference is a bug."* A replay driver puts a **reconstruction step** inside that claim — if the stored events do not pin a targeting choice or a switch unambiguously, a divergence might be the reconstruction rather than either engine, and the oracle weakens exactly where it is supposed to be absolute. Replay is worth adding *after* the harness is trusted, as a second corpus; it is the wrong thing to debug against first.

But "scripted" here does not mean a policy that plays well. It means a driver that plays to COVER. Will's phrasing is the specification — *every* move, item, ability and mon. So the action rule is:

at each decision, prefer the legal action that exercises the census mechanic furthest below its floor; break ties toward the entity (move / item / ability / species) least exercised so far.

This is §3.2's directed loop pulled inside the driver rather than bolted on around it. The swarm chooses *which teams take the field*; the driver chooses *which clicks happen* — and both are steered by the same coverage report. A skilful policy would be actively counterproductive: good play converges on a narrow set of strong lines, which is the uniform-sampling failure one level down.

Three consequences worth writing down before any code exists:

1. **A coverage-seeking driver will produce bad games, and that is correct.** It will click Quash into an empty slot and Trick Room on turn six. Nobody is scoring these games; they exist to make the two engines disagree.
2. **Legality is still the constraint.** The driver picks among *legal* actions only, or it tests a position the game cannot reach and any divergence is meaningless.
3. **Mode B keeps a separate driver.** Distribution comparison needs unbiased sampling over actions, not coverage-biased selection, or the rates being measured are the driver's and not the engine's.

And the fringe arrives in BUNDLES, which the driver must not break apart. Measured: **252 of 252 declared Quash carriers in the open-sheet corpus are Sableye with Prankster — 100%.**

Quash is priority 0 and therefore fails whenever the target has already moved; Prankster's +1 is the entire reason the move is playable. So "test Quash" without Prankster tests a configuration **nobody in 46,987 games has ever brought**. Real teams carry these bundles for free; a team generator would have scattered Quash across random bodies and produced only the version that does not exist. It is the strongest single argument in this document for §3's insistence that teams come from the store.

THE DRIVER AND THE SWARM COVER DIFFERENT SPACES, AND THE REPORT MUST NOT CONFLATE THEM. Will, following the Quash finding: *"i mean i guess a fast mon could use quash but we never see it."* True, and the reason is mechanical rather than fashion — **you Quash the thing that is faster than you.** If you were already faster you would simply move first and would not need the move, so the use case is inherently *"I am slower and must act first"*, which speed cannot solve by definition and Prankster's +1 can. The bundle is forced, not conventional.

The consequence is a trap: **a coverage-seeking driver will click Quash on a fast mon**, because it is legal and uncovered. That is *correct for engine testing* — the mechanic resolves and both engines must agree on it — but it means a coverage figure cannot be read as *"we tested what happens in real games."* The configuration it tested is one nobody has brought.

So the two instruments answer different questions and the artifact must report them separately:

	covers	answers
the driver (§3.3)	the mechanic space	does every move, item, ability and mon RESOLVE identically
the swarm (§3, §3.2)	the situation space	do the positions people actually reach resolve identically

A single line reading `Quash: covered` would be true and misleading at once — the same shape as the damage differential's 12% tolerance (§1) and the muzzled corpus contrast, both of which passed while proving nothing. Two columns, always.

4. Two modes, never blurred

This is the Csmith lesson applied literally: a comparison is only evidence where a difference is *definitionally* a bug.

Mode A — PINNED. Tests mechanics and sequencing. Tolerance zero.

Every die fixed identically in both engines: no accuracy misses, no crits, no chance secondaries, damage roll pinned to one index on both sides. In this mode the two engines are **deterministic functions of the same input**, so *any* difference is a bug — no tolerance, no statistics, no sampling error. This is where the WIRES will come from.

The one subtlety, already paid for once: pinning must be the **same** die. Showdown's `randomChance()` bypasses a `battle.random` override entirely and crits need `willCrit`; `tests/test-engine-diff.js` records both traps in its header and cost `engine/validate_damage_sim.js` two debugging rounds. Reuse that code rather than re-deriving it.

Mode B — ROLLED. Tests rates. Compares distributions.

Everything Mode A pins is exactly what Mode A cannot test: **the crit rate is 1/24, accuracy is a distribution, secondaries are a chance.** Pinning them off proves nothing about them. So Mode B runs many seeds and compares *distributions*, not individual games — a chi-square or a confidence interval on "how often did a crit land", not "did this turn match".

The two modes answer different questions and neither substitutes for the other. Reporting them as one number would be the 12%-tolerance mistake again, one level up.

5. Naming what is different — the harness must localise, not just detect

Will, 2026-08-06: "*make sure the harness identifies what exactly is different between showdown and medicham so we can isolate what is wired incorrectly.*" This is the requirement that separates a debugging instrument from a scoreboard, and it changes what gets compared.

Do not diff end-of-turn state. "Turn 4 HP differs by 12" is a symptom with a dozen possible causes — a damage multiplier, a missed residual, an item that should have triggered, a mis-ordered action — and it puts the reader back at the start.

Diff the EVENT STREAM, and use Showdown's own vocabulary as the common format. The official simulator already emits a structured, ordered protocol log:

```
|move|p1a: Incineroar|Fake Out|p2a: Garchomp
|-damage|p2a: Garchomp|154/175
|-ability|p1a: Incineroar|Intimidate|boost
|-unboost|p2a: Garchomp|atk|1
|-weather|Sandstorm|[upkeep]
|-damage|p1a: Incineroar|160/175|[from] Sandstorm
|-enditem|p2a: Garchomp|Focus Sash
```

That is not merely a log — it is a **step-level trace of every decision the authority made, already labelled with the mechanism that made it**. So: have MEDICHAM emit the same event shapes, align the two streams, and the **first differing event names the wired-wrong thing directly**. A missing `|-unboost|` is Intimidate. A `|-damage|` with the wrong number after identical preceding events is the damage math. An out-of-order `|move|` pair is turn order. An absent `|-enditem|` is the Sash.

This is the same reasoning that makes differential testing work at all (§2.1), applied one level down: the finer the granularity at which two implementations are compared, the more directly the difference points at its cause. McKeeman compared program *outputs*; comparing *traces* is strictly more informative when both sides can be made to emit one.

Each divergence is then classified and counted by class, not by instance. The report should read

turn order	12 games	3 distinct causes
residual damage	8 games	1 cause
item trigger	5 games	2 causes
stat stage	4 games	1 cause

because the point is to fix a class of wiring, and 12 instances of one turn-order bug is one WIRE, not twelve findings. This is how the arming pass worked: WIRE 129 was one root cause under six dead arms.

Then, and only then, the mechanical reduction:

- 1. First divergence per game only.** Everything after the first is downstream of it, and counting consequences as separate findings inflates the number and hides the cause.
- 2. The divergence, reduced.** Run `ddmin` over the game: drop turns, drop Pokémon, drop moves, drop items, keeping only what still diverges. Ship the 1-minimal repro, not the game — and because the trace already names the event, the reduction has a target to preserve rather than a vague "still fails" predicate, which is exactly the structure-aware reduction later delta-debugging work (HDD, Perses) exploits.

3. **The repro must be emitted as a runnable probe**, in the shape `tests/probe_red_demo.js` already takes, so a divergence converts directly into an armed census probe. The harness's output is then not a report anybody has to transcribe — it is the next test.
4. **The run's own mechanic coverage**. Which of the 235 census mechanics were actually exercised. **A run that never triggered Illusion has not tested Illusion** and must say so rather than reading as a clean sweep. This is the project's own standing rule — *a capability that cannot prove it ran is assumed broken* — pointed at the harness itself, and it is what makes Will's "every single mechanic" a measurement rather than an aspiration.
5. **The swarm composition**, per §3.
6. **A frozen engine release id**. It is a measurement; `engine/engine_release.js` and the photograph rule apply.

5a. STEP ONE IS BUILT — MEDICHAM emits the stream (3.58.0)

`engine/medicham2-browser.js` emits **36 protocol event types** on request (`battleInit(A, B, {trace: []})` , off by default; every emit site is a falsy test against a module-level sink that is null unless a caller asked for one).

The event list is **derived, not typed** — `engine/derive_protocol_events.js` reads Showdown's own `add()` call sites and writes `data/protocol-events.json` , whose `showdownEvents` , `emittedCount` , `notEmittedCount` and `partialCount` read **91 / 36 / 58 / 10** — every non-emitted event carries a written reason. Two gates fail the run: an INVENTED event (claimed and never emitted upstream) and an UNDECLARED one (emitted upstream and neither emitted nor explained here).

The scan covers `data/mods/champions/` **as well as** `sim/` , **and it had to.** `sim/battle-actions.ts:1800` emits `add('-supereffective', target)` — two fields — and `data/mods/champions/scripts.ts:271` overrides it with `add('-supereffective', target, Math.min(typeMod, 2))` — three. A trace built from `sim/SIM-PROTOCOL.md` alone would have carried the wrong shape on every super-effective hit **in this format**, forever, and the differ would have reported it as a divergence.

§6's worked example separates, and on one line. Measured, both arms, with the stat blocks aligned and both engines' dice pinned to the same die:

```
showdown  Intimidate 112/170   control 112/170   <- the crit ignored the -1
medicham2 Intimidate 130/170   control 111/170   <- it did not
    the two streams agree on all 15 lines before the divergence
    the FIRST differing line is the |-damage| itself
```

The control is the same scenario with Intimidate swapped for an inert ability, and each engine is compared against ITSELF across it. The obvious test — "do the two engines' numbers match" — is false for a reason that has nothing to do with crits, and finding that out is §5's first result:

MODE B IS NEEDED SOONER THAN §4 IMPLIES. *medicham2's damage range for Knock Off Incineroar → Snorlax is 57..67 — eleven integers, sampled uniformly — and Showdown rolls sixteen indices onto the same span, flooring each separately, so the multiplicities are unequal. The ENDPOINTS agree, which is all `tests/test-engine-diff.js` ever compares (`roll=0` against `min` , `roll=15` against `max`), so 149/150 is compatible with every one of the fourteen middle rolls being off by one or two **and** with every roll's probability being wrong. §1's table says the damage differential's 12% tolerance is free slack; this is the sharper statement — the comparison is endpoint-to-endpoint by construction and the interior is unmeasured.*

A second finding the state comparison cannot see. medicham2 resolves the knock-off, the resist berry and the contact punish **before** subtracting the target's HP, so its stream reads `-enditem`, the Rough Skin `-damage`, then the target's `-damage`; Showdown subtracts first. End-of-turn state is identical — which is exactly why `tests/test-game-diff.js` agrees on all five scripted games and the trace does not. That is §5's whole argument, demonstrated on the first night it existed.

Still to build: the comparison driver over two streams, the divergence classifier, `admin`, and the coverage report. This section records only that the stream exists and that it localises.

6. A worked example of what Mode A should catch on day one

Three parts of a critical hit are declared unmodelled in `engine/medicham2-browser.js`, confirmed against Showdown's `sim/battle-actions.ts:1683-1691`:

```
if (isCrit) { ignoreNegativeOffensive = true; ignorePositiveDefensive = true; }
const ignoreOffensive = !(move.ignoreOffensive || (ignoreNegativeOffensive && atkBoosts < 0));
const ignoreDefensive = !(move.ignoreDefensive || (ignorePositiveDefensive && defBoosts > 0));
```

A crit ignores the **attacker's negative** offensive stages, the **defender's positive** defensive stages, and screens. It does **not** ignore burn — those rules operate on boost *stages*, and burn is an `onModifyAtk` multiplier. Gen 2 ignored burn on a crit and Gen 3 onward does not, which is why that half is the one people misremember.

The first of the three is the expensive one here: **Intimidate is everywhere in this format**, so an Intimidated attacker landing a crit should hit at full Attack, and MEDICHAM prices it at `-1`. A `willCrit` move (Flower Trick, Storm Throw, Frost Breath) thrown by an Intimidated attacker in Mode A is a two-line scenario that separates the engines exactly — and no staged pair test we own is shaped to notice, because the census asks *does the mechanic fire, not is the number right afterwards*.

7. Cost, and what is honestly unknown

Building it is roughly a day: both engines already play games, `tests/test-engine-diff.js` already solves the die-pinning traps, and the store already holds the teams. Running it is cheap — MEDICHAM measures ~1,600–2,000 battles/sec, so 1,000 games is under a minute and the official engine is the slower side.

The fix cycle afterwards is not estimable, and saying otherwise would be a number to retract later. Nobody knows what it finds, because it is the first instrument that could see a sequencing bug at all. The base rate in this repository is bad: every deep instrument built so far found real bugs immediately — the arming pass found seven, the interaction matrix found twelve, the accuracy probe found that all four doors were dead. Csmith found 325 bugs in compilers that had been in production for decades and were tested far harder than this. Expect several WIRES.

Done looks like: divergences per 1,000 games is small, every survivor has a written reason, and the run's own coverage is high enough that "we tested it" is a measurement rather than a claim. Until then MEDICHAM's output is not trusted for decisions — which is exactly the bar Will set.

Sources

- McKeeman, *Differential Testing for Software* — <https://www.cs.tufts.edu/comp/15oFP/archive/bill-mckeeman/DifferentailTesting.pdf>

- Yang, Chen, Eide, Regehr, *Finding and Understanding Bugs in C Compilers* (PLDI 2011) — <https://users.cs.utah.edu/~regehr/papers/pldi11-preprint.pdf>
- Groce, Zhang, Eide, Chen, Regehr, *Swarm Testing* (ISSTA 2012) — <https://agroce.github.io/issta12.pdf>
- Groce, Alipour, Zhang, Chen, Regehr, *Generating Focused Random Tests Using Directed Swarm Testing* (ISSTA 2016) — <https://agroce.github.io/issta16.pdf>
- Zeller & Hildebrandt, *Simplifying and Isolating Failure-Inducing Input* (delta debugging / `ddmin`) — <https://www.computer.org/csdl/journal/ts/2025/03/10859156/23X97jMgYjm>
- Groce et al., *Randomized Differential Testing as a Prelude to Formal Verification* (ICSE 2007) — <https://agroce.github.io/icse07.pdf>
- Reducing failure-inducing inputs, The Fuzzing Book — <https://www.fuzzingbook.org/html/Reducer.html>