

The whole-game differential — design, and the research it is taken from

Version 3.71.0 · Last updated 2026-08-07

THE INSTRUMENT WAS MEASURING ANNOUNCEMENTS, AND THE HEADLINE IS NOW THE BOARD AT THE END OF TURN 1 (3.70.0). `engine/board_state.js` reads HP, status with its counters, items, all seven stat stages, aliveness, every field condition WITH ITS CLOCK and the persistent volatiles out of BOTH engines' live bodies at every turn boundary, after the whole residual phase. The ladder that stood here ran under `data/state-ladder.json`. **EVERY TURN-1 AND WHOLE-GAME BOARD-AGREEMENT FIGURE THAT STOOD HERE IS RETRACTED, 2026-09-06, AND IS NOT REPLACED.** Three measured reasons, any one sufficient: the artifact carries no `steering.driver_code`, so `engine/arms_comparable.js` answers UNKNOWN for every pair drawn from it and no work done today can repair that; its own `determinism.verdict` reads *"THE TWO BASELINES DISAGREE. Do not read the table below as a ladder."* — a red line no document quoting it has ever mentioned, and one the account below shows is PROBABLY an instrument false alarm that could not be settled, because the per-arm artifacts needed to settle it were not kept; and its sample cannot be re-taken, because the pool was drawn live from `data/games.bo3.jsonl` at `ff6529f6a6d7` and that file is a month of hourly appends further on. A re-run today would use a different pool, a 643-row census against its 252-row pin, a changed protocol skip list and four pinned arms where it had one — a different question, and a re-run answering a different question is worse than a withdrawal. Account: `docs/_reports/2026-09-06-comparability-restore.md`. The comparator proves itself first: 7 representation mappings red-demonstrated in both directions and 25 planted state divergences, each of which must be caught at the planted boundary and localised to the planted field — 25/25 on all fourteen arms.

THIS INSTRUMENT'S OUTPUT DOES NOT PREDICT THE LEAF, AND THAT IS NOW MEASURED (3.69.0). `data/leaf-engine-contrast.json` joins this file's first-divergence depth, per position, to MILTANK's per-position leaf error. **Every figure of that join is QUARANTINED — withheld, not annotated:** the artifact is downstream of MEDICHAM (`engine/leaf_engine_contrast.js` reaches `engine/medicham2-browser.js` through `require`), so the sample size, the rank correlations, the smallest detectable effect and the sign under the older engine are all absent, and it becomes quotable again when the gate opens AND this is re-run: `node engine/leaf_engine_contrast.js`. That is not an argument against the instrument — it measures conformance to the authority, which is worth having on its own terms and is what ADR-002 asks for. It is an argument against **spending further engine effort in order to move the leaf**, and it retires "line depth is the better steering number" as an unvalidated assumption: the two metrics behave identically here, and neither reads.

Two properties of this file were confirmed in the course of that measurement and belong here:

- **The driver's statefulness is load-bearing and is not a detail.** `CLICKS` and `COV_HITS` carry across games deliberately, so `playGame`'s divergence depth is a function of the position **and of every game played before it**. A per-position reproduction check re-running a 24-position slice in a fresh process disagreed on 16 of 24 — correctly, by construction. Any caller joining depth to a per-position quantity must control for run order or measure what it costs.
- **A reversed-order control quantifies that cost.** Same release, same positions, order reversed inside each shard: the two depth readings correlate but are not identical, so part of a depth reading is the run and not the position. How large that part is was measured on a superseded engine and is withheld until the control is re-run.

THE INSTRUMENT HAS BEEN TURNED ON ITS OWN HISTORY (3.68.o).

`engine/wire_ladder.js` replays every frozen release of the 2026-08-06/07 wire night through this driver under **one** pinned census and **one** team pool, so all nine arms are mutually comparable rather than only adjacent — `data/wire-ladder.json` . Three results bear directly on this design:

- **§5.3's median is the number that decides whether this approach terminates, and it did not move.** 1,995 games per arm, six wires, and the median game still parts after **one completed turn** at every rung. Whole-game agreement went 2 → 22 of 1,995. The turn count is too coarse to grade a rung at all; the useful unit is the PROTOCOL LINE, where the mean goes 15.01 → 23.97 and p90 goes 30 → 57.
- **A per-class delta is not readable on its own.** `-damage field 3 RISES 170 → 216` across WIRES 1-3 and then falls to 141: earlier wires push games deeper and expose divergences that were masked by an earlier stop. Fixing a bug can therefore make a class count go UP, and 141 of 1,995 games part EARLIER than the baseline after six correct fixes. §5's "a class is a wire, an instance is not" holds; "a smaller class count is progress" does not.
- **The determinism claim is now demonstrated over ten runs, not asserted.** The pre-WIRE-1 baseline ran first and last with eight arms between and reproduced byte for byte, including the per-game divergence depth, and the whole ladder was run three times end to end.

THE DIFFERENTIAL HAS RUN, AND MEGAS ARE IN IT (3.62.2). `engine/game_differential.js` plays a real stored team through MEDICHAM and through the official Showdown engine, step for step, against a stamped frozen release. **Read every figure from `data/game-differential.json` , never from this sentence** — the first version of this paragraph quoted a run that a later one replaced within the day, which is the drift this whole document set keeps having to correct.

At the time of writing it reports every measured game diverging, with the median parting after a single completed turn. **Mega bodies are now tested** (ROADMAP #31): no stone is stripped from the measured arm, and every mega choice Showdown offered was taken by both engines.

Two limits travel with any rate this instrument prints and must never be separated from it. Nothing past the first turn is exercised, because a game stops at its first divergence. And both sides are built Serious / 0 EVs / 31 IVs so the two engines compute the same stat line before *and* after a forme change — **this tests RULES, not the spreads the ladder actually brings.**

ROADMAP #68. Written 2026-08-06 from Will's specification: "*we would want to play n games with like a thousand different teams to really test every single mechanic in the game*", under his bar: "*i dont care if it takes a while, i just want it done correctly so we can trust it before using it to make all these decisions.*"

1. Why the instruments we already have cannot answer the question

Every part is proven and nothing proves the whole.

instrument	reads	what it establishes	what it cannot see
mechanics census	235 probed, 234 live, 235 armed, 0 unarmed, 0 directCall	each mechanic fires, and each probe was shown RED on a broken engine first	anything about sequencing
damage differential	149 / 150	damage matches Showdown on staged pairs	any error under its 12% tolerance , uncounted
interaction matrix	1,624 / 1,643	pairs of mechanics compose	still staged, still pairwise

instrument	reads	what it establishes	what it cannot see
data/game-diff.json	5 games	five hand-written scenarios, 6 turns each	its <code>not_compared</code> list excludes hp amounts, accuracy misses, chance secondaries, crits, any pair where one engine KOs and the other does not , the <code>protect</code> volatile, and Showdown-only volatiles

Two of those deserve their own note.

The damage differential's tolerance is not absorbing roll variance. `showdownDamage()` is called at `roll=0` and `roll=15` against MEDICHAM's min and max, so both engines emit a pinned 16-roll range and the comparison is endpoint-to-endpoint. **A correct engine scores 0.000 on every row.** The 12% is free slack, and the artifact publishes only `worst`, so whether the 149 passing rows sat at 0.000 or at 0.119 is recorded by nothing.

The 31.1-point win-probability disagreement quoted in `docs/SUMMARY.md` predates WIRES 123–132. `engine/rerun_list.js` classes it unquotable. We do not currently know what that gap is.

2. The research this design is taken from

Each citation drives one decision below. They are not decoration.

2.1 Differential testing itself — McKeeman (1998)

Differential Testing for Software, Bill McKeeman, Digital Technical Journal. [PDF](#)

The founding statement: when two or more implementations of the same specification exist, feed both the same mechanically generated input and compare. **The oracle is free.** This is a materially stronger position than fuzzing, which can only see crashes and hangs — a differential harness sees *wrong answers*, which is the entire class of bug this project keeps finding.

McKeeman's other contribution is the **input tier**: he tested C compilers across seven tiers, from random ASCII up to model-conforming C programs. Deeper tiers find deeper bugs. Our tiers are already built and we are standing on the third of four: staged pair → staged interaction → **real game** → real game with real decisions.

Decision it drives: the harness compares two engines rather than asserting against expectations. We already own the authority (`engine/champions_sim.js`, the official simulator) and ADR-002 already says Showdown is it.

2.2 The input must make a difference *definitionally* a bug — Csmith (PLDI 2011)

Finding and Understanding Bugs in C Compilers, Yang, Chen, Eide, Regehr. [preprint](#) · [ACM](#)

325+ previously unknown bugs over three years. **Every compiler they tested both crashed and silently generated wrong code on valid input.** The methodological core is not the random generator — it is that Csmith generates programs *avoiding undefined and unspecified behaviour*, because where the standard permits two answers, a difference is not evidence of a bug and the oracle collapses.

Our undefined behaviour is **the dice**. Where either engine may legally roll, a divergence proves nothing. This project has already paid for that lesson once: CHANGELOG 3.45.0 records that the interaction matrix's *"two pinned dice were not the same die, so every sub-100-accuracy move had been MISSING in the reference engine while medicham2 hit"* — a whole class of false disagreement from one unpinned roll.

Decision it drives: two modes, and never one blurred mode. See §4.

2.3 Diversity comes from OMITTING features, not adding them — Swarm Testing (ISSTA 2012)

Swarm Testing, Groce, Zhang, Eide, Chen, Regehr. [PDF](#) · [ACM](#)

This is the paper that turns Will's "*a thousand different teams*" from a good instinct into a technique, and it says something counter-intuitive enough to be worth stating carefully.

Swarm testing generates a large population of configurations **each of which deliberately omits some features**, rather than one configuration that enables everything. In a week of testing it found **42% more distinct ways to crash a collection of C compilers than the heavily hand-tuned default configuration** of the same random tester.

The mechanism is the part that matters to us, and the paper gives it in two halves:

1. **Some features actively suppress the behaviour you want to reach.** Their example is a stack: including `pop` calls prevents the tester from ever driving the stack into its overflow-detection logic. **Our `pop` is Protect.** It is on almost every competitive set, it is the single most common click in the format, and a turn spent Protecting is a turn in which no damage, no secondary, no ability trigger and no field interaction occurs. A swarm that samples real teams uniformly will spend an enormous fraction of its turns testing nothing.
2. **Features compete for space.** Every slot spent on one mechanic is a slot not spent on another, so a uniform sample explores each mechanic shallowly. Omission frees the depth.

Decision it drives: team sampling is a *swarm over feature omission*, not a uniform draw from the meta. See §3.

2.4 A divergence is not actionable until it is minimal — Delta Debugging

Simplifying and Isolating Failure-Inducing Input, Zeller & Hildebrandt (the `ddmin` algorithm). [retrospective](#) · [The Fuzzing Book's treatment](#)

`ddmin` systematically removes parts of a failing input and re-runs, isolating a **1-minimal** subset — one where removing any single remaining element makes the failure disappear. Worst case is $O(n^2)$; in practice far better, and later work (HDD, Perses) exploits input structure to do better still.

Why we need it: a divergence found in turn 5 of a 4-vs-4 game with eight Pokémon, thirty-two moves and a field state is a *report*, not a *bug*. Nobody can act on it. Reduced to "these two Pokémon, this one move, turn 1", it is a WIRE.

Decision it drives: the reducer is part of the harness, not a follow-up. See §5.

3. Team selection — a swarm, not a sample

Teams come from `data/games.ladder.jsonl` — the whole ladder store, real games with real items, spreads and movesets. Real teams matter because a generated team is a team nobody would bring, and the mechanics that matter are the ones people actually click.

But drawing 1,000 teams uniformly from the meta is the *heavily hand-tuned default configuration* that Swarm Testing beats by 42%. Every top team shares the same features — Protect on three of four, Fake Out, the same two weather setters — so a uniform draw tests those deeply and everything else never.

So: **sample a swarm of configurations, each omitting a feature class**, and draw teams that satisfy each. Candidate omission axes, each of which unblocks something it currently suppresses:

omit	what it lets the run finally reach
Protect	turns that actually resolve — damage, secondaries, contact abilities, item triggers

omit	what it lets the run finally reach
priority moves	the ordinary speed-order path, and dynamic re-sorting (WIRE 118)
weather setters	terrain, and the no-weather damage path
Intimidate	attack stages at their declared values, and the crit interaction in §6
KO-capable attackers	long games, which is where residual damage, sleep counters and field expiry live
single-target moves	spread damage, Wide Guard, redirection

OMISSION IS NOT THE ONLY AXIS, AND WILL NAMED THE CASE THAT PROVES IT — *"some moves hit thru protect, like feint or unseen fist."* A mechanic that exists only as an *interaction with* a feature cannot be reached by omitting that feature. Drop Protect from every team and you have made Feint (397 uses), Phantom Force (424) and Unseen Fist untestable, because there is nothing left for them to punch through.

So the swarm needs **paired configurations** alongside the omitting ones: Protect present *and* concentrated with its breakers, redirection present *and* concentrated with the moves that ignore it, weather present *and* concentrated with the abilities that overwrite it. Swarm Testing's own framing supports this — the paper's claim is about *variance* in which features are present, not a monotonic preference for fewer. Omission is what buys **depth** in the features that remain; pairing is what buys the **interaction** surface, which is precisely where this project's bugs have actually lived (the interaction matrix found twelve that no single-mechanic probe could reach).

The check on whether the swarm is working is §5's coverage report, not intuition: if a configuration never produced a turn where a Protect was broken, it did not test breaking Protect, and it must say so rather than being counted as a run.

Report the swarm's composition in the artifact. A configuration that produced no games is a configuration that tested nothing, and must say so.

3.1 A blind spot the swarm must not inherit: we only read the BASE form's ability

Will, 2026-08-06: *"you seem to miss a lot of mega abilities because you only read the base form ability."* That is the root cause of two separate findings from the same evening, and it is one broken assumption rather than a list of missing tags.

`uses` in `data/tags.json` counts **sheet-declared** abilities. A team sheet declares what a Pokemon starts with. A mega's ability exists only **after** it evolves. So every ability that is mega-only reads `uses: 0`.

That zero is what makes it silent: the coverage bar in `tests/test-medicham-coverage.js` is a usage-weighted THRESHOLD, not a reading out of the artifact, and a `uses: 0` entity can never clear it — so a mega-only ability is never scheduled to be wired, **by construction, silently, forever.** (2026-08-10: the threshold sentence was split off its own paragraph. It states a bar, and standing in a block that cites `data/tags.json` made `tests/test-docs-current.js` read it as a figure ATTRIBUTED to that artifact — which it never was. The check was right that the attribution was wrong; the number was not.)

Measured: of the 63 entries in `MEGA_ABIL`, twelve read `uses = 0` or are absent from `tags.json` entirely, and six have no mechanic derived at all — `trace` on Mega Alakazam and Mega Meowstic (207 declared uses, none of them the mega's), `berserk` on Mega Drampa, `stalwart` on Mega Skarmory, `unseenfist` on Mega Golurk, `piercingdrill` on Mega Excadrill. Fairy Aura on Mega Floette is the same bug and reaches **10.5% of ladder sides** (ROADMAP #64); Unseen Fist is ROADMAP #69.

This is CLAUDE.md's *prefer OBSERVED over DECLARED* failing **inside the instrument that decides what gets wired**, which is the worst place for it: the bar cannot report a gap it is structurally unable to see.

Consequence for this harness. Two, and the second is the one that generalises:

1. Team sampling must draw from **observed post-mega state**, which the store already records directly — `floettemega ab=Fairy Aura`, `golurkmega ab=Unseen Fist` — rather than from sheet declarations. A swarm built on declared abilities will faithfully reproduce the blind spot.
2. **The coverage report in §5 must count the ability a body actually had when it acted**, not the one its sheet declared. Otherwise the harness reports "Trace exercised: 0" when it megaed into Trace on turn one and used it, and reads as a clean sweep of something it never touched — the exact failure §5 exists to prevent, arriving through the door §3 left open.

3.2 The swarm alone will never reach the fringe — directed swarm testing

Will, 2026-08-06: *"we want swarms of all the different modes, but we also need to test the fringes, like upper hand, trick room (especially!) feint, ilusiion, disguise, i could go on and on."*

He is right, and the arithmetic is unforgiving. Usage below is the `uses` field of `data/tags.json` (sheet-declared counts); the per-team and per-game figures later in this section come from the open-sheet corpora `data/games.bo3.jsonl` and `data/games.ots.jsonl`, and the swarm's own composition is `data/diff-swarm.json`. **The two are different denominators and must not be mixed** — the same mistake was made and corrected on 2026-08-06 when `tags.json` counts were quoted against `tests/regulation_usage.js` counts, which disagree by up to 13× (ROADMAP #70).

Re-derived 2026-08-06 (3.62.2) from the current `data/tags.json` (`sheet_entries` 119,616). The column below used to hold a stale snapshot and `tests/test-docs-current.js` could only see two cells of it — Trick Room and Wide Guard, whose old values 7,423 and 3,353 appear in no artifact. **The other thirteen passed by coincidence:** 74,245, 13,292, 424 and 76 all happen to occur in unrelated files under `data/`. That is the honest reason this table is being rewritten whole rather than patched in two places.

tier	mechanic	uses	how rare, relative to Protect
dominant	Protect	80,328	the reference
	Fake Out	14,380	saturated
global	Trick Room	8,077	plenty of games, few of the interactions that matter
	Prankster	8,061	fine
	Sucker Punch	7,704	fine
	Wide Guard	3,568	fine
fringe	Phantom Force	465	~170x rarer
	Feint	437	~180x rarer
	Illusion	242	~330x rarer
	Trace	221	~360x rarer — and untagged , so 0
	Ally Switch	190	~420x rarer
	Quash	175	~460x rarer
	Instruct	173	~460x rarer
	Disguise	133	~600x rarer
	Upper Hand	80	~1,000x rarer

The last column changed meaning, and that is a correction rather than a cosmetic edit. It used to read *"what a uniform 1,000-game sample gets"* and gave counts — ~9, ~4, ~3, ~1.6. **The conversion factor from a declared use to an expected exercise is recorded nowhere in this document, in any artifact, or in any script**, and it is not `uses / games`, so those counts could not be reproduced by anyone including their author. They are replaced by a ratio that is arithmetic on the column beside it. The argument is unaffected and is in fact stated more plainly: **Upper Hand is three orders of magnitude rarer than Protect**, so a uniform swarm sized to exercise Protect will reach Upper Hand about once.

A handful of exercises is not coverage, it is an accident. And the fringe is where the bugs are: Upper Hand and Sucker Punch share a tag and have *different* conditions (ROADMAP #60); Illusion mis-attributes every move it disguises (#67); Instruct is absent by construction; Ally Switch fails on consecutive uses. Every one is a mechanic that decides a game outright when it fires.

TRICK ROOM IS A DIFFERENT CASE AND EARNS ITS EMPHASIS. At 8,077 uses it is not rare — a uniform swarm sees it constantly. It is dangerous because it is the only mechanic that **inverts the meaning of every speed comparison on the field for five turns**, and 3.49.0 made speed order *dynamic*, re-sorted before every action (WIRE 118). So Trick Room does not add one branch; it doubles the meaning of every existing ordering branch, and it multiplies with priority brackets (which it does **not** invert), with Prankster, with Quash and After You, with dynamic re-sorting, and with speed control that expires mid-turn. Sampling teams that *run* Trick Room is easy; reaching the ordering interactions *inside* it is the hard part, and no amount of uniform sampling does it on purpose.

The technique is directed swarm testing — Groce, Alipour, Zhang, Chen, Regehr, ISSTA 2016, *Generating Focused Random Tests Using Directed Swarm Testing* (PDF). Its result is the natural sequel to the 2012 paper: you can use the statistics a swarm run already produced to learn **which configurations correlate with hitting a rarely-covered target**, then bias generation toward those configurations to hit it far more often — without hand-writing a test for it.

So the harness runs a loop, not a pass:

1. **Swarm** for breadth — §3.
2. **Read the coverage report** — §5, which already counts exercises per census mechanic. This is not a new instrument; it is the one already required, used as the feedback signal.
3. **Direct** at every mechanic under its floor. Bias team sampling toward the configurations the swarm's own statistics show correlate with reaching it.
4. **Repeat until every mechanic clears its floor**, and report the floor each one cleared.

The floor is not uniform, and the rule for it already exists in this project. The coverage bar in `tests/test-medicham-coverage.js` carries Will's carve-out — *anything that can turn a CERTAINTY into a FAILURE, regardless of usage*. Upper Hand at 76 uses belongs to that class: a bot that thinks Upper Hand beats an Earthquake makes a confident, wrong, game-losing click. Usage sets the floor for ordinary mechanics; the carve-out sets it for the ones that flip outcomes.

This is also what makes Will's "*i could go on and on*" a non-problem. **The list must not be hand-written**, or it inherits every weakness of the hand-maintained ban list this project already replaced with a mechanism. The 235-row census IS the list; the coverage report says which rows are starved; the directed phase feeds them. Nobody has to remember Upper Hand.

3.3 DECIDED: the driver is scripted, and it is coverage-seeking rather than skilful

Will, 2026-08-06: "*yes i agree lets do scripted to make sure we test every move, item, ability, and mon.*"

The two candidates were **replay** (feed both engines the decisions real humans made, reconstructed from the store's turn events) and **scripted** (drive both from a policy this harness controls). Scripted wins for Mode A, for a reason that is about isolation rather than convenience: Mode A's claim is "*the engines are deterministic functions of the same input, so any difference is a bug.*" A replay driver puts a **reconstruction step** inside that claim — if the stored events do not pin a targeting choice or a switch unambiguously, a divergence might be the reconstruction rather than either engine, and the oracle weakens exactly where it is supposed to be absolute. Replay is worth adding *after* the harness is trusted, as a second corpus; it is the wrong thing to debug against first.

But "scripted" here does not mean a policy that plays well. It means a driver that plays to COVER. Will's phrasing is the specification — *every* move, item, ability and mon. So the action rule is:

at each decision, prefer the legal action that exercises the census mechanic furthest below its floor; break ties toward the entity (move / item / ability / species) least exercised so far.

This is §3.2's directed loop pulled inside the driver rather than bolted on around it. The swarm chooses *which teams take the field*; the driver chooses *which clicks happen* — and both are steered by the same coverage report. A skilful policy would be actively counterproductive: good play converges on a narrow set of strong lines, which is the uniform-sampling failure one level down.

Three consequences worth writing down before any code exists:

- 1. **A coverage-seeking driver will produce bad games, and that is correct.** It will click Quash into an empty slot and Trick Room on turn six. Nobody is scoring these games; they exist to make the two engines disagree.
- 2. **Legality is still the constraint.** The driver picks among *legal* actions only, or it tests a position the game cannot reach and any divergence is meaningless.
- 3. **Mode B keeps a separate driver.** Distribution comparison needs unbiased sampling over actions, not coverage-biased selection, or the rates being measured are the driver's and not the engine's.

And the fringe arrives in BUNDLES, which the driver must not break apart. (*Corpus sizes are deliberately not printed in this document. Two hand-typed store counts here had already drifted apart from each other and from the file — node engine/status.js prints the live figure.*) Measured: **252 of 252** declared Quash carriers in the open-sheet corpus are Sableye with Prankster — **100%**. Quash is priority 0 and therefore fails whenever the target has already moved; Prankster's +1 is the entire reason the move is playable. So "test Quash" without Prankster tests a configuration **that nobody in the entire ladder store has ever brought**. Real teams carry these bundles for free; a team generator would have scattered Quash across random bodies and produced only the version that does not exist. It is the strongest single argument in this document for §3's insistence that teams come from the store.

THE DRIVER AND THE SWARM COVER DIFFERENT SPACES, AND THE REPORT MUST NOT CONFLATE THEM. Will, following the Quash finding: "*i mean i guess a fast mon could use quash but we never see it.*" True, and the reason is mechanical rather than fashion — **you Quash the thing that is faster than you**. If you were already faster you would simply move first and would not need the move, so the use case is inherently "*I am slower and must act first*", which speed cannot solve by definition and Prankster's +1 can. The bundle is forced, not conventional.

The consequence is a trap: **a coverage-seeking driver will click Quash on a fast mon**, because it is legal and uncovered. That is *correct for engine testing* — the mechanic resolves and both engines must agree on it — but it means a coverage figure cannot be read as "*we tested what happens in real games.*" The configuration it tested is one nobody has brought.

So the two instruments answer different questions and the artifact must report them separately:

	covers	answers
the driver (§3.3)	the mechanic space	does every move, item, ability and mon RESOLVE identically
the swarm (§3, §3.2)	the situation space	do the positions people actually reach resolve identically

A single line reading Quash: covered would be true and misleading at once — the same shape as the damage differential's 12% tolerance (§1) and the muzzled corpus contrast, both of which passed while proving nothing. Two columns, always.

4. Two modes, never blurred

This is the Csmith lesson applied literally: a comparison is only evidence where a difference is *definitionally* a bug.

Mode A — PINNED. Tests mechanics and sequencing. Tolerance zero.

Every die fixed identically in both engines: no accuracy misses, no crits, no chance secondaries, damage roll pinned to one index on both sides. In this mode the two engines are **deterministic functions of the same input**, so *any* difference is a bug — no tolerance, no statistics, no sampling error. This is where the WIRES will come from.

The one subtlety, already paid for once: pinning must be the **same** die. Showdown's `randomChance()` bypasses a `battle.random` override entirely and crits need `willCrit`; `tests/test-engine-diff.js` records both traps in its header and cost `engine/validate_damage_sim.js` two debugging rounds. Reuse that code rather than re-deriving it.

Mode B — ROLLED. Tests rates. Compares distributions.

Everything Mode A pins is exactly what Mode A cannot test: **the crit rate is 1/24, accuracy is a distribution, secondaries are a chance**. Pinning them off proves nothing about them. So Mode B runs many seeds and compares *distributions*, not individual games — a chi-square or a confidence interval on "how often did a crit land", not "did this turn match".

The two modes answer different questions and neither substitutes for the other. Reporting them as one number would be the 12%-tolerance mistake again, one level up.

5. Naming what is different — the harness must localise, not just detect

Will, 2026-08-06: *"make sure the harness identifies what exactly is different between showdown and medicham so we can isolate what is wired incorrectly."* This is the requirement that separates a debugging instrument from a scoreboard, and it changes what gets compared.

Do not diff end-of-turn state. "Turn 4 HP differs by 12" is a symptom with a dozen possible causes — a damage multiplier, a missed residual, an item that should have triggered, a mis-ordered action — and it puts the reader back at the start.

THE PARAGRAPH ABOVE IS SUPERSEDED BY §5b (3.70.0) AND IS KEPT BECAUSE IT WAS ACTED ON. *It is right that a state diff is a poor localiser and wrong that it should therefore not exist. Taken as "compare the stream INSTEAD of the board" it left the project with no board comparison at all, and ten wires were aimed with an instrument that could not see whether a fix changed anything a search can act on. The correction is not to replace one with the other: the stream localises, the board decides whether there was anything to localise. Both run, on the same games, in the same process.*

Diff the EVENT STREAM, and use Showdown's own vocabulary as the common format. The official simulator already emits a structured, ordered protocol log:

```
|move|p1a: Incineroar|Fake Out|p2a: Garchomp
|-damage|p2a: Garchomp|154/175
|-ability|p1a: Incineroar|Intimidate|boost
|-unboost|p2a: Garchomp|atk|1
|-weather|Sandstorm|[upkeep]
|-damage|p1a: Incineroar|160/175|[from] Sandstorm
|-enditem|p2a: Garchomp|Focus Sash
```

That is not merely a log — it is a **step-level trace of every decision the authority made, already labelled with the mechanism that made it**. So: have MEDICHAM emit the same event shapes, align the two streams, and the **first differing event names the wired-wrong thing directly**. A missing `|-unboost|` is Intimidate. A `|-damage|` with the wrong number after identical preceding events is the damage math. An out-of-order `|move|` pair is turn order. An absent `|-enditem|` is the Sash.

This is the same reasoning that makes differential testing work at all (§2.1), applied one level down: the finer the granularity at which two implementations are compared, the more directly the difference points at its cause. McKeeman compared program *outputs*; comparing *traces* is strictly more informative when both sides can be made to emit one.

Each divergence is then classified and counted by class, not by instance. The report should read

turn order	12 games	3 distinct causes
residual damage	8 games	1 cause
item trigger	5 games	2 causes
stat stage	4 games	1 cause

because the point is to fix a class of wiring, and 12 instances of one turn-order bug is one WIRE, not twelve findings. This is how the arming pass worked: WIRE 129 was one root cause under six dead arms.

Then, and only then, the mechanical reduction:

1. **First divergence per game only.** Everything after the first is downstream of it, and counting consequences as separate findings inflates the number and hides the cause.
2. **The divergence, reduced.** Run `ddmin` over the game: drop turns, drop Pokémon, drop moves, drop items, keeping only what still diverges. Ship the 1-minimal repro, not the game — and because the trace already names the event, the reduction has a target to preserve rather than a vague "still fails" predicate, which is exactly the structure-aware reduction later delta-debugging work (HDD, Perses) exploits.
3. **The repro must be emitted as a runnable probe**, in the shape `tests/probe_red_demo.js` already takes, so a divergence converts directly into an armed census probe. The harness's output is then not a report anybody has to transcribe — it is the next test.
4. **The run's own mechanic coverage.** Which of the 235 census mechanics were actually exercised. **A run that never triggered Illusion has not tested Illusion** and must say so rather than reading as a clean sweep. This is the project's own standing rule — *a capability that cannot prove it ran is assumed broken* — pointed at the harness itself, and it is what makes Will's "every single mechanic" a measurement rather than an aspiration.
5. **The swarm composition**, per §3.
6. **A frozen engine release id.** It is a measurement; `engine/engine_release.js` and the photograph rule apply.

5a. STEP ONE IS BUILT — MEDICHAM emits the stream (3.58.0)

`engine/medicham2-browser.js` emits **44 protocol event types** on request (`battleInit(A, B, {trace: []})`, off by default; every emit site is a falsy test against a module-level sink that is null unless a caller asked for one).

The event list is **derived, not typed** — `engine/derive_protocol_events.js` reads Showdown's own `add()` call sites and writes `data/protocol-events.json`, whose `showdownEvents`, `emittedCount`, `notEmittedCount` and `partialCount` read **91 / 44 / 50 / 10** — every non-emitted event carries a written reason. (*Those four move. `-mega` and `detailschange` were on the declared-not-emitted list with the reason "mega evolution happens before `battleInit`, so there is no in-battle event to emit", and they moved into `TRACE_EVENTS` on 2026-08-07 when mega evolution became a mid-turn choice — read the counts from the artifact, never from this sentence.*) Two gates fail the run: an INVENTED event (claimed and never emitted upstream) and an UNDECLARED one (emitted upstream and neither emitted nor explained here).

The scan covers `data/mods/champions/` **as well as** `sim/`, **and it had to.** `sim/battle-actions.ts:1800` emits `add('-supereffective', target)` — two fields — and `data/mods/champions/scripts.ts:271` overrides it with `add('-supereffective', target, Math.min(typeMod, 2))` — three. A trace built from `sim/SIM-PROTOCOL.md` alone would have carried the wrong shape on every super-effective hit **in this format**, forever, and the differ would have reported it as a divergence.

§6's worked example separates, and on one line. Measured, both arms, with the stat blocks aligned and both engines' dice pinned to the same die:

```
showdown  Intimidate 112/170   control 112/170   <- the crit ignored the -1
medicham2 Intimidate 130/170   control 111/170   <- it did not
    the two streams agree on all 15 lines before the divergence
    the FIRST differing line is the |-damage| itself
```

The control is the same scenario with Intimidate swapped for an inert ability, and each engine is compared against ITSELF across it. The obvious test — "do the two engines' numbers match" — is false for a reason that has nothing to do with crits, and finding that out is §5's first result:

MODE B IS NEEDED SOONER THAN §4 IMPLIES. *medicham2's damage range for Knock Off Incineroar → Snorlax is `57..67` — **eleven integers, sampled uniformly** — and Showdown rolls **sixteen** indices onto the same span, flooring each separately, so the multiplicities are unequal. The **ENDPOINTS** agree, which is all `tests/test-engine-diff.js` ever compares (`roll=0` against `min`, `roll=15` against `max`), so 149/150 is compatible with every one of the fourteen middle rolls being off by one or two **and** with every roll's probability being wrong. §1's table says the damage differential's 12% tolerance is free slack; this is the sharper statement — the comparison is endpoint-to-endpoint by construction and the interior is unmeasured.*

A second finding the state comparison cannot see. `medicham2` resolves the knock-off, the resist berry and the contact punish **before** subtracting the target's HP, so its stream reads `-enditem`, the Rough Skin `-damage`, then the target's `-damage`; Showdown subtracts first. End-of-turn state is identical — which is exactly why `tests/test-game-diff.js` agrees on all five scripted games and the trace does not. That is §5's whole argument, demonstrated on the first night it existed.

Still to build: the comparison driver over two streams, the divergence classifier, `ddmin`, and the coverage report. This section records only that the stream exists and that it localises.

5b. THE BOARD AT THE TURN BOUNDARY — the correction to §5 (3.70.0)

Will: "ALL WE CARE ABOUT IS HP/ STATUS, ITEMS, MONS ALIVE AT THE END OF THE TURN? ANNOUNCING IT DOESNT REALLY MATTER RIGHT?", then "AND STAT BOOSTS/ DROPS", "AND FIELD CONDITIONS", "TRAPPED STATUS", "LEFTOVERS FIRING". Then, after a more correct engine failed to predict better: **"I ONLY CARE ABOUT TURN 1 TO START."**

The two instruments answer different questions and neither substitutes for the other.

	the protocol stream (§5)	the board (engine/board_state.js)
asks	do the two engines tell the same story	do they reach the same position
good at	naming the wired-wrong mechanism	saying whether anything a search can act on changed
blind to	whether a difference matters	which of a dozen causes produced it

WHAT IS COMPARED. Per active body: HP, status with the toxic stage and turns slept, item, all seven stat stages, fainted, species. Per side: every party member's HP and aliveness, hazards as LAYER COUNTS, each screen with its counter, Tailwind with its counter. Field: weather, terrain and Trick Room, each with its counter. Volatiles: Substitute with its HP, Taunt, Encore, Disable, Leech Seed, confusion, Perish, and MOVE trapping with its counter. **Counters are not optional — Tailwind with 3 turns left is not Tailwind with 1, and a search that cannot tell them apart plans the wrong turn.**

WHERE THE BOUNDARY IS DECIDES WHAT IS MEASURED. It is taken **after the entire residual phase**, at the instant the next turn's choices are made. Leftovers, burn and poison chip, Toxic's escalating stage, sand and hail, Leech Seed draining off two bodies, Perish counting down, Wish landing and every field counter ticking only touch the board there. A boundary taken earlier measures a board nobody plays from.

REAL STATE, NEVER NARRATION. medicham2 is read off `S.actA[i].curHP` , `.boosts` , `._sub` , `S.field.weatherT` , `S.sfa.sc` ; Showdown off `battle.sides[i].active[j]` , `side.sideConditions` , `battle.field` . **Neither engine's log is opened.** Deriving a board from the protocol would reproduce the original defect one level down: if our stream omits an `-enditem` , a stream-derived board keeps the item and reports a divergence the ENGINES do not have.

TURN 1 IS THE HEADLINE, AND THE MEDIAN WAS THE WRONG STATISTIC. Turn 1 is the only turn that begins from a board both engines agree on; every later turn begins from wherever the run had already drifted, so a pooled per-turn rate is contaminated by earlier error and hides which turn moved. Both are published — `agreement_by_turn` states its denominator at each turn, `turn_boundary_agreement` pools and is kept beside it so the contamination is visible. The median first-divergence turn read 1 at all ten rungs and was reported ten times as "nothing moved"; the distribution is **bimodal** and a median cannot move on one until half the mass crosses.

MAGNITUDE IS BUCKETED, BECAUSE OFF-BY-ONE AND OFF-BY-FORTY ARE DIFFERENT WORK. An HP off by one is the residue of §WIRE 4's fixed-point arithmetic; an HP off by forty is a missing mechanic. Lumping them is how `-damage field 3` stayed one opaque class for ten wires.

REPRESENTATION MAPPINGS ARE DECLARED AND RED-DEMONSTRATED. Two engines can hold one fact in two shapes — `at / atk` , `floette-eternal / floetteeternal` , a hazard as presence or as layers, an absent key as zero, `sand / sandstorm` , a faint as `fainted:true` or as `status:'fnt'` , sleep counted up or down. Each is declared with a pair it MUST collapse and a pair it must NOT, and both directions run before any board is read. A mapping that collapses its `distinct` pair is a **silencer** and the run says so. The weather and terrain translation is **the engine's own exported function**, not a second copy of it.

WHAT IS DELIBERATELY NOT COMPARED ships with every artifact as `NOT_COMPARED` : ability trapping (medicham2 stores no flag and evaluates the tag at switch time — a comparator would have to reimplement its rule and would then be checking its own belief), item `DISPOSITION`, PP, and the stall counter behind consecutive Protect. An absent field reads exactly like an agreeing one, so each is named with its reason.

THE COMPARATOR PROVES ITSELF FIRST. 25 planted state divergences, one per compared field family, written into the LIVE medicham board at a boundary the clean arm agreed at. Each must be APPLIED (an unapplied plant reads exactly like a comparator that found nothing), CAUGHT, at EXACTLY that boundary, and LOCALISED to the planted field. Gated by `tests/test-state-differential.js`, which also proves the party is keyed by species rather than index — Showdown reorders `side.pokemon` on every switch-in, and index matching manufactured 123 party divergences in 179 games.

6. A worked example of what Mode A should catch on day one

Three parts of a critical hit are declared unmodelled in `engine/medicham2-browser.js`, confirmed against Showdown's `sim/battle-actions.ts:1683-1691` :

```
if (isCrit) { ignoreNegativeOffensive = true; ignorePositiveDefensive = true; }
const ignoreOffensive = !(move.ignoreOffensive || (ignoreNegativeOffensive && atkBoosts < 0));
const ignoreDefensive = !(move.ignoreDefensive || (ignorePositiveDefensive && defBoosts > 0));
```

A crit ignores the **attacker's negative** offensive stages, the **defender's positive** defensive stages, and screens. It does **not** ignore burn — those rules operate on boost *stages*, and burn is an `onModifyAtk` multiplier. Gen 2 ignored burn on a crit and Gen 3 onward does not, which is why that half is the one people misremember.

The first of the three is the expensive one here: **Intimidate is everywhere in this format**, so an Intimidated attacker landing a crit should hit at full Attack, and MEDICHAM prices it at -1. A `willCrit` move (Flower Trick, Storm Throw, Frost Breath) thrown by an Intimidated attacker in Mode A is a two-line scenario that separates the engines exactly — and no staged pair test we own is shaped to notice, because the census asks *does the mechanic fire*, not *is the number right afterwards*.

7. Cost, and what is honestly unknown

Building it is roughly a day: both engines already play games, `tests/test-engine-diff.js` already solves the die-pinning traps, and the store already holds the teams. Running it is cheap, and the official engine is the slower side.

The throughput sentence here read ~1,600–2,000 battles/sec for MEDICHAM and is superseded. Re-measured 2026-09-08 on this box, 40 pinned team pairs, both engines interleaved in one process: MEDICHAM **134 whole games/sec** and Showdown's raw `Battle` **33.7** — so 1,000 games is about 8 seconds of play on MEDICHAM and about 30 on Showdown. Those are PLAY time only. This instrument also aligns two protocol logs per game, which is in neither figure, so do not read them as a runtime estimate for the differential itself. Artifacts and pins: `data/verification/speed-2026-09-08/`.

The fix cycle afterwards is not estimable, and saying otherwise would be a number to retract later. Nobody knows what it finds, because it is the first instrument that could see a sequencing bug at all. The base rate in this repository is bad: every deep instrument built so far found real bugs immediately — the

arming pass found seven, the interaction matrix found twelve, the accuracy probe found that all four doors were dead. Csmith found 325 bugs in compilers that had been in production for decades and were tested far harder than this. Expect several WIRES.

Done looks like: divergences per 1,000 games is small, every survivor has a written reason, and the run's own coverage is high enough that "we tested it" is a measurement rather than a claim. Until then MEDICHAM's output is not trusted for decisions — which is exactly the bar Will set.

Sources

- McKeeman, *Differential Testing for Software* — <https://www.cs.tufts.edu/comp/150FP/archive/bill-mckeeman/DifferentailTesting.pdf>
- Yang, Chen, Eide, Regehr, *Finding and Understanding Bugs in C Compilers* (PLDI 2011) — <https://users.cs.utah.edu/~regehr/papers/pldi11-preprint.pdf>
- Groce, Zhang, Eide, Chen, Regehr, *Swarm Testing* (ISSTA 2012) — <https://agroce.github.io/issta12.pdf>
- Groce, Alipour, Zhang, Chen, Regehr, *Generating Focused Random Tests Using Directed Swarm Testing* (ISSTA 2016) — <https://agroce.github.io/issta16.pdf>
- Zeller & Hildebrandt, *Simplifying and Isolating Failure-Inducing Input* (delta debugging / `ddmin`) — <https://www.computer.org/csdl/journal/ts/2025/03/10859156/23X97jMgYjm>
- Groce et al., *Randomized Differential Testing as a Prelude to Formal Verification* (ICSE 2007) — <https://agroce.github.io/icse07.pdf>
- Reducing failure-inducing inputs, *The Fuzzing Book* — <https://www.fuzzingbook.org/html/Reducer.html>