

Written at the end of the session that began with `docs/HANDOFF-2026-08-01.md` . Read this first, then `docs/ARTIFACT-ACCESS-RULES.md` for the rules that came out of it.

Branch: `refit-and-semantics-guard` , 19 commits, not yet merged. Suite: **54 passed, 1 failed, 0 skipped** (from 46 / 1 / 1).

Three items were asked for. All three are done. The refit — item 1 — is not the interesting part.

**The hypothesis it was meant to test is REFUTED.** `allyHit` did not move:  $-0.0187 \rightarrow -0.0126$ , about 0.2 standard errors. The immunity bug was not why it sat near zero. It sits near zero because humans genuinely do not avoid hitting their own partner, and it should be reported that way. It is one of four features whose interval contains zero.

**What the refit found instead was a much larger defect.** `fit_joint.js` matched a human's click by requiring the move's TARGET to match, and `board.js` builds a spread candidate with `targetMon: null` because Earthquake is not aimed. **No spread click could ever match.** Spread moves are 14.94% of all human move clicks; 1,393 of 1,397 were discarded. The pair fit was estimating 74 weights from **24,997 of 82,483 joint turns**, and the missing 70% was precisely the turns the pair model exists to learn from.

Nine of eighteen pair weights changed sign. `spreadFreeBesideAlly` went **-4.986  $\rightarrow$  +0.859**.

**The corrected vector beats the shipped one at 70.9% of decisive pairs** [66.0, 75.4], replicated on three disjoint seed blocks (66.7%, 65.9%, 70.9%) with `sprt.js` and `paired_h2h.js` agreeing to the decimal. The bar is 55%.

**Roadmap item 1's stated premise is void.** It read  $-5.054$  as evidence that the imitation objective prices good play badly. That number was a fitter defect. The item may still be worth doing; this is no longer evidence for it, and DODUO's 42.0% was measured on the contaminated vector.

`MC.mons` keys formes with a hyphen ( `rotom-wash` ); every lookup stripped punctuation first. **101 of 308 entries could never be found — 8.17% of all observed metagame usage.** For those, every damage-derived feature read zero: no kill odds, no threat, no risk. Silently, because "not found" is a legitimate answer.

Fixed with one shared resolver ( `engine/mc_key.js` ) plus a cosmetic-forme fallback that is deliberately strict — base stats AND types identical, so the substitution is exact. Gourgeist-Super (same types, different stats) correctly still resolves to null.

Degradation on the standard workload: **27.64%  $\rightarrow$  8.42%  $\rightarrow$  0.00%**.

*Every expensive bug this session had one shape: something degrades, it counts itself correctly, and nothing fails.*

| defect              | what the code already knew                                            |
|---------------------|-----------------------------------------------------------------------|
| the spread matcher  | printed "57,486 of 82,483 unmatched" in every run                     |
| the forme lookup    | <code>dmgFailures.unknownSpecies</code> counted every miss, for weeks |
| mega, four times    | the mega rate was 0%                                                  |
| MEDICHAM's backtest | every forme silently filtered from its sample                         |

The counters were right. The numbers scrolled past. **A counter without a budget is a number; a counter with a budget is a guard.**

tests/test-degradation-budgets.js is that guard. Ceilings are MEASURED ( --ratchet ), never chosen, and may only be tightened. Fitters are read from the artifact each one writes rather than re-run, because a fit takes half an hour and is the only honest reporter of what it saw.

The second pattern, from docs/ARTIFACT-ACCESS-RULES.md : **sharing an artifact is not enough; the ACCESSOR has to be shared too.** Every engine already read the same MC.mons from the same generated file. It broke anyway, because five places each wrote their own doorway into it — two correct, three wrong, all five written by Claude in five different sessions.

| guard                             | catches                                                            |
|-----------------------------------|--------------------------------------------------------------------|
| tests/test-feature-semantics.js   | a feature changing MEANING under a stable name                     |
| tests/test-degradation-budgets.js | any degradation rate exceeding its declared ceiling                |
| tests/test-mc-key.js              | a hand-rolled species lookup (ratcheted at 16, may only shrink)    |
| tests/test-artifact-keys.js       | a new artifact whose keys can be missed, with no declared accessor |
| tests/test-sprt-arm-sign.js       | the ship/no-ship verdict being reported with its sign inverted     |
| tests/test-dead-volatile.js       | re-applying a volatile the target already has                      |
| tests/test-team-preview-race.js   | committing the bring before the opponent's sheet arrives           |

**R7, learned twice in one day: a guard only guards what it exercises.** The semantics fixture missed the forme bug entirely because every species in it was hyphen-free. It now carries Rotom-Wash and Sinistcha-Masterpiece.

- 1. REDIRECTION IS NOT MODELLED, and it is probably most of a 23% drop rate.** board.js has the features ( redirectThenAttack , redirectThenSetup ) and medicham2-browser.js 's own header promises redirection — but nothing implements it. Consequences: MAG's threat model is wrong whenever Follow Me is up; the rollout engine ignores it; and fit\_policy.js 's own caveat names it as the leading cause of unmatched clicks. fit\_joint still drops **23.20%** of two-slot turns. This is a missing mechanic, not a new feature.
  - 2. The 18 pair weights have never been collinearity-audited.** Nine flipped sign today and redirectThenAttack sits at -0.405 for a textbook play. The 56-feature audit shows five features are "strong alone, lost in company"; the pair block is unmeasured. Cheap diagnosis.
  - 3. Self-play teams are not the metagame.** Total variation distance **54.8%** against real ladder usage, with every top species under-represented (Kingambit 16.1% generated vs 22.4% real). Anything trained on self-play faces a softer field than MAG actually plays.
  - 4. test-mag-page** — the page implements 21 of 56 features. The one red test. A project.
  - 5. Dead artifacts.** policy-weights-opppmodel/nomf/machamp.json have zero consumers; policy-weights-nopop.json is documented as a corrupted fit and is referenced only in comments. All git-tracked, so recoverable — but nopop and machamp are cited in docs/MODELS.md and two review docs, so retire the citations rather than silently breaking them.
  - 6. build/omnibus.py still exits 0 after printing FAILED.** PDFs were produced this session with headless Chrome instead. A build step that lies about succeeding is its own bug.
- 
- Ran a fixed 200,000-game H2H** when engine/sprt.js exists to stop early, having read its header the same day. Every comparison was decided inside ~420 games. SPRT's own words: *"would have saved 16,948 games (97.7% of the run)."*
  - Reported a head-to-head with its sign inverted**, then "fixed" sprt.js — which was right. Arm 1 is the challenger. Reverted; tests/test-sprt-arm-sign.js now forces both analysers to agree.

- **Chased a self-play side bias through four corpus regenerations.** It was my own invocation: the levers are PER ARM, `--greedy` arms only side A, and I armed one side. `validate_selfplay.js` now says so instead of blaming the pairing.
- **Claimed `vmmem` was thrashing.** It was a closed laptop.
- **Pinned six cores while Will was using the machine,** and had to be told to stop.
- **Nearly weakened a freshness guard** because it was inconvenient. Checked first; the guard was right and two site artifacts really did derive from self-play.

The recurring shape in that list: a confident diagnosis stated before the measurement that would have settled it. Three of the six cost real time.

---

- `SHOWDOWN_PATH=C:/Users/willj/Projects/Pokemon/pokemon-showdown` is still required.
- Never `git add -A` — `data/archetypes.json`, `kad-replays.js`, `live.js`, `conformance.json` churn.
- **6 processes maximum**, and ask before running wide while Will is at the keyboard.
- The levers are PER ARM: `--greedy` / `--greedy2`, `--joint` / `--joint2`. For a symmetric self-play corpus, arm BOTH.
- `validate_selfplay.js` 's mirror check uses `RandomPlayerAI`, so "mirror is 50/50" tests the SIMULATOR — not MAG and not mew's wiring.