

MEASURE — can we believe a number

Owens: engine/mew.js , engine/sprt.js , engine/provenance.js , engine/status.js , engine/backtest_winrate.js , engine/paired_h2h.js , engine/feature_engine_contrast.js , the noise floor, the corpus stamps, and the MAG refit.

Its one number: leaf calibration — when the leaf says 90%, is it 90%.

May not: change a policy, a search knob or an engine mechanic. This division builds the rulers; it does not compete on them.

<!-- GENERATED: engine/status.js -->

```
MEASURE — can we believe a number

leaf calibration: QUARANTINED — the figure is withheld, not annotated.
  data/winrate-backtest.json is downstream of MEDICHAM: its generator engine/backtest_winrate.js is in t
  MEDICHAM is not correct — 2 of 9 gate clauses fail (whole-game differential / BOARD-MATERIAL — games w
  it becomes quotable again when the gate opens AND this is re-run: node engine/backtest_winrate.js
engine correctness -> leaf: QUARANTINED — the figure is withheld, not annotated.
  data/leaf-engine-contrast.json is downstream of MEDICHAM: its generator engine/leaf_engine_contrast.js
  MEDICHAM is not correct — 2 of 9 gate clauses fail (whole-game differential / BOARD-MATERIAL — games w
  it becomes quotable again when the gate opens AND this is re-run: node engine/leaf_engine_contrast.js
provenance: 189 unsafe, 2 void (declared), 33 possibly stale, 30 ok, 0 missing
click censoring: QUARANTINED — the figure is withheld, not annotated.
  data/click-censoring-census.json is downstream of MEDICHAM: its generator engine/click_census.js is ir
  MEDICHAM is not correct — 2 of 9 gate clauses fail (whole-game differential / BOARD-MATERIAL — games w
  it becomes quotable again when the gate opens AND this is re-run: node engine/click_census.js
the weights are QUARANTINED — data/policy-weights.json and the joint weights were fitted on features con
REFIT OWED — weights fitted 2026-08-28 15:46
  feature_fixture --check FAILED:   or restamp with: node engine/feature_fixture.js --stamp <file> |  C
  moved after the fit: engine/medicham2-browser.js  2026-09-06 00:08
  moved after the fit: data/engine-data.js  2026-08-31 00:08
  moved after the fit: data/abra-tags.js  2026-09-04 21:38
```

stamped 2026-09-06 01:28

<!-- /GENERATED -->

SIXTY-SEVEN WITHHELD FIGURES WERE STILL BEING PUBLISHED IN THE LIVING DOCUMENTS, EVERY CITATION FAITHFUL — 104 → 37, AND THE 37 THAT STAND ARE THE RULER, NOT THE DEBT. 2026-09-06, CHANGELOG 5.261.0

WHAT MOVED. tests/test-docs-quarantine.js counted **104** figures across 14 living documents sourced from artifacts engine/quarantine.js withholds. **37 still fire.** Cleared to zero: docs/SUMMARY.md , docs/MEASURE.md , docs/MILTANK.md , docs/PRIORITIES.md , docs/EXTERNAL-EVIDENCE.md , docs/GAME-DIFFERENTIAL-DESIGN.md , docs/ENGINE-COVERAGE-PLAN.md ; docs/ABRA-whitepaper.md 15 → 2, docs/MODELS.md 23 → 5, docs/SEARCH.md 5 → 2, docs/WEB.md 4 → 2. Every removal uses status.js 's own sayHeld shape —

the figure is ABSENT, the artifact is named with why it is downstream, and the re-run command is given — because a strike-through, a "previously" and a PRE-CHANGE caption have each already been shown not to stop a number being quoted.

THE DECK AND THE TECHNICAL DOCS WERE CARRYING THE SAME FIGURES AND THE RATCHET COULD NOT SEE THEM. The clause fires on a figure that shares a PARAGRAPH with a citation of the withheld artifact. docs/ABRA-deck-plain-english.md states its numbers in plain English with the artifact named nowhere near them, so it scored zero offenders while republishing the leaf comparison and the censored-label counts. They are withheld now for the same reason as the whitepaper's, but the general point is the instrument's: **a document that cites nothing is invisible to a rule keyed on citations**, and the plain-English deck is exactly the document most likely to cite nothing.

THE 37 THAT STAND ARE NOT DEBT OF THIS KIND AND MUST NOT BE "FIXED" BY DELETING THE FIGURES. They divide in two, and both are the CLASSIFIER rather than the documents.

- **A COINCIDENCE ENGINE INSIDE** data/policy-weights.json . artifactNumbers walks strings, and that artifact carries a featureHashes.features map of integer HASHES. Six of them match a figure in a document by accident — the hashes of benchRisk , tgtBulk , koFirst , switchSurvives1 , pivots and statusBites — once the matcher's legitimate $\times 100$ and $\div 100$ rescaling is applied. Two fitted floats do the same: one weight rescales onto a percentage a document states, and one unweighted column onto another. The hash values are not written out here, because writing them out is itself a republication and this rule has no prose escape — it was demonstrated on this very paragraph, which fired three times before it was rewritten. Every one of those document figures belongs to a DIFFERENT and non-withheld source — the damage differential, the empirical-click driver's reach rate, a planned game budget, DODUO's unpaired win rate, an out-of-sample behaviour rate. This is the "registry that fires on every occurrence of a small integer" failure engine/docs_scan.js already records, arriving through a hash table.
- **INSTRUMENTS THAT MEASURE MEDICHAM RATHER THAN CONSUME IT.** engine/quarantine.js has the exemption mechanism — MEASURES_THE_ENGINE — and it holds exactly two modules, game_differential.js and derive_protocol_events.js . Everything else is withheld on pure require reachability, so the mechanics CENSUS (data/all-mechanics-fire.json), the rate runner that tallies MEDICHAM's dice against declared targets (data/million-run.json , -staged), the register audit (data/register-reality.json), the register copy (data/open-work.json) and the click-coverage probe (data/medicham-represented-clicks.json) are all withheld — while CLAUDE.md says in as many words that the census is NOT quarantined. Fifteen of docs/ROADMAP.md 's nineteen are this. **Withholding a legitimately quotable number is its own defect**, so none of them was touched here; the fix is an entry in MEASURES_THE_ENGINE , which is an ENGINE-side change and is owed.

WHAT IS OWED. (i) The two classifier repairs above. (ii) node engine/status.js --write was deliberately NOT run — an ENGINE agent was editing the simulator and rewriting data/game-differential.json throughout this pass — so the generated blocks in every division ledger are one pass behind. (iii) The seeded census in tests/test-docs-quarantine.js still reads 104 and now reports 67 lines to delete; re-seeding it to 37 belongs to whoever owns that file. Full account: docs/_reports/2026-09-06-withdraw-quarantined-figures.md .

THE PUBLISHED NUMBER AND THE MEASURED NUMBER ARE ONE NUMBER AGAIN — BOARD-MATERIAL 50 OF 961, PROTOCOL 151 OF 961, REPUBLISHED OFF A SETTLED TREE, AND BOTH CLAUSES NOW FAIL ON A COUNT RATHER THAN ON WITHHELD STALENESS. AND THE 1.53× THIS DIVISION PUBLISHED LAST

NIGHT WAS MEASURED AGAINST THE WRONG DENOMINATOR.

2026-09-06, CHANGELOG 5.258.0

THE PUBLICATION FIRST, BECAUSE IT IS THE ONLY THING HERE THAT CHANGES WHAT MAY BE QUOTED. `data/game-differential.json` was rewritten on release `db248fe67a5e` — 961 games, cap 20, arm `middle`, steering `empirical`, `--end-state`, census pin `data/verification/census-pin-9446a684709d.json`, pool `--team-store data/team-pool-frozen`. It holds **board-material 50 of 961 and protocol first-divergence 151 of 961**. It had carried a 46 measured on a superseded release for 1.3 days, and while it did, ****both whole-game clauses were failing on WITHHELD STALENESS — "MEASURED AGAINST A DIFFERENT ENGINE ... EVERY COUNT IN IT IS WITHHELD"** — rather than on anything measured. **A clause that fails for the right reason and a clause that fails because its input is old are the same colour and are not the same statement, and this division should be the last place that lets them read alike. 911 games never part a board; 4 of the 50 part a board with the protocol identical all game; 10376 of 10539 compared turn boundaries were identical. The settled-tree run reproduces the last fix step** to the byte** on `classes`, `first_divergences`, `state.first_board_divergences` and the by-cause summary, which is what makes this a publication rather than a sixth measurement of one quantity.

THE CORRECTION THIS DIVISION OWES IS ITS OWN: THE PROTECT AMPLIFICATION WAS A RATIO WITH THE WRONG DENOMINATOR UNDER IT. Published last night: the empirical driver over-clicks the protect family at **1.53×** its own input, cause renormalisation. **Renormalisation is about half of it. The denominator is the other half, and the denominator was never a defect at all.** Decomposed over the run's own 17,532 decisions, so that no step compares two populations:

step	value	factor
declared input, acts-weighted over the whole table	13.565%	
the SAME table's marginal, weighted by the decisions this arm took	16.209%	×1.195
renormalised over the legal candidate set — the weights the sampler was handed	20.257%	×1.250
realised — what the sampler returned	20.374%	×1.006
		×1.502

The sampler is faithful to ×1.006; all of the residual is in the weights. The first factor is not a driver defect: the arm plays a census-steered pool whose bodies carry the family at a higher marginal rate than the ladder's own click distribution does. **So 13.565% is the wrong ruler for this run** — the pool-matched same-table denominator is **16.209%**, and against it the arm reads **×1.257**. The general rule, and it is this division's business rather than the driver's: **an error expressed as a ratio is only as good as the population its denominator was taken over, and a run whose census pin moves takes the denominator with it.** Any future sentence of the form *"the arm reads X% against 13.565%"* carries the pool-matched marginal beside it or it is not a measurement.

AND THE SECOND HALF OF THE EARLIER DIAGNOSIS IS WITHDRAWN OUTRIGHT, WITH ITS SIGN INVERTED. Legality subsetting was named as a contributor on the strength of candidate sets averaging about 3.14 of four. Measured on this run the mean is **3.772**, and conditioning on it: **87.0% of decisions — 15,253 — already have all four moves and are the ones reading 21.724%**, while a body down to one legal candidate reads **8.134%**, because a body narrowed to one move is usually narrowed to its attacking move rather than to its Protect. **Legality subsetting is a small NEGATIVE contribution.** It is withdrawn rather than quietly dropped, because a cause that vanishes without a sentence is indistinguishable from a cause somebody stopped honouring.

AND THE PART THIS DIVISION SHOULD BE PROUDEST OF IS THE SMALLEST NUMBER IN IT: A NOISE FLOOR WAS MEASURED BEFORE THE EFFECT WAS BELIEVED, FOR THE FIRST TIME ON A CALIBRATION FIGURE HERE. The only place both a prediction and a ground truth exist is the human corpus, so the rule was scored there: 185,422 scored clicks, split half by game into **92,949 and 92,473**, the correction fitted on one half and evaluated on the other, so nothing is scored against what it was fitted on. The marginal reads **14.233%**, the driver's rule reads **16.228%**, humans clicked **14.757%**. **The half-vs-half spread of the OBSERVED rate is 0.002 points**; the rule over-predicts by **+1.644 and +1.646 points**, which is 800 times the floor. That is what makes the over-prediction an effect rather than a reading. A carriage correction removes about **72%** of it, held out, landing at **+0.405 and +0.504**. **Nothing here is tuned**: the repair changes the driver's declared input, and applying it in the same pass as five engine fixes would leave none of the six attributable. The expected landing point is written down instead — near **15.0 to 15.1%** on a ladder-shaped population and near **16.7%** on this census-steered pool, and explicitly **not** the 14 to 15% the earlier report predicted, because the pool is not the ladder.

TWO PREDICTIONS MISSED TONIGHT AND BOTH MISSES ARE ONE MISTAKE ABOUT A SAMPLE, WHICH IS A RULER FAILURE AND THEREFORE THIS DIVISION'S. The Leech Seed step called **58 / 157** and read **56 / 158**; the Fairy Aura step called **54 / 156** and read **51 / 153**. Both reasoned from `state.first_board_divergences`, which **is capped at 40 rows and is a SAMPLE, never the population** — the sowerless Leech Seed chip was in the other nineteen, and the aura's reach was mis-scoped because the cause string names the VICTIM rather than the mechanic. Three of five landed at their point estimates, including the staged-pin step's prediction that nothing would move. **A capped list is a sample and must be read as one**; that is the second time in two nights it has cost a point estimate here, and it is now written down rather than remembered.

THE INSTRUMENT REPAIR WAS PAID FOR RATHER THAN ARGUED, AND THE STAMP THIS DIVISION ADDED ONE VERSION AGO IS WHAT MADE THE BILL VISIBLE. Binding the staged pins by name moved `driver_code` from `e87506b2d737` to `0c1fc935a5fb` over eleven files, which breaks comparability with the **56 by construction**. It was re-measured paired on the same release, census pin and pool: board-material 56 to 56, protocol 158 to 158, with `classes`, `first_divergences`, the end state and every first board divergence identical. **A stamp that fires on your own repair and forces a paired re-run is the stamp working**, and this is the first time it has been honoured on a change this division did not make.

WHAT THIS DIVISION STILL OWES.

- **Leaf calibration — this division's one number — stays WITHHELD and was not run.** `data/winrate-backtest.json` is downstream of MEDICHAM and the gate is shut on the board-material clause. **No reliability curve is published in this version and none is implied by anything above.** The standing brief is not discharged by this document.
- **The noise floor for the WHOLE-GAME differential is still unmeasured, and tonight makes the gap sharper rather than smaller.** A floor now exists for the protect-rate ruler and it cost one split-half pass; nothing in this repository still states what a one-game or a two-game move in the board-material count is worth, so a step of that size is reported as an attribution — a named cause, a knob-cleared control, a closed row and no new row — or it is not reported at all.
- **The MAG refit stays OWED and it is a REFIT, not a restamp.** No fit was started and no fitted vector was written; `data/policy-weights.json` was not touched. The damage table under the fitted vector moved from 318 species to 322, so the feature FUNCTION's input changed and a restamp would write over the evidence for the refit rather than answer it.
- `node engine/status.js --write` **WAS run by this pass, last in the sequence**, so the `<!-- GENERATED -->` block at the top of this file is current and was not hand-edited.

THE PINS FROZE EVERY INPUT AND NEVER FROZE THE INSTRUMENT, AND THE COMPARABILITY CHECK SAID COMPARABLE ABOUT TWO RUNS PLAYING DIFFERENT DRIVER CODE. THE "NON-REPRODUCIBLE JOINT ARM" IS WITHDRAWN — SIX RUNS ON ONE SET OF PINS ARE BIT-IDENTICAL WITHIN EACH SIDE OF THE EDIT. 2026-09-05, CHANGELOG 5.257.0

THIS IS THIS DIVISION'S SUBJECT AND THIS DIVISION MISSED IT, SO IT IS RECORDED THAT WAY. A measurement here pins three things — the engine release, the census, the team pool — and **all three are INPUTS**. The code that reads them was never digested at all. `engine/game_differential.js` and `engine/empirical_driver.js` are read LIVE off the working tree by every run, and an edit to either one selects the sample just as surely as the tables do. The photograph rule in CLAUDE.md says nothing in frame may move; the instrument was standing outside the frame holding the camera.

AND THE CHECK BUILT TO CATCH EXACTLY THIS ANSWERED THE WRONG WAY, WHICH IS WORSE THAN NOT HAVING IT. `engine/arms_comparable.js` was asked about two of tonight's runs and printed *"COMPARABLE. Both arms selected their sample the same way, so a difference between their numbers is the change under test."* The two runs were playing different driver code. **Its own limits block had named the hole in prose the whole time** — *"THE DRIVER ITSELF ... no artifact records its digest. WIRE 4 asserted it by hand"* — which is the finding to keep: **a named limit is not a guard**. A file that describes its own blind spot reads, to anybody scanning it, exactly like a file that covers it.

THE CLAIM THAT IS WITHDRAWN, NAMED, BECAUSE IT WAS PUBLISHED THREE HOURS BEFORE IT WAS RETRACTED. This version's changelog note records that `joint-empirical-click/v1` gave **167, 167 and 138** on identical pins, that the cause was under investigation, and that until it was settled every joint figure of the night — 110, 53, 167, 138 — was one draw from an uncharacterised distribution and could not be quoted. **It is settled and it was never a distribution**. Six runs on one identical set of pins split perfectly cleanly at the protect fix, which landed at 02:27:

runs on the driver as it was	protocol / board-material	prefer_narrowed
three of them	121/34, 121/34, 138/53	20,507 / 20,507 / 20,353
three on the repaired driver	167/69, 167/69, 147/55	0 / 0 / 0

Each side is bit-identical within itself, down to `credit_events` **and** `shuffle_calls` . So the joint **53 STANDS** and is reproducible; what is withdrawn is the diagnosis, not the observation. **The general rule for this division: a measurement that will not reproduce is a code-movement hypothesis before it is a randomness hypothesis**, and the cheapest test is not more runs — it is a digest of the thing doing the running.

WHAT LANDED, MEASURED RATHER THAN ASSERTED. `engine/steering.js` digests the instrument's own require closure into `steering.driver_code` . `engine/game_differential.js` takes that digest at module load, takes it again at write time, and **VOIDS the run if it moved — withholding** `diverged` , `mid_void` **and** `state` **rather than captioning them**, which is the CLAUDE.md rule that a caption is not a quarantine applied to the instrument for the first time. `arms_comparable.js` now COMPUTES its limit line from the two artifacts in front of it, and `steering.comparable` treats **exactly one side carrying the stamp as a REFUSAL**, not a pass — because "nothing recorded it" has to read as a refusal or the stamp buys nothing on the first pair that matters.

THE SECOND CORRECTION, AND IT IS THE ORDINARY KIND: A FIGURE PUBLISHED TONIGHT NEEDS ITS INSTRUMENT NAMED. The empirical arm on the repaired driver reads **147 protocol / 55 board-material**, not 47. **47 is not wrong** — it is the same run read by the comparator BEFORE the leaf widening, at 40 leaves rather than 54, and it remains correct as the second leg of the paired protect measurement. Two comparators, two readings, one run. **This division's rule stands unchanged and is the reason the discrepancy was traceable at all: name the ARTIFACT and the ARM beside every whole-game figure.**

THE PROTECT MEASUREMENT ITSELF WAS PAIRED PROPERLY AND IS THE ONE THING HERE THAT NEEDED NO CORRECTING. 961 games each, release `688e696f00c8`, empirical arm, the driver rule the only difference, with the before leg **re-run under the knob** rather than lifted from a published artifact. **22.2% of decisions reached the sampler with exactly one candidate, 60% of them Protect**; with all four moves present the arm was already at **15.3%**, the human rate. Games that finish went **56% → 79%** (resolved **539 → 762**), and board-material went **34 → 47 — an increase that is the expected consequence of the repair rather than evidence against it**, because a game that ends reaches positions the old arm never played. The residual rate is **1.53×** its own input for a separate, measured reason: the move-prior table is a marginal renormalised onto the four moves a body carries (row mass 0.917 over 8 becomes 0.521 over about 3.1) and Protect always survives the subsetting. Named, not tuned.

WHAT THIS DIVISION STILL OWES.

- **Leaf calibration — this division's one number — stays WITHHELD and was not run.**
`data/winrate-backtest.json` is downstream of MEDICHAM and the gate is shut. No reliability curve is published in this version and none is implied by anything above.
- **The MAG refit stays OWED and it is a REFIT, not a restamp.** No fit was started and no fitted vector was written.
- **The noise floor for the whole-game differential is still unmeasured, across both arms.**
Tonight sharpens why it matters: two figures 8 apart on the same run were a comparator width, and nothing in this repository states what an 8 is worth.
- **Every artifact on disk predates the instrument stamp.** They were never checked on that axis and still are not; `steering.comparable` refuses a pair where only one side carries it, so the first honest instrument-checked comparison is a run yet to be taken.
- `node engine/status.js --write` **was NOT run by this documentation pass and no** `<!-- GENERATED -->` **block was hand-edited.** The block at the top of this file is stamped to an earlier pass.

THE RULER WAS HIDING THE DEFECT: A CHARGED MOVE STRUCK THE WRONG SLOT AND NO INSTRUMENT HERE COULD HAVE SEEN IT, BECAUSE THE OLD DRIVER AIMED EVERY FOE-AIMED CLICK AT ONE INDEX AND MADE RE-AIMING A NO-OP BY CONSTRUCTION. JOINT ARM 110 → 53, EMPIRICAL ARM 35 → 34, AND THE TWO MAY NOT BE PAIRED. 2026-09-05, CHANGELOG 5.256.0

THIS IS THIS DIVISION'S SUBJECT IN ITS PUREST FORM, AND IT IS WORSE THAN A BLIND SPOT IN AN INSTRUMENT — IT IS A BLIND SPOT THE INSTRUMENT MANUFACTURED. `vol.charging` 's release turn aimed at `live(foes)[0]` where the authority replays the `targetLoc` stored on the sub-volatile, so a Phantom Force charged at slot b struck slot a. The driver in use until this version resolved every foe-aimed click with the lowest live index **for both slots**, so an engine that

replays a charged move's stored target and an engine that always strikes the lowest live index play **identical games**. The defect was not unmeasured. It was outside the space the measurement could reach, and it would have stayed there for as long as the driver did. A search is worth what its model is worth, and a differential is worth what its driver is worth: **the driver is part of the ruler, not part of the subject**.

EVERY FIGURE IN THIS VERSION BELONGS TO AN ARM, AND `arms_comparable` REFUSES TO PAIR THEM BY DESIGN — THIS DIVISION SHOULD BE THE LAST PLACE THAT TRIES. They are different policies playing different populations. **Joint arm: board-material 110 → 53, protocol first-divergence 191 → 138, VOID 38 → 4. Empirical arm: board-material 35 → 34, protocol 120 → 121.** Charge moves are entirely absent from `unshared_address_shapes` after the fix, which is where the VOID collapse comes from — a move re-aimed differently in the two engines puts a die at an address only one side ever draws, and a game whose shared addresses fall below the overlap floor is discarded rather than scored. The census is level at 829 of 829 with 0 unprobed and the gate is back to 2 of 9 failing clauses.

AND A DRIVER CHANGE IS A POPULATION CHANGE, WHICH MEANS IT RESETS A BASELINE RATHER THAN IMPROVING ONE. THIS IS THE SECOND TIME IN ONE NIGHT. The coverage-to-empirical swap took board-material 0 → 135 by playing games that end instead of games that run to a cap; the focus-fire-to-joint swap has now exposed a targeting defect that a constant aim made unobservable. **Neither number was wrong; each was a statement about a population.** The rule this division enforces on everyone follows directly: a whole-game figure is quotable only with its artifact AND its steering policy, and a figure taken under one driver may never be differenced against a figure taken under another.

THE CONTROLS ARE THE PROOF, AND THEY ARE THE ONLY REASON A 57-GAME MOVEMENT MAY BE ATTRIBUTED AT ALL. Knob-cleared runs — same release, same pinned pool, same driver, the two fixes switched off — reproduce the earlier figures **exactly**: joint 110 and empirical 35, against 53 and 34 with the fixes live. `engine/engine_release.js` drift reports that only `medicham2-browser.js` moved between the two releases. **So the entire 57-game movement in the joint arm is these two fixes and nothing else.** An exact reproduction under a cleared knob is stronger evidence than any before-and-after on a moving tree, because it removes the one thing a before-and-after cannot: everything else that changed.

THE NOISE FLOOR IS STILL UNMEASURED, AND THAT IS WHY THE TWO ARMS ARE REPORTED DIFFERENTLY HERE. Nothing in this repository states the spread between two halves of a single arm. The joint arm's movement is carried by an exact control reproduction and a one-file drift report, which is attribution rather than an effect estimate and does not need a floor. **The empirical arm moved by one game and this division does not call that an effect** — it is smaller than any spread that has ever been measured here, because no spread has ever been measured here. Splitting one arm in half and measuring the difference between the halves remains owed, and it is owed more loudly now that there are two arms to do it for.

A GREEN TEST WAS A STALE ASSERTION THAT HAD ONLY EVER PASSED BY LUCK, WHICH IS THE FAILURE THIS DIVISION EXISTS TO CATCH AND DID NOT. `tests/test-pin-arms.js` selected its arms with a `!x.top` filter that swept in the `middle` arm, whose `chance` is a live uniform and whose own description reads *"moves miss at their printed accuracy"*. **It was green only when one hash landed under 0.01.** A green demonstration that has never been shown able to go red is not evidence; this one could not reliably go green either, and it had been standing as a pin on the arms for as long as it existed. Fixed, and all arms now pass.

THE PREDICTION CARD SCORED 3 OF 8 AND THE SHAPE OF THE MISSES IS THE FINDING, NOT THE SCORE. Both empirical misses were by one and are attributed. **All three joint misses ran in the same direction — the fix being much larger than called.** A one-sided miss set is a

systematic error in the prior, not noise in the measurement, and it is only visible because the misses are recorded beside the hits. A prediction card is a ruler only if the misses go on it.

WHICH ARTIFACT HOLDS WHICH WHOLE-GAME FIGURE, STATED BEFORE EITHER IS QUOTED. `data/game-differential.json` was again **NOT** republished and still holds board-material 46 of 961 — that is what `node engine/status.js` prints, and it is stale as a description of tonight's engine. The measurements above live in this version's verification artifacts under `data/verification/`, one per arm, each with its knob-cleared control beside it. The published figure is superseded by a settled-tree pass and never by a sentence.

WHAT THIS DIVISION STILL OWES.

- **Leaf calibration — this division's one number — stays WITHHELD and was not run.**
`data/winrate-backtest.json` is downstream of MEDICHAM and the gate is shut. No reliability curve may be published through an engine whose boards still part, and the standing brief is not discharged by this document.
- **The MAG refit stays OWED and it is a REFIT, not a restamp.** No fit was started and no fitted vector was written. The simulator moved again this version, so a restamp would write over the evidence for the refit.
- **The noise floor for the whole-game differential is still unmeasured, now across two arms.** Until it exists, a one-game move is reported as an attribution or not at all.
- `node engine/status.js --write` **was NOT run by this documentation pass and no** `<!-- GENERATED -->` **block was hand-edited.** The block at the top of this file is stamped to an earlier pass and says nothing about this one; read the gate, not the block.

A RELEASE ID IS THIS DIVISION'S STALENESS SIGNAL, AND IT FIRED WITH NOTHING BEHIND IT — FIVE HEAVY RE-RUNS SPENT ON 26 CONTENT-IDENTICAL SOURCES. THE ANSWER IS A DIAGNOSIS, NOT AN EXEMPTION: THE DIGEST IS UNCHANGED AND TONIGHT'S TWO CLAUSES STILL FAIL WITH EVERY COUNT WITHHELD. 2026-09-05, CHANGELOG 5.255.0

THIS IS A RULER FAILING IN THE DIRECTION THAT COSTS MOST, AND IT IS THIS DIVISION'S SUBJECT. Everything measured here is measured against a frozen release, and the release id is what tells a reader that a figure is owed again. An id that moves for a real reason is the whole point of the mechanism. An id that moves for no reason is a false alarm inside the staleness detector itself — and a false alarm there is not cheap, because the prescribed response is to re-measure. It was paid at full price: an agent found the engine gate at **7 of 9 failing clauses instead of 2 of 9** and re-ran **five heavy clauses** to restore it, and the cause was that `engine/medicham2-browser.js` had its **line endings** changed while all **26** frozen sources were content-identical — `diff --strip-trailing-cr` returns 0 differences. The standing case is live in the tree: one cited release differs from the working tree in exactly one file, **39,932 LF against 39,932 CRLF, zero characters edited.**

THE FIX IS A THIRD DOOR, AND THE TWO OBVIOUS ONES ARE BOTH CORRECTLY SHUT. Pinning the nine deliberately unpinning sources to LF **would rewrite those files, move every release id, and break** `tests/roster.js`, **whose red demonstrations match a carriage return against the simulator's own source** — a repair that strands every existing release to stop a nuisance is not a repair. And normalising the comparator is forbidden here on a stated ground this division should be the last to weaken: *the difference is already observable to an instrument*, so teaching the comparison not to see it

makes the instrument and the comparator disagree, which is the failure mode, not the cure. **So the drift is DIAGNOSED rather than excused.** `engine/engine_release.js` and `engine/pin_guard.js` classify a moved digest as CONTENT-CHANGED or EOL-ONLY, `tests/probe_release_drift_diagnosis.js` holds the demonstration, and the nine-source job stays filed — now visible every time it costs something instead of being rediscovered.

NOTHING IS EXCUSED, AND THE CLAUSES THAT FAIL TONIGHT PROVE IT BY STILL FAILING. The digest is unchanged and still hashes raw bytes. The id still moves on a line-ending change. Artifacts stranded by an aged-out release stay stranded — a stranded figure is withheld and re-measured, never resurrected. ****This version's two failing clauses are diagnosed CONTENT-CHANGED and still FAIL, with every count withheld and "A re-measurement IS owed" printed where a number would go.**** That is the withhold rule as this division wrote it: the figure is not annotated, it is absent. A diagnosis that made a red clause green would have been an exemption wearing a diagnosis's name, and it would have been this division's own signature failure.

THE CLASSIFIER IS A PROPERTY, NOT AN ENUMERATION, AND ONE IMPLEMENTATION SERVES BOTH READERS. It reads no filename, no path, no `.gitattributes` entry and not the frozen source list; one probe case diagnoses a file whose name exists nowhere in this repository, by calling the classifier on two bare buffers. `engine/status.js` and `engine/quarantine.js` **needed no edit at all**, because both already render the guard's `why` verbatim — the same discipline that keeps `status.js` shelling out to `provenance.js` rather than reimplementing its rules. Two implementations of one fact is the most expensive recurring failure in this repository, and a second drift classifier would have been that failure dressed as its own remedy. Measured on the current tree: of **36** cited releases, **34** CONTENT-CHANGED, **1** EOL-ONLY, **1** NO-DRIFT, **0** UNDIAGNOSABLE, at **37 ms** inside the gate.

AND IT CAUGHT A BUG IN ITSELF — A FALSE ALARM FROM THE THING BUILT TO STOP FALSE ALARMS — WHICH IS WHY IT WAS SHOWN RED FOUR WAYS BEFORE BEING BELIEVED. Its first line-terminator counter reported two identical files as differing. With the normaliser broken to the identity function the probe scores **9 of 16**, reproducing the pre-fix behaviour exactly; with an over-excusing break it scores **10 of 16**; repaired, **17 of 17**. A green demonstration that has never been shown able to go red is not evidence, and this division has already paid for that lesson once this week in the roster's plants.

A COMPARATOR THAT DOES NOT LOOK AT A LEAF AGREES ABOUT IT FOR FREE, AND THE DENOMINATOR MOVED 37 → 40 OF 56. This is a ruler change and belongs in this ledger even though the work is ENGINE's. Will's ruling is that the leaves come before chasing the count to zero: *board-material zero on 37 of 56 standing leaves is not the same claim as zero on all of them.* `lockon`, `minimize` and `noretreat` are wired, each traced to a real **write AND read** site before wiring — the Unburden check, because a leaf can look wireable on every derived column while the engine holds nothing under that name — and each staged in a real game to confirm it stands at the boundary. All three red first with a control; `lockon`'s clock compared as a clock. **Board-material stayed flat at 35, called exactly in advance**, because the pinned pool holds zero Lock-On and zero Dragapult: the lab moved and the pool correctly did not. Say which scoreboard a fix should move before running it, and this one said so.

WHICH ARTIFACT HOLDS WHICH WHOLE-GAME FIGURE, STATED BEFORE EITHER IS QUOTED. `data/game-differential.json` was again **NOT** republished and still holds board-material 46 of 961 — that is what `node engine/status.js` prints and it is stale as a description of tonight's engine. The measurement on tonight's engine is board-material 35 of 961 with protocol first-divergence 120 and VOID 4, in this version's verification artifact under `data/verification/`. Both are true, they are not the same question, and neither may be quoted without its artifact. The published figure is superseded by a settled-tree pass and never by a sentence.

STORE-REPLAY IS REFUTED AS A DRIVER AND THE REFUTATION IS STRUCTURAL, WHICH IS WORTH MORE THAN A MEASUREMENT THAT CAME OUT BADLY. A replay stops being a replay at the **first damage-dependent faint**, and **at least 24.8%** of frozen-pool games have one on turn 1 — because Champions sheets never publish the **66** stat points, so the spread is absent on 100% of sheet bodies measured across **47,856** of them and simulated damage cannot match the real game's. It is also a POPULATION change rather than a driver change: the differential pairs one game's team A against another game's team B, so **no recorded sequence exists for the matchups it plays**. A driver that cannot be replayed on the population under test is a different experiment, not a stricter one, and refuting it before it was built is the cheap order.

THE COORDINATION GAP WAS ALREADY QUANTIFIED IN THIS REPOSITORY AND NOBODY HAD ACTED ON IT. THAT IS A MEASUREMENT FAILURE, NOT A SEARCH ONE.

`engine/board.js:377` measures humans double-targeting **23.4%** of the time against roughly 50% under independent choice, and `engine/empirical_driver.js:56-64` declares that it has **no target model and no switch model**, with switches at **12.1%** of real slot decisions. The consequence is visible in the games it produces: a median of **11** turns against real VGC's **7**, with **49%** hitting the turn cap instead of ending. A figure that sits in the tree and changes nobody's decision is worth exactly what an unmeasured figure is worth — the same shape as the registered gate that was red, correct, and filed in a report.

THE RE-DIAGNOSIS OF THE STANDING LEAVES FOUND NO BIG BUCKET LEFT, AND A NEGATIVE RESULT STATED PLAINLY IS THE POINT. Of the 37 rows standing at the start of this version: about **12** damage-value rows already fenced by a filed row, **5** `stall` (refuted — a die, not a missing mechanic), **3** Cursed Body / Flame Body / castform (refuted), **2** Poison Touch (a die value; the ability is wired and both engines reach the draw at the same point), **2** `vol.charging`, and a one-row tail. **After fencing, the largest actionable group was two rows.** `vol.charging` is a confirmed defect that cannot be probed today: the scripted chooser in `engine/game_differential.js` falls back to a target field a locked move's request does not carry, so the authority refuses the choice and no staged scenario here has ever played a two-turn release turn. Reported, not fixed — and it can move the artifact's `directed` block, so it is a currency risk rather than only a coverage one.

PREDICTIONS: 4 OF 4 ON THE LEAF BATCH, AND ONE MISS RECORDED RATHER THAN EXPLAINED AWAY. All four exact, with the flat board called in writing before the run. On the mechanics batch the VOID call was wrong — **6** predicted against **4** measured — because the bounced Parting Shot restored a shared accuracy address, which fell from **8** to **5**. The miss is recorded. A prediction card is only a ruler if the misses go on it.

WHAT THIS DIVISION STILL OWES, UNCHANGED BY ANY OF THE ABOVE.

- **Leaf calibration — this division's one number — stays WITHHELD and was not run.** `data/winrate-backtest.json` is downstream of MEDICHAM, the gate is shut, and no reliability curve may be published through an engine whose boards still part. The standing brief to point the backtest at the current leaf and publish the curve cannot be honestly discharged tonight, and it is not discharged by this document.
- **The MAG refit stays OWED and it is a REFIT, not a restamp.** No fit was started and no fitted vector was written. The damage table under the feature function has moved, so a restamp would write over the evidence for the refit. **The drift diagnosis does not touch this:** it reports why an id moved, and the inputs to this fit changed in content.
- **The noise floor for the whole-game differential is still unmeasured.** Board-material moved by two games this version and by three the version before; nothing in this repository states the spread between two halves of one arm, so a two-game move is reported as an attribution — by id, with zero new — rather than as an effect. That is the right report and it is not a substitute for the floor.

- **The nine unpinned sources are filed and not fixed**, deliberately, for the reason given above.
- `node engine/status.js --write` **was NOT run by this documentation pass and no** `<!-- GENERATED -->` **block was hand-edited**. The block at the top of this file is stamped to an earlier pass and says nothing about this one; read the gate, not the block.

A REGISTERED GATE WAS RED, NAMED THE EXACT PAIR IN 1.56 SECONDS, AND WAS FILED IN A REPORT — SO A GENERATED ARTIFACT SHIPPED A LIVE ENGINE DEFECT IN HEAD FOR SIX DAYS. THE FIX IS A PLACEMENT, NOT A SIXTH COMPARISON.

2026-09-05, CHANGELOG 5.254.0

THIS IS THIS DIVISION'S SUBJECT AND IT IS THE WORST VERSION OF IT. The failure was not an absent instrument, a wrong threshold or an unstamped artifact. `engine/artifact_audit.js` check G is a PROPERTY — it derives every self-declaring bundle under `data/` from that bundle's own `GENERATED` by `<builder>` from `<source>` header and spawns the builder's own `--check`, so it needs no named pair and covers pairs nobody has thought of. It exited 1, printed `1 GAP(S) FOUND`, named `data/abra-tags.js` against its source, and took 1.56 seconds. It is registered in `tests/run-all.js`. **It went unacted-on because only a full suite runs it, and because the failing result was written into a session report annotated "this check got BETTER across the session."** That is *a check nobody acts on is not a check*, verbatim, and it is the same normalisation as "one of the two known failures."

THE CORRECT LEVER WAS THE MOMENT, NOT ANOTHER RULER. `.githubhooks/pre-commit` now runs the audit **above its scope guard**, because that guard exits 0 for a commit staging only `data/` and *"regenerate the tags, commit the data"* is the shape of all five instances of this drift class. **No new comparison code was written**, which is the part this division should defend: two implementations of one fact is the most expensive recurring failure in this repository, and a sixth comparison would have been that failure dressed as its own remedy. Shown RED then GREEN on the real pair against an isolated `GIT_INDEX_FILE`.

WHICH ARTIFACT HOLDS WHICH WHOLE-GAME FIGURE, STATED BEFORE ANYTHING ELSE IS QUOTED. `data/game-differential.json` was **NOT** republished this pass and holds board-material 46 of 961 with protocol first-divergence 141 — that is what `node engine/status.js` prints, and it is stale as a description of tonight's engine. The measurement on tonight's engine is board-material 37 of 961 with protocol 122, and it lives in `data/verification/fix-batch-7.json` on release `316669459d67`. Both are true; they are not the same question, and neither may be quoted without its artifact. The published figure is superseded only when a settled-tree pass rewrites it, never by a sentence.

A SUITE THAT HAS NEVER BEEN SHOWN ABLE TO FAIL IS NOT EVIDENCE — FOUR ROSTER PLANTS COULD NOT GO RED. Every shipped roster artifact is written **without** `--reds`, so the red list is empty and the clause passes whatever the plants do. Four plants had drifted off their sites: two carried a literal carriage return against an LF engine file, one anchored a signature that grew a fourth parameter, one matched two sites. All four were re-aimed and then verified **independently, one rule at a time**, rather than read off the artifact that had been hiding them — which is the only acceptable order when the instrument under suspicion is also the reporter. The shipped stages now carry 18 / 29 / 35 reds, 0 not-ok, 0 unaimed plants.

AN OOM THAT LEAVES BOTH THE CONTENT AND THE MTIME UNCHANGED IS A STALENESS BUG WEARING A CRASH'S CLOTHES. `engine/tag_dex.js` was exhausting the default heap, and its only write is near the end of the file — so a death left `data/tags.json` holding old content with

an old mtime, nothing on disk recorded it, and every consumer read a stale tag set that looked freshly generated. ABRA-HEAP: 3072 is therefore a SAFETY declaration rather than a performance one, and the rebuild was **verified by the stamp rather than the exit code**: the tag artifact's own generated stamp moved from 2026-08-29T13:37:50Z to 2026-09-05T00:34:28Z. This is the class engine/provenance.js exists for, reached through a door provenance cannot see — an artifact that did not change cannot be caught by comparing it to anything.

FIVE CANDIDATES REFUTED BEFORE ANY EDIT, AND ONE REFUTATION MADE

PERMANENT. The Dire Claw sleep clock is correct in both engines; stall — five games, the largest single named leaf — is a die and not a missing mechanic; castform.species was not reachable in the shapes staged; Flame Body and Cursed Body fire. **The Dire Claw refutation now lives as an arm of its own probe**, which is the difference between a negative result that holds and one that has to be re-argued the next time somebody reads the leaf list. Predictions were recorded pre-run in data/verification/2026-09-04-sun-frz-prediction.json and went 4 of 4, then 4 of 4.

WHAT THIS DIVISION STILL OWES, UNCHANGED BY ANY OF THE ABOVE. Leaf calibration stays WITHHELD and is not annotated: data/winrate-backtest.json is downstream of MEDICHAM and the gate is shut at 2 of 9, both failing clauses reading the unreleased whole-game artifact. **The MAG refit stays OWED and is a REFIT, not a restamp** — and this release adds a reason rather than removing one, because data/abra-tags.js is now named among the inputs that moved after the fit, and it is frozen into every release a measurement is taken against. node engine/status.js--write was not run by this documentation pass and no <!-- GENERATED --> block was hand-edited; the block at the top of this file was stamped at the end of the engine pass that produced this version, and it agrees with the prose here — it prints the 2-of-9 gate, the REFIT OWED clause, and data/abra-tags.js among the moved inputs.

A ROW ASSERTING A LIVE DEFECT READ CLOSED TO THE GATE BECAUSE OF HOW ITS CELL PARSED — THE REGISTER IS A RULER AND IT HAD A READING ERROR, IN BOTH DIRECTIONS.

2026-09-04, CHANGELOG 5.253.0

NO MECHANIC CHANGED, NO DIFFERENTIAL RAN, AND NO WHOLE-GAME FIGURE WAS RE-MEASURED. data/game-differential.json still holds board-material 46 of 961, exactly where the previous release published it. Nothing in this block is evidence about the simulator, and no sentence in it should be read as the engine having moved.

THE DEFECT IS A READING ERROR IN AN INSTRUMENT, WHICH IS THIS DIVISION'S ENTIRE SUBJECT. The shipping status reader takes the text after the LAST pipe in a register row. #175's status cell begins open – engine DEFECT ; a code span inside that cell carries a pipe, the cell was split, and an appended closure narrative was read as the verdict. So a row asserting a live defect reported **CLOSED** to the gate while saying the opposite in its own first words. It is open and asserting breakage again, and that is the entire +1 in the gate's defect column.

THE CLASS WAS TWO DEFECTS AND THE BIGGER HALF HAD NO NAME. The pipe half hid 8 rows. The emphasis half hid nine: the cell's leading-whitespace skip is \s*, which does not skip **, so a cell authored | **CLOSED 2026-08-13** | fails to match with no pipe anywhere in the row. Eighteen rows were repaired, notation only — 18 insertions and 18 deletions, one line each. Open rows fell 237 → 222 and open-and-asserting-breakage rose 50 → 51. A whole-register replay through the shipping detectors confirms **exactly 17 verdicts moved and no eighteenth by accident**, and the two detector functions are byte-identical to HEAD.

THE GUARD IS WRITTEN AS A PROPERTY AND IT IMPORTS THE DETECTOR IT CHECKS.

One invariant per row: *the gate-visible verdict must be the same whether the status cell is read as the shipping detector reads it, or as the column the author wrote and then rendered to plain text*. It imports the shipping closed-row detector rather than reimplementing it — a third copy of that detector once disagreed with the canonical one on **24 of 292 rows in both directions**, and a verifier that re-implements the rule it is checking is a documented failure class in this repository. It was shown red twice: on seven synthetic doors, each with a repaired twin that goes quiet, **four of them doors nobody here has ever used**; and then on real data, where pointing it at the pre-repair register names **15 rows and exits 1** — the repair pass's own list, reproduced by an instrument that had never seen it. A mixed-corpus arm and a lift arm fail if the comparison is deleted or if the shipping reader ever stops cutting, so it cannot go green by asking nothing. Current register: 506 rows, 0 verdict failures, 0.17s, deterministic, writes nothing.

REPORTING RATHER THAN RATCHETING IS A JUDGEMENT, AND IT IS THE KIND THIS DIVISION IS SUPPOSED TO MAKE OUT LOUD. 15 cells are still cut by a pipe whose verdict is unchanged either way. They are printed as a list with no bar and **deliberately not ratcheted**: a count that may only go down invites the next author to argue their row is the exception, and the verdict clause already covers every cut that has a consequence. Ninety rows and 631 pipes were left alone, every one checked and none with a verdict at risk — 13 closed, 2 correctly open. Churn across ninety rows for no behavioural change is not worth the diff, and saying so is part of the work.

THE ESCAPE THAT LOOKS LIKE THE FIX WAS MEASURED AND IS NOT ONE. Escaping a pipe changes nothing: the capture uses a negated-pipe character class, which stops at an escaped pipe exactly as it stops at a bare one. Proven on a synthetic row through the shipping detectors — the authored form and the fully-escaped form extract identical wrong text. Only removing the pipes recovers the status.

A PARSE REPAIR MUST NOT LAUNDER ITSELF INTO A VERIFICATION, AND SIXTEEN CLOSURES SAY SO IN THEIR OWN CELLS. Sixteen rows were made READABLE without being RE-VERIFIED, and each records that in the cell rather than in a report somebody has to find. Cheap artifact checks did pass for seven of them; the weakest four rest only on the named instrument existing on disk, and say that too. A row that becomes legible is a row whose verdict can now be READ. It is not a row whose verdict was re-measured, and keeping those two apart is why this division exists.

WHAT IS OWED HERE.

- **Leaf calibration — this division's one number — stays WITHHELD and was not run.** `data/winrate-backtest.json` is downstream of MEDICHAM, the gate has not opened, and no reliability curve may be published through an engine whose boards still part.
- **The MAG refit stays OWED and it is a REFIT rather than a restamp.** No fitted vector was written and no fit was started; the damage table under the feature function has moved, so a restamp would write over the evidence for the refit.
- **The sixteen readable-but-unverified closures are owed a verification**, and it is not scheduled here. Their cells carry the admission; the register does not carry a re-measurement.
- **The 15 cut-but-harmless cells are carried, not fixed**, by the decision recorded above.
- **The defect-clause split this ledger printed one release ago — 40 rows asserting breakage with no instrument, 7 answering nothing usable, 3 naming a green one — predates this repair and was NOT re-derived.** Read the printed work list, not that split.
- **This documentation pass did not run** `node engine/status.js --write` **and did not hand-edit inside a** `<!-- GENERATED -->` **block.** The block at the top of this file was restamped by the ENGINE division's own pass while this text was being written, so it is current to THAT pass and says nothing about this one. It carries no register figure, so it neither confirms nor contradicts anything above — and it now

prints a smaller MEDICHAM clause count than the dated block below, which is that division's result and is not claimed here.

THE PUBLISHED NUMBER AND THE MEASURED NUMBER ARE THE SAME NUMBER FOR THE FIRST TIME — BOARD-MATERIAL 77 OF 961 → 46 OF 961 (4.8%) — AND THE GATE WENT FROM UNMEASURED TO MEASURED-AND-RED. 2026-09-04, CHANGELOG 5.252.0

THE TWO-NUMBER PROBLEM THIS DIVISION HAS REPORTED FIVE TIMES IS CLOSED, AND IT CLOSED BY REPUBLICATION RATHER THAN BY ARGUMENT. `data/game-differential.json` was rewritten on release `0dec37ff5ad9`: board-material 46 of 961, protocol first-divergence 141 of 961 raw and 140 after one declared, 961 games PLAYED of a 1200-PAIR budget, 7 void, 1 thrown. There is no longer a correctly-measured value of this quantity living outside the artifact the gate reads. ****The standing caveat — *the published clause still holds 77* — is DISCHARGED****, and it is written here as discharged rather than deleted, because a caveat that vanishes without a sentence is indistinguishable from a caveat somebody stopped honouring.

THE FAILURE CHANGED STATE, AND THIS DIVISION EXISTS TO INSIST THAT IS THE HEADLINE. For most of the night the gate failed on artifacts that could not prove which engine produced them — **state (a), in which nothing is known**. Five of the six failing clauses were in that state. All five pinned artifacts were regenerated against one release and **no clause is in state (a) now**. The one remaining FAIL is **state (b): a named instrument genuinely RED, on the current engine, over 46 real games**. An unmeasured clause and a measured red clause are not the same finding, and reporting them under one integer is how a ruler stops being one. The gate reads `CLOSED - 1 of 8 GATING clauses fail`.

A NOTE ON THE TWO RENDERINGS OF THAT COUNT, BECAUSE BOTH ARE IN THIS REPOSITORY AND THEY LOOK LIKE A DISAGREEMENT. The gate's own banner prints `1 of 8 GATING clauses fail`; the `<!-- GENERATED -->` blocks in these ledgers print `2 of 9 gate clauses fail` and name both whole-game clauses. They are the same verdict counted two ways: there are 9 clauses of which 8 gate, and 2 fail of which 1 gates. The non-gating failure is NARRATION, which is a separate gate by Will's 2026-08-22 call. Neither figure is wrong and neither may be quoted without saying which population it counts.

STATE (c) IS REPORTED SEPARATELY AND HOLDS NOTHING SHUT, WHICH IS EXACTLY WHY IT HAS TO BE NAMED. Inside the defect clause that PASSES: 40 open register rows assert breakage with no instrument that decides them, 7 name an instrument that was asked and answered nothing usable, and 3 name a green one. Folding the 40 into the gate's verdict would report a broken simulator where what exists is an unmeasured one — the same error, in the opposite direction, as reporting an unmeasured clause as a passing one.

THE PREDICTION WAS WRITTEN BEFORE THE RUN AND SCORED 4 OF 4, ALL EXACT — AND THE PREMISE IT RESTED ON WAS INCOMPLETE, CAUGHT IN ADVANCE RATHER THAN AFTERWARDS. Called: board-material 46 in a band of 44 to 48, protocol 141 in a band of 138 to 145, void 7, thrown 1. Read: 46, 141, 7, 1. The brief's premise was that the only change since the 46 was counter declarations. Byte-diffing the two frozen snapshots showed four of 26 files had moved and the fourth was not a counter: `data/smogon-priors.json` advanced from Smogon month 2026-07 to 2026-08 and from 284 species to 283, and it reaches the run through `engine/champions_sim.js` into `engine/set_priors.js` and `engine/smogon_priors.js`. **It was named as the falsifier before the run and then measured inert —**

`classes` and `first_divergences` byte-identical, every differing field a stamp or a live corpus count. That is the difference between a prediction that held and one that was lucky, and it is the only reason the 4-of-4 is worth recording. Running record: 2-of-3, 4-of-4, a one-game protocol miss, 4-of-5, now 4-of-4.

A FIGURE MEASURED AGAINST BYTES NO COMMIT CONTAINS CANNOT BE REPRODUCED, AND ONE WAS. The release that first produced the 46 was cut from a working tree reachable from no commit. It is superseded by one cut from committed bytes, and every figure above names the latter — but the comparison between them had to be made by diffing frozen snapshots rather than by `git diff`, and if the snapshot directory is ever pruned that measurement is gone. Recorded now rather than discovered in a month.

WHAT IS OWED HERE.

- **Leaf calibration — this division's one number — stays WITHHELD.** `data/winrate-backtest.json` is downstream of MEDICHAM and the gate has not opened. It was not run. The standing brief to point the backtest at the current leaf and publish a reliability curve cannot be honestly discharged while the board-material clause is non-zero: a calibration curve measured through an engine whose boards part at the rate now published in `data/game-differential.json` would be a number wearing a receipt.
- **The MAG refit stays OWED and it is a REFIT, not a restamp.** The damage table moved from 318 species to 322, so the feature FUNCTION's input changed. No fitted vector was written and no fit was started; MAG is paused by its owner's decision.
- **The narration baseline restamp is available and unrun**, so direction of travel on protocol divergence stays withheld. Two pins are two instruments.
- **The register's coverage gap is unchanged:** 40 open rows assert breakage with no instrument, 7 answer nothing usable, 3 name a green one.
- `node engine/status.js --write` **DID run**, last in the sequence and for the first time in three releases, so the `<!-- GENERATED -->` block at the top of this file is current and was not hand-edited. `docs/WEB.md` carries no generated block, so four of the five ledgers are stamped.

MORE THAN HALF THE UNVERIFIABLE DEFECT CLAIMS WERE ALREADY FALSE — OPEN ROWS 261 → 237, OPEN-AND-ASSERTING-BREAKAGE 74 → 50. NO WHOLE-GAME FIGURE WAS RE-MEASURED, SO THE BOARD DID NOT MOVE. 2026-09-04, CHANGELOG 5.251.0

READ THIS RELEASE AS THE REGISTER AND THE COUNTERS, AND NOT AS ENGINE PROGRESS. No mechanic changed, no differential was re-run, and nothing in this block is evidence about the simulator. `data/verification/fix-batch-M6instr-defog.json` still holds board-material 46 of 961 and protocol 141; `data/game-differential.json` still holds 77 of 961 and is what the gate clause prints. Both are carried forward unchanged and neither is re-derived here.

THE UNIT OF THIS PASS IS THE CLAIM, AND THE MAJORITY OF THE POPULATION WAS ALREADY FALSE. The register carried 43 open rows that asserted a defect while nothing in the repository could confirm or refute them — the gap this division reported as the real one in the previous release. Worked row by row: **24 were already fixed** and **6 were duplicates of another row**. Open rows fell **261 → 237** and open-and-asserting-breakage fell **74 → 50**. A backlog whose majority is false is not a backlog, it is an unmeasured ruler, and every judgement of the form *what is still broken* was being taken over a population that did not exist.

NONE OF THE TWENTY-FOUR WAS CLOSED ON "THE TRIAGE SAYS SO", AND THAT IS THE ONLY REASON THE NEW INTEGERS ARE WORTH ANYTHING. Every closure carries its own evidence inside the cell — a live census row, a roster verdict, a tag value, or a knob-and-counter pair — dated, with the prior status carried forward verbatim rather than overwritten. **22 of the 24 were re-verified against the newer commit** rather than against the one the triage had read, because a commit landed between the two passes and every artifact had to be re-staged. A triage is a hypothesis about the register; a closure is a measurement. Closing on the hypothesis would have produced identical integers and none of the warrant.

THE REPAIR THIS DIVISION WAS HANDED DOES NOT WORK, AND THAT WAS MEASURED BEFORE ANYTHING WAS TOUCHED. The specified fix was to escape the pipes inside a roadmap cell as `\|`. It changes nothing: the status-cell capture takes the text after the LAST pipe through a negated-pipe character class, which stops at `\|` exactly as it stops at `|`. Proven on a synthetic row driven through the shipping detectors — as authored and fully escaped both extract the same wrong fragment, and only removing the pipes extracts the real status. **The pipes were removed and the shared closed-row detector was left alone:** it is exported so that one detector answers *is this row closed* for everything, and mutation testing showed it can be replaced with `return true` while all **159** of its assertions still pass. Editing an unproven detector to cure a symptom in its caller is how a ruler gets quietly broken.

THE SWEEP FOUND A CLASS, NOT AN INSTANCE, AND THE CLASS WAS REPORTED AND LEFT. 98 rows carry **669** unescaped pipes inside inline code and **22** have their status cell cut by one. Only the row that started this also asserted breakage. **Nine of the remaining 21 read OPEN against an authored closed or PART DONE cell**, and one parses empty. They were not closed here: none reaches the gate, and inferring a closure from a parse artifact is precisely the error this session had already paid for once. Their owners must restate them. A parse artifact is a fact about the ruler and never about the world it measures.

FOUR COUNTERS WERE NaN, AND ONE WENT OUT AS null IN TWO PUBLISHED ARTIFACTS. A field incremented into an object that never declared it reads `NaN` forever. `MEDSEEN.retaliateWhenLowered` was published as `null` in `data/million-run.json` and in `data/million-run-150k.json`: the capability fired and the counter recorded nothing, twice. `roostRiderNoPrimary` was declared inside the wrong object literal while being incremented on another, so one read zero forever and the other was `NaN` — the documented *compares undefined to zero and can never go red* trap, live in the tree. **No || 0 was added at any increment site;** that hides which object owns the counter, and ownership was the defect. A capability that cannot prove it ran is assumed broken, and a counter reporting `NaN` is worse than an absent one, because it is a capability wearing a receipt.

THE BRANCH THAT FIRES WHEN AN ARTIFACT WAS MEASURED AGAINST A DIFFERENT RELEASE WAS UNREAD. `engine/quarantine.js` asserted that the pin guard had proved it ran, quoted the founding rule in its own comment, and then named five of six branches — omitting the one that is the entire reason the guard exists. It was built by this same session, hours earlier. The branch list is now DERIVED from the guard object's own keys rather than typed, so a sixth branch is covered with no edit. Selftest **216 → 218 passed, 0 failed.** One correction to the audit, made in the code rather than in a report: the branch WAS already being driven; what did not exist was any reader of the counter.

THE RULE BEHIND THE FOUR IS A PROPERTY, AND IT NEEDS NO EXEMPTION LIST. `tests/test-counter-init.js` asserts that every increment targets a declared field. Across **507** files and **602** counter literals it finds exactly **4** violations with zero false positives, so nothing had to be excused — which is the strongest available evidence that a rule is the right shape, and the opposite of every check here that had to

be taught its own exceptions. It was **shown RED on all four first**, naming each with a real line number, it carries a synthetic red-proof arm so it cannot pass by asking nothing, and it names the **6** computed-key increments it honestly cannot decide.

FIVE COUNTERS THE AUDIT CALLED UNREAD ARE READ, THROUGH `Object.assign`, **WHICH A NAME GREP CANNOT SEE.** Recorded as a correction to the audit rather than acted on: they reach published artifacts by spread, so the counter's name never appears beside a reader. Two more were declared unread in the source with their reasons rather than given a guessed assertion, and the second tier was decided per counter instead of mass-wired. `tests/test-no-silent-failure.js`'s `++` rule was deliberately NOT changed — it asks a different question, and it has already been corrected four times for exactly this over-reach.

THE HYGIENE FIGURES THAT SAY NO MARKER WAS INVENTED. The census was untouched at **829** live. `NOT A DEFECT` was written zero times across the **24** closures. The `VERIFIED BY` count reads **138** before and **138** after. Two of the triage's own claims were transcription errors and are recorded as such rather than absorbed, and one row's figure was corrected in its own cell to four causes across six games.

WHAT IS OWED HERE.

- `node engine/status.js --write` **was NOT run, for the third release running.** It computes through `engine/quarantine.js`, which was uncommitted in the worktree all night; running it would read a file mid-edit and write the result into every ledger's generated block. The `<!-- GENERATED -->` block at the top of this file is three passes behind and was not hand-edited. Read the gate, not the block.
- **The 47-entry counter-reader registry was deliberately NOT built.** Without `engine_counters_zero` measured first its entries would be guesses, and a registry of guesses is a ruler nobody can audit.
- **98 rows still carry unescaped pipes and 22 still have a cut status cell.** Reported, left, and owed to their owners to restate — nine of them read OPEN against a cell their author wrote as closed.
- `data/game-differential.json` **is still not republished.** The published clause reads 77 of 961; the current measurement is 46 and lives in the verification artifact, and every sentence quoting either names the artifact that holds it.
- **Leaf calibration stays WITHHELD.** `data/winrate-backtest.json` is downstream of MEDICHAM and the gate has not opened. The refit stays OWED, no fitted vector was written, and no fit was started.
- **Nothing was committed in this pass.** It wrote documentation only.

THE NUMBER COMES BACK DOWN — BOARD-MATERIAL 53 OF 961 → 46 OF 961, PROTOCOL 154 → 141 — AND THE SEQUENCE THIS DIVISION PUBLISHES IS 77 → 61 → 50 → 53 → 46, RISE INCLUDED. 2026-09-04, CHANGELOG 5.250.0

A SEQUENCE WITH THE UNWANTED STEP REMOVED IS AN EDITED SEQUENCE, AND THIS DIVISION IS THE ONE THAT HAS TO REFUSE TO EDIT IT. The rise to 53 was the previous pass correcting an address that had been hiding four broken games behind a coincidence; the fall to 46 is this pass closing the other half of the same defect and a Defog that swept the wrong side. Both steps are engine work on identical pins — release `7ffc58da8ef8` → `252025cfddc`, same census pin, same frozen team pool, same empirical steering driver — with VOID at 7 throughout and the census level at 829 of 829. Writing the evening as *77 down to 46* would be true and would also be a smoothed curve. It is written in full.

THE SAME TWO-NUMBER PROBLEM, ONE PASS FURTHER ON. The gate clause reads board-material **77 of 961** and that is the PUBLISHED number; the re-measured **46 of 961** lives in `data/verification/fix-batch-M6instr-defog.json`, together with protocol **141**. There are now five correctly-measured values of one quantity in this session's record and exactly one of them is published. The rule stands as written at 5.245.0: a figure travels with the artifact that holds it, in the same sentence, every time.

A CAUSE COUNT IS NOT A GAME COUNT, AND THE 13-AGAINST-7 GAP IS THE WHOLE LESSON OF THIS PASS. The confusion defect was fourteen CAUSES, not fourteen games. Thirteen were damage-value disagreements and all thirteen closed, with that class falling **22 → 9**; the fourteenth is a stream-position defect carrying a confusion line at the join, a different family, untouched and unmovable by this fix. Board-material nevertheless fell by only **7**, because **6 of the 13 games carry a second, already-counted divergence**. Every other class is unchanged to the game, so those six were already in another class rather than having moved into one — which class, per game, is unmeasured. **A report that closes N mechanisms and claims N games has made this exact error**, and this division's job is to make the unit explicit before the number is quoted.

THE SCOREBOARD WAS CALLED IN WRITING BEFORE THE RUN, IT SCORED 4 OF 5, AND THE MISS IS THE HEADLINE FIGURE. Board-material was predicted at **39**, range 39–43, and came in at **46**. Protocol was called at 140–148 and read 141; VOID at 7 or lower and read 7; the census unchanged and was unchanged; the Defog fix was called to move the pinned pool by **exactly zero** and did — an ally-aimed Defog is a play nobody makes, and the pinned pool replays real human clicks, so the lab was the only scoreboard that could move. The miss has one cause and it is not a new one: *a game with a second cause stays counted* was named IN the prediction and then under-weighted when the integer was chosen. Running record on called scoreboards: 2-of-3, 4-of-4, a one-game protocol miss, now **4-of-5**. It is recorded because a prediction whose misses are not counted is not a prediction.

THE COMPARISON SPANS A CHANGED INSTRUMENT AS WELL AS A CHANGED ENGINE, AND THAT DISCLOSURE IS THIS DIVISION'S. `PIN_DIGEST` moved `ccb365985023` → `bc38e47d94f`: half of the confusion fix is the differential's own wrapping of the authority's `getConfusionDamage` draw, so the ruler and the measured thing moved in one pass. **The 13 → 0 cause count survives it** — a cause either names the same event on both sides and disagrees or it does not. **The board-material integer does not, and it is not presented as if it were one-variable.** Where an integer cannot be attributed to a single change, the caveat IS the result; a cleaner number here would be a false one.

****THE facts are global RULE WAS CORRECTLY NOT APPLIED, AND THAT IS WORTH AS MUCH AS A FIX.**** Three `sweepField` call sites were reported as three copies of one selection. They are not: Tidy Up is `target: 'self'` and names the mover's own side plus the foe sides with conditions, so consolidating it into the Defog repair would have collapsed its two bags into one and left the opponent's hazards standing; the damaging branch never reads the argument at all. Two sites that look identical can be asking different questions. Both memberships are DERIVED from the tag corpus and the authority text on every run rather than asserted, so a new carrier fails by name.

FOUR UNCOMMITTED DECLARATION ROWS WERE DESTROYED BY A `git checkout --` DURING THE BATCH, AND IT IS REPORTED AS A LOSS RATHER THAN A REPAIR. Two were restored verbatim from a printout taken before the loss. Two are RECONSTRUCTIONS, labelled as such in the file itself, each quoting the original answer verbatim from the expired key's own text and asking to be re-read by its author. The code they cover has not moved a byte and the census verifies that by digest, so the substance survived. `git checkout --` **on a file another session has edited and not committed is unrecoverable — git holds nothing to restore from.** The symptom was an undeclared count rising with no engine change, and it was chased as instrument drift before it was recognised as self-inflicted, which is the same diagnostic trap as *suspect the instrument before the engine* running in the opposite direction.

REGISTER COVERAGE, MEASURED, AND THE HONEST NUMBER IS THE ONE NOBODY ASKED FOR. Of **219** open rows: **13** name a marker the classifier admits, **43** declare `INSTRUMENT OWED` in writing, **0** carry a marker the classifier refuses (nine the night before). **163 of the 219 carry neither, and 43 of the 74 rows that assert breakage are among them** — for those rows the clause holding the open-defect gate shut rests on prose that nothing can confirm or refute. ROADMAP #381 measured 35 of 49 on 2026-08-23: the declared half has improved and the undeclared half has grown with the register. Both are true and reporting only the first would be the flattering half.

WHAT IS OWED HERE.

- `data/game-differential.json` **IS STILL NOT REPUBLISHED.** The published clause reads 77 of 961; the current measurement is 46 and lives in the verification artifact. Every document quoting 46 names that artifact in the same sentence.
- `node engine/status.js --write` **was NOT run in this pass.** The `<!-- GENERATED -->` block at the top of this file is one pass behind and was not hand-edited. Read the gate, not the block.
- `tests/staged_board.js` **exits 1 on a pre-existing red** — 25 of 25 scenarios board-identical, and the `fakeout-flinch` proof plant's anchor is absent at `HEAD` as well as in the working tree. Its own 2026-08-19 header files the remaining half in `engine/engine_release.js`, which is this division's file. Reported, not patched, and not filed as a known failure.
- **Two premature-close candidates are filed with evidence and NOT reopened.** Their rows and their shared instrument are byte-identical to the tree that produced the verdict, and the instrument was not run because its whole-repo pass scans files a live agent owns. A third row is flagged open-and-asserting-breakage while its unchanged instrument reads green.
- **Four closures rest on markers the register has never once executed.** A closure resting on an unrun marker is a closure resting on prose.
- **M2 stays FENCED rather than split,** and the fence is deliberate: the proposed split does not survive the arithmetic.
- **The surviving ordering cause has no probe,** and six of the thirteen games whose confusion cause closed remain board-material on a second cause that is not attributed per game.
- **Leaf calibration stays WITHHELD.** `data/winrate-backtest.json` is downstream of MEDICHAM and the gate has not opened. The refit stays OWED; `data/policy-weights.json` was not written and no fit was started.
- **Nothing was committed.** This pass wrote documentation only.

THE HEADLINE NUMBER GOT WORSE AND THAT IS THE FIX WORKING — BOARD-MATERIAL 50 OF 961 → 53 OF 961, PROTOCOL 150 → 154 — AND THE LARGER RESULT IS THAT 9 OF 124 REGISTER MARKERS WERE BEING SILENTLY REFUSED.
2026-09-04, CHANGELOG 5.249.0

A RISE IS NOT AUTOMATICALLY A REGRESSION AND THIS DIVISION IS THE ONE THAT HAS TO SAY WHICH IT IS. The confusion self-hit drew both of its values from one address where the authority uses two — `data/conditions.ts` writes `this.activeTarget = pokemon` between the roll and `getConfusionDamage`, so the two draws have different addresses whenever the confused body clicked at a foe. **THE DEFECT'S TRUE SIZE IS 14 GAMES. FOUR HAD BEEN HIDDEN BY A SHARED COIN LANDING THE SAME WAY ON BOTH ENGINES,** and correcting the address removed the coincidence. So the count rose because the instrument can see more, not because the engine got worse.

THE BAR FOR ACCEPTING THAT SENTENCE IS ATTRIBUTION, NOT PLAUSIBILITY, AND IT WAS MET IN ADVANCE. Three things carry it, in the order that matters. The direction was written down BEFORE the run, in `data/verification/2026-09-04-M6-address-prediction.json`, as neutral-to-slightly-worse with the arithmetic; protocol missed the stated band by one, which is stated rather than smoothed. Exactly one divergence class moved — `-damage field 18 to 22` — and nothing else in the run changed by a single game. And eight of the eight moved causes carry `[from]confusion`. A rise that were a real regression would be spread across classes; this one is confined to the class the edit touches. **THE FIX IS KEPT, NOT REVERTED.** The revert is one environment variable and it is the owner's call, named rather than taken.

THE SAME TWO-NUMBER PROBLEM, ONE PASS WORSE. The gate clause reads board-material 77 of 961 and that is the PUBLISHED number; the re-measured 53 of 961 lives in `data/verification/fix-batch-M6-sidesel.json` and protocol 154 with it. `data/game-differential.json` was not rewritten again in this pass. Both are correctly measured on identical pins and only one is published — so every clause in this ledger names its artifact in its first words, every time.

AND NOW THE PART THAT IS PROPERLY THIS DIVISION'S: THE RULER WAS LYING ABOUT ITS OWN COVERAGE. A `VERIFIED BY:` marker is a register row's claim that an instrument decides it. `engine/register_reality.js` carried a `SAFE` predicate that refused 9 of 124 markers outright — the row claims an instrument, the tool declines to run it, and nothing anywhere says so. **2 of the 9 rows that both assert breakage AND name an instrument were among them, so 22% of this register's live instrument coverage was fictional.** Worse than the count: a refused marker was published under the SAME verdict a gate produces when it ENOENTs or is killed — so *my ruler would not read this marker and this instrument is broken* were one number and one label.

ONE NUMBER WAS HIDING TWO DEFECTS AND THE SPLIT IS THE FIX. In the last published artifact **all 27 rows counted as an unrunnable instrument were `SAFE` rejections** — a defect in the RULER and a defect in the WORLD summed under one heading. `engine/quarantine.js` now reports `MARKER REJECTED` (nobody was asked) apart from `unrunnable` (asked, answered nothing usable), each with its own count, its own sentence and its own per-row verdict. Selftest **210 → 216 passed, 0 failed**, all six new arms shown RED first on four separate deliberate breaks, and `tests/test-divergence-composition.js` stayed green **without being edited**.

THE REPLACEMENT IS A PROPERTY, NOT ANOTHER LIST OF SPELLINGS, BECAUSE THE LIST WAS TRIED AND FAILED. ROADMAP #521 fixed this exact class once by enumerating one wrong form, and nine markers walked past it. `classifyMarker()` reasons about the shape instead: there is no shell in the path (`execFileSync`, never `shell:true`), so node reads only the tokens BEFORE the entry point, that region **fails closed**, the entry must resolve inside the repository, and post-entry tokens are inert by construction. `--arm middle` and `--arm=middle` are now one fact, and 22 hostile strings are still refused. The evidence that this is a widening and not a hole is the full-population regression: **124 markers, 115 identical argv, 0 moved, 0 lost**. Selftest 58 → 73.

AND THE GUARD THAT WAS REMOVED WAS ITSELF FICTIONAL, WHICH IS WHY THIS IS NOT A RELAXATION. Its comment claimed bare values were refused to keep multi-minute game-playing runs out of the register pass. **Measured on the pre-fix bytes, that is false:** `--arm=middle`, `--stage=moves` and `--games=1200` were already admitted, one character away from the forms it rejected. It guarded spelling, not cost. A guard whose stated justification has never been measured is indistinguishable from one that works — which is the same shape as the staleness gate that hashed a field no row carries.

THE LIVE rejected COUNT IS 0 AND THAT IS HONEST RATHER THAN A FIX THAT DID NOTHING. `data/register-reality.json` is the 2026-08-27 artifact and predates the new vocabulary, so the affected rows are still spelled with the old label in it. The mechanism is in the tool; the artifact has not caught

up. **Regenerating it is blocked on an undecided question and that decision is FILED, not half-guarded:** four newly-runnable markers are heavy and two of them rewrite gate inputs, so the full register pass must not be run beside a live agent until an owner decides.

THE BEST JUDGEMENT IN THE PASS WAS A REFUSAL TO WRITE A MARKER. ROADMAP #318 is an OPEN row asserting breakage. The obvious repair — dropping the unexpanded environment prefix, exactly as was done on four sibling rows — would have produced a marker that **reports the row GREEN while the defect it names is untouched:** `tests/roster.js:9167` passes `learnsetMode: 'report'`, so 632 learnset refusals go to a printed bucket and never reach the count that becomes the exit code. No marker was written; the row's `INSTRUMENT OWED` declaration is its honest coverage. **A green instrument that cannot see the defect is worse than a declared gap,** and this is the second time in two passes that the cheap repair would have manufactured coverage.

THE SIDE-SELECTION ALARM WAS THE MEASURING TOOL AND NOT THE ENGINE — CORRECTLY DIAGNOSED, AND WITH A REAL DEFECT UNDERNEATH IT.

`engine/side_selection_census.js` read undeclared 84 against a ratchet of 81 with four sites outside every hunk of the session's diff. All four are byte-identical to code classified on 2026-08-29 — same expression, same site digest — and what moved is the census's ANCHOR: a new branch test above two of them, and the enclosing function drifted **1,660 lines against a 1,500-line window** for the other two. Three declarations expired for lines that never changed. All four were re-derived from the authority and found correct; undeclared **84 → 80**, ratchet met, exit 0. The new `ANCHOR-DRIFT` verdict was shown against a control that flags exactly those four and none of the other eighty. **Suspect the instrument before the engine** held again — and it did not become an excuse, because proving the fourth site innocent surfaced a live Defog target-side defect at `engine/medicham2-browser.js:26471` whose old note had named the wrong line.

WHAT IS OWED HERE.

- `data/game-differential.json` **IS STILL NOT REPUBLISHED.** The published clause reads 77 of 961; the current measurement is 53 and lives in the verification artifact. Every document quoting 53 names that artifact, which is a caption and not a fix.
- `node engine/status.js --write` **was NOT run in this pass.** The `<!-- GENERATED -->` block at the top of this file was last stamped earlier today and is not restamped here — read the gate, not the block.
- **The instrument half of the confusion defect is open.** `engine/game_differential.js` reads that draw un-inverted where this engine mirrors it; ENGINE cannot flip it alone without the corner arm parting on every self-hit. It is worth all 14 games and it moves the pin digest, so it needs its own before/after.
- `data/register-reality.json` **is stale and is WITHHELD, not annotated.** It is the 2026-08-27 artifact, taken against 416 register rows and 112 markers where the register now carries more of both; its `INSTRUMENT UNRUNNABLE` bucket is the pre-split vocabulary.
- **Leaf calibration stays WITHHELD.** `data/winrate-backtest.json` is downstream of MEDICHAM and the gate has not opened. The refit stays OWED; `data/policy-weights.json` was not written.
- **80 side selections remain undeclared,** and three of the sites examined are three copies of one selection.
- **The Defog target-side defect** at `engine/medicham2-browser.js:26471` is ENGINE's and is not closed here.
- **Nothing was committed.** This pass wrote documentation only.

THE GATING QUANTITY MOVED A SECOND TIME IN ONE SESSION — BOARD-MATERIAL 61 OF 961 → 50 OF 961 — AND THIS

DIVISION STILL HAS NOT PUBLISHED IT. 2026-09-04, CHANGELOG 5.248.0

THE SAME TWO-NUMBER PROBLEM, ONE PASS WORSE. The rule stands as written at 5.245.0: each clause names its quantity in its first words and a figure travels with the artifact that holds it. So — **the gate clause reads board-material 77 of 961 and that is the PUBLISHED number**; the re-measured 50 of 961 lives in `data/verification/fix-batch-M5M7M8.json` and is NOT published. There are now three correctly-measured values of one quantity in this session's record (77, 61, 50) and exactly one of them is what the gate prints. Republishing is this division's job.

THE COMPARISON IS PINNED, WHICH IS WHAT MAKES ELEVEN GAMES ATTRIBUTABLE. Release `f3504e5f88d6` → `9b449a41c865` with every other pin identical, 961 games, the empirical steering driver. Protocol first-divergence — the clause that reports without gating — is `161` → `150`. VOID did not move and the mechanics census was level throughout, so neither can be offered as a side effect.

ALL FOUR PREDICTIONS HIT. THAT IS A RESULT ABOUT THE METHOD AND IT IS NOT THE BASE RATE

The previous batch scored two of three and this ledger recorded the miss in its own words. This one scored four of four. **Both entries stand together or neither means anything** — the reason to write a prediction down before a run is that the record is kept when it is unflattering, which is the same discipline as never reading an interim SPRT and reporting the bound's verdict rather than a p-value.

A PROBE WAS VACUOUSLY GREEN AND A COUNTER CAUGHT IT. THE BOARDS COULD NOT HAVE

M8's first fixture had the holder clicking Protect, so the contact move never reached the handler on EITHER engine, and the two boards agreed about a mechanic that never fired. **A comparison cannot detect its own silence.** This is the noise-floor discipline in a different costume: before believing an agreement, measure that the thing under test ran at all. It is also the ninth entry in this repository's command-that-succeeds-having-done-nothing class — `--out` without `--write` writes no file and exits 0 — and that one was written by the coordinator into the brief, not found in the tree.

THE INSTRUMENT BEHIND THE RETRACTED UNBURDEN CLAIM NOW MEASURES THE ENGINE, AND REFUTES THE CLAIM ON ITS OWN FIXTURE

`tests/probe_leaf_widening.js` compared its own stand-in against the authority; that is what produced the claim retracted at 5.246.0. It now calls `effSpeed` out of the same frozen release the game is played on, twice per body, against the authority's `getStat('spe')` modified versus unmodified. Both arms sit in one game: a non-carrier that lost its item reads `effSpeed 122<122 x1` and does not register, and the carrier reads `344<172 x2` and does. **It was shown red twice before being trusted** — injecting the retracted claim prints `DISAGREE`, and forcing `effSpeed` to throw prints `COULD-NOT-MEASURE`, which is an instrument declining to make a claim rather than making a weak one. ROADMAP #535's engine half is measured at `effSpeed 244<122 x2`, so the doubling follows the CURRENT ability; its authority half stays `INSTRUMENT OWED` and no citation was invented for it.

~60 UNRUN CHECKS WERE TRIAGED AND ZERO WERE WIRED, WHICH IS THE CORRECT OUTCOME

RUNNABLE NOW 0 · NEEDS A BASELINE CHOICE 65 · NOT A CHECK 1, with unaccounted moving 66 → 63. **63 of the 66 load a simulator and an agent was editing that simulator throughout.** Certifying a red arm and a green arm against a moving tree is not a certificate; it is the photograph rule broken at the

instrument instead of at the run. The PENDING-WIRE audit reads 0 of 36 wireable today, 34 blockers verified still true and 2 carrying stale text, corrected.

OWED

- `data/game-differential.json` **IS STILL NOT REPUBLISHED.** The gate clause prints 77 of 961; every document quoting 50 must name the verification artifact. This is MEASURE's job and it is now one pass further behind than it was.
- `node engine/status.js --write` was NOT run — the game-playing slot was held — so the `<!-- GENERATED -->` block at the top of this file is one pass behind. Read `node engine/status.js`, not the block.
- **CUTTING THE RELEASE STRANDED PINNED ARTIFACTS** — `engine-diff`, the roster stages and `all-mechanics-fire` name a release the tree has moved past. They are WITHHELD until regenerated. That is the pin guard working; a stranded artifact is re-measured, never resurrected.
- `engine/side_selection_census.js` **EXITS 1**, undeclared 84 against a ratchet of 81, with all four new sites arriving in earlier commits and nothing having run the check.
- **9 of 124** VERIFIED BY: **MARKERS ARE REFUSED BY** `engine/register_reality.js` 's OWN **SAFE PREDICATE AND READ AS** `NOT_STARTED`. The row claims an instrument, the tool declines to run it, and the row is reported neither verified nor unverified — an unaccounted check shows RED and these show GREEN. ROADMAP #521 fixed this class once and fixed only the `-r <preload>` case. This is a ruler defect and therefore this division's.
- **Leaf calibration is unchanged and still WITHHELD**, as is everything downstream of MEDICHAM. The gate did not open, and the standing `data/winrate-backtest.json` item is untouched by this pass.
- Full accounts: `docs/_reports/2026-09-04-fix-batch-M5M7M8.md`, `docs/_reports/2026-09-04-probe-leaf-widening-repair.md`, `docs/_reports/2026-09-04-unrun-checks-triage.md`, and `docs/_reports/2026-09-04-living-docs-5248.md`.

THE GATING QUANTITY MOVED ON ENGINE WORK FOR THE FIRST TIME IN THREE VERSIONS — BOARD-MATERIAL 77 OF 961 → 61 OF 961 — AND THIS DIVISION HAS NOT PUBLISHED IT. THE 61 IS IN `data/verification/fix-batch-M1M3M4.json` ; `data/game-differential.json` **WAS NOT REWRITTEN, SO THE GATE STILL PRINTS 77. 2026-09-04, CHANGELOG 5.247.0**

TWO CORRECTLY-COMPUTED NUMBERS, ONE PUBLISHED. THAT IS THIS DIVISION'S PROBLEM AND IT HAS ALREADY COST THREE RECONCILES IN ONE SESSION. The rule stands as written at 5.245.0: both clauses name their quantity in their first words, and a figure travels with the artifact that holds it. So — **the gate clause reads board-material 77 of 961 and that is the PUBLISHED number**; the re-measured 61 of 961 lives in the verification artifact named in the heading above and is not published. Neither is wrong. They are not the same question, and pooling them or quoting the better one without its artifact is the failure the two-clause split was built to stop.

THE COMPARISON IS PINNED AND THAT IS WHAT MAKES SIXTEEN GAMES

ATTRIBUTABLE. Release `8ad06030e129` → `f3504e5f88d6` (engine digest `5391ae37d4d8`), census `9446a684709d`, pool `0d103fb9fa87`, `--steering empirical`, 961 games — the pins are IDENTICAL across the two runs, so the three fixes are the only variable. A photograph rule that pins only the code would not have carried this claim: the store moves hourly and the census is steered, so both were pinned too. Protocol first-divergence — the clause that reports without gating — is 168 → 161, and VOID is 8 → 7.

THE PREDICTION WAS WRITTEN DOWN BEFORE THE RUN AND SCORED TWO OF THREE. THE MISS IS THE ENTRY THAT MATTERS

Board-material called **60–70**, landed **61**. Protocol called **about 162**, landed **161**. VOID called **0–2** and **DID NOT MOVE — it stayed at 7. That is a MISS**. The stated reason is that M1 removed one divergence shape entirely while the games carrying it retain OTHER unshared addresses, so they stay void. **A prediction whose misses are not counted is not a prediction, it is a summary written afterwards** — the same failure shape as reading an interim SPRT and stopping at the number you like. Two of three is the score; it is not "the prediction held".

A REGRESSION WAS CAUGHT BY THE CENSUS AND NOT BY THE PINNED POOL, WHICH IS THE TWO-SCOREBOARD RULE PAYING FOR ITSELF

The first form of the M4 fix freed a lock on a fixture body clicking off its own move list, taking the mechanics census **829 → 828**; narrowed to *did the move list CHANGE*, it returned to 829. **The pinned pool measured identically on both forms of the fix**. The pool is usage-weighted by construction and answers *does this matter*; the lab stages one scenario per entity and answers *is this correct*. Only one of them could see this, and it was not the one carrying the headline.

A GATE THAT DIES MID-AUDIT REPORTS A PARTIAL VERDICT, AND A PARTIAL VERDICT IS NOT A VERDICT

`tests/test-artifact-rerunnable.js` audits 91 stamped artifacts across 30 releases and died at node's default heap PARTWAY THROUGH — after a STRANDED line and before its clean summary — so it **blocked commits on a partial verdict that read WORSE than the truth**. With the heap set it is all green, 5 checks. **The real fault was not the file**. `.github/pre-commit` ran its whole gate loop with bare `node` and ignored `ABRA-HEAP`, so every gate in that loop carried the same trap; the hook now reads each script's own declaration the way `tools/lownode.cmd` does. This is a ruler defect and therefore this division's: an instrument that answers with whatever it managed to say before it died is worse than one that fails cleanly, because the partial answer is well-formed and gets believed.

OWED

- `data/game-differential.json` **IS NOT REPUBLISHED**. Until it is, the gate clause prints 77 of 961 and every document quoting 61 must name the verification artifact. Republishing is this division's job, not ENGINE's.
- `node engine/status.js --write` was NOT run — an engine agent held the game-playing slot — so the `<!-- GENERATED -->` block at the top of this file is one pass behind and still reads `5 of 8 gate clauses fail`. Read `node engine/status.js`, not the block.
- **The per-hit damage differential's currency is unknown, not clean**. `tests/test-engine-diff.js` was not re-run and the simulator moved, so unlike the previous two passes the reason cannot be *nothing that feeds it moved*. `node engine/provenance.js` compares content and decides it; this ledger does not restate the per-hit result.
- **Leaf calibration is unchanged and still WITHHELD**, as is everything else downstream of MEDICHAM. The gate did not open; a smaller number of divergent games is not an open gate, and the standing winrate-backtest item is untouched by this pass.
- Full account: `docs/_reports/2026-09-04-fix-batch-M1M3M4.md` and `docs/_reports/2026-09-04-living-docs-5247.md`.

RETRACTED — "EVERY BODY IN THIS ENGINE THAT LOSES AN ITEM GETS UNBURDEN'S SPEED DOUBLING" IS FALSE, AND THE

INSTRUMENT WAS WRONG BEFORE THE ENGINE WAS. 2026-09-04, CHANGELOG 5.246.0

THIS RETRACTS A SENTENCE THIS LEDGER PUBLISHED HOURS EARLIER, in its own OWED list and in four living documents. The dated section below stands as written and is superseded from here. The OWED bullet that carried it is a WORK ITEM rather than evidence, so it is corrected in place with its old wording quoted — **a false open item sends someone to fix a non-defect**.

WHAT THE CODE DOES. `engine/medicham2-browser.js:14770` is the ENTRY GUARD (`m._hadItem && !m.item`). The multiplier is pushed on the next line, `:14772` , gated on `TAGS.param('ability', m.ability, 'speedOnItemLoss')` . Walking the ability block of `data/tags.json` returns exactly one key carrying that parameter: `unburden` .

arm	expectation	census reading
ability none	must NOT move	187,187
Unburden	must move	187,374

That control was already on record before the claim was published. It is the knob-cleared pair this division asks for by rule, and it separates the two readings on its own.

THE MECHANISM, WHICH IS THE DURABLE HALF. `tests/probe_leaf_widening.js:277` holds `const stand = m => (m && m._hadItem && !m.item) ? 1 : 0` and reports THAT as the engine's side. It never read `effSpeed` . Its agreeing reading therefore asserted one thing only — *both bodies had lost an item* — and could not have detected the multiplier, present or absent, in either direction. **An instrument that compares its own stand-in against the authority can only ever confirm its own stand-in.** A noise floor cannot save that: the floor tells you whether an effect is bigger than the spread, and it says nothing at all when the instrument is not measuring the quantity written on its label.

AN INSTRUMENT CAN BE WRONG BEFORE THE ENGINE IS, AND IT WAS THREE TIMES IN ONE NIGHT — a probe printing `undefined<>undefined` as its receipt, a selftest exercising a five-line copy of the rule it was checking, and this one. Two were caught by the agents that owned them; this one was relayed as fact without being verified at the line.

WHAT IS ACTUALLY OWED IS NARROWER: ROADMAP #535. The doubling is recomputed from the body's CURRENT ability rather than held as the volatile the authority grants at the moment the item is lost. So an Unburden acquired AFTER the hand is already empty doubles here and does not in the authority; Skill Swap is the reachable door. Filed `INSTRUMENT OWED` — nothing in the tree decides it, and the row says so rather than citing a probe that does not exist.

FIVE OF THE SIX ROWS FILED THIS PASS SAY `INSTRUMENT OWED` **AND ONE CARRIES A** `VERIFIED BY` (register `497 → 503` rows, `251 → 257` open). That ratio is the honest state of the register rather than a shortfall to be tidied: a row that names the instrument it does not have is auditable, and a row that cites a probe nobody ran is not. `NOT A DEFECT` **was used in none of the six**, because `engine/quarantine.js:1040` matches that string in a status cell and treats it as a ruling that overrides the derived verdict — a casual note in a register cell is executable, and it has previously subtracted three live turn-order divergences from the gate.

NO FIGURE IN THIS LEDGER MOVES ON A RETRACTION. Nothing was re-measured in this pass, no artifact was regenerated, and no quarantined figure becomes quotable. `node engine/status.js --write` is still OWED, so the `<!-- GENERATED -->` block above is one pass behind — read the gate, not the block.

THE GATING CLAUSE WAS COUNTING THE WRONG QUANTITY WHILE THE RIGHT ONE SAT UNREAD IN THE SAME ARTIFACT.

BOARD-MATERIAL 77 OF 961 (8.0%) GATES; PROTOCOL FIRST-DIVERGENCE 167 OF 961 REPORTS. 2026-09-04, CHANGELOG 5.245.0

THIS IS A REVERSAL OF WHAT THIS LEDGER AND EVERY LIVING DOCUMENT HAVE BEEN PUBLISHING. The dated sections below stand as written and are superseded from here, not rewritten. Any sentence that says the whole-game clause counts protocol first-divergence, or quotes **167** as *the* whole-game number, is retired by this section.

clause	quantity	reading	gates
whole-game differential / BOARD-MATERIAL	board_material_games	77 of 961 (8.0%) — state.games 961 less state.games_board_never_diverged 884	yes
whole-game differential / NARRATION	protocol_first_divergence_games	167 of 961 — 168 raw, less the one declared row	no — REPORTS

THE TWO FIELDS ARE READ, NOT DERIVED. engine/game_differential.js writes both; the clause subtracts them and does nothing else. There is no second implementation of *did a board part* — engine/board_state.js decides that once and the differential records the answer, which is this division's own rule about a fact having one owner.

WHY THIS IS NOT A RELAXATION, AND THE ARGUMENT IS MEASURED RATHER THAN ASSERTED. Of the 168 protocol divergences, protocol_diverged_board_never_did is **102**: those games write no differing board leaf at any compared turn boundary. They are real work and they are narration. Going the other way, **11 of the 77 part a BOARD while the protocol never diverges at all** — derived from four artifact fields as material - (protocol_diverged_games - protocol_diverged_board_never_did) = 77 - (168 - 102), named as derived, and deliberately NOT clamped, so a negative reads as *the artifact contradicts itself* rather than as a clean bill of health.

THOSE 11 ARE THE REASON THIS SPLIT IS A TIGHTENING. Under the single clause a game whose board went wrong with nothing in the narration pointing at it left the count altogether — no cause, no class, no shape, nothing to grep. **A fix that closed a protocol divergence without repairing the board would have improved the headline without improving the engine.** That is the failure mode this division exists to catch, and it was live inside our own headline number.

NOTHING MAY BE SUBTRACTED FROM THE BOARD COUNT, AND THAT IS STRUCTURAL RATHER THAN STRICT. Both subtraction mechanisms — DECLARED_DIVERGENCE and data/decision-impact.json — attribute by protocol CAUSE over classes[].causes[] . A parted board is a leaf PATH and carries no cause, so there is no mapping to subtract through. The clause publishes a RAW count and says so; the one declared row subtracts from NARRATION, where it belongs, and cannot open this clause.

NARRATION IS A GATE THAT REPORTS, NOT A BACKLOG. It has a row, a count, a place in failing, and node engine/quarantine.js --narration exits 1 while it is red. The gates: false flag is applied by a WRAPPER to every return path including all the refusals — in the first draft only the final verdict carried it, so a stale or torn artifact would have made narration come back WITHHELD, default to gating, and hold the gate shut on the exact quantity the 2026-08-22 ruling took off the critical path. Two arms assert both refusal families keep the flag.

BOTH CLAUSES NAME THEIR QUANTITY IN THEIR FIRST WORDS AND CARRY A quantity FIELD. Two correctly-computed numbers printed side by side with only one published has caused three reconciles in a single session here. Never state either figure without naming which quantity it is.

GATE: CLOSED — 1 of 8 GATING clauses fail . Seven pass on artifacts that can now prove what produced them; the one failure is the engine backlog.

THE ARMS TESTED A COPY, AND A DELIBERATE BREAK COST ZERO OF THEM

Putting narration back on the gate should have gone RED and did not. The selftest arms were exercising a five-line reimplementation of the filter rather than the shipping one — **the verifier re-implementing the rule it is checking**, caught on itself. The rule is now `gateVerdict()` , exported, and the arms drive it. Selftest **186 → 210 passed, 0 failed**.

FIVE PINNED ARTIFACTS CAN NOW PROVE WHAT PRODUCED THEM, AND NO MEASURED FIGURE MOVED

`engine-diff` , `roster.{items,abilities,moves}` and `all-mechanics-fire` were regenerated and each carries **26** `source_digests` **plus** `showdown_commit` , verified by reading the fields rather than by trusting the run. `engine/pin_guard.js` is the one door that withholds on an absent, wrong or unverifiable pin. **Nothing measured moved:** `engine-diff` 0/6000 at the midpoint and both corners; `roster` 0 DIFFER / 0 DID-NOT-FIRE on all three stages; `mechanics` verdicts identical row for row. `Roster` ran without `--reds` , so `reds: []` carries forward — deliberate, to keep the regeneration a single-variable change.

THREE OF THE FIVE WERE IMPOSSIBLE UNTIL A GUARD ADDED HOURS EARLIER WAS REPAIRED, and the failure would have read as *skipped*: `game_differential.js` 's coverage-arm refusal ran at MODULE LOAD off the whole process's argv, so `roster.js --write` and `all_mechanics_fire.js --write` died at **exit 2 — the SKIP code** — before playing a game. `require.main === module` fixes it; both arms re-verified.

THE LEAF WIDENING MEASURED FLAT, AND THE PREDICTION WAS WRITTEN DOWN FIRST

`Board-material` 77 → 77, protocol **168 → 168**, on release `8ad06030e129` — **unchanged**, because `engine/board_state.js` is the comparator and not one of the 26 frozen sources, so the engine bytes were byte-identical and the comparator was the only variable. The call made in advance was *rise or stay flat, it cannot fall*. **Flat does NOT prove the three leaves clean** — that artifact records divergences, not per-leaf agreements, and `tests/probe_leaf_widening.js` is the separator.

THREE "PRE-EXISTING FAILURES" WERE NEVER FAILING, AND THAT WAS PROVED BEFORE IT WAS BELIEVED

A working-tree test written against the new `pin_guard.js` was run against a `git show HEAD:` copy of `quarantine.js` that predates it. A **VINTAGE GAP** — reproduced exactly by injecting HEAD's bytes, then disproved by finding the same arms green in the pre-rename working tree. The "separate undiagnosed defect" filed against them is RETIRED. A rename that silently turns something green is indistinguishable from one that hides it, which is why this was checked before it was committed. `tests/test-divergence-composition.js` is repointed to `narrationClause` — it composes divergence CLASSES, which is the narration quantity — and reads **18 checks, all pass**, was 17 with 10 red.

`git stash` **WAS REFUSED AS A DIAGNOSTIC HERE, CORRECTLY**. With `core.autocrlf=true` and 24 files of other agents' uncommitted work in the tree, a stash/pop can rewrite line endings — the same mechanism that has twice moved a release digest with no code change. A sparse HEAD worktree was used instead and removed afterwards.

OWED

- `node engine/status.js --write` was not run. The `<!-- GENERATED -->` block at the top of this file is one pass behind and **still says** 5 of 8 gate clauses fail ; the gate reads CLOSED — 1 of 8 GATING clauses

fail . Read the gate, not the block.

- **RETRACTED 2026-09-04 (CHANGELOG 5.246.0) — THIS OWED ITEM WAS FALSE AND IS REPLACED BY A NARROWER ONE.** It read, verbatim: *"Unburden is an ENGINE defect, not a comparator gap, and needs a register row. engine/medicham2-browser.js holds no state under that name; effSpeed recomputes the doubling from _hadItem && !m.item (:14770), so every body in this engine that loses an item gets the speed doubling."* **That is false.** :14770 is the ENTRY GUARD only; the push on the next line, :14772 , is gated on TAGS.param('ability', m.ability, 'speedOnItemLoss') , which exactly one ability key in data/tags.json carries. The evidence rested on tests/probe_leaf_widening.js:277 comparing its OWN stand-in rather than this engine's Speed. What is owed is ROADMAP #535 — the doubling is recomputed from the CURRENT ability instead of being held as the volatile granted when the item went — filed INSTRUMENT OWED . Full account in the retraction section at the top of this ledger.
- The weights restamp was NOT run, by Will's decision. data/policy-weights.json untouched, MAG paused. That the measuring path is free of it was verified in code: game_differential , roster , all_mechanics_fire , test-engine-diff and steering carry zero references to MAG or the policy weights, and the empirical driver samples data/move-priors.json — recorded human clicks, no model in the path.
- Full accounts: six reports under docs/_reports/2026-09-04-*.md .

THE REPOSITORY ALREADY KNEW AND NOTHING ACTED — SO THE INSTRUMENT THAT READS THE INSTRUMENTS IS THIS DIVISION'S. 2026-09-04, CHANGELOG 5.244.0

THIS IS ONE FINDING WEARING FIVE COSTUMES. engine/status.js was printing the row driver policies the gate quotes – 1 of 2 on line 103 of a 253-line output. engine/durable-ingest.js:464 explained the archive drift in a comment and instructed *"RUN THIS BEFORE ANY REPARSE"*. tests/run-all.js was red on its own unaccounted-checks clause. In each case the fact was DERIVED, was WRITTEN DOWN, and no consumer read it. This division's whole argument is that state is printed rather than typed — and printing is not reading. **A number that nothing acts on is not a measurement, it is a decoration.**

SO EVERY FIX IN THIS PASS CARRIES A LABEL, AND THE LABEL IS THE OWNER'S ACCEPTANCE TEST. CLASS-level means the failure becomes unreachable; INSTANCE-level means one door is shut. The question to ask of each: *would this catch a second instance, spelled differently, arriving through another door?* An INSTANCE-level fix is not wrong — it is incomplete, and saying which one you shipped is the difference between a repair and a re-arming.

CLASS-LEVEL — engine/sweep.js : WHAT EACH INSTRUMENT FAILS TO CHECK ABOUT ITSELF. Five derived sections, about three seconds, exit 1 on any finding. First run:

section	first run
checks nothing runs	60
gate clauses blind to their own staleness	3 of 8
published figures out of date	12
counter fields nothing reads	783 of 1,048
store rows with no raw log	614

IT WAS SHOWN RED BEFORE IT WAS TRUSTED, on every CANNOT DERIVE path, under SWEEP_BREAK=1. . 5 . That is not ceremony here: a sweep whose derivation silently fails reports a spotless repository, and this division has already paid for a ruler that printed a stable number because it could not see.

CLASS-LEVEL — MUTATION TESTING, PILOTED ON THE GATES AND NOT ON THE ENGINE. The engine has a strong external oracle in the official simulator. **The gates have none**, which is exactly how a gate reading `8 of 8 PASS` survived being false and why the pilot was pointed here first. `tools/mutate-gates.js` with `stryker.gates.conf.json` : **83 mutants, 57 killed, 26 survived, 28s**. A surviving mutant is a clause whose failure nothing observes — the same claim `tests/test-wiring.js` makes about a capability that cannot prove it ran, applied to the rulers.

CLASS-LEVEL — THE DAMAGE-TABLE DIGEST NOW DERIVES ITS HASHED FIELDS, AND 5.242.0's REPAIR IS REVERSED AS SUFFICIENT. That version repaired two named terms and this ledger recorded it as done. It was INSTANCE-level: an enumerated list means a field added tomorrow goes unhashed exactly as typing did, so closing the two instances re-armed the class. The hashed set is now derived from the rows' actual keys minus a declared exclusion map of six entries. **The derivation then found a second blind class the list could not have:** an ABSENT field hashes identically to a PRESENT AND EMPTY one, so deleting a body's entire moveset moved nothing — live on **4 rows for mv and 10 for item**. The digest value is **unchanged at** `9d289cf77e24`, so this adds NO third reason to restamp; the two reasons recorded at 5.242.0 stand, still fused by any single restamp, still the owner's call.

INSTANCE-LEVEL AND SAID SO — `tests/test-artifact-keys.js` **HAD FIVE SILENT LIMITS, NOT ONE.** It sliced to 8 keys, capped depth at 3 where the deepest real object is 10, never descended arrays, iterated a hard-coded two-name list, and **exited 0 when the engine data failed to load**. Compounded, `MC.priors` — **230 keys** — was never inspected once, inside the file written to catch hand-typed lookups. **The cost argument that justified the truncation was false when measured: 733ms full against 596ms truncated.** It now walks fully, FAILS rather than truncates at a budget, and prints its own out-of-scope set by name (**5,466** `.json` files in **12** directories). **13 undeclared tables** are newly visible and are REPORTED, not fixed.

ATTRIBUTION WAS MEASURED, NOT INFERRED, AND THAT IS THE PART THIS DIVISION OWNS. `tests/run-all.js` reads **143 passed, 28 failed**. **Exactly one failure was caused by this session**, established by re-running each red test at the prior commit in a control worktree rather than by reasoning about what had been touched. Five assert a stale fact, three were heap ceilings, twenty were already red. **The two strongest medicham2 hypotheses —** `test-pin-arms` **and** `test-game-differential` **— were both REFUTED by the control**, and neither may be cited as evidence about the engine.

THE COUNT WAS MISREPORTED TWICE BEFORE IT SETTLED: 23, THEN 30, AND THE TRUTH IS 28. The first was a TORN READ — a line count taken on the runner's output while it was still being written. CLAUDE.md documents that hazard for artifacts; this is the same hazard applied to a log, and it produced a well-formed, plausible, wrong number exactly as the artifact version does. **And a command ran for twenty minutes having done nothing:** the tell was not the exit code but **2 seconds of CPU**, because output piped through a buffering stage makes a stalled run and a working one identical.

THREE CHECKS COULD NOT RUN AT ALL — exit 134, out of heap, no declared ceiling. With one, `tests/test-engine-release.js` gives **71 passed / 0 failed**, `tests/test-set-realism.js` gives **6 passed / 0 failed**, and `engine/validate_selfplay.js` gives a real verdict plus a finding it was too dead to report. A check that cannot start and a check that passes are the same exit code to anybody not reading.

THE COMMENT CENSUS GIVES A BASE RATE THIS PROJECT HAS NEVER HAD — 495 files, 17,519 comments: **0 confirmed** contradictions of their own code across **1,371** candidates, **2** TODO/FIXME (both false positives), **37** references to things that no longer exist, **4** imperatives with no executor and **6** whose executor is RED. The standout is this section's own thesis: `engine/engine_release.js:653` heads *"PRUNING, AND WHY IT IS SAFE"* and reasons from **nine releases holding 23 MB** against **523 releases holding 2.5 GB** today. `ROADMAP #144` is cited in code and has never been a register row — **nine such ids across 30 citations**.

WHAT IS OWED. Every one of these was withheld on purpose; none is a finding waiting to be collected:

- `node engine/status.js --write` — NOT run. The `<!-- GENERATED -->` block at the top of this file is still one pass behind and still quotes the pre-5.243.0 gate.
- `node engine/feature_fixture.js --stamp` on the fitted policy weights — **WILL'S CALL, NOT TO BE RUN**. Unchanged from 5.242.0 and 5.243.0. The weights were not written and MAG stays paused.
- The owner's decision on moving the whole-game clause from protocol first-divergence to board-material. Unchanged from 5.243.0.
- `npm install` — NOT run. The mutation dependency is declared and uninstalled, so the pilot's numbers above are reproducible only after it is.
- **26 surviving mutants** to triage. Each is a gate clause whose failure nothing observes.
- ROUTED AND NOT FIXED: **89 duplicate ids** in the self-play store; one unattributed finding in `engine/conformance.js` ; the 13 undeclared key tables; a raw-log census artifact still asserting the old subset relation with no generator to correct it.

Full accounts: eleven reports under `docs/_reports/2026-09-04-*.md` .

THE GATE CERTIFIED THE ENGINE ON A POPULATION WHERE THE GAMES DO NOT END, AND THIS DIVISION OWNS THAT RULER. 2026-09-03, CHANGELOG 5.243.0.

THE READING IS REVERSED. `engine/quarantine.js` read **8 of 8 PASS**. It now reads **GATE: CLOSED — 1 of 8 clauses fail**, and that is the correct outcome. **Nothing was tuned to reach it**. The failing clause is state (b) — a named instrument actually failing on the same engine — not a stale release and not a register row without an instrument.

THE DEFECT IS THE ONE THIS DIVISION EXISTS TO CATCH: A CORPUS STAMP NOBODY CHECKED BEFORE ATTRIBUTING AN EFFECT TO IT. The artifact answering the whole-game clause was produced by the `census-coverage-seeking/v1` driver, whose objective is census coverage rather than winning. On that arm **17 of 961 games reached a result (1.8%)**, 944 stopped at the 12-turn cap with both sides standing, and **0 boards parted**. A board cannot part at the end of a game that has no end, so the clause's zero was a fact about the driver and not about the simulator.

THE CONTROL IS EXACT — ONE FREE VARIABLE, EVERYTHING ELSE PINNED.

pinned	value	
engine release	8ad06030e129	
turn cap	12	
frozen team pool	0d103fb9fa87	
census pin	9446a684709d	
games played / budget	961 of a 1200 PAIR budget	
the only difference	the driver	

driver	reached a result	boards parted
census-coverage-seeking/v1	17 of 961 (1.8%) — 944 stopped at the cap	0
empirical-click/v1	474 of 961 (49.3%)	77

THE CORRECTED FIGURES ARE TWO QUANTITIES AND THEY ARE NOT INTERCHANGEABLE. Board-material divergence is **77 of 961 (8.0%)**. Protocol first-divergence is **168**. Both are correct, only one is the bar, and every statement of either must name which it is.

THIS RETRACTS A PUBLISHED POSITION AND THE RETRACTION IS THE POINT. Every clean or near-clean whole-game reading in the ledgers, the white paper, the deck, the technical documentation and the summary was taken on the coverage arm. Those blocks are dated history and stand as written; they are superseded here, not rewritten. What they cannot support any more is the inference that was drawn from them — that the MEDICHAM quarantine was close to lifting. It is not. Nothing this division withholds becomes quotable, and the re-run list is unchanged.

WHY IT WAS INVISIBLE, AND WHAT NOW SEES IT. `engine/status.js` was already printing the row `driver policies the gate quotes - 1 of 2`. Nothing was wrong with the run; the wrong arm was wired to the clause, and **a gate that cannot tell which driver produced its input is the defect — the clean reading was only the symptom.** Three changes, in the order they close the door:

- `data/game-differential.json` **IS** the empirical arm now. The clause was NOT repointed at `data/verification/`, because two files that both claim to be the published quantity can drift.
- `engine/game_differential.js` refuses at second zero to publish a coverage run into that path: `--write` without `--out` now requires `--steering empirical`. That is the exact door this came through — `--out` had diverted the empirical run to `data/verification/` while the published file kept its coverage numbers.
- `wholeGameClause` reads the artifact's own `steering.policy` and withholds its figures from anything that is not `empirical-click/v1`, including an artifact carrying no steering block at all: verdict `MEASURED ON THE WRONG POPULATION`, `clauseExit 2`.

SHOWN RED BEFORE IT WAS TRUSTED. The refusal was demonstrated against the real coverage artifact before that artifact was replaced, and re-verified afterwards by feeding a preserved byte copy of it back in. Selftest 159 passed, 0 failed. A byte copy of the superseded coverage artifact was taken BEFORE the overwrite and kept at `data/verification/game-differential.coverage-2026-09-04T0141.json`; nothing was deleted.

ONE MISMATCH IS DECLARED AND WAS DELIBERATELY NOT CLOSED. The clause prints **167** — protocol first-divergence less the one declared case — and never computes `state.games - state.games_board_never_diverged`. CLAUDE.md records the owner's 2026-08-22 call that the bar is BOARD-MATERIAL, with narration as its own separate gate afterwards, which makes 77 the figure the clause should read. **Changing 167 to 77 in the same pass that turned the gate red is indistinguishable from tuning**, and this division does not compete on its own rulers. It is his call and it is owed below.

WHAT IS OWED. Every command here was withheld on purpose. None of them is a finding waiting to be collected:

- `node engine/status.js --write` — restamps the `<!-- GENERATED -->` blocks in the division ledgers. NOT run, deliberately: the five ledgers must not be stamped with a lifted quarantine that was never real. Until it runs, the generated block at the top of this file is one pass behind and still quotes the pre-change gate.
- The owner's decision on moving the clause from protocol first-divergence to board-material.
- `node engine/feature_fixture.js --stamp data/policy-weights.json` — **WILL'S CALL. NOT TO BE RUN.** Unchanged from 5.242.0. `data/policy-weights.json` was not written and MAG stays paused.
- `node tests/test-web-status.js` (12 red) and `node tests/test-web-quarantine-loaders.js` (2 red) — RED and PRE-EXISTING, stated rather than filed. Their payload was built 2026-08-25 declaring `open:false` against an OPEN gate; they were red before this change and WEB is under a declared pause.

Full account: `docs/_reports/2026-09-03-gate-reads-empirical-arm.md`.

THE RULER THIS DIVISION OWNS WAS BLIND TO TWO OF THE FIELDS IT CLAIMED TO HASH. 2026-09-03, CHANGELOG 5.242.0.

`engine/feature_fixture.js` 's `tableDigest()` is the function that answers *has the damage table under this fitted vector moved*. It hashed `m.ty` for typing. **The table carries no `ty` : the field is present on 0 of 322 rows, and typing lives in `t`**. So that term evaluated to a constant `null` on every row of every run since it was written, and a TYPE change has never been visible to the digest. `m.wt` was not hashed at all, and weight is a live damage input in this format — Low Kick, Grass Knot, Heavy Slam, Heat Crash, Heavy Metal and Light Metal. Both are hashed now. Term order is append-only, so the diff reads as one repaired term plus one new one. `nature` and `sp` stay unhashed on purpose, because they reach the damage formula only through `st`, which is hashed; so do the four provenance fields. Full account: `docs/_reports/2026-09-03-table-digest-blind-fields.md`.

THIS IS LESSONS §5 IN THIS DIVISION'S OWN HOUSE, AND IT IS THE HARDER VARIANT. The usual shape is an instrument accusing a correct engine. This one accused nobody: it printed a stable digest, which is exactly what a correct digest prints when nothing has changed. **A silent ruler and a ruler that cannot see are the same output.**

AN ABSENT FIELD HASHES AS `null` EXACTLY LIKE A FIELD THAT IS PRESENT AND EMPTY. That is why reading the code was never going to settle it — the comment three lines above the term said typing was covered, and nothing in the output disagreed. It was settled by MUTATION, with controls in both directions:

mutate one row's...	before the fix	after
<code>t</code> (typing)	digest does NOT move	moves
<code>wt</code> (weight)	digest does NOT move	moves
<code>mv</code>	moves — the control , the harness can see a change	moves
<code>st</code>	moves — the control	moves
an INVENTED <code>ty</code> field	moves — the term was live; the data never populated it	n/a

The second control is the one worth keeping. Without it, "the digest did not move" is equally consistent with *the term is dead* and with *the term is live and this field is empty everywhere*, and those two have different fixes. `tableDigest` is now exported so the probe calls the canonical function instead of copying it — the same rule `status.js` follows in shelling out to `provenance.js`.

MEASURED. The damage-table digest moved `1bda9df11d73` → `9d289cf77e24` **with no change to the table itself.** `tests/test-feature-semantic.js` : 24 passed, 0 failed. `engine/status.js` parses this gate's printed OUTPUT rather than its code, so the shape was deliberately left alone — it still matches `FEATURE SEMANTICS CHECK FAILED`, still exits 1, and its `how`: string is byte-identical to the one already stamped in the `<!-- GENERATED -->` block at the top of this file. Its verdict is unchanged.

THE STAMP IN `data/policy-weights.json` IS NOW STALE FOR TWO INDEPENDENT REASONS, AND ONE RESTAMP WOULD FUSE THEM. The damage table was regenerated (318 → 322 species) *and* the ruler that measures it changed. A single `--stamp` absorbs both causes and cannot separate them afterwards, so this note and the CHANGELOG entry are the only places the distinction survives. Whoever runs the restamp inherits that: after it, "the table moved" and "the ruler moved" are one undifferentiated green.

NO RESTAMP AND NO REFIT WERE RUN, BY THE OWNER'S EXPLICIT DECISION. THAT IS A DECISION, NOT AN OMISSION, AND IT IS NOT THIS DIVISION'S TO OVERRIDE. MAG stays paused until MEDICHAM is correct; MEDICHAM is upstream of the weights, so a refit under a simulator still

known wrong would only have to be repeated. `data/policy-weights.json` was not written. The RESTAMP-not-refit verdict on the table regeneration stands unchanged from 5.241.0 and waits for him.

WHY IT WAS SAFE TO FIX THE RULER WITHOUT THE RESTAMP, WHEN 5.241.0 SAID THEY SHOULD LAND TOGETHER. The deferral guarded against clearing a gate ahead of the verdict it exists to inform. That risk lives entirely in the RESTAMP, not in the fix: with nothing restamped, the table gate still fires, the evidence for the refit verdict is not written over, and the only thing that moved is a hex value. The two halves are still coupled in the other direction, which is what the paragraph above records.

WHAT IS OWED. Both commands below were withheld deliberately in this pass, and neither is a finding waiting to be collected:

- `node engine/status.js --write` — restamps the `<!-- GENERATED -->` blocks in the division ledgers. Not run: light mode. Until it runs, the generated block at the top of this file is one pass behind and its `feature_fixture --check FAILED` line quotes the pre-fix run.
- `node engine/feature_fixture.js --stamp data/policy-weights.json` — **WILL'S CALL. NOT TO BE RUN.** It is the restamp that fuses the two staleness causes above, and MAG is paused.

THREE OF THE FIVE CARRIED REDS WERE NEVER RED. TWO WERE THE HARNESS AND ONE WAS THE RULER. 2026-08-29, CHANGELOG 5.222.0.

Five checks had been reported red — correctly, and A/B-verified as pre-existing — through eleven test batches. CLAUDE.md bans "known failure" as a status, so each is closed with a verdict rather than carried again. Full account: `docs/_reports/2026-08-29-five-reds.md` .

check	verdict
<code>tests/test-resolution-order.js</code>	FIXED — the harness. <code>tools/lownode.cmd</code> did not honour <code>ABRA-HEAP</code>
<code>tests/test-engine-diff.js</code>	INSTRUMENT — exit 3 was <code>publish_guard.js</code> refusing a shrink
<code>tests/probe_shield_refusal_line.js</code>	INSTRUMENT — a blanket expectation accused a correct engine
<code>tests/probe_random_target_address.js</code>	INSTRUMENT — it assumed one address per die call
<code>tests/probe_instruct_shield.js</code>	FILED — a real engine defect, ROADMAP #532

THE PATTERN IS THIS DIVISION'S OWN, AND IT IS THE ONE LESSONS §5 IS ABOUT. Four of the five were the measurement, not the game. That is now roughly thirty instances, and the two new shapes are worth naming because neither is a probe staging the wrong body:

- **A RUNNER THAT DOES NOT HONOUR WHAT THE CHECK DECLARES.** `tests/run-all.js` reads `ABRA-HEAP` out of a child's own header and even annotates an OOM as *"a memory ceiling, not a verdict about the game"*. `tools/lownode.cmd` — which CLAUDE.md mandates for every heavy run, so it is the command actually typed — did not, and the check died at exit 134 for eleven batches. The declaration was right, the check was right, and nothing in between carried it.
- **AN EXIT CODE THAT NO ENGINE STATE COULD HAVE MADE GREEN.** `tests/test-engine-diff.js` defaults to `--n 150` against a published 6,000, so `engine/publish_guard.js` refused every discovered run and set exit 3 — correctly. A permanently-red check is not a finding about the simulator, and it was read as one.

AND ONE THING THIS DIVISION SHOULD READ TWICE. `probe_shield_refusal_line.js` printed `<-- FAIL, the knob did not reach the driver's module` while printing, two fields to its left, the stamp that proves the knob DID reach it. A wrong diagnosis is worse than none: it aims the next reader at the loader instead of at the expectation. The message is now derived from which clause failed.

THE GENERALISATION, MEASURED RATHER THAN EXTRAPOLATED. 79 files in this repository spawn a node child and **exactly one** — `tests/run-all.js` — **derives the child's heap from the child's own source**. Today's exposure is nonetheless small and is stated as such: only two files declare `ABRA-HEAP`, and the second (`tests/probe_endturn_clock_order.js`, 4096) was measured **passing** at the default. `engine/register_reality.js` honours no declaration and is named by two rows that run that probe — not bitten today, one arm away from it, and deliberately not patched with no red to show. The sharper half needs no run at all: `engine/tag_dex.js` 's single write is at line 9833 of 9854, so an OOM leaves `data/tags.json` with its previous content **and its previous mtime**, and nothing on disk records that the regeneration did not happen.

ROADMAP #533 IS THIS DIVISION'S OWN INSTRUMENT AND IS FILED UNREPAIRED.

Repairing the length check in `probe_random_target_address.js` let it run to completion for the first time since 2026-08-27, and what it then printed was `0 of 0` — ROADMAP #478's engine-side half has moved the draw out of the bucket clauses 2–4 scan, so its negative control is a control that cannot fail against a clause that cannot see anything. Clause 1 is unaffected and reproduces, so nothing published is retracted. Re-aiming the clauses is a re-measurement and gets its own pass.

WHAT IS OWED. `engine/register_reality.js` has not been re-run, so #532's marker is not yet in `data/register-reality.json` and `engine/quarantine.js` 's no open, known engine defect clause still reads PASS. It will read FAIL once that runs, which is the correct state and is recorded here in advance so it is not read as a regression. It was not run in this pass because it executes every marker in the register and rewrites its artifact unconditionally with no `--out`.

THE ROSTER'S MOVE ROWS PASSED ON THE ANNOUNCEMENT, AND THE AUDIT THAT FOUND IT WAS WRONG ON 2 OF THE 11 ROWS IT NAMED. 2026-08-29, CHANGELOG 5.215.0.

`engine/all_mechanics_fire.js` credits a move when the authority's segment carries any `-` event outside `NOT_A_CONSEQUENCE`. `-singleturn` is not in that set; the move ladder stops at the first resolving rung; and on that rung the receiver clicks `inert`, which resolves to **Agility**. Eight guard moves therefore read RESOLVED with nothing ever thrown at them, and no board comparison could notice: every one is a declared `duration: 1` leaf, ended in the residual before the boundary at which the board is sampled.

Every named mechanism verified. The count did not. Focus Punch and Beak Blast were in the audit's list of eleven and do not belong there — both are attacking moves, so their segment carries `-damage`, and their `-singleturn` is emitted by `priorityChargeCallback` at the top of the turn, before the `|move|` line, so it is not even inside the segment the verdict reads. The systemic half stands for all eleven: the roster covered the leaf's CREATION and not its FUNCTION.

`-singleturn` **stays OUT of** `NOT_A_CONSEQUENCE` **and the reason is at the declaration.** It is not bookkeeping; it is the authority saying a state was created, and for a move whose whole function is to create a state, creating one IS resolving. Demoting it makes Protect report *"produced no consequence line at all"* — a false accusation against a correct engine — and singles out one spelling of arrival among `-start`, `-sidestart`, `-fieldstart` and `-singlemove`. The gap was never that the announcement is counted; it is that nothing separately asked whether the effect ran.

The derivation over-matched 60 of 500 on the first try, and the ANCHOR was wrong after that

Rule 1 — *any event a non-arrival handler emits* — matched **60 of 500 moves**, 41 of those markers being `-end / -sideend / -fieldend off onEnd`, which is a leaf EXPIRING rather than working. Narrowed to INTERCEPTION handlers only — those that can run only because an incoming move reached the leaf — the match is **11 of 500**.

Then the anchor. The first version asked for the leaf's DISPLAY NAME on the marker line. Showdown announces the whole guard family generically:

```
|move|p1a: Toxapex|Baneful Bunker|p1a: Toxapex
|-singleturn|p1a: Toxapex|move: Protect      <- not "move: Baneful Bunker"
|-activate|p1a: Toxapex|move: Protect        <- the block, under Protect's name
|-status|p2a: Feraligatr|psn                 <- the Bunker's own punish
```

Three of the first run's seven reds were the ruler. The anchor is now the literal third argument of the handler's own `this.add`, read out of the authority's source. Suspect the instrument before the engine — twice in one pass.

VOLATILE_THEN_WHAT had never fired once, in either arm

	entries	thenWhatFor non-null	reached by a VOLATILE key
abilities (the only arm that called it)	201	18	0
items	148	1	0
moves (never called it)	500	32	7

All seven keys are volatiles written by MOVES. Cute Charm, the one ability that infatuates, is tagged `punishesAttacker` with no `statusInflict` param at all. Seven entries wired to the one arm where none of their keys exist, producing no error and no zero counter. `setupFor` now calls the same producer, and the 37 verbs the move arm cannot execute are counted rather than dropped.

The result, over all 500 rows on release 4b67526d29d8

11 rows declare an effect marker. **7** demonstrate the refusal on BOTH engines (`leafEffectSplit: {}` — no engine split). **4** do not and each states why on the row. **76** further rows write a leaf that prints nothing when it fires and belong to a counter comparison.

No row flipped RED and no engine defect was found. Dified row-by-row against `data/all-mechanics-fire.json` digest `acb84bbf62ad`: **0 of 500** changed `resolved`, `medicham_resolved`, `diverged`, `rung`, `turns` or `board.verdict`.

`endure` cannot be checked at all here — its `-activate` needs a LETHAL hit and every body in this fixture is at x6 HP by design, so nothing anybody clicks is lethal. That is `shapeUnbuildable`, a **fixture limit and not an engine finding**, and it is printed rather than passed.

The counter comparison, costed and not built

`GD.REL.require('engine/medicham2-browser.js').seen` already returns **607 live counter keys** (`MEDSEEN` declares 790 in source), including `sideGuardBlocked`, `sideGuardPierced`, `protectPierced`, `enduredLethalHit`, `preTurnShieldAnnounced`, `helpingHandBP` and the four flinch keys. **Nothing in engine/ or tests/ reads it against a Showdown-side count.** The Showdown side is now free — `leafEffectMarkers` returns the derived `{ev, label}` pairs. What it costs is: a per-game DELTA (MEDSEEN accumulates across 617 games in one run); a PAIRING table, the one part that cannot be derived; a per-pair SEMANTIC fixture, because a

counter and a protocol line need not count the same event; and extracting `leafEffectMarkers` into a data-only module, because `all_mechanics_fire.js` runs on require. **Start with the shield family, not flinch** — half of that pair is already written and unit-tested.

Full account: `docs/_reports/2026-08-29-roster-effect-check.md` .

THREE THINGS THE GATE'S VERDICTS DO NOT COVER, NOW PRINTED. 2026-08-29.

`engine/coverage.js` exists to print what a clean verdict is NOT a claim about. Yesterday produced three facts of exactly that shape and none of them was in it. All three are now derived at run time; each prints `NOT DERIVED` with a reason if its source stops parsing, shown on a deliberate break before being trusted.

1. THE DIFFERENTIAL'S SPREADS ARE SYNTHETIC — `differential` bodies on a `REAL` spread 0 of 17536 . An open team sheet carries no spread, so `game_differential.js` ASSIGNS one from the body's slot index. The row reads the rule off the driver's own constants rather than retyping it: 66 points, a 32 cap, a descending Speed ladder `[32, 22, 11, 0]` by slot, the remainder to the higher attacking stat then spilling to `spd` then `def` , and **o into HP** (Champions' Showdown line adds the investment plus 75 for HP and `medicham2`'s L50 line has no HP term, so HP points would diverge silently on every body). The `NATURE` is real — `--nature real` , 17,440 bodies from the sheet's own and 96 fallen back to `Serious` — and both engines are handed the same invented spread, so **the run is internally consistent and its damage is not metagame damage**. The artifact already declared `spreads_absent` ; what was missing was what got put there instead, and that it is a construction.

2. THE COMPARATOR ONLY READS AT A TURN BOUNDARY, SO THE CEILING IS 56, NOT 80 — `board` leaves compared 34 of 56 . The denominator was the population and a reader takes a denominator for a target. Of the 80 leaves a legal mechanic can write, 4 are declared uncomparable, **18 carry a declared duration of 1** and are ended in the residual, and **2 are removed inside their own action** (`volatile:fling` , `volatile:sparklingaria`). None of those 24 can be standing when the board is read, so 80 is not reachable and the widening work is the **22** that can. The most-written leaves in the hole — `flinch` (20 writers), `protect` — are among the permanently uncomparable, so the remaining work is smaller AND worth less than 34 of 80 suggested.

AND THE SELF-REMOVAL RULE HAD TWO PRODUCERS, DISAGREEING. It lived only in `tests/probe_leaf_name_map.js` , so `derive()` — the function `status.js` and `coverage.js` read — did not know about it and would have published a ceiling of **58** where that probe printed **56**. The rule and the boundary call-site count moved into `tests/probe_uncompared_leaves.js` ; the name-map probe now calls them and its printed output is unchanged. The ceiling holds only while `BS.snapshot` has one caller, so that is counted every run (1 call site, `stateCheck` ; 0 elsewhere) and printed beside the number it justifies.

3. WHICH DRIVER A WHOLE-GAME FIGURE WAS TAKEN UNDER — `driver` policies the gate quotes 1 of 2 . On one set of pins (release `e129bca605e3` , cap 12, pool `0d103fb9fa87` , 961 games each) the coverage-seeker reaches a result in **17 games (1.8%)** with **o** whose board diverged, and the empirical arm reaches **459 (47.8%)** with **135**. The gate reads only the first. The row prints both, states that the arms share their pins so the difference is the driver and nothing else, and quotes `engine/arms_comparable.js` 's actual refusal rather than paraphrasing it — the pair is refused on `policy` , so they are two instruments and not a before/after. It also carries the empirical arm's own limit, derived: **42 of 961 games (4.4%) truncate on a forced switch medicham2's placement cannot express to Showdown, so 47.8% is a lower bound**. Six older artifacts in the same family sit on other pins and are counted, not printed — a different cap, release or pool is a different question.

Full account: `docs/_reports/2026-08-29-coverage-scope-lines.md` .

THE WHOLE-GAME DIFFERENTIAL'S CLEAN SHEET IS A CLAIM ABOUT GAMES THAT DO NOT END. 2026-08-29.

Built: a second driver arm, `empirical-click/v1` , **beside** `census-coverage-seeking/v1` . Selected by id (`--steering coverage|empirical`), default unchanged. It draws the action from `data/move-priors.json` — $P(\text{move} \mid \text{species})$ over real recorded ladder clicks, with the turn-1 `lead` table on turn one — and takes a voluntary switch at the **9.98%** conditional rate measured off the raw logs of both human stores (`data/rollout-switch-census.json` , 58,639 finished games). Everything else is unchanged: same swarm teams, same Mode A pinned dice on both sides, same comparators, same census credit rule, same target rule. Only the action selection moved, which is what keeps every divergence a RULE.

THE MEASUREMENT. Two legs, one session, one binary, release `e129bca605e3` , census pin `9446a684709d` , `--team-store data/team-pool-frozen` , `--games 1200` -> 961 games, arm `middle` , cap 12, both diverted to `data/verification/` . The published artifact was not written.

	coverage-seeker (control)	empirical-click
both engines ended the battle	17 (1.8%)	459 (47.8%)
cut off by the turn cap (12)	944 (98.2%)	454 (47.2%)
protocol parted	6	248
board-material causes / games	0 / 0	114 / 128
DIFFERENT-END-STATE	0	99
band 1 DIFFERENT-WINNER	not reachable	0 of 459 resolved games
elapsed	102.8 s	252.9 s

THE CAP WAS NEVER THE BINDING TERM. THE DRIVER WAS — 1.8% -> 47.8% at an identical cap, release, census pin and byte-identical team pool (`0d103fb9fa87`). The control leg reproduces the published headline exactly, which is the proof the default arm did not move.

AND THE CLEAN SHEET DOES NOT SURVIVE A DRIVER THAT PLAYS THE GAME. `0` board-material becomes `128` on the same 961 games, and the two cause sets are **disjoint** — not one of the control's 6 appears among the empirical arm's 225. This is a RE-BASELINE, not a regression: `arms_comparable.js` refuses the pair on `policy` , correctly, and the 128 live in a population — turns 5-12 of games that resolve, forced switches, post-KO replacement, endings — that the coverage arm has never once visited. **The quarantine clause reads board-material zero on games that do not end.** That is now measured rather than assumed.

BAND 1 IS REACHABLE FOR THE FIRST TIME AND READS ZERO. Of the 99 games whose final boards differ, none disagree about who won. It is a result on one sample, not a proof about the winner rule: ROADMAP #362's simultaneous-double-wipe defect is rare and 459 resolved games may simply not contain one.

RECEIPTS, PRINTED INCLUDING THE ZEROS. 77,611 decisions; **0 on an unprofiled species** (derived, not lucky — 336 of the format's 347 legal species hit a prior row on `species.id` and the other 11 resolve through the base forme); 11 base-forme resolutions; 1,820 (2.3%) draws where the table held none of the body's legal moves; realised switch rate **9.675%** against 9.98% measured; **0 repeated driver addresses.** `ban_narrowed` reads 0 and that is correct — `diff_swarm.js` already selects `omit-*` teams by a predicate that excludes the feature, so the driver's ban has nothing left to remove.

THE SAMPLER IS DUPLICATED AND THE DUPLICATION IS PINNED BY A TEST. `rollout_leaf.pickByPrior` could not be called: that file requires a LIVE medicham2 and board.js at module load, and the differential reads a frozen release. `tests/test-empirical-driver.js` runs 720 draws of the same

rows through both implementations and asserts they agree — 720/720 — so the day they diverge is a red test rather than a silent second opinion. `engine/rollout_leaf.js` gained an EXPORT only; no require edge, no SOURCES change, no release stranded.

LIMITS, STAMPED IN THE ARTIFACT UNDER `not_modelled`. The priors carry no target model (the existing first-live-foe rule is unchanged against a real 23.4% double-target rate) and no model of WHICH body a switch sends (uniform over the legal bench). **42 of 961 games (4.4%) stop on an unmirrorable forced switch** after the boards had already parted, so 47.8% is a lower bound.

Full account, cause tables and the commands owed: `docs/_reports/2026-08-29-empirical-driver-arm.md`.

A FIGURE WHOSE GENERATOR PRINTS AND DOES NOT SAVE HAS NO SOURCE, AND THE ONLY HONEST FIX IS TO WITHDRAW IT. 2026-08-29.

`tests/test-docs-current.js` read **22 passed, 1 failed** on clause 3b(c) — the census of figures with no artifact behind them anywhere — naming ONE surviving figure in two living documents: the PORY two-feature material baseline that `engine/pory_baseline.py` published on 2026-07-25. It is now **23 passed, 0 failed**, census **37 → 35**, with `data/docs-currency-baseline.json` untouched and both ratchets where they were (3b(a) at 8, 3b(b) at 65). **The figure was withdrawn, not sourced.**

THE GATE WAS RIGHT, AND THE THREE REASONS ARE INDEPENDENT.

- **The generator writes nothing.** `engine/pory_baseline.py` holds no `json.dump`, no `open(..., 'w')` and no `write()`; it prints a five-arm table to stdout and exits. There was never a file to check the documents against. An independent walk of **5,288 JSON files under `data/`** finds nothing that rounds to the value at any scale — and its companion number in the same parenthetical passes the census only by **coincidental collision** with an unrelated usage share in `data/meta-usage.json` and `data/bring-bias.json`. That is the ENGINE agent's hypothesis confirmed in the half of the pair that is still standing: a figure can sit "traceable" for a month on an accident.
- **It was scored on the wrong population.** `git show e39329de:engine/pory_baseline.py` — the version that produced the table — has **no clean-data filter**. One landed on 2026-07-30, five days later, and the script's own comment says the unfiltered archive is mostly bots, forfeits and stubs, so *"the comparison that is supposed to keep this project honest is itself measuring the wrong population."* The clean corpus moves every arm by more than the gap the pair was reporting.
- **The CLAIM is superseded, not just the sourcing.** Both documents said PORY **LOSES** to the two-feature baseline. `data/pory-eval.json` answers the same question on the clean corpus with PORY's own shipped estimator, **paired and clustered by game** instead of two unpaired point estimates: both arms at **0.623623**, difference **+0.000001** (positive = PORY worse), 95% CI **[-0.000026, +0.000029]** over **1,177** held-out games of 5,883. That is a **TIE**, which `docs/MODELS.md` and this file have recorded since 2026-08-04. The two corrected documents were the last places carrying the stronger claim.

RE-SOURCING WAS AVAILABLE AND WOULD HAVE BEEN THE WRONG ANSWER.

`pory_baseline.py` could be given a `json.dump` in ten minutes and the raw archive is on disk. It would not have rescued this figure: the corpus and the filter both moved, so a re-run produces a DIFFERENT number, and writing the published one into an artifact to satisfy a gate is typing rather than measuring — the `1.256 / 1.544` P1 class one section down, arriving through a different door. **Sourcing a figure means finding what measured it, not making a file that contains it.**

AND THE CHANGELOG ROUTE WAS REFUSED ON PURPOSE. `changelogHas()` makes any figure written into `CHANGELOG.md` traceable everywhere, so one line would have turned the clause green with nothing measured. `engine/docs_scan.js` warns about that loop in its own header; the ENGINE agent declined it and so does this pass. The withdrawn numerals appear in neither the changelog entry nor either corrected document — only in `docs/REVIEW-2026-07-25.md`, which measured them and now carries a SUPERSEDED note saying what moved underneath it.

PORY IS NOT QUARANTINED AND THAT WAS CHECKED, NOT ASSUMED. `engine/pory.py` imports `json`, `os`, `math`, `random`, `numpy` and reads `data/games.ladder.raw-logs.jsonl`; it never loads the simulator and reads no rollout. **PORY — the logistic value net — is a different model from PORYGON2, the nearest-neighbour value function**, and PORYGON2 is the name on the quarantine list. Had it been the same model the fix would have been to WITHHOLD the figure rather than correct it, which is a different action; the two are easy to confuse by name alone.

REPUBLISHING AN ARTIFACT BREAKS EVERY DATED DOCUMENT THAT QUOTED IT, AND THE FIX IS A LABEL, NOT A NEW NUMBER. 2026-08-28.

`tests/test-docs-current.js` went red on two clauses after WIRE 158 landed: six version headers stale at the previous CHANGELOG top, and the cited-artifact clause up from its baseline of 65 to 71. Both are today's work and both are now green at **23 passed, 0 failed**, with the mismatch count back DOWN to 65 and **no entry added to** `data/docs-currency-baseline.json`.

THE SIX NEW MISMATCHES WERE ALL IN DATED-HISTORY BLOCKS, AND NOT ONE OF THEM WAS A WRONG CLAIM. The census went 780 → 782 and the roster's items stage 139 → 140, so `data/mechanics-census.json` and `data/all-mechanics-fire.json` stopped containing values that four documents correctly recorded as what those files SAID on an earlier release. Rewriting the figures would have falsified the record; `engine/docs_scan.js` already provides the right answer, which is that a block *about* a superseded reading is skipped (`QUALIFIED`). So each of those blocks now says so in its own words, and the current numbers live in a new version block above them.

THIS IS A TREADMILL THE FILE ALREADY DIAGNOSED FOR ONE CLAUSE AND NOT THE OTHER. `docs_scan.js`'s `changelogHas()` exists because *"the moment an artifact is republished, every figure it used to hold becomes 'in no artifact'"* — a gate that fires no matter what anyone does. That escape covers the census clause 3b(c) and does NOT cover the cited-artifact clause 3b(b), which has only the `QUALIFIED` prose test. Every future engine change that moves a headline count will therefore re-open 3b(b) against the blocks that recorded the old one. Labelling is the correct action each time; it is worth knowing it is structural rather than carelessness.

A DEDUPLICATION IN THE CLAUSE HIDES HALF THE WORK. Its key is `doc|figure|cites`, so two rows in `docs/SUMMARY.md` quoting 780 against the same artifact are ONE entry. Fixing the reported line left the clause red until the second, unreported line was found by grep. Read the count, then grep the document — the printed list is a set, not a census.

THE GATE NOW PRINTS ITS OWN COVERAGE, AND TWO OF THE FOUR EXAMPLES THAT PROMPTED IT WERE OVERSTATED.

2026-08-28.

Every unpleasant surprise of 2026-08-28 had one shape: **a verdict printed without its coverage**, where the number was correct and quietly narrower than it read, and **the scope was already recorded somewhere nobody looks** — in a field beside the pass count, in a probe nobody runs, in a clause tail, in a per-row `note`. That is a REPORTING defect, not an engine defect, and it is why finishing MEDICHAM has felt endless: the verdict is read, the scope is not, and the next person to look somewhere new produces another surprise.

`engine/coverage.js` **is the fix, and** `engine/status.js` **is its only shipping caller**. Every gate clause now prints, under its verdict, the AGE of the artifact it was drawn from and what that artifact's own denominator EXCLUDES; and a `COVERAGE` block states the finish line as a set of counts so *"is MEDICHAM done"* is one command rather than a judgement. Read the numbers from `node engine/status.js`; none is repeated here, because a figure typed into prose is the fourteen stale handoffs in a new costume.

THE FINDING THIS DIVISION SHOULD KEEP IS THAT TWO OF THE FOUR MOTIVATING EXAMPLES DID NOT SURVIVE BEING CHECKED. They were given to me as fact and I checked each against the artifact before wiring it, which is the only reason the coverage lines are not themselves wrong:

- **"the roster compares multi-hit moves only ever at 2 hits" — REFUTED.** `tests/roster.js` 's own `move/multihit` `why` block already records (2026-08-27) that this sentence was FALSE for nine days. A `[2,5]` move never reaches the pinned range form: `data/mods/champions/scripts.ts:441` draws it with `battle.sample` over a twenty-entry table, `PRNG#sample` is the one-argument `random(n)`, and the arms answer that `top ? n-1 : 0` — so the two arms reach the two ENDS and the INTERIOR is what nothing reaches. **The per-row `note` in `data/roster.moves.json` still prints the refuted sentence**, typed from `e.multihit[0]`; that is an open reporting defect owned by whoever holds `tests/roster.js`, and it is exactly how the wrong claim reached me.
- **"the damage differential's multi-hit skip is a hidden field" — OVERSTATED.** The raw skip COUNT was already on the differential line in `status.js`. What was genuinely missing is that the skip is a whole FAMILY of moves the volley loop has never run once — derived through the same door the instrument uses to build its skip set, the `multiHit` tag, so the two cannot part.

THE COVERAGE NUMBER MUST NOT BE THE NEXT INCOMPLETE VERDICT, SO THE MECHANISM IS JUDGED ON WHETHER IT FINDS A FIELD NOBODY TOLD IT ABOUT. Three mechanisms, in decreasing generality: a NAME VOCABULARY that reports any non-zero `skipped_*` / `*_dropped` / `did_not_fire` / `unreachable` -shaped field in any artifact, with no list of artifacts or fields; an ARITHMETIC RESIDUAL that catches an exclusion with **no name at all**, by comparing a population-shaped key to an accounted-shaped key in the same object; and DECLARED RANGES, which reports any tag param carrying `range: [lo,hi]` with `hi > lo`. Each was shown behaving on a deliberate break before being trusted, **including the negative control that proves the stated limitation is real**: an exclusion named `volleysNotRun` is MISSED, because the vocabulary cannot match a word nobody used. There is no mechanical defence against that; the partial answer is that `node engine/coverage.js --audit` prints the unmatched field names so a reader auditing a new instrument sees what exists rather than only what a regex recognised.

One producer per fact, enforced by import rather than by care. The leaf split comes from `tests/probe_uncompared_leaves.js` — which now exports `derive()` and whose CLI is a renderer over it, proved byte-identical before and after — and the tag split moved out of `status.js` into `coverage.js`, because

two implementations of one fact is the breach that had the closed-row detector disagreeing with itself on 24 of 292 rows in both directions.

And the gate now prints how old each artifact is. A stage run *without* `--write` prints a complete report and exits 0 while its artifact never moves; `data/roster.items.json` published a `DEFERRED-BY-OWNER` row that had been fixed hours earlier because of exactly that. The age is read from the artifact's own `generated` stamp and never from its `mtime` — "*newer than its source*" is no evidence at all — and the `mtime` is read for one purpose only: to say out loud when a file changed in the last minute, because a torn read is a plausible, well-formed, completely fictitious answer rather than an error.

Full account: `docs/_reports/2026-08-28-gate-coverage.md` .

MEDICHAM SPEED: THE LEAF COSTS 1.14–1.27 s AT n=200, AND I PUBLISHED TWO WRONG CLAIMS BEFORE THAT ONE. 2026-08-28.

Will asked how fast we can play games on MEDICHAM, for rollout budgeting. **Timing only — this makes no accuracy claim and does not touch the quarantine.** `data/rollout-cost.json` (R2 leaf cost) was not read; every figure was measured fresh under release `5f3f7141227c` with the pool pinned to `data/team-pool-frozen` .

The answer is QUARANTINED — the figures are withheld, not annotated. `rolloutWinProb` was timed at the shipped `miltank.js` `DEFAULTS` — `n=200`, `explore=1.0`, `uniform`, `maxTurns 14` — and every cost in that answer came out of `data/medicham-speed.json` , whose generator `engine/bench_speed.js` is in the play layer and reaches `engine/medicham2-browser.js` through `require` . This is R2, leaf cost, which `CLAUDE.md` names in the quarantine list by name. MEDICHAM is not correct — `node engine/status.js` names the failing clauses. No ms per leaf call, no playouts per second, no leaf calls per decision, no whole-game rate and no ms per turn is carried in its place, and the K the budget funds is therefore unstated here rather than restated. It becomes quotable again when the gate opens AND this is re-run: `node engine/bench_speed.js` (consolidated by `engine/bench_speed_consolidate.js`). Method and full account: `docs/_reports/2026-08-28-medicham-speed.md` .

THE FINDING THIS DIVISION SHOULD KEEP IS NOT THE NUMBER. It is that I got two things wrong first, in the two shapes this repository keeps paying for.

1. The artifact and the report disagreed, and only a reader comparing them could have noticed.

`engine/bench_speed.js` writes one file per run, so the published `data/medicham-speed.json` held ONE run while the report quoted the best of several. It reported a 14% slower engine and the *opposite* scaling conclusion. Caught in review, not by a gate. The artifact is now `DERIVED` from every run file and carries `RANGES` with each endpoint attributed to its file and conditions, so the two cannot drift again. **One run cannot carry a range, and the range is the honest answer here.**

2. "MEDICHAM does not scale across processes" was wrong, and the evidence for it was arithmetically impossible. `data/_bench-scaling.json` reads `1w 167.1 → 4w 143.2 → 6w 144.3` aggregate playouts/sec. Four processes cannot do less total work than one. Those were collapsed legs and I read them as a finding. Re-measured warmed AND with worker counts **interleaved** (`1,2,4,1,2,4`), it scales sub-linearly: **1.44x at 2 workers, 2.35x at 4** — while the two 1-worker legs of that same run differ by **7.3x**. The consolidator now detects an impossible row mechanically. Confidence stays `LOW`; plan against one worker.

The instrument has a noise problem and it is now characterised rather than assumed.

- **Within one arm, steady state, 2.5%** — reps 3–9 of a 10-rep probe. **I quoted that as though it described the spread between RUNS. It does not;** between-run spread is 8–28% for arms that

reached steady state and up to 342% for arms that did not.

- **V8 tier-up is 5.6x over ~4,000 playouts.** The harness warm-up plays ~36. A short run publishes a partly-cold number: cap 20 as a first arm read 42.7/sec against 188.7 elsewhere.
- **Sporadic collapses to a third or a seventh**, on an idle box, no reproducible trigger. Every figure is fastest-of-N because of this.
- **Priority is not it.** Like-for-like BelowNormal vs NORMAL: 170.3 vs 170.9 (cap 14), 158.5 vs 151.5 (cap 60) — **under 5%, sign flipping.** An earlier "0.4%" quoted only the cap-14 half and was too tight.
`tools\lownode.cmd` costs nothing detectable at a ~5% bound.
- **Part of the spread is unattributed and says so.** Two runs with cap 6 as the first arm read 94.3 and 240.1. Two candidate heuristics were tried against the data and both were contradicted.

Rule for anyone timing this engine: one cap per process, as the first timed arm, at least 6 reps, and every condition repeated at least twice in interleaved order. The sweeps published first violated all three.

One thing worth carrying to ENGINE, not acted on here. A release does not freeze `data/rollout-switch-census.json` — `rollout_leaf.census()` resolves it under `data/releases/<id>/data/` and does not find it, so a release-pinned leaf silently falls back to **switchRate 0, horizon 0 and cannot switch at all**. That is `engine_release.js` `requireClosure` 's declared `fs.readFileSync` gap. The bench passes the census explicitly; nothing else does.

THE DOCUMENTS ARE UNFROZEN AND FIVE GATE CLAUSES WENT BLANK BETWEEN 09:58Z AND 10:06Z. THE CAUSE IS A LINE ENDING. 2026-08-28, CHANGELOG 5.205.0.

Will paused MEDICHAM development and asked for the docs to be unfrozen. The living-docs deferral of 2026-08-10 is discharged: 274 sprint rows folded into the white paper, the deck, the technical docs, SUMMARY, MODELS and DAMAGE-STAGES; all six restamped to 5.205.0; the six `MEDICHAM SPRINT PAUSE` version pins retired from `data/docs-currency-baseline.json`; `docs/MEDICHAM-SPRINT-NOTES.md` deleted, which is what re-arms the full rule in `engine/docs_scan.js` and `.github/hooks/pre-commit`.

THIS DIVISION'S FINDING IS NOT THE DOCUMENTS. It is that the gate read 7 of 8 clauses PASS a few hours ago and reads 5 of 8 now, with NO ENGINE BYTE CHANGED. The three deliberate-roster stages, the whole-game differential and the staged-mechanics comparison all record engine release `5f3f7141227c` and were written at 09:56–09:58Z. The tree now hashes to a different release, so `engine/status.js` correctly calls every count in them *"an answer about other bytes"* and withholds it.

MEASURED, NOT ASSUMED — AND THE DISCRIMINATOR IS THE STRONGEST ONE AVAILABLE. Exactly ONE of the twenty-six frozen sources moved: `data/tags.json`, mtime 10:06Z. The release keeps a COPY rather than a checksum, so the two files can be compared directly:

<code>data/releases/5f3f7141227c/data/tags.json</code>	799,498 bytes
<code>data/tags.json</code>	842,110 bytes
raw equal?	false
CRLF-normalised equal?	true
JSON deep-equal?	true

A checkout under `core.autocrlf = true` rewrote a generator's LF output as CRLF between the runs and now. `tag_dex.js` writes LF; git hands back CRLF; the sha256 moves and the content does not.

THIS IS THE SECOND OCCURRENCE AND IT IS ALREADY IN THE RECORD. docs/ENGINE.md documents the identical event on 2026-08-26 — *"the reason is a line ending, which is worth writing down because it will happen again to whoever next regenerates data/tags.json"* — where the remedy was to cut the release over the bytes a checkout actually produces. It happened again. **The hazard is standing, not closed:** git diff currently warns *"LF will be replaced by CRLF the next time Git touches it"* on every one of the living documents as well, so any tracked LF file that a generator writes is one checkout away from stranding whatever measured against it.

WHAT THIS DIVISION DID NOT DO, DELIBERATELY. It did not restore the file to LF to make the gate read 7 of 8. Editing an input so a ruler prints the wanted number is the failure this division exists to prevent, and it would have produced a release id that no checkout reproduces. **The five clauses are WITHHELD and a re-run is OWED.** No rate, no diverged count and no roster column from those five artifacts was written into any document this pass — checked afterwards against data/quarantine-stamp.json's citation_sites ratchet, which is unmoved at three entries.

WHAT WAS PUBLISHED, WITH ITS BOUNDS STATED IN ALL SIX DOCUMENTS.

figure	artifact	bound written beside it
780 probed / 780 live / 0 missing	data/mechanics-census.json	a LAB — one staged scenario per mechanic, usage-blind; answers <i>is this correct</i> , never <i>does this matter</i>
6000 compared / 6000 agreed / 0 disagreed, and 0 at each of the sixteen band indices	data/engine-diff.json	its own scope is damage only — no items, no abilities, no turn order, no status duration, no switching — and skipped_multihit 134 with skipped_ability_multihit 17 means it has never applied a multi-hit move
1696/1696 exact, 0 at the wrong stage	tests/test-damage-stages.js , re-run green this pass	the stage chain only; the crit die and the event die are the battle loop's
33 compared / 4 declared / 43 in neither list, 25 live at the boundary	tests/probe_uncompared_leaves.js over 500 moves, 201 abilities, 148 items	"the boards match" is a claim about 33 leaves of 80 — recorded as the largest caveat in the document set

AND THE VOID RULING IS STATED IN ALL SIX RATHER THAN CAPTIONED. Every figure that passed through the event die before 2026-08-27 is VOID, not stale — two engines agreed over a narrower slice of outcome space than the comparison claimed, so an agreement is not evidence of one. The old paragraphs stay in place as dated history, under one governing sentence in each document that none of them may be cited as current. The alternative — a per-figure asterisk — is the thing CLAUDE.md already forbids.

OWED AND NOT RUN THIS PASS: the three roster stages, engine/game_differential.js and engine/all_mechanics_fire.js (a re-run is the only thing that lifts the stranding); the MAG refit, which engine/status.js still reports OWED with feature_fixture --check failing on both the fixture identity and the damage table; and the standing item — data/winrate-backtest.json , leaf calibration, which stays QUARANTINED and unmeasured behind the gate.

A CHECK THAT READS GREEN WITHOUT RUNNING IS WORSE THAN ONE THAT READS RED. 2026-08-28, CHANGELOG 5.201.0.

ROADMAP #521 , #522 , #523 , #524 , #525 . **This division builds the rulers, and tonight the ruler was the thing that was wrong — for the sixth time in a day.**

THE HEADLINE IS A NEGATIVE RESULT AND IT IS STILL THE ANSWER. Three probes were counted as ACCOUNTED FOR by tests/run-all.js , in prose that asserted engine/register_reality.js ran their markers. It could not: SAFE required a marker to begin node <script> and permitted flags only, and all

three begin `node -r ./tests/_live_release.js` . Run at last, **all three are GREEN**, each with a knob-cleared control that moves the arm it names. There was no hidden defect. The accounting was wrong and the engine was not — which is worth exactly as much as a finding, because it is the difference between a ruler that is trusted and one that is trusted correctly.

THE DECISION THAT MATTERED WAS WHICH SIDE OF THE FENCE TO EDIT, AND IT WAS MEASURED RATHER THAN ARGUED. `tests/_live_release.js` is load-bearing:

`engine/game_differential.js:196` CUTS A REAL RELEASE at require time when no `--release` is given, repointing `data/engine-release.json` under whatever else is measuring. All three probes detect the preload themselves and refuse with exit 2 without it. Rewriting the markers to satisfy the old regex would have bought three refusals and one moved release pointer per pass. So the REGEX moved.

AND IT MOVED NARROWLY, WHICH IS THE PART TO KEEP. `SAFE` still refuses BARE VALUES.

That is not laziness: widening it would silently admit `tests/roster.js --stage moves` , `engine/all_mechanics_fire.js --kind abilities` and two `game_differential.js --team-store` markers — multi-minute game-playing runs, three of which REWRITE artifacts other readers hold. A gate that runs those on every pass is a gate that rewrites the corpus it is auditing. `probe_trace_list.js` learned the `--cells=60` spelling instead; the parser moved, not the gate.

THE PREFIX WAS NOT DECORATIVE ON ALL FOURTEEN, AND THE ONE EXCEPTION IS THE POINT. The brief said thirteen were one deletion from running and warned that one probe running with the variable empty is evidence for that probe. Checked one at a time:

`tests/probe_endturn_clock_order.js` never required `engine/showdown_path.js` , so it asked the raw variable and exited 2 with `NOT RUN` under exactly the environment thirteen siblings ran to completion in. Worse, a literal paste of its marker sets the variable to the three characters `...` , which `showdown_path.js` HONOURS by design — so the documentation was not merely useless, it was a working way to break the probe.

A REFUSAL IS STILL PUBLISHED AS A PASS, AND THE WORSE HALF IS AT EXIT ZERO — NOT EXIT 2. The source report ranked the `ABRA-EXIT` gap by counting exit-2 paths. Exit 2 is UNDECLARED, which `register_reality.js` reads as NOT A VERDICT and counts as BAD: noisy, but honest. The unmeasured half is the staging refusal spelled `process.exit(0)` — 12 paths in 4 of 74 `tests/probe_*.js` — which the register reads as VERDICT-GREEN and a CLOSED row reads as CONFIRMED. **A COULD-NOT-STAGE is a claim about the FIXTURE, never about the mechanic,** and at exit 0 that claim is published as a clean bill of health. Five converted and shown red by forcing the guard; the rest are filed, not swept.

#496 IS A PREMATURE CLOSE AND IT IS FILED AS OBSERVED, NOT AS A REGRESSION.

`probe_trace_list.js` exits 1 on the live tree, reproduced **3 of 3** with byte-identical counters across two different pool digests. The row closed on a PINNED sample (`44bd49403231`) and this run drew a LIVE pool while `engine/medicham2-browser.js` was being edited — so the corpus stamps do not match and the claim stops at OBSERVED. Naming that boundary is this division's job; attributing it is not.

ONE THING WENT WRONG AND IT IS RECORDED RATHER THAN TIDIED AWAY.

`data/engine-release.json` was found holding an automatic `game differential mode A` cut, I judged it my own leakage and reverted it to HEAD, then put it back ~90 seconds later. It has since been superseded by ENGINE's deliberate cut `0415c53255a9` ("the absorb answers arrival one..."), so the net effect is zero. **The rule the error breaks is the one this division enforces on everyone else: a measuring agent does not write a file it does not own, and on a moving tree it cannot attribute one either.** The correct action was to report the drift and leave it.

OWED, NOT RUN. `node engine/status.js --write` — this pass was denied the game slot, so every `<!-- GENERATED -->` block in the ledgers is stamped one pass behind.

A SUBTRACTION THE GATE MAKES MUST NAME WHAT IT SUBTRACTED. 2026-08-28, CHANGELOG 5.196.0.

ROADMAP #520 . **This is a REPORTING change and it moves no clause.** Measured rather than argued: `HEAD:engine/quarantine.js` and the working copy were compiled in ONE process against the SAME on-disk artifacts and returned identical verdicts and identical counts on all eight clauses.

THE RULE THIS ENFORCES IS ALREADY THIS DIVISION'S, AND THE GUARD WAS BREAKING IT. A filter may only ever subtract from a number a reader can still see. The DECLARED register prints every row that MAY subtract, with its ruling, whether or not it fired — and the owner's closet, sitting one line away in the same clause, printed the integer 4 shelved by the owner and nothing else. No names, no dates, no rulings. A fifth entry could have appeared and nothing would have said so.

SHELVED BY THE OWNER now names each row with its carrier, its cause, its board verdict and the dated ruling; publishes `owner_shelved` / `owner_shelved_summary` / `owner_shelved_rows` ; prints at zero as well as at four; and **compares the derived rows against the artifact's own summary** rather than assuming they agree — a derived set is not a fact until something compares it to its source.

NAMING IT FOUND SOMETHING ON THE FIRST RENDER. Of the four shelved rows on release `aea838766e7f` , three are `ANNOUNCEMENT-ONLY` and `item:metronome` **is** `board_verdict: STATE` — 859/960 against 868/960, a board that genuinely parted. Will's Metronome ruling is explicitly cost-based and stands; the point is that **the shelf is not uniformly a no-board-effect shelf**, and until this run nothing said which of its rows were which.

THE COMPANION FINDING IS A REFUSAL. Bitter Malice and Night Daze were proposed as the CLOSETED kind's first two entries and were refused: both already carry `deferred = ILLUSION_SHELF` , derived from `GD.CLOSET_SPECIES` off the ABILITY rather than a name list, and `classifyMechanics` skips a `deferred` row **before** it asks `declaredMatch` — so a declaration could not have been reached even if it matched. `game_differential.js` additionally drops all 43 Illusion-carrying teams from the pool, so the whole-game clause holds zero zoroark causes. Writing the rows would have registered a permanent exemption that fires on nothing.

TWO METHOD NOTES THIS DIVISION SHOULD KEEP.

- 1. A SELFTEST THAT READS A LIVE ARTIFACT IS NOT A SELFTEST.** The first version of the printer assertion drove `mechanicsClause()` off disk and went RED mid-session for a reason that had nothing to do with the code: another division cut a release, `data/engine-release.json` moved `aea838766e7f` -> `b035aa665740` , and the clause correctly took its MEASURED-AGAINST-A-DIFFERENT-ENGINE early return. A green/red signal another agent can flip is noise, and noise is precisely how a red test becomes "one of the two known failures". `mechanicsClause` gained the `inject` door the file already uses twice.
 - 2. THE HAND-WALKED DERIVATION WAS WRONG AND THE VALIDATOR WAS RIGHT.** Walking `prevo/ baseSpecies` chains reported Zoroark-Hisui as a Night Daze learner. `TeamValidator` over the 347 legal species of the format refuses it — **Bitter Malice -> Zoroark-Hisui, Night Daze -> Zoroark, one legal learner each.** When a legality question has an authority, ask the authority; a chain walk is a reimplement of one, and this repo names its own casualty for that.
-

THE WHOLE-GAME AND BOARD-MATERIAL BASELINES WERE RESET ON PURPOSE ON 2026-08-27. DO NOT SUBTRACT ACROSS IT.

ENGINE fixed the `middle` arm's die (ROADMAP #489, CHANGELOG 5.176.0): bare FNV-1a has no diffusion after its last round and the last field of every address is `nth` , so an indexed draw only TRANSLATED the previous value — max circular shift **0.0351571** against ~0.5, and **1.75** distinct damage buckets from a ten-hit address instead of 7.56.

	before (<code>01be9daf14ee</code> , pins <code>2efbc9ed1946</code>)	after (<code>f9d6be635d34</code> , pins <code>f646b0163bc0</code>)
whole-game, counted	3 of 961 (8 raw, less 5 declared)	14 of 961 (19 raw, less 5 declared)
board-material	1 of 961	12 of 961

THE RISE IS THE INSTRUMENT SEEING WHAT IT WAS BLIND TO, NOT A REGRESSION, and it was predicted before the run. Attribution was checked rather than assumed: the two frozen releases differ in exactly one SOURCES file and its only executable difference is the three finaliser lines.

MEASURE'S PART IS THE REFUSAL, AND IT IS ALREADY WIRED. `DICE_MODEL` now rides in `PIN_DIGEST` (`engine/game_differential.js`), which `engine/arms_comparable.js` compares through `mode` . The gate prints, unprompted:

DIRECTION OF TRAVEL WITHHELD — the baseline was stamped under

A/middle/pins:2efbc9ed1946/... and this run is A/middle/pins:f646b0163bc0/... . One pin is one corner: those are two instruments, and subtracting one rate from the other invents a trend.

Without that, a pre-fix run and a post-fix run would have been tabled together as comparable and 3 -> 14 would have entered the record as an engine regression. **A number that quietly spans a reset is worse than no number.**

RE-STAMP DELIBERATELY OR NOT AT ALL. `node engine/quarantine.js --stamp-whole-game` moves the baseline to this pin. It has NOT been run here: whether `f646b0163bc0` is the pin to hold going forward is a MEASURE judgement, not an ENGINE one, and the withheld direction-of-travel line is the correct state until somebody makes that call.

Why the refit lives here, not in ENGINE or SEARCH

The refit is the expensive event on the one expensive edge, and it invalidates seven artifacts: counterplay, winrate-backtest, opponent-calibration, weight-multiplicity, then the mag / mew / scoreboard bundles. `provenance.js` derives that set rather than carrying a typed list of it.

The division that owns *knowing when a number stopped being true* is the one that should be pulling that trigger.

A restamp is only valid if the feature FUNCTION is unchanged. Damage table moved → refit. Not a restamp. There is no version of this where the shortcut is fine.

Open — in priority order

THE CLOSET IS A GATE EXEMPTION NOW, AND IT SHIPS EMPTY — 2026-08-27

Will's ruling, 2026-08-26: "no if i put things into the closet it should not be gated — like illusion", and 2026-08-27: "things in the closet shouldnt block a gate if we know why they fail and choose to accept it." Implemented as a third `DECLARED_KINDS` entry, `CLOSETED` , in `engine/quarantine.js` . ROADMAP #508 . Full

account: docs/_reports/2026-08-27-closet-gating.md .

THE GATE DID NOT MOVE AND IT MUST NOT BE REPORTED AS THOUGH IT HAD. 5 of 8 clauses PASS before and after, measured on the same artifacts (release f3d423e19e88 , 961 games, read at 23:47 EDT); whole-game 1 of 961 either way. A gate that improves because the rules changed is not an engine improvement, and this one did not even improve.

THE ONE ROW THIS KIND WAS BUILT FOR WAS NOT WRITTEN. Will closeted Tailwind's expiry order on 2026-08-24; #493 FIXED it on 2026-08-27. The current differential holds six causes and none is a -sideend / tailwind pair, so a declaration would have registered a permanent exemption for a divergence that no longer occurs. #355 is closed with that evidence and the refusal is recorded as a comment where the row would have gone.

WHAT THE DOOR ASKS, AND WHY IT IS A SCHEMA AND NOT A SENTENCE. Nine fields, refused at declaredMatch if any is missing: the owner, the date, his own quoted words and the register row; the instrument, release, date and finding of the measurement behind the no-board-effect claim; and what would make the entry wrong. This is the deliberate inverse of NOT A DEFECT , which is a regex over a register cell — *"a sentence somebody typed as a note"* is the failure mode this division was sent to avoid reproducing.

AND IT IS RE-CHECKED. Four declarations in this project have been refuted — speed ties, Tailwind's coin, Moody, and a fainted body in an active slot — and the die fix of 2026-08-27 voided every measurement taken before it. closetEvidenceStale compares the entry's evidence release against the release of the artifact it is excusing and prints EVIDENCE NOT RE-CHECKED naming both. It still subtracts; the owner ruled. An unstamped run reports nothing and never reports it fresh.

THE RECEIPT ON THE OTHER DOOR OVERSTATED ITSELF 8x, AND THAT IS FIXED AS A DISPLAY BUG. DEFECT matches inside the phrase NOT A DEFECT , so the open-defect clause reported every row carrying the phrase as excused. Measured over 432 register rows (205 open): **8 carry it, exactly 1 (#252) would have counted without it.** No verdict moves — the open set, the red set and the clause pass/fail are byte-identical. The nine rows named in docs/_reports/2026-08-27-open-defect-clause.md are now **eight**; #344 was refuted and closed the same night, and the list was re-derived rather than inherited.

OWED. #336 (the sleep draw ADDRESS, never checked) is still excused by the phrase and is the one of the eight worth re-opening: it claims nothing and carries a WHAT WOULD DECIDE IT , but *"the distribution is right and whether the draw is addressed identically has never been checked"* is the exact shape of the four refuted declarations. It is not suppressing anything today — the receipt now proves that — so it does not hold the gate.

THERE IS NOW A SPEED BASELINE, AND THE BOX IS PART OF IT — 2026-08-27

Full account: docs/_reports/2026-08-27-speed-baseline.md . Will: *"i dont care about the old one i just want an honest baseline of how fast medicham can play a game out"* — so no prior figure is quoted or compared against; a number taken on a different engine is a different measurement.

One complete game: wall p50 7.4–9.7 ms, p90 10.6–14.1 ms, p99 14.4–19.1 ms, CPU 11.5–14.3 ms. Do not quote a max from it — in all twelve repeat runs the slowest game WAS the cold first game. n = 2,831 games in the headline run, all played to natural completion by a side wipe, **zero** hitting the 200-turn safety cap. Turns p50 7, p90 10, max 21. Release 6a845424c450 (HEAD, 0 of 26 files moved), --team-store data/team-pool-frozen . Composable primitive: **~1.1 ms per turn**; the turn loop is 98.7% of a game and turn count explains 60% of the variance in its length.

The cap does not bind. Only 4.4% of games reach turn 12 at all, so the whole-game differential's 12-turn cap truncates 1 game in 23 and is a no-op on the rest.

Nobody is searching in that number. It is `battleTurn(S, rng)` with no action arrays — medicham2's own greedy policy on both sides, the same call `battle()` and `previewLeaf` make. A MILTANK-driven game will run longer and cost proportionally more. Scale by your own turn count, never reuse 8 ms.

AND THE HEADLINE IS A RANGE FOR A MEASURED REASON, WHICH IS THIS DIVISION'S PROBLEM AND NOT AN ASIDE. The within-run noise floor (LESSONS §9, one arm split by interleave) is **0.006–0.152 ms**. On that floor almost anything is significant. It is the wrong floor: **twelve byte-identical repeats drifted 8.04 → 9.74 ms p50, monotone, 140× that floor**, and an idle gap took it back to 7.82 before it began climbing again. `cpu/wall` held flat at 1.40–1.45 throughout and free RAM did not move, so it is neither contention nor memory — it is a laptop shedding boost clock.

So: on this box, any A/B whose arms run at different times, on a difference under ~20%, is measuring the thermal state of the laptop. Interleave the arms or do not run the test. That binds MEASURE's own future work first.

Also measured and worth carrying: **CPU per game is 1.4× wall** (V8 GC/compile threads), so six concurrent processes want ~11 of 16 cores at ~500 MB RSS each — the documented six-process cap is tighter than it looks. And `tools\lownode.cmd` (BelowNormal) cost **0.069 ms**, below the noise floor: free, exactly as its header claims.

THE FIXTURE-LEGALITY GATE WAS RED AT FIVE AND THREE OF THE FIVE WERE THE GATE — 2026-08-26

Full account: `docs/_reports/2026-08-26-fixture-legality.md` . ROADMAP #266. CHANGELOG 5.151.0.

5 FAILED → **ALL GREEN**, baseline **22 → 15 verdicts / 23 → 15 pairs**, no allowance added. **Two real illegal entities** (`Incineroar` can't learn `Knock Off.` , one commit old; `Milotic` can't learn `Calm Mind.` , a ternary branch that can never be taken and typed the name anyway) and **eight stale allowances** that ENGINE had already repaired in `24fe4c5c` — proved a repair rather than a scanner that stopped looking, since `24fe4c5c^` declares all eight, HEAD declares none, and the file is still the sweep's largest contributor at 204 declarations.

The other three failures were the ruler. Helper detection judged a top-level `const` on a flat 500-character window that ran off the end of the declaration, so `const ok = (cond, label, extra)` inherited the `const stage = rows => rows.map(...)` beneath it and **every assertion in two files was scanned as a set declaration** — `'medicham='` read as the species `Medicham`. Fixed at the window; measured at **875 → 872 declarations, diff exactly the three phantom rows**, strays 5 → 0.

AND THE CHECK DOES NOT CATCH THE CLASS, WITH A COUNT. Three shapes walk past it. The largest is a positional row whose moves array is not in slot 2 — `stage(rows)` in thirteen files writes moves LAST — hiding **124 distinct sets, 21 of them illegal, 15 verdict sentences not on the baseline**. It was measured and **deliberately not armed**: those repairs change what their scenarios measure and belong to the divisions that own them, which is the same order matcher (C) followed. Second is `engine/validate_damage.js` 's golden master (two species and a move, all scalars, no array) — the file where a human, not this gate, found Choice Band, Choice Specs and an Amoonguss on 2026-08-25; audited by hand at **36 rows, 0 problems**, clean and still unseen. Third is rule 1's normalisation, which lets a punctuated fragment "name itself".

No game number moved and none could — no engine byte, no census, no differential, no release. Neither edited fixture was re-run; this batch may not play a game, so that is a PREDICTION and is recorded as OWED in the report.

CLOSED. ONE DECLARED LIST, ONE READER — THE MECHANICS CLAUSE COUNTED A ROW THE WHOLE-GAME CLAUSE HAD ALREADY DECLARED — 2026-08-26

Full account: docs/_reports/2026-08-26-declared-list.md . ROADMAP #464. CHANGELOG 5.147.0.

DECLARED_DIVERGENCE in engine/quarantine.js was consulted at exactly one site, inside wholeGameClause . That clause declared the Supreme Overlord fallenundefined line AUTHORITY-WRONG and subtracted its 5 games; tests/test-mechanics.js carried a live probe asserting we refuse the line deliberately; and classifyMechanics counted it as a defect on the same run, filtering on !r.diverged || r.deferred and never looking at the list.

It was never two artifacts that needed reconciling — it was one grammar with one reader. Both artifacts write <cls> :: |lineA <> |lineB from the same comparator, so the matching rule needed no loosening and was not loosened. declaredMatch(cause, ev, threw) is now the one door; the whole-game clause's inline loop was deleted rather than duplicated.

MEASURED, HEAD beside the working copy in one process against MD5-frozen artifacts

(release 667278050dcf): mechanics clause 9 → 8 of 16, exactly one row leaving —

ability:supremeoverlord , 112 teams in 13,116 open-sheet games, the most-played of the sixteen, so the reach filter would never have removed it. Whole-game clause **unmoved at 10 of 961** with a byte-identical why string; board-material **unmoved at 2 of 961**; census **unmoved at 750 live / 753 probed / 3 missing**; shelf list identical; declared + counted + shelved + unknown + cleared = 16 = rowsSeen = summary . **No engine byte touched.**

Nine selftest assertions added (**100 → 109 passing**), pushing a SYNTHETIC declaration rather than asserting anything about Supreme Overlord, because a gate built from an instance catches that instance and not the class. Shown red first: breaking the mechanics door (const dec = null) turns 4 red and leaves the whole-game one green.

What still walks past the door is named in the code's own header — differentialClause (numbers, no cause string), the three roster stages (our two engines, DEFERRED-BY-OWNER), coverageClause and openDefectClause (no cause string), and orderProbeClause , which DOES carry a cause and does not read the list. That last one is inert today (0 pairs probed, no ordering declaration) and is the nearest thing to the next instance of this bug. SHOWDOWN-ONLY is likewise a verdict no clause reads; those eight ability rows are handled correctly BY ACCIDENT. Both are recorded, neither is fixed here.

CLOSED. THE SWITCH INDEX WAS NOT THE BUG — 2026-08-25

Full account: docs/_reports/2026-08-25-switch-index-instrument.md .

ENGINE handed over the last two board-material "a chosen switch the authority performs and medicham2 does not" games with the harness named as the suspect: engine/game_differential.js sends switch N against a side.pokemon array Showdown **reorders** (sim/battle-actions.ts:118-132), so a cached index would not name the same body twice. Right first suspect — in this repo the ruler has been the culprit five times in two days — and **refuted**.

The instrument resolves that index off the LIVE array, by species, immediately before battle.choose . It now says so with its own counter rather than by being read: release 2ecd3bdc274b , 961 games, **63,258 switch indices sent, 43,125 of them against an already-permuted party, 0 MISADDRESSED**. The counter has been shown RED — MEDI_SWITCH_BY_INITIAL_INDEX=1 restores the cached-index bug and produces 4,932 misaddressed, 1,796 choices refused by Showdown and 901 of 961 games thrown. **Zero of the 23 first divergences on this release are a switch-addressing artefact**, so every board-material figure taken on this instrument stands.

The real mechanism is **ENGINE's and is one defect for both games**: `medicham2` evaluates `preventsSwitch` at switch-EXECUTION time while Showdown evaluates it at CHOICE time and never re-asks, so a Gengar-Mega arriving on an earlier switch in the same turn retro-cancels a switch this engine had already been told to make. `tests/probe_trap_timing.js` isolates it in five arms with the authority's own refusal as the positive control. Predicted effect of the fix, stated before it is taken: 23 → 21 raw, gate 18 → 16 of 961, board-material 10 causes → 8, honest range 21–22 / 16–17 / 8–9.

A second instrument defect was found on the way and fixed: `freshBodies` dropped `_switchKey`, so it was `undefined` on every body this instrument has ever played and the medicham switch lookup has always fallen through to the mutable display name. CLAUDE.md already names that cause (Morpeko) and it was still live. Predicted before the run — misses stay 3, diverged stays 23 — and measured exactly so. A dead safety net looks exactly like a working one.

CLOSED. THE PLANTED-STATE PROOF WAS FAILING ON THE FIXTURE, NOT THE COMPARATOR — 2026-08-25

Full account: `docs/_reports/2026-08-25-planted-state-proof.md`.

`planted_state_proof_ok` has been **false on every committed artifact back to 24 August**, and `game_differential.js` exits 1 on it — so every board-material figure carried an unexamined caveat about whether the state comparator can detect anything at all.

It can. Thirteen of forty-two plants failed on the artifact's proof pair, and all thirteen were aimed at bodies that could not carry them: at the plant boundary side B is two corpses and all four benched bodies on both sides are dead. `board_state.js` holds the post-faint group — `item`, `status_counter`, `boosts`, `ability`, `vol`, `stall` — on a body both engines call dead, so seven volatile plants written into a fixed `S.actB[n]` moved no compared leaf and six bench plants correctly refused to plant at all. **The control is what settles it**: the same seven mutate functions handed a body that is standing are caught, localised, at the same boundary.

The pass/fail was a function of `--games`. `diff_swarm` picks by a deterministic stride, so the SECOND team of the proof pair moves with the sample size; the identical code at `--games 45` passes all 42 plants and at `--games 1200` fails thirteen.

Three fixture fixes in `engine/game_differential.js`, none of them a weakening: the volatile plants find a standing body and return the slot they used (a tighter localisation than four of them asserted before); `applied` now means the board MOVED where the comparator looks, asked of `BS.compare` rather than of the callback's return value; and a plant that cannot land at the last agreeing boundary walks back through the earlier ones — **on applied only**, never on `caught`, because retrying a plant that landed and was missed would hide the one thing this proof exists to expose. 42 of 42 CAUGHT+LOCALISED afterwards on the same pair, and `tests/test-state-differential.js` passes with 42/42 `applied` on all six pairs.

The published board-material figures stand. ROADMAP #314's consequence line — that DIFFERENT-END-STATE is a lower bound wherever quoted, including the 83 in CHANGELOG 5.54.0 — rested on those thirteen being indeterminate. They are the fixture, so that caveat is withdrawn. The committed artifact still carries the false flag until the next `--state` run regenerates it.

CLOSED. THE BROWSER RULEBOOK DRIFTED FROM THE NODE RULEBOOK, TWICE, AND NOTHING COMPARED THEM — 2026-08-25

Full account: `docs/_reports/2026-08-25-tags-drift.md`.

`data/tags.json` is what the node engine reads; `data/abra-tags.js` is the browser copy of it, frozen into every engine release. A `tag_dex` run at 2026-08-24 23:37 rewrote the first and left the second at its 22:59 content, **and both were committed that way** — so HEAD carried two engines disagreeing about three moves (Bug

Bite and Pluck stealing any item instead of only a berry, Thunder Wave ignoring the type chart). Both params have live consumers in `medicham2-browser.js`. An earlier instance of the same drift ran for **two days**. Neither was caught by a check; the second surfaced because a release cut happened to look.

No new gate. `build/build_tags.js.js` gained a `--check`, copying the pattern `build/build_guru.js.js` has had since 2026-08-04, and `engine/artifact_audit.js` — already a gate, already the file for *"a derived artifact is not a fact until something compares it to its source"* — gained **check G**, which derives the pairs from the `GENERATED` by `<builder>` from `<source>` headers under `data/` and runs each builder's own `--check`. Nothing is re-implemented and nothing is typed.

It was **RED on the real drift** and green after `data/abra-tags.js` was regenerated — that regeneration is the only artifact this pass changed, and it changed the file to agree with a `tags.json` it should always have agreed with.

`build/build_browser_data.js` got the same `--check` and is now compared too: `data/move-effects.js` (frozen in every release, carries move priority) and `data/mega-formes.js` were derived from the Champions dex with **nothing standing between them and it** but an mtime test. Shown red on a deliberate break in a scratch copy. **Six bundles remain uncovered and check G prints them by name on every run** — the list is derived, so it is never quoted from here.

oooooooooooooooooooo. MOODY IS DECLARED INCOMPARABLE; SPEED TIES WERE REFUSED, BECAUSE THIS HARNESS ALREADY SHARES THE TIE DIE — 2026-08-23

Full account: `docs/_reports/2026-08-23-declared-moody-ties.md`. Will ruled on both by name (*"yeah some things we can just quarantine as a known failure with a quoted reason (speed ties, moody, etc)"*), then *"yes we can have two things, impossible to compare, and too difficult and irrelevant to add at the moment"*). One of the two landed.

MOODY LANDED. `data/abilities.ts:2691-2716` picks the stat with `this.sample(stats)` twice, over the five main stats only, and `medicham2-browser.js:25831` implements the same rule over `['at','df','sa','sd','sp']` with a real draw — so the RULE is not in dispute and only WHICH STAT the die names is. The middle arm addresses a draw by `[seed, turn, category, move, target]` plus an occurrence index, and a residual `sample()` has no named category on either side, so both engines take it off the generic `any` stream at indices each fills with its own unrelated draws. 8 of 961 games. The row states that this is **the instrument's addressing, not a law**: give the residual pick a named stream on both sides the way ROADMAP #290 did for `tie`, and the row must be DELETED rather than kept.

SPEED TIES WERE REFUSED, AND THAT IS THE FINDING. The derivation handed over — `sim/battle.ts:429` `speedSort` ends in `prng.shuffle`, so a tie is a coin flip — is true of the GAME and false of this HARNESS, and only the harness's number reaches the clause. Showdown's shuffle is a **no-op** in every shipped arm (`sdShuffleReverses false`), `medicham2`'s tied-group key has had **its own stream since 2026-08-20** (`RNG_STREAMS` includes `tie`) which the middle arm neutralises with `o.tie = () => 0`, and 3.74.0 fixed the tie at the root — which is why the two `tie-second` arms were RETIRED for "breaking a correct one". **The sixth die the declaration would have claimed does not exist is already shared.** So the 3 causes whose own probe reads `speed_tied: true, speed_gap: 0, same_priority: true` are a REAL turn-order disagreement, and declaring them would have subtracted a live defect under a heading reading "nothing to fix" — the `medicham2-browser.js:17440` failure with a better-sounding reason. They stay UNDECLARED and are proposed as a roadmap row.

THE TWO KINDS PRINT APART AND ARE NEVER SUMMED. `INCOMPARABLE` ("the authority makes a random draw we have no shared address for — no defect") and `AUTHORITY-WRONG` ("matching it would make us less correct") get separate headings and separate counts, plus `declared_by_kind` on the clause result. The

clause moved `declared 5 -> 13`, `undeclared 43 -> 35` on 48 raw over 961 games; `diverged` and `declared` both still print, so the allowance is legible rather than absorbed. **It is still RED and nothing here could have opened it.**

THE THIRD KIND — DEFERRED — IS DELIBERATELY NOT BUILT, AND THE HOLE IS NAILED SHUT. A DEFERRED row asserts the opposite of the other two: that there IS a defect.

`DECLARED_KINDS` is a whitelist, so a row typed with any other kind is NOT subtracted, counts as UNDECLARED and is NAMED on the run. The selftest pushes a `kind: 'DEFERRED'` row with `match: () => true` and asserts `declared=0, undeclared=1, ok=false`. The design for it, if it is ever wanted, is in the report.

THE NEGATIVE CASES ARE THE PROOF, NOT THE POSITIVE ONE. Nine mutations of the Moody cause — each a real defect wearing the same shape — were injected through the shipping clause and every one fell through to UNDECLARED: accuracy in the pool, evasion in the pool, `+3` instead of `+2`, `-2` instead of `-1`, our line unattributed, one of two occurrences unattributed, the authority attributing the boost elsewhere, the boost following that slot clicking a move, and a different slot on each side. Speed ties are asserted undeclared at gap 0 **and** gap 40, so re-adding the row goes red. Selftest: 108 passed, 0 failed.

oooooooooooooooooooo. THE IDENTITY GATE IS A RUNTIME TRIPWIRE NOW, AND IT FOUND ON ITS FIRST RUN A DEFECT THE COUNT HAD BEEN GREEN ON FOR THREE WEEKS — 2026-08-23

Full account: `docs/_reports/2026-08-23-identity-gate-permanent.md`. Will: *"you do what you think is best just make it a permanent solution."*

What `tests/test-effective-identity.js` **asserts now.** Every mon a `Board` switches in gets a recording accessor on `ability`, `baseStats`, `weighthg` and `weightkg`. Every active on the test board holds a mega stone whose forme ability differs from every ability its base forme can have — swept out of the dex and the damage table, nothing named — so on that board **the declared ability is wrong for all four actives** and a raw read is a defect by construction. The board is then driven through the live decision path and every read is recorded with its stack. **Exactly one call site may see the raw field:** `board.js effective()`. That is Fowler's SELF ENCAPSULATE FIELD, which this file's header has cited since 2026-08-02, stated as an executable assertion instead of as a count.

IT FOUND A REAL ONE IMMEDIATELY. `engine/position_features.js:249` reads `f.mon.ability` off a live stone-holder to build the defender list for `priorityRefusedAbove()`. `engine/board.js:3520` is the same function one file over and **it resolves** — `effAbility(f.mon, dex)`. One fact, two implementations, which is the failure CLAUDE.md names as *FEATURES ARE PER-MODEL, FACTS ARE GLOBAL*. **The retired count ratchet was GREEN on that file and had been since 2026-08-02**, because `position_features.js` sits exactly at its per-file number of 5 and a count only ever asks whether the number went up. Exposure is zero today and the gate **re-derives that on every run** rather than asserting it: the value is consumed only by the priority bar, and no legal mega in this format gains or loses a `blocksMove` ability. If one ever does, the gate goes red on that entry by name. MEASURE does not own `engine/position_features.js`, so it is declared with the guard and proposed as a roadmap row rather than edited.

SHOWN RED ON SIX PLANTED SHAPES AND ON ONE REAL BREAK. `ABRA_EI_PLANT=all` compiles six stale reads through `vm` — so they exist only while the knob is set and no scanner can ever count them as debt — and the gate named all six and exited 1: `mon.ability`, `const {ability} = mon, ({ability}) => ...`, `{...mon}`, `Object.assign({}, mon)` and `mon[k]` with a computed key. **The last three were conceded by the old header as undetectable by any text scan.** Separately, `engine/board.js:3520` was edited to read raw, the gate named `engine/board.js:3520` with the offending source line, and the file was restored from a byte copy taken first (SHA-1 verified identical, `90ae57fe3bcaf886a29f8affc273bf437dd9d2af`) — not `git checkout`.

WHAT IT PROVABLY CANNOT DO, and this is stated in the gate's own header rather than here. Its coverage is EXECUTION coverage. It drives `board.js` `candidates()+featuresFor()` and `foeActionDistribution()` for all four actives, `position_features.js` `positionFeatures()` both sides, and `rollout_leaf.js` `rolloutWinProb()` both sides — the MAG feature path and the leaf. **The drive list is printed on every run** so a reader can check whether their consumer is in it. It does not cover `magnemite.js` 's live loop, the fitters, or the differential harnesses. It is also blind to a value copied out of the sheet before a Board existed and then treated as live: that is the old scan's `Object.assign` hole **moved rather than closed**, and it is written down as such.

THE COUNT RATCHET IS RETIRED, NOT RESTAMPED. `data/effective-identity-baseline.json` keeps every number it ever asserted under `last_count_baseline` (generated 2026-08-11, count 1198, 80 files) with the reason beside it; `--update` and `--propose` now refuse and exit 2. **Nothing was adopted.** The scan still runs as a printed inventory that asserts nothing, and the 32 walked-file notes are kept because each is somebody's line-by-line account of a file and that is the expensive part. One narrow static assertion survives — `baseSpecies(...).baseStats` at zero — because it is the one text shape a whole-repository walk named as dangerous and it reaches files the tripwire never executes. It is a supplement and the file says so.

The 18 behavioural assertions are untouched, including the three that build a body, mega-evolve it inside a real turn, and assert the effective ability is what gets read. The run is 24 passed, 0 failed, exit 0.

oooooooooooooooooooo. THREE OF THE FIVE FAILING GATE CLAUSES WERE STALENESS, NOT A BROKEN SIMULATOR — FULL REFRESH 2026-08-23 ON RELEASE
`0faabe2a3f1b`

Full account: `docs/_reports/2026-08-23-gate-refresh.md`. The gate went **6 of 8 clauses failing to 3 of 8** without a single line of engine code being touched, because the three roster clauses had been measured against release `c36782953dee` and were correctly withheld as *an answer about other bytes*. Re-run on one fresh release: **items 0 DIFFER / 0 DID-NOT-FIRE, abilities 0 / 0, moves 0 / 0**. Named, because a row number is not a finding — the previously red entities *Big Root*, *Greninjite*, *Dragon Cheer*, *Fake Out*, *Matcha Gotcha*, *Psych Up* and *Transform* all now match.

THIS IS THE CASE THE PASS / FAIL-BECAUSE-BROKEN / FAIL-BECAUSE-UNMEASURED SPLIT EXISTS FOR. Five clauses read red this morning and three of them carried no information about the engine at all. A gate that cannot tell those apart trains people to read *closed* as *broken*, which is the same normalisation failure as "one of the two known failures".

AND ONE CLAUSE WAS BLIND FOR A THIRD REASON THAT LOOKED LIKE THE OTHER TWO. *"THE REACH FILTER CANNOT BE APPLIED"* was neither staleness nor breakage: `engine/all_mechanics_fire.js` defaults to `--kind moves`, so the published `data/all-mechanics-fire.json` held `rows.moves` and nothing else, and `quarantine.js:718` refuses to filter unless all three populations have per-entity rows — falling back to counting every divergence unfiltered. Re-run `--kind all` and the clause answers for the first time: **29 of 36 diverging mechanics played and uncleared** (moves 22, abilities 12, items 2; 7 below the reach shelf), worst by reach *Cursed Body* 2,177 teams, *Toxic Debris* 1,840, *Disable* 1,799 clicks, *Regenerator* 1,596, *Poltergeist* 1,383, *Mental Herb* 967. **22 → 29 is a RE-BASELINE, not a delta** — different population and a filter that could not previously run.

The remaining two failures are real. Whole-game: 961 games on the `middle` arm, 73 parted less 5 declared = **68 undeclared (7.1%)**, of which the end-state comparison says **21 DIFFERENT-END-STATE and 52 wording-only**; direction of travel stays WITHHELD because the stored baseline is stamped under `top-tie-first` and this run is `middle` — one pin is one corner, and `--stamp-whole-game` was deliberately not run. Register: five open rows name a RED instrument (#218, #224, #241, #258, #273).

BLOCKED, NOT OWED: `engine/wire_ladder.js` **CANNOT RUN AT ALL.** It exits 4 and writes nothing — correctly. All **fifteen** arms pin releases frozen before `engine/medicham2-browser.js` exported `natureL50`, `rngStreams`, `spreadL50`; that is all fourteen distinct ids, and `compat` reports 168 of 351 releases predate an export. LESSONS §12 at ladder scale — `data/wire-ladder.json` stays UNSAFE, the release-ladder figure stays WITHHELD, and the unit of work is re-pinning the arms, never a re-run.

And leaf calibration — this division's one number — is still QUARANTINED and still withheld.

A reliability curve measured through an engine that parts from the authority on 21 of 961 games is a claim about the wrong engine, and publishing it with a caveat is the bug. It becomes re-runnable when the gate opens, not true.

oooooooooooooooooooo. THE MUTATION ARTIFACT WAS SIXTEEN DAYS STALE, AND WHAT IT WAS HOLDING WAS NOT WHAT IT SAID — RE-RUN 2026-08-22 ON RELEASE

`6fb9ebd3b704`

`data/mutation-coverage.json` had not moved since 2026-08-06 and was quoted as *"163 class-A rows, tag never read"*. Re-run against a fresh cut of the current tree. **The count is 148 and the count is the least useful thing in the file.**

It could not run at first, twice, and both refusals were correct. The triage calibration hand-decided `taunt / forbidsStatusMoves` as class A against release `032b4a2979dd`; ENGINE has wired it since, so the rule returned D and the sweep refused. The planted `speedMult` stub anchored on `s*+=_sm.mult`, which is now `_mods.push(+_sm.mult)`, and `plant()` threw rather than planting nothing. A hand answer is a claim about a BUILD; neither constant recorded which. Both repaired, `decidedAgainst` added, and the A branch moved to a synthetic that no fix can close — ROADMAP #326.

The two counts are not comparable and the per-key story is. The tag corpus grew 182 -> 292, so the battery scope moved `620f24df16ff -> c5c9acf2cc9d` and the ratchet correctly declined to compare them. Of the 163 old class-A operators: **10 NOW-LIVE** (Taunt, Knock Off's `failsIfNone`, Substitute, Poltergeist, Pollen Puff, Shed Tail — closed by ordinary engine work), **28 reclassified** to B/C/D, **11** became UNREACHED-BY-THIS-BATTERY, **3** downgraded, **6** no longer emitted, and **105 still class A**. 43 of the 148 are rows the old battery never had.

CLASS A IS NOT A DEFECT LIST, AND SAYING SO IS THE FINDING. 31 of its 36 carrier x tag rows carry an ARMED census probe that PROVES the mechanic works. I hand-read the other five: `move:leechseed / immunityGate` is graded A with *"Nothing in the simulator implements this fact"* and the Grass immunity is implemented at `medicham2-browser.js:18616` through a SIBLING tag the classifier cannot see. ROADMAP #323. Three real gaps survive — Trick vs Sticky Hold (**o corpus uses**), `punishesMinimize` (Minimize is **32 of 198,840 sheet entries**), item Metronome (**27**) — total reach under 0.3%, filed LOW as ROADMAP #327 so they are not re-discovered rather than so they are fixed.

THE VEIN IS CLASS B AND C: 119 VALUES IN `data/tags.json` **THAT NOTHING READS.** The engine consumes the tag's membership and substitutes its own number. Protect's `shieldsUser / stallCounterChecks / targetClass` at **67.75%** of sheet entries, Fake Out at 12.44%, **Life Orb's** `costsPerAttack` **at 10.69%** — the row rediscovered the expensive way on 2026-08-21 after sitting in this file for sixteen days — Sucker Punch 7.11%, Flare Blitz and Wave Crash `recoil.readFrom`, Knock Off's `mult = 1.5`, the drain `num / den` on Matcha Gotcha and Giga Drain. None is a behavioural defect today; that is the point. It is *A DERIVED ARTIFACT IS NOT A FACT UNTIL SOMETHING COMPARES IT TO ITS SOURCE* at 119 live sites with no comparator. **The unit of work is one comparator, not 119 fixes** — ROADMAP #324.

And what is NOT a defect, said out loud because the counts grew fastest here. 98 UNREACHED-BY-THIS-BATTERY (was 32), 44 UNSTAGEABLE tags (was 25), 9 NO-CARRIER. The fixture did not keep up with a tag corpus that grew 60%. A COULD-NOT-STAGE verdict is a claim about the fixture, never about the mechanic, and none of these may be reported as breakage.

Pinned: engine release 6fb9ebd3b704 (26 files). The battery is NOT census-steered — censusCrossRef() reads data/mechanics-census.json only to ORDER class-A rows, so no census pin is needed for the sample; the census was digested 80e648f34d56 and was not regenerated. ENGINE cut 136a9894af62 while this ran, which is exactly what a release is for.

0000000000000000. TWO OF THE THREE ROWS HOLDING THE THIRD GATE CLAUSE SHUT ARE FIXED AND CLOSED; THE THIRD IS RED FOR A NAMED REASON — 2026-08-19

engine/board.js (stampSky , jointFeaturesFor), engine/magnemite.js (_candsFor), engine/seed_source_audit.js , engine/gate_weather_guard.js , engine/gate_seed_source_audit.js , and twenty-four catch blocks across eighteen files. ROADMAP **#286 CLOSED, #287 CLOSED, #258 still open**. The clause no open, known engine defect went from naming **six** rows to naming **four** (#218, #241, #258, #290).

#286 — A GUARD THAT HAD NEVER BOUND, ON A FIELD NOTHING WROTE.

weatherSetupHelpsPartner read norm(A.__weather || '') !== w , and __weather occurred exactly once in the live tree — that read. So the comparison was '' !== w and the "only if the weather is not already up" clause had never bound in the history of the file. board.stampSky(cands, board) is the assignment, and **both** candidate producers call it: board.candidates for the fitter, the fixtures and the rollouts, and magnemite._candsFor , which builds its menu from the live REQUEST and does not route through the first. Stamping one of the two would have been the fitting/playing mismatch this project has already paid for twice. The value is board.weather , the expiry-aware accessor #276 built.

THE SECOND DEFECT THE ROW RECORDED IS ALSO FIXED, AND THE EXISTING ARM COULD NOT SEE IT. The loop runs [[A,B],[B,A]] and asked A.__weather on both passes. Reverting that left gate_weather_guard.js at **exit o** — because both candidates of a real pair come off the same board and carry the same stamp, so the wrong-half read agrees by accident. A fourth pass was added: with only the SETTER's field populated, the feature must answer the same in either argument position. That is exact rather than sampled, so it cannot false-positive — the bar the row's rejected static scan failed. Both halves were shown RED on a deliberate break before being trusted.

REACH, MEASURED WEIGHT-FREE, BECAUSE THE WEIGHTS ARE QUARANTINED. Over **2,000 corpus games, 21,757 joint decision points and 1,211,091 pair evaluations**, through joint_rows.build with a ZERO ranker and K above the longest candidate list so that no pair is dropped by a quarantined number, and with both arms taken from ONE process so ENGINE editing the simulator underneath could not separate them: the feature fired **344 times before the repair and 219 after — 125 suppressed, 36.3% of every firing of this column was a redundant weather setup** — on 64 of 21,757 decision points (0.29%). At 300 games the same quantity reads 49.2% of 59 fires, which is what a sample that small is worth. **THE ARGMAX-FLIP RATE IS WITHHELD RATHER THAN UNMEASURED:** each flipped pair's score moves by exactly the fitted weight for this column, and publishing a rank change would be quoting the quarantined vector.

AND THE LIVE HALF IS PROVEN, NOT ASSUMED. A 4-game self-play through the joint path reports skyStamped=6779 , jointSkyUnstamped=0 , skyStampNotAnArray=0 . tests/test-wiring.js caught the first attempt outright — _dropDeadVolatiles returns {cands, user, doubles} , not an array — which is precisely the class of defect that file exists for.

A SEARCH-OWNED FILE WAS EDITED UNDER A MEASURE BRIEF. `engine/board.js` and `engine/magnemite.js`, on the router's instruction. It is on the record here rather than implied.

IT MOVES A FITTED JOINT COLUMN, AND NO REFIT WAS RUN. `weatherSetupHelpsPartner` means a different quantity after this than before it, so it joins the set the refit owes. The refit was ALREADY owed and is gated behind MEDICHAM rather than behind compute, so this adds a column to a debt that exists; it does not create a new one. Nothing was fitted, and nothing downstream of the weights was re-quoted.

#287 — AN AUDIT THAT SUBSTRING-MATCHED A HAND-TYPED LIST. `STUB_HAS_NO` is gone. `seed_source_audit.js` derives the class by calling `board.unmodelledBasePower(dex, board)` — asking each callback to produce a number and recording the ones that cannot — over a union of an empty board, its own staged board and every canonical fixture board, with the board count recorded and a fixture that will not build REPORTED rather than swallowed. Nothing reads a callback's spelling any more. The derived class is **four: avalanche, beatup, payback, ragefist**; the row's earlier seven, including Triple Axel, is exactly the figure a derivation corrects with no edit, because #283's `mv.hit = 1` made Triple Axel answerable after that list was typed.

THE ARTIFACT'S FALSE CLAIM IS GONE AND ITS STATE IS DERIVED. `openAndNotFixed` keeps its AUTHORED caveat text and reads OPEN/CLOSED from `docs/ROADMAP.md` through `quarantine.roadmapRowIsClosed`, imported and never copied. #244 reads CLOSED, so the caveat is DROPPED with its reason recorded rather than annotated; an unreadable register WITHHOLDS every caveat instead of publishing on an unchecked premise.

THE ROW'S REASON FOR NOT REGENERATING WAS MEASURED AND IS FALSE. It said regenerating today would bake ENGINE's in-flight `data/tags.json` into the artifact. With `engine/tags.js` and `data/tags.json` both made unreadable by a preload, the regenerated artifact is BYTE-IDENTICAL apart from its timestamp. The file's header already declared it does not read tags; that declaration is now checked rather than trusted, which is the difference between a claim and a measurement.

AND THE GATE ITSELF HAD THE LESSON. Arm 2 read `.row` off a single object. When the repair turned `openAndNotFixed` into `{checkedAgainst, rows, dropped}` the arm found ZERO assertions and went GREEN — a check that stops looking and calls the silence a pass, inside the gate written to catch exactly that. It was found by breaking the artifact deliberately and noticing the arm did not shout. It now accepts three shapes and REFUSES anything else with exit 2.

#258 — RE-MEASURED ON THE DAY, AND IT IS STILL RED. Both figures the row carried were stale. Today: **827 catch blocks, 291 silent (35%), 102 manufacturing; 201 baselined, 1 fixed, 91 NEW, 48 of those manufacturing.** All **24 of the new-manufacturing blocks reachable this session are fixed**, taking NEW 91 -> 67 and new-manufacturing **48 -> 24**. Each repair keeps the conservative value and adds the reason — a counter, a named `console.error`, or both — because inventing a different answer is a second defect rather than a fix. The most consequential were `mod_audit.js` x3 (three reads whose EMPTY answer read as *"Champions changes nothing"* inside the tool whose whole job is that question), `champions_sim.js` (a `checkCanLearn` exception became a LEGALITY CLAIM about this format) and `quarantine.js` (a declared-divergence matcher that threw silently moved its cause into UNDECLARED and INFLATED the rate this file publishes).

THE 24 THAT REMAIN ARE NAMED, AND EVERY ONE IS IN A FILE ANOTHER DIVISION HELD THIS SESSION: `engine/game_differential.js` (6), `engine/tag_dex.js` (2), `engine/medicham2-browser.js` (1) and fifteen across `tests/`. **This is not a filed failure.** The gate is RED, the reason is stated, the reason is ownership, and the next holder of those files closes it. NOT re-baselined.

oooooooooooooooo. THE HEADLINE WITHHOLDS, THE FIGURE WAS RE-TAKEN, AND FIVE ROWS THAT ASSERTED BREAKAGE NOW HAVE GATES — 2026-08-18

engine/quarantine.js (--whole-game , clauseExit), engine/gate_fail_and_silent.js , engine/gate_weather_guard.js , engine/gate_seed_source_audit.js , engine/gate_offfield_target.js . ROADMAP #298 / #241 / #286 / #287 / #218 / #224, and #299 / #300 / #301 filed.

#298 — THE MOST-QUOTED ENGINE NUMBER IS WITHHELD ON A RELEASE MISMATCH, AND CLAUDE.md HAD ALREADY DECIDED THAT. wholeGameClause was the only one of four clauses that compared no releases at all; mechanicsClause , decisionImpact and orderProbeClause have all refused an artifact cut against other bytes since they were built. It now refuses too, and on a mismatch **nothing measured comes back** — no rate, no diverged , no games , no class composition, and none of those numbers inside the verdict string either, because a withheld figure still sitting in the prose is the same bug with a different word in front of it. What comes back is which release it wanted, which it got, and the command that repairs it. It refuses a MISMATCH and not an ABSENCE: an unstamped artifact still answers, exactly as orderProbeClause allows one. The exit mapping is now clauseExit , one implementation shared by --whole-game and --order-probe , so a withheld verdict cannot exit 0 through one command and 2 through the other. node engine/quarantine.js --selftest is **87 cases, 0 failing**, four of them new and two of them asserting the red case twice over.

AND THEN IT WAS RE-MEASURED RATHER THAN LEFT WITHHELD.

engine/game_differential.js on the tree's own release: **696 of 995 = 69.9%** (data/game-differential.json , release 978ca8fe72c9 , arm A/middle, 700 raw less 4 declared, 0 cleared on decision impact, planted-divergence proof caught, 0 threw). node engine/quarantine.js --whole-game prints it and exits 1. **READ THAT DENOMINATOR AS WHAT IT IS: every game played, including the ones where the instrument itself desynced.** It is not the rate among games the comparison was valid for, and #299 below is that finding.

IT IS NOT A BEFORE/AFTER AGAINST THE 690 of 1230 THIS WAS FILED ABOUT, AND THE PAIRED ARM IS WHY RATHER THAN A HEDGE. The driver steers off a census that is regenerated and draws pairs from a team pool read live from the store, so two runs a day apart do not play the same games — the same request for 1,230 games returned 995. --release 6875c8ace00e replays the OLD engine bytes under TODAY's steering, which is the comparison that isolates the engine, and the two arms are the same run to the game: identical primary-arm counts, identical class composition, identical event missing from medicham2 block. **The engine moved by one game in one tie arm and by nothing at all in the primary one.** Every difference between 56.5% and 69.9% is the sample.

#299 — FILED, AND IT IS THE FINDING THIS RE-RUN ACTUALLY BOUGHT. The driver's own rule is that a game whose per-category draw counts diverge is VOID, not a divergence. Its stdout says so in three numbers; the artifact persists two of them. By the run's own arithmetic **72 of the 700 diverged games are themselves void**, and on the games where the instrument is valid the engines disagree about **628 of 645 = 97.4%** rather than 69.9%. The honest figure is worse, not better, and neither one can be published today because the void count lives only in a console line nothing parses. The clause is green only at zero either way; what is wrong is that the published ratio is over two populations that are not the same one.

THE FIVE UNGATED ROWS, EACH WITH ITS INSTRUMENT AND ITS EXIT CODE. Every gate refuses a stale artifact with exit 2 rather than passing, and every one was shown red on its own selftest before being trusted.

row	instrument	exit	what it measured
#241(3)	node engine/gate_fail_and_silent.js	1	30 causes / 51 games, green only at zero
#286	node engine/gate_weather_guard.js	1	2 of 2 staged pairs still fire with the weather already up

row	instrument	exit	what it measured
#287	node engine/gate_seed_source_audit.js	1	both directions plus the false claim about a closed row
#218	node engine/quarantine.js --whole-game	1	the gating half, runnable at last
#224	node engine/gate_offfield_target.js	0	the defect is ABSENT on a current artifact

#241 — THE PIN WAS SEEDED AT 3 AND RE-SEEDED AT 30 THE SAME DAY, AND THAT IS THE LESSON RATHER THAN A CORRECTION. The gate exited 3 and said *this got WORSE*. It had not: the paired `--release` arm shows the class byte-identical across both releases under one sample. The population moved, not the engine. So the pin now carries the census digest and the team-pool digest out of the run's own steering block, and **a REGRESSION verdict is WITHHELD whenever they differ** — still red, still LIVE, but no movement attributed to a lever nobody measured. A pin without its sample stamp is a coincidence, and one integer is exactly as good at hiding that as the bare `25` in #295 was.

#286 — THE FUNCTIONAL ARM, AND THE STATIC SCAN STAYS REJECTED. The gate finds the pairs the shipping `jointFeaturesFor` itself scores as a weather synergy with no weather up, then re-asks the same boards with `board.setWeather` through the Board's own door and the candidates rebuilt through `B.candidates`. It names no move, no type and no weather — it reads them off the candidate it was handed. A third pass populates the guard field BY HAND and requires the feature to drop, which is not a pass condition and not a fix: it is the proof the gate is not stuck red, because a check no repair can clear is decoration.

#224 — CLOSED ON A MEASUREMENT, AND GATED ANYWAY. It had closed on the right kind of claim and then nothing re-took it for six days while the simulator was rewritten nightly. A measurement taken once and quoted afterwards is prose about a number. The gate reads the engine's own `traceBodyOffField` counter as well as the artifact text, and **does not count a stale artifact**: `data/divergence-turns.json` is stamped five releases back and is named and excluded rather than contributing an old zero.

TWO MORE ENGINE FINDINGS, FILED AND NOT TOUCHED. #300: the catch that reports a failed Showdown wrap assigns a `let` declared five hundred lines below it, so a bad `SHOWDOWN_PATH` kills the run with a temporal-dead-zone error naming neither the path nor the require — the failure REPORT is the part that does not work. #301: `switch` lookups that MISSED: `medicham 22, showdown 0`, on a line the run itself marks MUST READ 0, identical on both releases, caught today only by a human reading stdout.

NOT TOUCHED, DELIBERATELY: #220 (with ENGINE), and #273, whose `tests/probe_red_demo.js` exits 1 and reads PREMATURE CLOSE in `register_reality.js` — `tests/` is ENGINE's. `register_reality.js` exits 1 today for that row and that row alone.

oooooooooooo. THE REACH SHELF IS ONE ANCHOR AND A RATE, AND THE REGISTER WIRE HAD NEVER CARRIED A ROW — 2026-08-18

`engine/quarantine.js ( --reach , --order-probe , --selftest )`, `engine/register_reality.js`, ROADMAP #295 / #296 / #297 / #298.

#295 — THE SHELF WAS TWO THRESHOLDS WEARING ONE NUMBER, AND FIXING IT MOVED NOTHING. The literal 25 was compared against CLICKS in a 64,846-game population and against TEAMS in a 13,116-game one, so one integer meant **0.0020% of clicks** or **0.095% of teams** — a move cleared the bar at roughly **48x lower relative usage** than an ability. Nobody decided that. The anchor is unchanged and now carries its unit in its name (`REACH_SHELF_CLICKS = 25`, from `tests/roster.js:1517` 's `USAGE_SHELF_BELOW`, applied at :1522 as `clicks >= USAGE_SHELF_BELOW` — a clicks threshold, ruled on as a clicks threshold; Will, 2026-08-18: "leave it at 25"). The common unit is **occurrences per stored game**, which both instruments already express: `clicks / click-counts.store_games` (64,846) and `teams / sheet-usage.sheet_games` (13,116). The teams threshold is derived at the anchor's rate — $25 \times 13,116 / 64,846 = 5.06$ **teams**, counting from 6 — and the comparison is integer cross-multiplication, so the move shelf is preserved bit-for-bit at `n < 25` and no float boundary can move a row.

	before	after
counted (played, uncleared)	34	34
shelved by reach	15	15
rows that moved		none, in either direction

IT REPRODUCES AND THAT IS STILL A NEW THRESHOLD. No diverging ability sits between 6 and 24 teams — the lowest counted one is `angerpoint` at 25 — and every shelved row is a move decided by the unchanged anchor. So the equality is a property of today's population, not of the change: an ability diverging at 6–24 teams used to be shelved and now COUNTS, which is the direction the defect demanded. **Anchoring the other way does not reproduce:** take the ability figure as the anchor and the move shelf becomes 124 clicks, shelving twelve currently-counted moves from `move:ficklebeam` (112) down to `move:attract` (30). Same fix, opposite result — which is why the anchor is named rather than assumed. Measured on `data/all-mechanics-fire.json` release `488fd1bf3f7c`.

THE DENOMINATOR NOW TRAVELS WITH EVERY COUNT. `reachOf` carries `denom / denomLabel`, the clause prints `2177 teams/13,116 open-sheet games`, and `--reach` prints the population, the rate per 10k games and the applicable shelf on every single row. That was the actual defect: a bare integer beside a name cannot carry its unit, and one header line is not a fix. Two further literals were untangled — `coverageClause` reads the CLICKS anchor because its loop walks `click-counts.json.moves` and nothing else, and the decision-impact sample floor became `DECISION_POINTS_FLOOR`, the same 25 in a THIRD unit.

#296 — THE GATE READ THREE KEY NAMES THE VERDICT ARTIFACT HAS NEVER WRITTEN. The open-defect clause read `rr.rows`, `v.command` and `v.exit`; `engine/register_reality.js` writes `results`, `cmd` and `green`. **Zero rows had crossed that wire since the day it was built.** Every open row fell into the DEBT bucket, `withRed` was structurally empty, and the clause printed *"no open row names an instrument that is RED"* for the reason that should have made it loudest — while #258's instrument was exiting 1. That is `merge_mega_into_engine.js` to the letter: matched on nothing, reported success. The fix is a declared shape (`REGISTER_REALITY + registerRealityRows`) that the WRITER writes through, not a tolerant reader that would hide the next rename; an artifact carrying no recognised array now returns NULL and says which keys it found, because *no rows* and *I cannot see the rows* are the two answers this bug confused. **Consequence, stated plainly: the clause now FAILS.**

#297 — THE CLOSED-DETECTOR LET PROSE OUTRANK AN open STATUS CELL. `roadmapRowIsClosed` honoured a cell saying `closed` and ignored a cell saying `open`, so a row whose narrative contains `CLOSED 2026-08-11` about a part that IS closed read as a closed ROW. Measured over all 217 register rows before changing it: **exactly five verdicts move, all closed → open, three of them asserting breakage — #218, #220, #224.** #220 is the expensive one: its `-fail` vs `-singleturn ... protect family` is **238 games of the 695** in `data/game-differential.json`, the largest single family in the whole-game differential, sitting in a row the gate could not see and `open_work.js` was not printing. The open-defect population went **5 → 8**.

#298 — FILED, NOT FIXED. `wholeGameClause` publishes its rate with no release check while `mechanicsClause`, `decisionImpact` and the new `orderProbeClause` all refuse an artifact cut against other bytes. Today's `690 of 1230` was measured on release `6875c8ace00e`; the tree moved twice while this was written. Making the headline clause refuse withholds the project's most-quoted engine figure and is a judgement with an owner.

ROADMAP #290 — THE ORDER PROBE IS A GATE NOW, AND ordering 229 IS NOT ONE POPULATION — 2026-08-18

node engine/quarantine.js --order-probe is the instrument #290's INSTRUMENT OWED specified: it FAILS while any order_probe row carries speed_tied: false AND same_priority: true . Exit 1 = the defect is present, exit 2 = the clause CANNOT ANSWER (no artifact, no probe, or a probe cut against other bytes), 0 = clean. Both non-zero codes are RED to register_reality.js deliberately: a VERIFIED BY exiting 0 on a missing artifact would close a live defect.

ON THE LAST ANSWERABLE ARTIFACT (release 6875c8ace00e , 1,230 games, arm A/middle): 26 move-vs-move pairs probed, 15 tied, 11 CARRYING THE CONJUNCTION across 11 distinct games. Speed gaps of 10, 20, 30, 32, 42, 56, 64, 76, 76, 147 and 213, every one at identical priority with every die pinned on both sides. Nine are Tailwind, three are Protect.

THE TWO NUMBERS ARE NOT THE SAME POPULATION AND THEY ARE ONE APART BY ARITHMETIC. They are two different groupings of one run: classes[.cls === 'ordering' is 228 games, while the gate's composition ordering is divergence_shape.shapeOf(cause) summed across ALL classes, which is 229. Cross-tabbed:

	games
shape ORDERING and class ordering	223
shape ORDERING inside class event missing from medicham2	6
class ordering that shapes as EMISSION	5

223 + 6 = 229 and 223 + 5 = 228. The counts differ by 1 while the sets differ by 11 games. Reading them as one number was about to attribute a whole class to the wrong cause.

WHAT IS ESTABLISHED IS 11 GAMES, NOT 228. The probe covers move-vs-move pairs only, so 11 is a FLOOR on the ordering class and never a count of it. Extrapolating 11/26 across 228 gives ~96 and is an extrapolation wearing a measurement's clothes; it is not published. The evidence is also one release stale — the differential ran at 04:09Z on 6875c8ace00e , release 488fd1bf3f7c (#294, the per-target accuracy roll) was cut at 05:05Z — so the gate exits 2 today and the count is WITHHELD rather than annotated.

THE FIVE UNMEASURED ROWS, WORKED — 2026-08-18

row	outcome	instrument	exit
#258	CONFIRMED — open, red instrument	node tests/test-no-silent-failure.js	1
#290	CONFIRMED — gate built, defect demonstrated on the last answerable artifact	node engine/quarantine.js --order-probe	2 (cannot answer: artifact one release stale)
#241	CONFIRMED, re-measured — part (3) is 3 causes / 3 games on the current artifact, not the stale upper bound of 8	INSTRUMENT OWED (a ratchet, which the re-run now makes possible)	—
#286	CONFIRMED by measurement — __weather has 1 read, 0 assignments, 265 frozen copies	INSTRUMENT OWED (the functional arm; the static scan was built, measured at 7 false positives of 8, and rejected)	—
#287	CONFIRMED by measurement — both directions reproduce, and the artifact still asserts closed #244 is open	INSTRUMENT OWED (compare against board.unmodelledBasePower , which exists at board.js:3239)	—

A RED TEST FOUND IN PASSING AND NOT FILED AS A STATUS: register_reality.js reports #273 as a PREMATURE CLOSE — the row is closed and node tests/probe_red_demo.js exits 1. tests/ is ENGINE's and was live while this ran, so it is reported here by name rather than touched.

oooooooooooo. THE MEDICHAM GATE'S FINISH LINE IS DECISION-EQUIVALENCE NOW, AND THE FILTER IS DERIVED FROM THE STORE — 2026-08-18

`engine/quarantine.js` (`--reach` , `--selftest`), `data/click-counts.json` , `data/sheet-usage.json` , and a contract for a `data/decision-impact.json` that does not exist yet.

THE BAR IS WILL'S. *"medicham needs to be good enough that when miltank uses it, it wouldnt affect any decisions"*, with the worked example *"so like if trick or treat isnt working, that doesnt matter because no one uses gorgeist."* Two clauses were asking for something else. The mechanics clause counted every diverging entity whether or not anybody plays it; the whole-game clause demanded literal zero over a unit nothing was allowed to filter. Neither is a weaker bar now — they are a different question, and the answer to it can actually be delivered.

FILTER 1 — REACH, at the threshold the project already had. A diverging mechanic below **25 observations in the store** is shelved: still staged, still played, still printed with its usage, and it stops counting. 25 is `tests/roster.js`'s `USAGE_SHELF_BELOW` and the coverage clause's shelf, not a new number; `coverageClause` used to carry its own literal 25 and now reads the same constant. The unit is whatever the store can observe for that kind — **clicks** for moves (`engine/click_counts.js` , 64,846 games, 1,259,717 clicks), **teams** for abilities and items (`engine/sheet_usage.js` , 26,232 teams from 13,116 open-sheet games). Both denominators print on every run because 25 teams and 25 clicks are not the same exposure and must not be read as if they were.

`tests/roster.js`'s header says **no honest store-derived ability usage exists. That was true on 2026-08-10 and** `data/sheet-usage.json` **was generated on 2026-08-11 —** *"The first honest store-derived ability usage this project has had."* The shelf reaches abilities and items because the instrument now exists, not because the rule was relaxed.

UNKNOWN IS NOT ZERO, and that half is what makes the filter honest. An entity the usage instrument STRUCTURALLY CANNOT SEE is not below the shelf: it counts, and it is named in its own column. `data/sheet-usage.json` declares that set itself — a team sheet reveals the PRE-MEGA ability, so an ability that lives only on a mega forme can never appear in it. Absence from an instrument that covers the whole class (every move click in every stored game) is an observed zero and is a different thing. If a usage artifact is missing outright, every row of that kind is UNKNOWN and nothing is shelved.

THE FILTER HAD TO NOT USE `tags.json` , **AND THIS ALMOST WENT WRONG.** The nine entities handed to me as unplayed carried `tags.json.uses` figures. Against `engine/click_counts.js` : bittermalice reads **0 there and 519 real clicks**, attract 2 against 30, fairylock 1 against 11. A reach filter run off `tags.json` would have shelved a 519-click move as unused — ROADMAP #70 landing on a live decision for the second time.

FILTER 2 — DECISION IMPACT, wired as a contract and refusing everything today. For a mechanic people do play, the question is whether fixing it moves the argmax, and that is a measurement, not something a clause may guess. `engine/argmax_paired.js` (ROADMAP #278) is the instrument. The clause reads `data/decision-impact.json` and clears a row only when **all** of: the artifact exists; `null_demonstrated: true` (its identical-dice control reported exactly 0 flips); its `engine_release` equals the tree's; and the row reports `flips: 0` over `paired >= 25` decision points and names the arm the fix landed in. The 95% upper bound from the row's own n (rule of three) prints beside every cleared row — **0 flips in 25 is an 11% bound, a floor and not a zero.** The same contract serves the whole-game clause through `cause: rows`, which is the only filter that clause gets: a game is not an entity, 690 of its divergences are `ordering` and `field` classes naming no mechanic, and thresholding those on usage would be a fabrication.

NOTHING IS CLEARED TODAY. There is no `data/decision-impact.json` , and the clause says so in the same sentence as the count — an exemption mechanism silent about being empty makes "o were cleared" and "we did not check" the same line.

THE drag: a different body DECLARATION IS WITHDRAWN. It was declared the same morning as "not a rule, a bench index", and the mechanism was right: `sim/battle.ts` `getRandomSwitchable` `samples` `side.pokemon` order, the authority takes one index under the pinned die and this engine takes another. The CONSEQUENCE was wrong for this bar. A different body arriving on the field is a different position, and every decision after it is taken against a board the authority does not have — the largest decision impact a divergence can have, not the smallest. It was never a claim that the authority is wrong, which is the only thing `DECLARED_DIVERGENCE` is for. **Supreme Overlord's** `fallenundefined` **stays:** an ungarded `onEnd` emitting a literal broken string on a `[silent]` line is a typo, and reproducing a typo is not correctness.

AND THE RECEIPT FOR KEEPING ALL THREE BARS NARROW. `medicham2-browser.js:17440` carried a declared divergence whose stated reason was COST — "*a far larger change than this wire is buying*". It hid the largest real defect in the engine: Showdown rolls spread accuracy per target, this engine rolled once per move, so Rock Slide (18,122 clicks) and Heat Wave (11,121) could never hit one body and miss the other. At 90 accuracy the exactly-one case is 18% of outcomes and did not exist here at all. **That one clears a reach filter without breaking stride.** Reach may never be the only filter, and no row may reach any of the three axes by a cost argument.

MEASURED, on `data/all-mechanics-fire.json` **generated 2026-08-18To4:06Z (release** `6875c8ace00e` **; the tree is** `29ddfe81594a` **, so the clause refuses these rows as stale and** `--reach` **prints them anyway — a listing is not a verdict):**

diverging mechanics, closet excluded	49 (moves 35, abilities 14, items o)
below the reach shelf	15, all moves, o–22 clicks
no usage figure	o — every diverging row has a store-derived number
cleared on decision impact	o — no run exists
counted: played and uncleared	34

THE WORKED EXAMPLE DOES NOT CLEAR THE SHELF AND THAT IS REPORTED, NOT PAPERED OVER. Trick-or-Treat is **70 real clicks**, not the 10 `tags.json` reports, so at a shelf of 25 it would hold the gate. Gourgeist-Super is **13 of 26,232 teams (0.05%)**, comfortably below. The move outlives the body that brings it. Raising the shelf to cover the example needs a number ≥ 71 and that is a decision with an owner; picking one here to fit one anecdote is not.

Shown RED on a deliberate break before being trusted, all three: unknown-reads-as-zero, `null_demonstrated` ignored, and the shelf boundary moved to `<=` (an ability sits at exactly 25 teams today, so that off-by-one would have excused a live row). `node engine/quarantine.js --selftest` is 56 cases, 0 failing.

oooooooooooooooo. ROADMAP #258 — THE SILENT-CATCH RATCHET: MEASURE'S OWN FILES ARE CLEAR, THE FLOOR FELL 216 → 201, AND 75 REMAIN IN OTHER DIVISIONS' FILES — 2026-08-17
`tests/test-no-silent-failure.js` , `data/silent-catch-baseline.json` .

THE GATE IS STILL RED AND I AM NOT FILING IT. The row is CONFIRMED against its own instrument by `engine/register_reality.js` , which is the honest state: the row asserts breakage and the gate exits 1. What changed is that the remaining population is now entirely outside this division.

MEASURE TOOK EVERYTHING IT OWNS. Fifteen silent catch blocks in `provenance.js` (8), `sprrt.js` (4), `mew.js` (1), `stamp.js` (1) and the ratchet itself (1) now speak, and `--update` locked the fixes in: **baseline 216 → 201, ten keys removed, nothing laundered.** `node tests/test-no-silent-failure.js --in`

<file>... lists what is left in any named file, so a division can work its own without grepping a repo-wide list — the second-scanner shape CLAUDE.md names.

ONE OF THEM WAS HIDING A LIVE DEFECT IN THE STALENESS CHECKER, WHICH IS THE ROW'S WHOLE ARGUMENT. `provenance.js` 's corpus classifier follows one level of `require('./x')` to decide whether an artifact's game count should be judged against the LADDER or the OPEN-SHEET ceiling. It resolved every such token against `engine/`, so for any generator living in `tests/` it resolved nothing at all — the read threw, the catch returned `''`, and the `games.(ots|bo3).json1` test ran against the empty string and returned the same `'ladder'` it returns for a generator that genuinely reads the ladder. **A capability absent, reporting success**, inside the tool whose job is catching that. The moment the catch was made to speak it named six files. Fixed by resolving against the requiring file's own directory and by reading the comment-stripped source (`engine_release.js` documents its API with `require('./x.js')`, which was being resolved as a dependency). **Verdicts before and after are byte-identical on today's tree** — the fix removes a hole rather than moving a number.

THE RATCHET COULD NOT SEE ITS OWN BLIND SPOT. A source file it failed to read was skipped in silence, so `files scanned 333` and a clean bill of health for 333 files were indistinguishable from a clean bill for 332 plus one nobody looked at. It now counts, names and FAILS on it. Shown red on a deliberate break before being trusted.

WHAT IS LEFT, RANKED, AND IT IS A HANDOVER RATHER THAN A LIST TO WORK THROUGH. Re-measure on the day — `node tests/test-no-silent-failure.js --all` — never this paragraph; the population moved twice while this was being written as SEARCH added files. As of 2026-08-17 it is **75 above the floor, 41 of them manufacturing a value**, and the manufacturing ones go first because each hands a made-up answer to whatever reads it. By file, manufacturing-first:

owner	file	mfg	the shape
ENGINE	<code>tests/roster.js</code>	5	<code>buildableSpecies</code> returns false when <code>mcKey</code> throws, so <i>unbuildable</i> and <i>not in this format</i> are one answer — the <code>buildMon("Scizor")</code> shape; and a precondition that THROWS is reported as <code>COULD-NOT-STAGE</code> , a claim about the fixture
ENGINE	<code>engine/mod_audit.js</code>	3	<code>return null</code> / <code>return []</code> — an empty override set reads as <i>Champions changed nothing</i>
ENGINE	<code>engine/tag_dex.js</code>	2	<code>return_ptShape</code> / <code>return null</code> on the tag derivation
ENGINE	<code>engine/game_differential.js</code>	2	<code>base = null</code> , <code>rt = null</code>
ENGINE	<code>engine/mega_census.js</code>	2	<code>return null</code> — the comment already says null must never read as <i>nothing is a mega</i>
ENGINE	<code>engine/champions_sim.js</code>	1	<code>ok = false</code>
ENGINE	<code>engine/medicham2-browser.js</code>	1	<code>rows = null</code>
ENGINE	<code>engine/merge_mega_into_engine.js</code>	1	<code>priors = null</code> — this is the builder whose 67 writes once all missed
ENGINE	<code>engine/conformance.js</code> , <code>engine/scenario_catalogue.js</code> , <code>engine/fixture_legality.js</code> , <code>tests/staged_board.js</code> , <code>tests/probe_volatile_leaves.js</code> , <code>tests/test-switch-carry.js</code> , <code>tests/test-forme-assert.js</code>	1 each	fixture and probe paths; <code>test-forme-assert.js:113</code> returns false from a <code>buildMon</code> throw
SEARCH	<code>engine/million_run.js</code>	2	<code>return false</code> , <code>k = null</code>

owner	file	mfg	the shape
SEARCH	engine/rollout_switch_census.js , engine/rollout_switch_probe.js , engine/click_counts.js	1 each	return null from a census reader
MEASURE-adjacent, unclaimed	engine/leaf_engine_contrast.js (2), engine/diff_swarm.js (2), engine/million_targets.js (2)	6	these sit on the leaf/target path and no division ledger names them; routed on request rather than taken unasked
WEB	tests/test-web-quarantine-loaders.js , tests/test-web-quarantine.js	1 each	w.ABRA_STATUS = null , age = '(mtime unreadable)'
OPS/infra	tests/test-lownode.js , tests/test-farm-ram-guard.js , tests/test-roadmap-register.js , tests/test-unmodelled-clicks.js	1 each	three of these four are false positives of the detector: sawFailure = true; code = e.status , threw = true and log = null all record the failure and are asserted on the next line. The detector spots a recorder by NAME (/fail\w*\s*(/), which noteReadFailure and sawFailure do not match. Not widened here: a regex guarding 801 catch blocks is not the place to buy four rows

oooooooooooooooo. ROADMAP #241 / #276 / #283 — THE THREE ROWS NOTHING DECIDES, NOW SAYING SO — 2026-08-17

engine/register_reality.js , data/register-reality.json .

Every one of the three is OPEN, asserts breakage, and had no instrument named — so the MEDICHAM gate was being held shut by three sentences. Each had a plausible-looking candidate gate that decides something else, and pointing a VERIFIED BY at any of them would have made a live defect read CONFIRMED-and-green, which is worse than the prose it replaced:

- **#241** — engine/game_differential.js MEASURES the missing -fail emissions and deliberately exits o on them ("a divergence is a FINDING", its own header), and tests/test-game-differential.js fails only when the INSTRUMENT is wrong. Nothing decides the row by exit code.
- **#276** — tests/test-seed-clock.js is GREEN and decides **#270**, the SEED's clock. #276 is the BOARD's own weather field, which no gate reads.
- **#283** — tests/test-rollout-fallen.js is GREEN and decides **#244/#245/#246**, the SEED's roster. #283 is board.movePower 's stub, which no gate reads.

So register_reality.js gained a second marker, INSTRUMENT OWED: , which records that NOTHING decides a row and names what would have to exist. It is counted and printed **separately from the verified figure**, because a declared debt is not a measurement.

oooooooooooooooo. ROADMAP #285 — THE DOCS-CURRENCY CENSUS WAS EXEMPTING FIGURES ON THE STRENGTH OF ITS OWN COMPLAINT ABOUT THEM — 2026-08-15

engine/docs_scan.js , tests/test-docs-current.js , data/docs-currency-baseline.json .

THE GATE IS RED AND I AM NOT FILING IT. docs/ABRA-whitepaper.md: 12 untraceable figures, was 11 . The whitepaper has not been edited since 2026-08-10; the regression came from underneath it. Three instrument defects were found chasing it, all three fixed and shown red; **the figure itself is not resolved** and #285 says exactly what the next actor needs.

1. **SELF-REFERENCE — a gate that could satisfy itself.** allArtifactNumbers unions every data/*.json , and the baseline IS one. It records the census's own findings as strings, and the artifact walk reads numbers out of strings, so **writing a figure down as an offender made that figure traceable on the next run.** Demonstrated by building it worse: recording the untraceable SET dropped the

whitepaper from 12 to 5 in one run. The loop had been live through `citation_mismatches` since that list was created. **And the same loop arrives through the REGISTER**, caught in the act within the hour: `data/open-work.json` is a copy of ROADMAP.md's prose, so any figure quoted in a register row becomes traceable. The row filed about this defect quoted the offending value while describing it, and the clause went green with the real regression still underneath. **Writing down "this number has no source" must never become that number's source.** It was masking **nine figures across three documents**; closing it takes `docs/MEDICHAM-SPRINT-NOTES.md` from 25 to 32 and restores `docs/ROLE-FAMILY.md`. Those are a DISCOVERY, not a regression, and the counts were **not** raised to match — that split is ROADMAP #188.

- 2. `--update` **WAS #258's LAUNDERING BUTTON, IN A SECOND FILE.** It turned a failing clause green *and* adopted every new offender into the floor, and the write gate was `if (F === 0 || UPDATE)` — so a red run could record its own regression. Monotone now, and proved: `--update` leaves the artifact byte-identical and exits 1.
- 3. **IT RATCHETS A COUNT IT CANNOT EXPLAIN.** The baseline said `11` and not *which* eleven, so the twelfth is not recoverable from the tree. A count also launders by SWAP — fix one, add one, count unchanged, gate green.

ONE HAND EDIT TO A RATCHET, AND THE ONLY REASON THAT IS EVER ALLOWED: the self-referential run had written the counts DOWNWARD (whitepaper 5), which is stricter than the truth and would have left the gate permanently red against a floor no correct run could meet. They are restored to the values the file held at `2026-08-15T03:11:13.624Z`, with the reason written into the artifact beside them.

oooooooooooo. ROADMAP #266 — THE FIXTURE-LEGALITY RATCHET COULD BE LAUNDERED, IT WAS UNDER-COUNTING, AND ONE ENTRY WAS A DIFFERENT DISEASE — 2026-08-14

`tests/test-fixture-legality.js`, `engine/fixture_legality.js`, `engine/champions_sim.js` (the shared oracle), `data/fixture-legality-baseline.json`, and the fixtures repaired in `tests/test-protocol-trace.js` and `tests/test-click-censoring.js`.

THE COUNT, AND WHY THE ROW'S 41 WAS NOT IT. Read from `data/fixture-legality-baseline.json` and `engine/fixture_legality.js`, never typed:

verdicts open at the start of this pass	32 (the row's 41, minus the ten repaired earlier the same day)
illegal DECLARATIONS behind those 32	34 — <code>pairOrigin.count</code>
verdicts now	22 — <code>count</code>
declarations now	23 — <code>pairCount</code>
repaired in this pass	10 verdicts / 11 declarations
UNREACHABLE — no legal carrier anywhere in the regulation	1

THE VALIDATOR STOPS AT THE FIRST PROBLEM PER POKÉMON, SO THE RULER WAS SHORT. A Snorlax declared with Swords Dance, Iron Defense, U-turn and Roar — four moves it cannot learn — produces the single sentence *"Snorlax can't learn Swords Dance."* Two declarations were hidden that way, and the cost is worse than a low count: repairing the first illegal move in a set makes the second appear as a **NEW** verdict, so a repair reads as a regression. Every declared move now also goes through `TeamValidator#checkCanLearn` and the baseline ratchets those pairs beside the sentences.

THE CLASSES ARE NOT ONE PROBLEM. Of the 34: **0** named a species this format does not contain, **1** named an entity with no legal carrier at all, and **33** were a legal body holding something a legal body somewhere else can hold — re-aimable, which is what the repairs did. The one that matters most is the smallest: **Spore has ZERO legal carriers in Champions Reg M-B**, derived twice (whole-roster

checkCanLearn , and an independent raw learnset walk through the prevo and base-forme chains; both answer 0), and tests/staged_status_counters.js stages its two sleep-counter scenarios on it. That is the shape that manufactured four phantom engine defects in one session — a probe fails, it reads as an ENGINE DEFECT, and the engine was correct not to model a thing the format cannot produce. **The repair is not a rename:** every sleep move that does have a carrier here is sub-100 accuracy (Hypnosis 60, Sleep Powder 75, Sing 55) and both scenarios run on the top-tie-first arm, where every sub-100 move MISSES by construction. Milotic can legally learn Hypnosis, so the mechanic is reachable on the bottom arm and nowhere else. **Held, not repaired** — that file belongs to the owner of the differential and moving a scenario between pin arms changes what it measures.

THE LAUNDERING BUTTON, THE SAME SHAPE AS #258's --update . "Repair a verdict" and "excuse a new one" were the same edit: a line out, or a line in. The label PRE-EXISTING is a **claim about time** and nothing checked it. The 41 sentences that existed when the check was added are now written down as a closed origin set, and a PRE-EXISTING entry that is not one of them fails by name. A genuinely new illegal set can still be declared — as DELIBERATE, with a reason, which is a visible and arguable act — and an UNREACHABLE entity may never be called DELIBERATE, because there is no body to stage it on. **Shown red on the full laundering attempt:** the offender planted in a fixture AND added to both verdicts and pairs AND marked unreachable , and the gate still failed by name and exited 1.

A POPULATION HOLE, MEASURED AND CLOSED. The sweep found a declared set two ways — a helper call, or an object literal keyed species: — and a set written as a **positional row** ['species', [moves], ability, item] was neither. Twelve such rows drive 200 games in tests/test-protocol-trace.js and **five of the twelve were illegal, including a Toxapex holding an item that does not exist in Gen 9**, while the gate reported the tree clean. Repo-wide that shape occurs in exactly one file (measured), so the rows were repaired first and the matcher turned on after: the population grew 627 → 639 declarations and the verdict count did not move at all.

WHAT IS STILL OPEN: 22 verdicts / 23 declarations, all in files this division does not own. tests/staged_board.js (16) and tests/staged_status_counters.js (7) are the differential's staged scenarios, filed to their owner; tests/test-switch-carry.js (2) needs a filler POLICY rather than a rename; and tests/test-click-censoring.js (1) is Farigiraf/Encore, where the repair is a re-cast rather than a swap — **derived: no legal body in this regulation learns Encore, Roar and Follow Me together**, and that fixture's one p1a body plays all three roles.

THE ORACLE IS SHARED, BECAUSE FACTS ARE GLOBAL. champions_sim.legalRoster / abilityCarriers / moveCarriers / canLearn / unreachable — one derived answer to "can anything in this format carry this", in the module that already owns checkLegal . tests/probe_red_demo.js still carries its own abilityUnreachable with its own dex walk; that file is ENGINE's and adopting the shared one is theirs to do. Two implementations of "is this legal here" will diverge silently, because both keep working.

oooooooooooo. ROADMAP #258 — THE SILENT-CATCH RATCHET HAD A LAUNDERING BUTTON, WHICH IS WHY IT SAT RED FOR A WEEK — 2026-08-14

tests/test-no-silent-failure.js (reworked), engine/provenance.js , engine/quarantine.js , engine/status.js , engine/open_work.js , data/silent-catch-baseline.json (lowered, never raised).

THE NUMBERS, MEASURED RATHER THAN QUOTED. The register row said 78 new and 101 manufacturing; an earlier note said 81 and 101. Both were stale and neither was today's:

catch blocks	787
silent — say nothing at all	296 (38%)
...of which MANUFACTURE a value	97
...of which only skip or continue	199

baselined floor	216 (was 220)
NEW since the baseline	80 — 41 manufacturing, 39 skipping

THE TWO POPULATIONS ARE NOT ONE PROBLEM AND THE REPORT NOW SAYS SO. A `catch { continue; }` over a torn JSONL line is correct silence. A catch that RETURNS A PLAUSIBLE VALUE is this project's named failure mode in its purest form — a capability that could not prove it ran, reporting success. The new list is grouped by file, manufacturing first, instead of 25 flat names and "... and 57 more".

THE INSTRUMENT'S OWN DEFECT WAS THE REASON THE DEBT COMPOUNDED. `--update` rewrote the floor to the CURRENT silent set, so locking in a fix and laundering every new offender were THE SAME COMMAND. The row says as much — *"NOT RE-BASELINED, because `--update` would launder a week of it into the floor"* — and the consequence was a gate red for a week with three separate agent reports correctly observing they had not caused it. The sum of those correct observations is a check nobody acts on, which is "known failure" wearing the word *pre-existing*.

`--update` is now MONOTONE: `min(baseline, current)` per key. It removes what was fixed and cannot add what is new. The only door into the floor is `--accept <file> "reason"`, one file at a time, with the reason written into the artifact beside the keys, refused outright if either argument is missing. **Demonstrated:** two `--update` passes took the floor 220 → 216 and the gate stayed RED at 80 NEW, where the old behaviour would have written 303 and reported zero.

FIXED HERE — seven, all MANUFACTURE, all in files this division owns:

file	what the silence cost
<code>provenance.js</code> ×3	<code>declaredWriter</code> , <code>keysOf</code> and <code>mtime</code> each turned <i>corrupt / permission-denied / mid-write</i> into the same answer as <i>absent</i> , inside the tool whose whole job is saying whether an artifact can be believed. They call <code>logUnreadable</code> now and the count prints every run, including at zero
<code>quarantine.js</code> ×2	an unreadable artifact was reported to the gate as <code>NO ARTIFACT — run the instrument</code> , which is a wrong diagnosis handed to the clause that decides whether MEDICHAM is correct; and the open-defect clause failed loudly without ever saying WHY it could not read the register
<code>status.js</code> ×1	<code>metaPathFor</code> throwing silently disabled the sidecar clause, so a figure whose stamp is rotten would print as clean — the one thing that block exists to prevent
<code>open_work.js</code> ×1	an unreadable instrument dropped out of the measured half of the report that exists to print what is open

AND THE FIRST DRAFT OF THE `quarantine.js` FIX WAS ITSELF THE LESSON. It reported on `.js` bundles that are handed to a JSON reader on purpose, and printed six lines of noise on a clean run. Narrowed to `.json` paths: a ratchet that flags code for doing what it asked is how a ratchet gets ignored, and that is the fourth correction of exactly this shape in this repository.

A second detector blind spot is now STATED in the file rather than worked around: `blank()` erases template literals, so a reason travelling in `${e.message}` is invisible while `' (' + e.message + ' )'` is seen. Fixing that would mean changing the brace scanner every other check in the file depends on; the call sites were changed instead, and the limit is written down. It costs a false POSITIVE, never a false negative.

STILL OPEN, AND NOT THIS DIVISION'S TO EDIT. The remaining 80 sit in `tests/roster.js` (5 manufacturing), `engine/mod_audit.js` (3), `engine/tag_dex.js`, `engine/game_differential.js`, `engine/million_run.js`, `engine/diff_swarm.js`, `engine/mega_census.js`, `engine/leaf_engine_contrast.js`, `engine/board.js`, `engine/medicham2-browser.js` and others owned by ENGINE and SEARCH. **The population moved by eight blocks during one hour of measuring it**, because both divisions were writing at the time — which is worth knowing before anyone quotes a single figure off this gate as though it were a constant. `node tests/test-no-silent-failure.js` prints the current per-file list; do not type one.

00000000000. ROADMAP #257 — A VERIFICATION RUN CAN NO LONGER REPUBLISH A SMALLER MEASUREMENT — 2026-08-14

engine/publish_guard.js (new), tests/test-publish-guard.js (new), data/published-samples.json (new), tests/test-engine-diff.js (four write sites).

FIRST, THE FIGURE, BECAUSE THAT IS THE HALF THAT MATTERS: IT IS TRUE.

data/engine-diff.json carries requested 6000, compared 6000, agreed 6000, disagreed 0, generated 2026-08-14T06:36Z. The white paper, the deck, the technical docs and the site cite 6,000 and are backed by a real run of that size; tests/test-docs-current.js is green on all 21 clauses. Nothing had to be restated and nothing is withheld. **One caveat that is not this row's to close:** engine/medicham2-browser.js moved at 03:19 local, 43 minutes AFTER that run, so 0-of-6000 is a photograph of the 02:36 build. That is provenance's question and status.js prints it.

SECOND, THE DEFECT, WHICH IS THE --write AND NOT THE 150. The stated mitigation on the row — "a verification run passes --out or omits --write" — was measured to be a no-op: the file has neither flag and wrote data/engine-diff.json unconditionally on every run. A rule you can only follow by remembering it is worth what the ban list of four was worth. So it is a mechanism:

- **A publish below the published sample is REFUSED.** The run still completes and its output goes to data/verification/<name>.n<sample>.json, so a verification run loses nothing except the ability to republish. process.exitCode = 3, because a run that did not publish what its own output describes must not read as a pass — that is exactly how the original quick check looked successful.
- --republish-smaller **is the deliberate door**, and it stamps sample_shrunk {from, to, at} INTO the artifact, so the next reader sees the decision in the file rather than in a closed terminal.
- **A high-water mark per artifact** lives in data/published-samples.json, so the refusal survives a hand edit, a merge, or a writer that never called the guard at all.
- **The three conformance sections now amend the path publish() actually wrote.** They used to re-read data/engine-diff.json BY NAME — so a --plant run, which goes out of its way to write beside the artifact instead of over it, stamped three sections onto the published one anyway.

SHOWN RED ON THREE DELIBERATE BREAKS BEFORE BEING TRUSTED: the refusal disabled (clause A fails on 5 assertions), the artifact shrunk behind the guard (clause C names data/engine-diff.json and prints ON DISK compared=150, RECORDED 6000), and a bare write restored in the caller (clause B names the file and line). And on the real defect: tests/test-engine-diff.js --n 150 now exits 3, leaves the 6,000-comparison artifact untouched, and writes data/verification/engine-diff.n150.json.

The guard is deliberately general — publish({file, artifact, sampleKey, argv}) and amend(path, sampleKey, fn) — and only tests/test-engine-diff.js uses it today. **27 other top-level artifacts in data/ carry a numeric sample field and are still unprotected** (counted, not listed — the shape is compared / requested / n / games / pairs / decisions); node engine/publish_guard.js record <artifact> <key> is how one joins, and each belongs to the division that writes it.

00000000000. ROADMAP #242, FIRST HALF — THE RESIDUAL ORDER TABLE WAS A SUBSET PRESENTED AS A POPULATION: 42 ROWS AGAINST 90 — 2026-08-14

engine/residual_order.js (rewritten), tests/test-residual-order-population.js (new), data/residual-order.json (regenerated, 42 → 90 rows). engine/medicham2-browser.js, tests/test-mechanics.js, engine/board.js, engine/rollout_leaf.js were **not touched** — the placement half of #242 is ENGINE's and reads this table rather than re-deriving it.

THE NUMBER. The table published 42 walk participants. The authority walks 90. Measured against Battle#fieldEvent in gen9championsvgc2026regmb, filtered to the regulation:

walk participants	90 (was 42)
that expire from inside the walk	58
...of which own no residual handler at all	48
...of which announce a protocol line	28
...silent but still order-bearing	30
that declare no order and sort at the 4294967296 sentinel	29

WHY IT WAS 42. `fieldEvent` sets `getKey = 'duration'` for the Residual event, and every collector admits an effect on `callback !== undefined || state[getKey]`. This file enumerated on the first clause only, so a `if (!hook) return;` guard threw away every effect that is in the walk purely because it has a live duration — Tailwind, Trick Room, the screens, Safeguard, Protect, the guards, the two-turn moves.

AND THE DIAGNOSIS IN THE ROADMAP ROW WAS WRONG IN THE WAY THAT MATTERS, INCLUDING MINE. The row says these effects have no order of their own. **They do, and it was in the format the whole time.** `tailwind.condition` declares `onSideResidualOrder: 26, onSideResidualSubOrder: 5` and simply declares no `onSideResidual`; `resolvePriority` reads `effect[callbackName + 'Order']` whether or not the matching callback exists. The order was never missing — the guard threw the effect away before anything read it. That is a better outcome than the row assumed: 28 of the 58 expiries have an exact published position, not a heuristic one.

PROVED BY RUNNING IT, NOT BY READING IT. One staged walk on the official simulator, four bodies, every family live at once:

-weather none	sandstorm expiry	order 1
-heal ... [from] Grassy Terrain	terrain heal	order 5
-heal ... [from] item: Leftovers	Leftovers	order 5
-damage ... [from] psn	poison	order 9
-end p1a: ... move: Taunt	Taunt expiry	order 15
-sideend p1: A move: Light Screen	screen expiry	order 26 sub 2
-sideend p1: A move: Tailwind	Tailwind expiry	order 26 sub 5
-fieldend move: Trick Room	Trick Room expiry	order 27 sub 1

Twelve seeds, identical sequence. It is not a coin flip: Light Screen precedes Tailwind because their declared subOrders differ, not because a tie shuffled.

FOUR PUBLISHED VALUES WERE WRONG, AND ALL FOUR WERE THE SAME MISTAKE — A RULE COPIED INSTEAD OF CALLED. The old file owned a literal transcription of `resolvePriority`'s subOrder defaults:

- `psn`, `tox`, `brn` were published at subOrder **2**. Their effectType is `Status`, which is not in the authority's map at all — the real value is **0**.
- `wish` was published at subOrder **2**. It is a slot condition, so the authority gives it **3**. The old file *had* a clause for that and its own de-duplication defeated it: the row was already emitted by an earlier loop, so the slot-condition patch never reached it.
- the map's flat `Condition: 2` is not what the authority does at all. `resolvePriority` refines by **where the effect is attached** — `state.target instanceof Side` → 4, `isSlotCondition` → 3, `instanceof Field` → 5 — which no effectType can express.

None of the four changed a grouping medicham2 currently consumes, so nothing downstream moved. They were wrong in the published artifact regardless, and they are the argument for the fix: **this file no longer owns a copy of the sort rule. It calls** `Battle.prototype.resolvePriority` **and**

Battle.prototype.comparePriority **on the real objects.**

THE GATE FOUND AN ERROR IN MY OWN REPLACEMENT WITHIN A MINUTE OF EXISTING, which is the strongest thing I can say for it. I shaped the `state` argument by hand — `{ target: Object.create(Field.prototype) }` — and published Fairy Lock at subOrder 5. Live, it resolves to 2, because `field.addPseudoWeather` **writes no** `target` **at all**: the field-condition refinement branch is unreachable in this Showdown build. The generator now stages a throwaway battle and keeps the real state object each attach method produces. A derivation that models the authority instead of asking it is wrong exactly where nobody thought to model.

THE GATE. `tests/test-residual-order-population.js`, auto-discovered by `tests/run-all.js`, 14 checks. It stages one battle, calls the authority's OWN `findSideEventHandlers` / `findFieldEventHandlers` / `findPokemonEventHandlers` with the same `getKey: 'duration'` that `fieldEvent` passes, and asserts every handler that comes back has a row whose `(order, subOrder)` equals what `resolvePriority` just produced live. **Shown RED first**: reverting the enumeration to handler-only makes it name 41 live participants with no row, plus 31 duration-carrying conditions from an independent lower bound.

AND THE FIXTURE HAD TO BE MADE INDEPENDENT, BECAUSE THE FIRST DRAFT STAGED FROM THE TABLE IT WAS AUDITING. With `ROWS.filter(site === 'side')` as the stage list, a table that had dropped every side condition also staged no side condition, and the membership clause reported one missing row instead of forty-one. The ids now come from the moves. This is the same defect as the row itself, one level up, and it is only visible because the break was run.

WHY `tests/test-residual-order-observed.js` **COULD NOT HAVE CAUGHT ANY OF THIS.** It stages a turn and checks the chips come out in the table's order — but it only ever looks at effects the table already contains, so a table missing a whole class of participants agrees with it perfectly. It is green today and was green throughout. **A check that watches a copy of the population cannot report that the population is short**, which is the sixth instance of that shape this week.

THREE MORE FINDINGS THAT ARE NOT ABOUT ORDER.

1. **Lucky Chant and Mist are** `isNonstandard: 'Past'`. The roadmap row's hand-derived list of 19 announcers names both. They are not in this regulation and cannot appear in a Champions turn. The list also omits the four weathers, `partiallytrapped`, Syrup Bomb, Encore, Perish Song and Uproar. The measured count is **28**, and Heal Block belongs on it — the *move* is `Past`, but the volatile is reachable in the regulation through **Psychic Noise**.
2. **Grassy Terrain is in the walk TWICE, eleven orders apart.** Its heal arrives through the per-active field collection at `onResidualOrder` **5** carrying the healed body's speed; its expiry arrives through the field collection at `onFieldResidualOrder` **27** with no speed at all. One row cannot say both, which is why a row is now an `(effect, attachment site)` pair rather than an effect. The 42-row table said only the first.
3. **A duplicate published key would have silently unplaced a medicham2 step**, and the first run of the rewrite produced one: both Grassy Terrain participants wanted `field:grassyterrain`, and medicham2 builds a `Map` from these rows, so the last would have won and its `terrain` step would have moved from order 5 to order 27 **while still reporting itself placed**. The generator now refuses to publish a duplicate and the gate asserts it. All 42 legacy keys still resolve; only `condition:grassyterrain` is gone, and it was a duplicate of `field:grassyterrain` that nothing read.

WHAT ENGINE NEEDS FROM THIS TABLE FOR THE SECOND HALF — a read, not a re-derivation:

- `rows[].site` is the load-bearing field. It decides the callback name, hence which `*Order` is read, hence where the effect lands. `ns:id` is the stable published key.

- `route` — `handler`, `duration`, or `handler+duration`. A `handler+duration` **row decrements first and, on the turn it reaches zero, runs `end` and SKIPS its own residual callback** (`continue` in `battle.ts`). Ten rows: the four weathers, Encore, Perish Song, Uproar, Syrup Bomb, `partiallytrapped`, `lockedmove` .
- `announces` + `announceLine` — 28 rows emit protocol at their sorted position. **30 more expire silently and are still order-bearing**, because the `duration--` spends a position in the same walk. An engine that skips the silent 30 will place the loud 28 correctly and still diverge.
- `order: null` means no declared order: `resolvePriority` stores `false` and `comparePriority` substitutes **4294967296**, so those 29 sort LAST — after Trick Room — by holder speed then subOrder.
- `speedFrom` — a side or field condition's expiry has **no speed** (its holder is a Side/Field, which has no `getStat`), so it sorts as 0 and lands behind every body-held effect at the same order.
- The blank-player-name defect on `| -sideend|p2: |move: Tailwind` is untouched here. The authority emits `| -sideend|p2: B|move: Tailwind ;` that is a protocol-rendering bug in `medicham2`, not an ordering one.

NOT DONE, AND IT IS THE HALF THAT MOVES A GAME. `medicham2` still places these expiries wrongly — 21 games of `data/game-differential.json` regroup to `ordering :: | -damage|<slot>|H/H| [from]sandstorm <> | -sideend|<side>|tailwind` , the largest single shape on the board. That is ENGINE's, it is live in that file, and the table it needs now exists.

oooooooooooo. ROADMAP #266 — TEN OF THE FORTY-ONE ARE REPAIRED, AND ONE OF THEM WAS NOT A LEGALITY DEFECT AT ALL — 2026-08-14

`tests/test-volatile-duration.js` , `tests/test-perish-song.js` , `tests/test-protocol-trace.js` , `tests/test-game-diff.js` , `tests/test-speed-tie.js` , `tests/test-nature-differential.js` , `tests/probe_volatile_leaves.js` , `data/fixture-legality-baseline.json` . `engine/medicham2-browser.js` , `tests/test-mechanics.js` , `engine/board.js` , `engine/rollout_leaf.js` , `engine/game_differential.js` , `tests/staged_board.js` and `tests/staged_status_counters.js` were NOT touched.

THE NUMBER: 41 distinct verdicts → 32, 55 rejected sets → 45, and 1 stray literal → 0.

`tests/test-fixture-legality.js` is ALL GREEN on the shrunk baseline, including its stale-allowance clause — every removal came off because the set became legal, and every one is named with its cost in the `repaired` block of `data/fixture-legality-baseline.json` . Population unchanged at 627 declarations and 227 distinct sets, so nothing came off because the scanner stopped looking.

verdict removed	what replaced it	what it cost
literal "dampro" (the stray)	damprock	nothing — the body is benched all scenario
Toxapex's item Black Sludge does not exist in Gen 9	Leftovers	nothing — benched body, and EXISTENCE is never deliberate
Clefable can't learn Perish Song	Azumarill sings, in BOTH files that share the singer	Clefable's Unaware leaves the field; nothing in either scenario boosts, so it could never have fired
Weavile can't learn Encore	Alakazam is the user	Pressure. Named below.
Snorlax can't learn Pound	Round / Terrain Pulse (duration), Round (perish, never clicked)	both Special and 100-accurate, so they are paid into SpD rather than Def
Snorlax can't learn Agility	Recycle — this repo's own derived no-op	strictly quieter: Agility was moving the foe's Speed two stages a turn
Corviknight can't learn Pound	Body Press	nothing — the anchor only ever clicks Protect
Archaludon can't learn Body Press	Iron Defense	nothing — the slot Protects for six turns
Whimsicott can't learn Stealth Rock	the hazard and the screen swapped bodies	the screen guards the special side now, and the rocks lose Prankster's +1

verdict removed	what replaced it	what it cost
Ninetales can't learn Dazzling Gleam	Psyshock	the click is single-target instead of spread

THE WORST ENTRY ON THE LIST PRODUCED NO VALIDATOR VERDICT AT ALL. `tests/test-protocol-trace.js:131` gave Politoed the item `'dampro' . dex.items.get('dampro')` returns a row that does not exist whose `.name` is the empty string — **and an empty item is a LEGAL item** — so `checkLegal` was asked about a Politoed holding nothing and answered honestly. The fixture built a Politoed holding nothing and every check went green. That is *"a capability was absent and everything reported success"*, inside a test, and it is exactly why `fixture_legality.js` reports stray literals in their own section rather than folding them into the illegal-set list: nobody is being accused of an illegal pairing, the string simply means nothing. The intended item is spelled correctly sixteen lines further down the same file.

A REPLACEMENT BODY IS DERIVED FROM THE FORMAT, NEVER RECALLED, AND THE DERIVATION IS THE ANSWER TO "WHY THAT ONE". Exactly **seven** legal bodies in this regulation learn Taunt, Encore and Disable together (Alakazam, Salazzle, Gardevoir, Gallade, Banette, Chimecho, Sableye) and exactly **six** learn Perish Song and Protect together (Gengar, Altaria, Absol, Politoed, Primarina, Azumarill). Alakazam is the fastest of the seven at 120 base Speed — the nearest thing to Weavile's 125 — and equally frail, so "the user must survive three turns of being hit" stays a real constraint instead of being quietly removed. Azumarill is the slowest and bulkiest of the six and the only one whose ability neither sets weather nor copies anything: Politoed's Drizzle would put rain on every board, Altaria's Cloud Nine would take it off, Absol brings Pressure into a PP-compared board.

AND THE COST OF THAT ONE IS A CAPABILITY THIS FILE NO LONGER STAGES, SAID OUT LOUD RATHER THAN ABSORBED: NONE OF THE SEVEN CAN HAVE PRESSURE. Weavile's Pressure made Snorlax pay double PP for every click into the user's slot, and PP is a compared board leaf. `tests/test-volatile-duration.js` still stages PP — both sides spend it every turn — but it no longer stages the DOUBLED rate. That rate is staged by `tests/staged_board.js`, whose Corviknight/Pressure bodies are actually targeted, so the coverage is not lost from the repo; it is lost from this file. Same shape as the Weather Ball case (#265): no legal body can carry the combination, so the honest move is to name what went with it.

THE FIXTURE AUDIT IN BOTH RED GATES IS NOW GREEN, AND `tests/test-perish-song.js` IS GREEN OUTRIGHT. Its two scenarios — four bodies dying at once, and the pivot that escapes the count — pass against the official engine with an Azumarill singing.

`tests/test-volatile-duration.js` **PASSES ITS AUDIT AND STILL PARTS FROM SHOWDOWN ON HP, AND THAT IS NOT THIS REPAIR. IT WAS MEASURED RATHER THAN ASSUMED.** The audit calls `process.exit(1)` before a single game runs, so nobody had seen this file's engine comparison since the learnset clause landed on 2026-08-12. The **pre-repair fixture was replayed body-for-body** — Weavile / Snorlax / Clefable, Pound and Shadow Punch, the same three-turn scripts — and it parts in all four scenarios by the same magnitude and in the same direction (`weavile sd=72 us=76`, `snorlax sd=199 us=204`). The legality repair moved the numbers; it did not create the disagreement.

WHAT THAT DISAGREEMENT IS, IS NOT MINE TO SAY TONIGHT, AND SAYING SO IS THE POINT. Isolated one click at a time it reads like a DAMAGE ROLL: `medicham2` answers the same number every time while Showdown's varies with the scenario id, which is the signature of one engine on a pinned corner and the other on a die. `tests/test-speed-tie.js` shows the same shape on bodies this pass never touched — its `no-tie-CONTROL` case parts on `-damage` 75 vs 69 with Weavile and Volcarona. But `engine/medicham2-browser.js` was written at **02:26 tonight, while these runs were going**, and `engine/game_differential.js` carries in-flight middle-arm RNG work belonging to another agent. **A**

measurement is a photograph and nothing in frame may move. So the absolute verdict is WITHHELD rather than annotated; what is published is the paired claim, which is valid because both arms were measured minutes apart against the same bytes: **pre-repair and post-repair part identically.**

WHAT IS DELIBERATELY NOT REPAIRED, AND WHY EACH ONE IS A SEPARATE PASS. 32 verdicts remain, and the reason is written per entry in the baseline rather than summarised here. Three blocks: `tests/staged_board.js` (15) and `tests/staged_status_counters.js` (7) are the staged-board library, which ROADMAP #266 routes to `@engine` and whose illegal clicks are load-bearing FILLER BOOSTS — a body clicking Iron Defense at +2 Def while a damage arm reads it is not a rename, it is a different board; `tests/test-protocol-trace.js` keeps 10, every one of which moves an event-coverage board (Cleable's Trick Room, Slowking's Perish Song and Haze, Incineroar's Knock Off and Politoed's Scald are all CLICKED to reach a protocol line, and the file's own verdict cannot be re-measured while ENGINE is live in the simulator); and `tests/test-switch-carry.js` keeps 2 because repairing only the two the static sweep can SEE would leave the identical defect in the Milotic and Vaporeon fillers it cannot see — those are declared through identifiers, so they are outside the population, and the fix is a filler policy rather than two edits.

ONE SITE CAME OFF WITHOUT A VERDICT COMING OFF, AND IT IS COUNTED AS A SITE AND NOT AS A REPAIR. `tests/test-nature-differential.js:152` gave Incineroar Knock Off, which it cannot learn here; it is now Darkest Lariat, its legal Dark physical click. *"Incineroar can't learn Knock Off"* is still produced by five other sites in two files, so the verdict stays on the baseline — a ratchet keyed on the sentence only shrinks when the last site producing it is gone.

oooooooo. ROADMAP #265 / #266 — THE FIXTURES WERE DECLARING TEAMS THE GAME WOULD REFUSE, AND THE RULER FOR THAT DID NOT EXIST — 2026-08-13

`engine/fixture_legality.js` (new), `tests/test-fixture-legality.js` (new), `data/fixture-legality-baseline.json` (new), `engine/feature_fixture.js` (six repairs + a body digest), `tests/test-feature-semantics.js` (three new clauses). `engine/medicham2-browser.js`, `tests/test-mechanics.js`, `engine/game_differential.js` and its artifacts were not touched.

THE HEADLINE IS A POPULATION, NOT A VERDICT. 334 .js files, 627 set declarations across 26 files, 227 distinct sets, 55 REJECTED by `TeamValidator`, producing 41 distinct verdicts. Six of those were in MEASURE's own `engine/feature_fixture.js` and are repaired here; the other 41-minus-those are baselined, named, and owed as #266.

THE VERDICT IS THE AUTHORITY'S AND NOTHING HERE RE-IMPLEMENTS IT. Will, mid-task: *"i thought we were using showdowns team validator as the ultimate legality test."* The brief this pass started from said to check the species with `isNonstandard`, then the ability against the species' list, then the item, then each move with `checkCanLearn` — **four hand-written rules standing in for a table Showdown maintains**, which is FACTS ARE GLOBAL broken inside the file written to enforce a rule. A piecemeal check only finds the rules somebody thought to write. `validateTeam` carries the 66-point SP budget, the 32-per-stat cap, the level, the item clause and every ban the Champions mod adds next. So `fixture_legality.js` calls `champions_sim.checkLegal` — one shared validator instance, five padding slots that are themselves proved clean — and prints the validator's own sentences rather than a summary of them. **It is also less code.**

THE SCANNER IS DERIVED AND IT WAS WRONG FOUR TIMES BEFORE IT WAS RIGHT, WHICH IS WHY EACH RULE IS WRITTEN DOWN WHERE IT LIVES. A hand-listed set of fixture files is the shape this repository has been wrong about most often, so the sweep walks the tree and finds sets by shape:

the rule	what it was reporting before it existed
a literal is an entity only if it normalises to that entity's OWN id	<code>dex.species.get('p2')</code> answers Porygon2 ; <code>p2</code> in this tree is a SIDE
moves come from ARRAY literals only	<code>hit('vaporeon','icebeam')</code> clicks the move AT that body — three false accusations
helpers are taken at TOP LEVEL only	<code>test-mechanics.js</code> redeclares <code>run / hit / at</code> in dozens of probe callbacks; a file-global name set paired an ability stamped on the ATTACKER with the DEFENDER's name
a set is declared by an object literal or by MUTATION	requiring a <code>species</code> key hid 64 real sets in <code>test-protocol-trace.js</code>

507 CONSTRUCTION SITES CARRY NO LITERAL SET AND ARE OUTSIDE THE POPULATION, WHICH IS CORRECT AND IS PRINTED RATHER THAN IMPLIED. A fixture that builds its bodies from the format — `tests/roster.js` walks the regulation, `engine/all_mechanics_fire.js` reads the tag dex — cannot type a name that does not exist. A fixture that builds through `buildMon(key)` and assigns the click LATER, which is most of `tests/test-mechanics.js`, has no set to validate at the point the body is made; the sweep says so instead of guessing, and that is the honest limit of a static sweep.

THE GATE IS A RATCHET, SHAPED LIKE `data/fixture-learnset-baseline.json` **AND A SUPERSET OF IT.** Keyed on the **validator's own sentence**, not on `file:line`, so the same illegal declaration moved to a new scenario is the same defect. Six clauses, and two of them exist because a green gate is worth nothing on its own: a **population floor** (a scanner that finds nothing passes everything below it — it must find ≥ 150 distinct sets, against 227 today), and a **self-check that the ratchet still discriminates**, run on every invocation rather than demonstrated once in prose. The baseline **may only shrink**: a baselined verdict that stops being produced FAILS, because a repair that leaves its allowance behind is a hiding place for the next illegal set. Every entry carries a `kind` (`DELIBERATE` or `PRE-EXISTING`) and a written reason, and an entry with neither fails.

SHOWN RED THREE WAYS BEFORE BEING TRUSTED, and the first is the real historical instance: re-planting Rocky Helmet on Venusaur reports `[EXISTENCE] Venusaur's item Rocky Helmet does not exist in Gen 9`. naming `engine/feature_fixture.js:87`; two invented baseline entries fire the stale-allowance clause by name; one carrying a bad `kind` fires the reason clause. The tree was restored byte-for-byte from a copy taken before the plant.

THE SIX REPAIRS ARE ALL IN MEASURE'S OWN FIXTURE AND COVERAGE WAS MEASURED ACROSS EVERY ONE. Rocky Helmet → Miracle Seed (4 sites), Electric Seed → Mental Herb, Throat Spray → White Herb, Grimmsnarl's Thunder Wave → Taunt, Hippowdon's Weather Ball → Rock Slide, and a new board.

	before	after
boards	10	12
candidates / pairs	352 / 1,407	384 / 1,534
columns that never fire	0	0
columns that LOST a firing	—	<code>statusBites</code> 5 → 4, and nothing else

TWO OF THOSE REPAIRS WERE NOT RENAMES AND THE COST IS RECORDED RATHER THAN ROUNDED AWAY. (a) **No body in this regulation can carry Weather Ball in sand.** Derived over the format, the move has **80 legal carriers and not one is a Rock or Ground body or sets sand**, so `sand-is-up` had been claiming the ROCK type-flip path through a move Hippowdon cannot learn. Repairing it IN PLACE was tried first and measured: swapping in Heliolisk costs that board its Earthquake and its Yawn and thinned `allyHit` **4 → 3**, `abilityBlock` **3 → 2**, `deadStatus` **2 → 1**, `statusBites` **5 → 2** — the first two being exactly the pair this file was built to protect — while **every column stayed non-zero, so**

the coverage check would have waved it through. The flip moved to its own board, `sand-weather-ball`, which is what this file's own record of two bad edits says to do. *(b) Grimmsnarl has no legal primary-status move at all*, only secondary-chance ones, so Taunt cannot carry `deadStatus`; the click moved to Torkoal's Will-O-Wisp into an already-paralysed Incineroar. That restores `deadStatus` and leaves `statusBites` one short, because the only body that could repay the last one is the Clefable whose fourth move is the single 4x hit anywhere in the fixture. Trading `eff4`'s only firing for one of five is a worse board, so it is not made. **4/384 against 5/352, declared.**

AND THE FINDING THAT OUTLIVES THE FIX IS THE GUARD'S OWN, FOR THE EIGHTH TIME. `feature_fixture.verify()` decided "*is this the same fixture?*" from `ROUND` and the list of scenario **LABELS**. A label is not a board. Edit a species, an item, an ability, a nature, a move or a pre-set state inside a scenario that keeps its name and the identity check passes — after which every moved column is announced as "*these features changed MEANING since the weights were fitted*", which is **FALSE**, and which accuses `board.js` of a change it did not make. That is the refit/restamp distinction this whole file turns on, collapsed. It now carries a **body digest** covering species, item, ability, nature, moves, bench and pre-set state, reported as "*the fixture's BOARDS changed while its scenario labels stayed the same*" and explicitly **NOT** as evidence about `board.js`. A stamp written before the digest existed is treated as an OLDER STAMP and not as staleness, matching the convention the `table` block already set.

WHAT THIS DOES NOT SAY. It says nothing about whether a fixture's board is a good test — only that the team could be brought. It cannot see a set assembled at runtime. And the 41 baselined verdicts are **not repaired**: they are named, costed per file, and routed as **#266**. Nothing was re-baselined into silence, and the gate is what enforces that.

REPORTED, NOT FIXED, NOT SILENCED. `tests/test-mechanics.js` was run WITHOUT being edited: **567 live / 0 missing, exit 0** — but ENGINE was working in the same window and took the census 565 → 567 under #264, so that number is ENGINE's and is quoted here only as evidence that nothing in this pass moved it. `tests/test-protocol-trace.js:131` gives Politoed the item `'dampiro'`, which names nothing in this format: the fixture builds a Politoed holding **nothing** while its source says otherwise, and reports success. That is *a capability was absent and everything reported success*, and it is filed rather than repaired because that file has twelve other verdicts and they move as one batch.

oooooooo. ROADMAP #254 — THE HAZARD SIDE. THE FIT DOES NOT MOVE, THE LIVE BOT DOES, AND THE GUARD COULD NOT SEE EITHER — 2026-08-13

`engine/board.js` (`sideFor` + three call sites), six sweep sites, `engine/feature_fixture.js` (one new board), `tests/test-hazard-side.js` (new), `tests/test-feature-antics.js` (one new clause). `engine/medicham2-browser.js` was not touched.

THE HEADLINE IS THE REFIT VERDICT, AND IT IS NO. `board.js` recorded every side condition on the MOVER's side. Seven of the eleven legal side-condition moves in this regulation are `target: allySide` and were right by accident; the four `foeSide` ones — Stealth Rock, Spikes, Sticky Web, Toxic Spikes — landed on the side they can never be on. That is a feature the FIT reads, so the question this division has to answer is whether the weights now describe a different quantity.

THE WHOLE FIT CORPUS, REPLAYED THROUGH `FP.decisionsFor` UNDER BOTH BOARDS: 9,226 games, 237,052 decisions, 1,731,851 candidate vectors, 0 errored, identical decision key set, and 0 decisions, 0 games and 0 feature vectors move. The pre-#254 board was compiled in memory and injected into `require.cache` under `board.js`'s own resolved path; the tree was never reverted, so both arms ran against the same everything-else.

THAT ZERO HAS TWO POSSIBLE MEANINGS AND THEY ARE OPPOSITE, SO BOTH WERE MEASURED. Thin exposure would make the zero uninformative. Exposure is small but not nil: **24 hazard clicks in the fit corpus** (Stealth Rock 19, Toxic Spikes 2, Sticky Web 2, Spikes 1) against **9,911 allySide clicks**, and **258 of 237,052 decisions (0.109%) carry a hazard as a legal candidate**. The zero is therefore not "nothing to move".

IT IS AN ISOMORPHISM, AND IT WAS CONSTRUCTED RATHER THAN ARGUED. `noteMove` wrote to the mover's side and both readers read the mover's side, so the old board is the new one with the two sides RELABELLED — every read that was derived offline lands on the same set. Same probe, all eleven moves, both boards:

how the condition arrived	pre-#254	fixed		
OFFLINE — <code>noteMove(..., worked=true)</code> , which is every corpus replay	<code>deadSide p1=1, p2=0</code>	<code>p1=1, p2=0</code> — identical		
LIVE — <code>`\</code>	<code>-sidestart\</code>	<code>, which is engine/magnemite.js:647`</code>	p1=0, p2=1	<code>p1=1, p2=0</code>

SO THE DEFECT WAS ONLY EVER IN PLAY, AND IT IS THE WORSE HALF. The live path takes the side straight off the protocol and calls `noteMove(..., worked=false)` , so the offline derivation never runs and nothing relabels the read. On the old board the LAYER read `deadSide=0` after laying Stealth Rock — it would re-lay it every turn believing it was not up — and the VICTIM read `deadSide=1` and refused to lay its own. Both exactly backwards. **MAG was fitted in a world where `deadSide` was effectively right and played in one where it was inverted**, which is *fitting environment and playing environment must match* broken in a new place — the same shape as the sheet-visible fit and the broken-mega champion — and it is closed here in the one direction that costs no refit.

A REFIT IS NOT OWED BY THIS ROW. The weights describe the same quantity on the corpus they were fitted on, to the vector. A refit remains OWED for the reasons `status.js` already prints (`medicham2-browser.js` , `board.js` , `engine-data.js` and `abra-tags.js` all moved after the fit) and nothing here changes that; no fit was run and no weight file was restamped.

ONE CONSEQUENCE TO NAME RATHER THAN DISCOVER LATER. Adding a fixture board takes the scenario count 10 → 11, so `feature_fixture.js --check` now answers *"the fixture itself changed ... Old hashes cannot be compared"* for every stamped file, which SUPERSEDES the per-feature message it printed before. That branch is correct and it is also less informative, and it is a restamp-or-refit decision that belongs to whoever next moves the weights — which is Will, who is reworking them.

THE FINDING THAT OUTLIVES THE FIX IS THE GUARD'S, AND IT IS THE REASON THIS ROW CAME HERE. `tests/test-feature-antics.js` is the guard against a feature changing MEANING under an unchanged NAME. Run under both boards with the fixture as it stood, **0 of 76 hashed columns moved** — `engine/feature_fixture.js` contained **zero hazard clicks** and the only side conditions any of its ten boards pre-set were `reflect` and `tailwind` , both `allySide` , which is exactly the half `sideFor` leaves alone. The guard would have certified this silently. R7 for the seventh time, and the new part generalises: the earlier misses were a SPECIES the boards did not stand on and a FIELD STATE they never entered; this one is a MOVE CLASS nobody clicks.

A NEW BOARD, hazards-already-up , AND IT IS DELIBERATELY ONE-SIDED. `stealthrock` and `spikes` up on p2, `stickyweb` and `toxicspikes` up on p1, so the flip is exercised in BOTH directions on one board — setting a hazard on both sides makes every reader answer 1 before and after and moves nothing. `deadSide` now moves `b984c210828d` → `d1be4f95d589` and no other column does. Every set is learnset-checked against `Dex.forFormat` , and the board is added rather than an existing one edited, per that file's own record of two attempts that were inert and one that regressed coverage.

AND ONLY `deadSide` **IS OBSERVABLE, WHICH CORRECTS THE ROADMAP ROW RATHER THAN REPEATING IT.** The row said a write-only fix would credit `setupTurns` every turn. Measured: none of the four hazards carries a `condition.duration` (Reflect 5, Tailwind 4, the hazards none), so `dur()` returns 0 and `setupTurns` is 0 for them under every variant. The `alreadyUp` site is routed through the helper for the one-fact-one-function rule and binds today only on the seven `allySide` moves, where `sideFor` is the identity.

NINE CALL SITES, NOT THREE. The write, the two reads, and six verbatim copies of the same mover's-side derivation that ENGINE's sweep found outside `board.js` — `fit_policy.js:793` (the fit itself), `joint_rows.js`, `branch_recall.js`, `corpus_shift.js`, `feature_coverage.js`, `redirect_audit.js`. Fixing only `board.js` would have left the fit deciding the fact for itself. `tests/test-hazard-side.js` holds all six to reading `B.sideFor`, and reverting one of them fails that clause **by file name**.

SHOWN RED THREE WAYS BEFORE BEING BELIEVED, none of them by editing the tree.

`tests/test-hazard-side.js` was **5 passed / 3 failed at HEAD** and is 11/0 after; under the injected pre-#254 board the semantics guard's new clause names all six wrong candidates in both directions; and reverting one sweep site fails by name. `tests/test-mechanics.js` is **565 live / 0 missing**, unmoved — the whole-game differential exercises `medicham2`, not the feature board, so it was not expected to move and did not.

REPORTED, NOT FIXED, NOT SILENCED. `engine/selftest.js` is RED (exit 1) on *"every raw reader of the ladder store declares why — 9 file(s)"*: it is a MISSING DECLARATION, not a corrupt read — the gate demands each raw reader carry a `RAW-STORE-OK`-style reason and nine do not. Pre-existing and untouched here. `tests/test-fragility.js` is RED on two Storm Drain assertions and is **identical under both boards**, so it is not this change. And `engine/feature_fixture.js` gives Venusaur a **Rocky Helmet**, which this format banned on 2026-08-04 — left alone deliberately, because changing a fixture body moves every stored hash.

oooooooo. GROWING A `need` **LIST SILENTLY RETIRED OLD MEASUREMENTS. IT NOW COSTS A NAMED ARTIFACT — 2026-08-12**

`tests/test-artifact-rerunnable.js` (new) + the `provides` recorder in `engine/engine_release.js`. Nothing else was touched.

Will: *"should i be concerned we suddenly cant run old things"* → *"you need to fix it and make it so it doesnt happen again i dont know the solution"*.

A release freezes the ENGINE and not the READER. Every symbol a caller adds to its `need` list retroactively strands every release cut before that symbol existed: the snapshot still verifies, still holds its bytes, and simply stops being openable. ROADMAP #222 split the RNG streams and `rngStreams` appeared in 30 snapshots; `spreadL50` stranded two more. **Both changes were right. The bug is that paying their cost was invisible.**

THE HEADLINE, AND IT CONTRADICTS THE FIRST READING OF THE SAME FACT. 41 artifacts name a release across 21 releases: 29 RE-RUNNABLE, 1 STRANDED, 11 UNKNOWN-PRODUCER, 0 retired. The first scan reported that essentially nothing could be re-run. That scan unioned every `need` list in `engine/` into one 24-symbol set and held every artifact to it — so `game_differential.js`'s artifacts were failed for lacking `hitChance`, which only `million_run.js` has ever asked for. **An artifact is judged against the caller that PRODUCED it**, read from its own `by` field, and the requirement table comes from `engine_release.js`'s `callerNeeds()` rather than a second scanner. The one genuine stranding is `data/nature-arms.json`, whose release predates `rngStreams` and `spreadL50` that `game_differential.js` now needs.

AND "168 OF 200 RELEASES ARE UNOPENABLE" IS ONE CALLER'S NUMBER WEARING THE STORE'S NAME. Asked per caller over all 201 releases on disk, against `engine/medicham2-browser.js`:

caller	releases that can serve it	predate an export	pruned / absent / unloadable
engine/game_differential.js	28 of 201	168	5
engine/million_run.js	97 of 201	99	5
engine/replay_differential.js	140 of 201	56	5
engine/speed_vs_pokeenv.js	196 of 201	0	5
the UNION, which is what census() asks	28 of 201	168	5

The union is the strictest caller and nothing else, so `data/release-census.json` 's `runnable` is a lower bound on every caller but one. It is not wrong — its own comment says it asks the hardest question any live caller asks — but it must not be read as "the store is 86% dead". **Filed:** `census()` **should report the per-caller column beside the union.**

UNKNOWN-PRODUCER IS A BAND, NOT A PASS. Eleven artifacts either record no `by` or name a producer with no `REL.require(file,{need})` site — a `.py` script, a driver that shells out, or a caller that reaches a snapshot through `REL.read / REL.path`. Their requirement cannot be `READ`, so they are not accused; they are still held to the one requirement that needs no guess, that the release opens at all. `data/all-mechanics-fire.json` — the 93 move divergences — is in that band, and the cheapest way to move it out is for its generator to stamp a `by`.

THE PARSER WAS DELETED RATHER THAN PATCHED, AND ITS FAILURE IS THE SAME ONE THREE TIMES. The check originally hand-rolled a `module.exports` reader and reported that a release cut MINUTES earlier lacked `fails` and `hitChance`. `medicham2-browser.js` interleaves block comments with the keys inside that literal, so a comma split glues each comment onto the key after it: **11 of 78 exports lost, and 4 more invented out of prose — one of them the word** `deliberately`. That is `provenance.js` 's `writesNear` hole and `callerNeeds` 's stripped-comment rule, for the third time. `engine_release.js` already answers the question by `LOADING` the frozen module (`surface()`), at 18ms, so the verdicts use the loader and this file contains no parser.

THE RECORDED `provides` **FIELD IS CORRECTED AND IS NOW AUDITED RATHER THAN TRUSTED.** It is kept because it is the only thing that survives a prune — once the bodies are gone `surface()` can answer nothing. It is stamped `provides_by` so a legacy list is distinguishable from a current one, and the check compares every recorded list against the loader on every run. The one manifest written by the old recorder is reported, not rewritten: a release record is immutable and is not edited to make a check green.

SHOWN RED BY REPRODUCING THE REAL EVENT. Adding `rngStreams` to `engine/replay_differential.js` 's `need` list — exactly what #222 did to the differential — turned its five artifacts `STRANDED` and named all five as `NEW` against the ratchet, while every artifact of every other caller stayed put and `replay-differential-freezes.json` (no `by`) correctly stayed `UNKNOWN-PRODUCER` rather than being accused. The caller was restored byte-identically (sha256 `a21d2158...`) and the check returned green.

The ratchet is `data/artifact-rerunnable-baseline.json`; it may only fall, and a new stranding fails by name.

Filed, not fixed: `callerNeeds()` reads `REL.require` sites only, so the file requirements of `tests/roster.js` and `tests/mutation_harness.js` — which reach a snapshot through `REL.read` and `REL.path` — are invisible to it and their ten artifacts are judged on openability alone. That under-reports and never over-reports. `callerNeeds()` also hard-codes an `engine/` prefix on the caller name, so a `tests/` scan is mislabelled and the label is repaired at the call site. Both belong in `engine_release.js`, in a pass that is not running beside two other divisions.

000000. THE END-STATE COUNT BECAME A SEVERITY LADDER, AND THE WORST BAND IS NINE PERCENT — 2026-08-12

engine/end_state_severity.js (new) + engine/game_differential.js end-state reporting + tests/test-end-state-severity.js (new). Nothing in medicham2-browser.js or tests/test-mechanics.js was touched; the run reads the frozen release 6155acc0fb26 , the same one the standing 31.9% / 36.6% bar and the published end-state table use.

WHY. DIFFERENT-END-STATE was a COUNT. A game in which a healthy body is killed by a move it cannot be hit by — after which a replacement comes in and every later line is a different game — weighed **exactly the same as a three-HP rounding residue**. On the same day, five defects were verified against Showdown's own source, staged, shown red, fixed, and the whole-game count moved by **+1 and +3**; all five were narration. Will then read twenty-five battles by hand and found three wrong OUTCOMES. **No instrument here could have separated those two kinds of finding, because a count has no order on it.**

THE HEADLINE — 2,300 games per arm, frozen team pool, release 6155acc0fb26 , turn cap 12, 0 threw, planted proof green, data/game-differential-endstate.json . Bands are over the DIFFERENT-END-STATE games, which is the parted ones PLUS the games whose narration never parted and that still ended apart (31 top, 18 bottom — a class the protocol instrument is structurally blind to).

band	what it means	top-tie-first	bottom-tie-first
1	DIFFERENT WINNER	0	0
2	DIFFERENT BODIES ALIVE — dead in one engine, standing in the other	25 (9.0%)	27 (9.1%)
3	HP differing by more than a typical hit	47 (16.9%)	59 (19.9%)
4	different species / typing / ability on a LIVE body	86 (30.9%)	95 (32.0%)
5	a status, item, hazard, screen, weather, PP or volatile	51 (18.3%)	47 (15.8%)
6	HP under one hit, or a stat stage — plausibly rounding	69 (24.8%)	69 (23.2%)
	banded	278	297

SO ROUGHLY ONE DIVERGED GAME IN ELEVEN ENDS WITH A DIFFERENT SET OF POKEMON ALIVE, and about three in ten with a body whose identity we have wrong while it is still standing. Neither was visible before; both were inside one number.

THE BAND 3 THRESHOLD IS MEASURED AND THE RULER IS THE AUTHORITY'S, NOT OURS. Every DIRECT -damage event Showdown narrated, as a fraction of the struck body's own maximum HP; the median is the threshold. **2,092 hits, median 25.9% of a health bar, IQR 14.1–47.3** (top); **1,808 hits, 28.8%, IQR 17.5–53.1** (bottom). Cut per arm and never pooled — the arms sit at opposite corners of the damage roll.

TWO THINGS THE FIRST VERSION OF THAT RULER GOT WRONG, BOTH FOUND BY READING ITS OWN OUTPUT. (a) Residual chip — sandstorm, burn, Leech Seed, hazards, Life Orb — was counted as a hit and put the median at **12.7% with quartiles 6.2 and 34.2**, the middle of a bimodal mixture, describing neither half. Showdown's own line says which is which ([from] ...), so the filter is read off the protocol rather than guessed from a size cutoff, which would be circular. **4,838 and 4,947 residual events are excluded and counted.** (b) 0 fnt carries **no denominator**, so a parse requiring \d+/\d+ discarded **every killing blow** — 13 direct hits narrated in one measured game and 8 reaching the ruler, the five missing being the knockouts. A ruler built only from hits too small to kill anybody. It now resolves against the body's carried maximum; hits went 1,612 → 2,092 and 1,321 → 1,808.

THE SHAPE PRIOR HOLDS AS A RATE AND IS WRONG AS A HEADLINE. The prior was that ORDERING-shaped divergences land harmless and RULE-shaped ones land severe. Share of each shape's games reaching band 2:

shape	top	bottom
RULE	7/47 = 14.9%	11/68 = 16.2%
EMISSION	16/137 = 11.7%	15/126 = 11.9%
ORDERING	1/14 = 7.1%	1/21 = 4.8%
UNPARSED	0/48	0/57
protocol never parted	0/31	0/18

RULE is the most dangerous shape per game and ORDERING the least, so the prior is right. **But EMISSION supplies 16 of the 25 and 15 of the 27 band-2 games** — the majority of the worst outcomes — because it is by far the largest bucket. "Fix the rule-shaped ones first" is right per game and wrong per evening. **And ORDERING is not zero: one game in each arm has an ordering-shaped first divergence and a different set of bodies alive at the end.** That is the more interesting half and it is not dismissed here.

THE WORKLIST, RANKED BY BAND FIRST AND BY CORPUS USAGE SECOND (teams containing the body in `data/meta-usage.json`, read LIVE — it is not in the frozen release, and it is **generated 2026-08-04**, so the ranking rests on an eight-day-old corpus model):

- `sinistcha` **is the largest single body in band 2 — 6 games (top) and 7 (bottom), 2,668 corpus teams**, and the same shape every time: `0/146 for us against 146/146 for the authority`. A body at FULL health in Showdown and dead here. Filed to `@engine`; not diagnosed here.
- `pair-redirect-priority` **supplies 14 of 25 and 13 of 27 band-2 games**, `pair-protect-bust` 8 and 7. That is a fact about the steered sample as much as about the engine and both readings are open.

BAND 1 IS ZERO AND THAT IS NOT A CLEAN BILL — THE HARNESS CANNOT CURRENTLY MEASURE IT. A battle that does not resolve cannot have a different winner, and **2,275 of 2,300 games stop at the 12-turn cap**. Raised to 19 (`data/game-differential-endstate-turn19.json`, 983 games, proof green) it is still **938 of 983 at the cap, band 1 = 0 in both arms**.

AND ABOVE TWENTY TURNS THE COMPARISON COLLAPSES, FOR A REASON THAT WAS THE HARNESS. `battleOver` is `S.turn >= (S.maxTurns || 20) || ...` and this driver had never set `maxTurns`. Measured at `--turns 40`, 983 games: **943 ENDED-APART, 937 of them "ONLY medicham2 ended the battle"** — 96% of a run produced by one hard-coded default, reading exactly like a catastrophic engine disagreement. (`data/game-differential-endstate-turn40.json` is that run, kept as the receipt; it was taken under the pre-fix driver and its band-3 figures are not to be quoted.) The driver now sets `S.maxTurns = max(MAXTURNS + 1, 20)`, so **every 12-turn run is byte-identical to before** and only the already-broken configurations move.

A 30-TURN RUN THEN REFUSED TO PUBLISH, CORRECTLY. With the horizon lifted, one of the state comparator's own planted defects — *PP spent on a slot NOTHING has touched* — can no longer be staged, because in a 30-turn game every slot has touched every move. The run printed `THE STATE COMPARATOR FAILED ITS OWN PROOF` and declared its state numbers worthless. That is the instrument working. **So the honest position is that DIFFERENT-WINNER has no denominator in this harness yet**, and closing it needs a plant that survives a long game, not a bigger sample.

SHOWN RED FIRST, AND SHOWN NOT FIRING. `tests/test-end-state-severity.js`: a planted death reaches the top rung and is localised; **a planted three-HP residue must NOT** — a ladder that answers "severe" to everything passes the first half alone; a one-sided forme change must land on the identity rung and not read as a death, because the party is keyed by species and a rename is indistinguishable from a corpse in the diff list; all six rungs reachable; a board carrying a death, a boost and a burn bands as the death; identity on a body dead in BOTH engines does not reach the identity rung; `severity()` REFUSES to run without a measured threshold; the `|split|` duplicate is stepped over; and the whole path through the real driver, a planted faint surviving to the last board against the same pair and seed run clean as a control.

THE ONE PLACE THE BRIEF WAS THE WRONG SHAPE, AND IT IS RECORDED RATHER THAN QUIETLY OBEYED. The brief ranked *a different set alive* above *a different winner*. First-match-wins over that order makes the winner rung **structurally empty** — a different winner is always also a different set alive — and a band that can never fire is not a band. They are swapped, and the containment is printed on every run. A fifth rung was also added: folding a burn, or a layer of Spikes, into "plausibly rounding" would repeat in miniature the exact weighting bug the ladder exists to fix.

AND A RED THAT WAS OPEN ON ARRIVAL IS CLOSED, WITH THE ATTRIBUTION CORRECTED. docs/ENGINE.md recorded tests/test-end-state.js PART 3 failing and attributed it to staleness in the pass that wrote it, having restored both medicham2-browser.js and game_differential.js to their HEAD bytes and reproduced it. That control was sound and it held the wrong variable still. **Same frozen release, one flag apart: live team pool → 2 FAILURES;** --team-store data/team-pool-frozen → **ALL GREEN.** diff_swarm.js reads the pool LIVE from a file OPS appends to, so pairsFor returns a different first pair as the store grows, and PART 3's item plant — legitimately undoable by a Knock Off or a berry — eventually lands in a battle that undoes it. The test now pins the pool by default and a caller can still override it.

WHAT THIS DOES NOT SAY. A band is exactly as strong as what board_state.js compares; its NOT_COMPARED list is published with the artifact. SAME-END-STATE remains a claim about where the two engines ARRIVED and not about the turns in between. ENDED-APART (4 and 6) is a THIRD answer and is counted beside the ladder, never inside it.

ooooo. TURN 1 IS THE PRIMARY READING NOW, AND THE CONTAMINATION THEORY IS HALF RIGHT — 2026-08-10

Will: "lets only do turn 1's so we know nothing from the previous turn is messing up. then we can begin to look into later turns". --turn1-only refused to write an artifact, correctly, because one of the three planted defects is **structurally impossible on turn 1** — every body starts at full HP, so an hp event written at the top of turn 1 puts 100 over 100 and there is nothing left to detect.

THE REFUSAL WAS NOT WEAKENED. THE GATE IS STATED OVER CLASSES OF DEFECT INSTEAD OF ARM NAMES. An arm declares its class; an arm the mode cannot carry declares itself INAPPLICABLE with the reason; and the run refuses unless every class has an arm that RAN and was CAUGHT. Dropping an arm because the mode cannot carry it is the same bypass as never running it, one level up from the --selftest flag this file already refuses to hide behind.

A FOURTH PLANT COVERS THE PRE-TURN-BOARD CLASS ON TURN 1: A SPECIES SWAP. A lead body is replaced by one whose FASTEST legal Champions spread — taken over every weather — is slower than the SLOWEST legal spread of the body the record shows it outspeeding, so the recorded resolution order becomes something no legal spread can produce. It is planted at **every suitable site up to 12 and every one must be caught**; planting once and stopping at the first success is how a detector that works on one board in twenty passes a gate.

plant	class	all turns	turn 1 only
a damage figure cut to a quarter	outcome	runs	runs (restricted to turn 1)
a freeze the game cannot produce	unreachable	runs	runs (restricted to turn 1)
a wrong pre-turn HP under a recorded KO	pre-turn board	runs	INAPPLICABLE — declared, not skipped
a species swap on the pre-turn board	pre-turn board	runs	runs

TWO CANDIDATE PLANTS WERE MEASURED AND REJECTED FIRST, AND THE FIRST ONE IS A FINDING. A species swap into a **type immunity** is invisible to this instrument: dmgRange returns an EMPTY roll array for an immune matchup and damageVerdict turns that into unresolved — dmgRange returned no rolls rather than a divergence. So "the record shows damage on a body our engine says cannot

be touched at all" is REFUSED, not accused. **FILED, not fixed** — it would move the headline on the same day the mode changed. A swap into a bulkier body was rejected for a different reason: the attainable interval is ~60 points of max HP wide, so it fires or does not fire depending on the sample, and a red proof that is a coin flip is not a proof.

SHOWN RED TWICE, DELIBERATELY, BEFORE BEING BELIEVED. Sabotaging the order comparator's disjointness test produced 8 of 8 PLACED PLANTS WENT UNNOTICED ; marking the species arm inapplicable produced A WHOLE CLASS OF DEFECT HAS NO ARM THAT RAN AND WAS CAUGHT: preturn . Both refused to write.

THE HEADLINE: 1,249 of 19,715 turn-1 units diverge = 6.34%, release 30b21c5a335a , 20,000 games read and 285 skipped (1.43%), 0 exceptions. bot-v-human 6.71%, human-v-human 5.62%.

THE MOVE FROM THE PUBLISHED 5.36% IS ENTIRELY SAMPLE, AND THAT IS MEASURED RATHER THAN ASSUMED. On the same first 3,000 games the rebuilt instrument reproduces 158/2,947 = 5.36% exactly, at the old release 70794711fe6d AND at the new one — so the instrument moved nothing and the engine moved nothing on this arm. Games 3,001–20,000 run 6.51%, and the gap survives inside both population strata rather than being the rising bot share alone.

TURN 1 AGAINST LATER TURNS, SAME 20,000 GAMES, SAME RELEASE, ONE RUN OF EACH MODE. --turn1-only reproduces the full run's turn-1 arm to the unit (1,249/19,715 both ways), so the mode changes the denominator and not the measurement.

	turn 1	turns >= 2
turns compared	19,715	106,558
diverged	1,249 (6.34%)	8,185 (7.68%)

Per 1,000 turns, by family — and the split between "everything that fed this was known" and "something was never revealed" is the whole point:

family	turn 1	later	later / turn 1
damage, all inputs known	17.40	25.08	1.44
damage, attacker's item never revealed	34.90	37.53	1.08
turn order	5.02	8.39	1.67
status, source ability known or no source	1.07	2.97	2.79
status, source ability never revealed	8.17	7.56	0.93
weather	0.00	0.17	—

THE STATUS BUCKET: THE RAW COUNT BARELY MOVES AND THE CLEAN ARM COLLAPSES. 182 status divergences on turn 1 (9.23 per 1,000) against 1,123 later (10.54 per 1,000) — a 12% reduction, which does NOT support "the 188 status divergences are largely false" as a blanket claim. **88% of the turn-1 status rows (161 of 182) are rows where the body that clicked into the status had an ability the record never revealed**, and that split did not exist until this pass: the damage family has carried [the attacker's item was never revealed] since it was built and the status family was quoting two different claims as one number.

Split, the answer is sharp. **SLEEP IS THE CARRIED-OVER-STATE CONTAMINATION AND IT IS ISOLATED:** 130 cause-known sleep divergences on later turns (1.22 per 1,000) against ONE on turn 1 (0.05) — a Yawn or a sleep counter from an earlier turn, exactly as suspected, and it cannot survive a turn-1 arm. **FREEZE IS THE OPPOSITE** and is the one clean status candidate for ENGINE: 10 of the 14 turn-1 freezes have a known cause and every witness is a Blizzard or an Ice Punch.

THE LARGEST SINGLE TURN-1 STATUS WITNESS IS NOT AN ENGINE DEFECT: sneasler Fake Out -> psn , **57 times in 20,000 turn-1s.** That is POISON TOUCH, and data/engine-data.js carries Sneasler's modal observed ability, which is UNBURDEN — in a closed-sheet game, 98.3% of this store, the body is built with it and cannot poison on contact. **A general lesson rather than a Sneasler one:** the modal-ability prior is doing the same damage to the status family that the unrevealed item does to the damage family.

AND THE SPREAD-MOVE INGEST DEFECT IS WORSE ON TURN 1, NOT BETTER. The unresolved rate is **39.56% of 41,762 damage units** against 35.30% over all turns, because **10,304 of the 16,519 unresolved rows (62%) are the spread-move conflation** — turn 1 is where both foes are most likely to be alive, so it is exactly where a spread move hits two bodies and the store collapses them into one row. Turn-1-only removes the invisible-carried-state class and removes none of this one. sinistcha Matcha Gotcha -> brn is the top turn-1 burn witness, which is that defect showing up in the status family too.

THE JOINT 66-POINT BUDGET IS NOW ENFORCED AND IT BUYS NOTHING — 0 of 28,982 corners clamped. Every corner this instrument evaluates pushes at most TWO stats of one body (defensive stat + HP pool = $2 \times 32 = 64 \leq 66$), so the constraint never binds and no interval moves by a point. The clamp is live rather than decorative and starts counting if the cap ever changes. **The over-width that is real is a different one:** every EVENT is evaluated at its own corner, so one body can be the max-offence spread while it attacks and the max-bulk spread while it defends, in the same turn, on the same 66 points. Tying a body to one spread across a game is a joint solve, and it errs toward ACCUSING the engine — not done here, and not on the day the mode changed.

Filed, not fixed: the immune-matchup hole above; readGames holds the whole sample in memory, which is what caps the sample rather than time (20,000 games is 35s in turn-1 mode); and the status family still has no split for a self-inflicted orb or a residual, which land in [source ability KNOWN or no preceding move] beside the genuine candidates.

0000. ROADMAP #68 — THE ENGINE NOW GETS MEASURED AGAINST GAMES THAT ACTUALLY HAPPENED — 2026-08-10

engine/replay_differential.js → data/replay-differential.json and data/replay-differential-freezes.json . New file; nothing existing was edited.

The authority is data/games.ladder.jsonl , **not Showdown.** tests/test-engine-diff.js is one damage calculation per row with no turn loop, and at its default n=150 it reported zero disagreements for weeks while n=20,000 finds 19. engine/game_differential.js has a turn loop and plays SYNTHETIC games against the library. Neither asks whether our engine describes the game people actually play.

THE HEADLINE, 3,000 games, release 70794711fe6d . 2,947 replayed, 53 skipped (1.77%), all for the same reason — the row has no turns. **18,421 turns compared, 1,008 diverged = 5.47%.** On turn 1, where no invisible state can exist yet, **5.36% of 2,947. 0 exceptions.** Bot games are included and split rather than filtered: bot-v-human 5.95%, human-v-human 5.06%.

THE BOARD IS REBUILT FROM THE RECORD AT THE START OF EVERY TURN, so each turn is an independent unit and one 6-turn game is ~6 tests rather than one fragile chain. The price is enumerated in cannot_see and the turn-1 arm is the one with no unmeasured confounder in it.

ROLL RECOVERY DOES NOT WORK ON THIS CORPUS, AND THE REASON IS MEASURED.

Will's proposal — roll all 16 and identify which one was played — is the right shape, and it turns the damage figure from an input into a test. It cannot run here: **Champions team sheets do not declare SP** (884 of 52,089 stored games carry a sheet; every one has evs: null), and the record states damage as an integer percent of an unknown maximum. The **median attainable damage interval is 60.1 points of max HP**, so one roll step is **3.76 points** — the legal-spread envelope is several times wider than the entire 16-roll band.

matched fires **2 times in 37,177**. The test is therefore inverted into the one the record supports: all 16 rolls, at the observed crit state, at both corners of the legal SP envelope, and is the observed value inside. **719 no-match, 13,359 ambiguous, 13,533 unresolved (36.4%), 9,564 KO-clamped**. The 36% is printed at the top because a rate that size decides whether the headline means anything.

THREE SOURCES OF RANDOMNESS ARE RECOVERED FROM THE RECORD RATHER THAN SIMULATED — a miss is `miss:1`, a crit is `crit:1` and its ABSENCE is knowledge because Showdown always announces one, and a secondary is an `x` or `b` event. **The resolution order is the event order (ROADMAP #43), and nothing in this repository had ever read it:** 3,081 turns where our speed calculation is forced and agrees, **159 where no legal Champions spread can produce the order the record shows**, 9,312 refused as inside the envelope, 5,016 refused for a Prankster/Gale Wings/Triage carrier.

EVERY DIVERGING TURN IS A READABLE FROZEN BOARD, not a counter — Will diagnoses by reading boards. `data/replay-differential-freezes.json` carries the board before, the reconstructed clicks with a per-click confidence, the board after from the log against what our engine could reach, and the raw events. **Four instrument bugs were found by reading its own output**, each of which had been sitting at or near the top of the mechanic table: a mega not applied before the turn it happens in, a Weather Ball priced in clear skies because the sun was set three events earlier in the same turn, a KO clamp read off a reconstructed HP instead of the record's own figure (56 of 158 divergences), and weather with no expiry doubling Swift Swim and Chlorophyll speeds for the rest of the game.

THE RED PROOF RUNS ON EVERY PUBLISHED RUN AND THE ARTIFACT REFUSES TO EXIST WITHOUT IT. Three planted defects in a real stored game — a damage figure cut to a quarter, an impossible freeze, a corrupted board HP that puts a recorded KO out of reach. A `--selftest` flag somebody has to remember is the same failure one level up. It has been **shown red twice for real**, both times correctly.

FILED TO ENGINE, from the mechanic table. These are candidates and not verdicts; where the attacker's item was never revealed the row says so IN ITS KEY and must not be quoted with the others. Largest with everything known: `x0.5-ish` 72, `x0.25-ish` 56, `x2-ish` 32, `x4-plus` 17, and the status family — `slp` 59, `brn` 57, `psn` 31, `par` 18, `frz` 15 — where the record applied a status no pin of our engine can produce.

FILED TO OPS — three ingest defects this instrument had to work around. Each is a store change, not an analysis change. (1) `durable-ingest.js` records **no cant event at all**, so flinched, fully paralysed, asleep, frozen and recharging are one indistinguishable absence — Showdown emits `|cant|p2a: X|flinch` and it is dropped. (2) **A spread move's damage is conflated:** every `-damage` is attributed to the last `m` event with `dmg = max(dmg, delta)` and `tgt = tgt || <slot>`, so one row carries the maximum of two deltas against the first target's species and the last target's `tgthp`. That refuses **8,360 of 37,177** damage units — the single largest unresolved bucket. (3) **Everything before |turn|1 is dropped**, including a lead's entry weather, so a Drought lead's sun is invisible.

NO GROUND TRUTH FOR THE CHOICES EXISTS. The 22 stored games with `willhoop` as a player carry the same extracted schema as every other row — no record of what was clicked. So the reconstruction cannot be validated against known answers and rests on the per-click confidence in the freeze dump. The live bot logging its own choice string would make this instrument exact.

DEFERRED, DESIGNED, AND BEHIND `--rates` : **the aggregate secondary-rate test.** Pinning an outcome to what was observed makes a wrong PROBABILITY agree with itself — the same trap as pinning damage. The counter-test compares each move's observed secondary rate across the corpus against the chance our tag declares. It is not part of any figure above and needs a far larger corpus before an interval on a 10% secondary means anything.

ooo. THE ABILITIES CLAUSE WAS GREEN OVER 27% COVERAGE AND FIFTEEN UNMEASURED ROWS — ROADMAP #120, #121, #122 — 2026-08-10

Will: "check the abilities clause for the same hole". data/roster.abilities.json read 84 FIRED-AND-BOARDS-MATCH, 217 COULD-NOT-STAGE, 15 CONTROL-NOT-QUIET and engine/quarantine.js printed clean: 84 fired and matched and PASSED. Three separate faults sat inside that sentence.

THE CLAUSE NOW FAILS, AND IT FAILS ON SEVEN ATTRIBUTED DEFECTS.

THE ENGINE MOVED UNDER THIS AND THE TWO CAUSES ARE SEPARATED RATHER THAN BLENDED. A release was cut by another division at 04:21 while this was being written, so the published run reads a different simulator from the artifact it replaces. The new instrument was therefore re-run pinned to the OLD release a4c7f898ad0e , and the two engine snapshots differ on exactly one row: Magic Bounce, DID-NOT-FIRE on the old and FIRED-AND-BOARDS-MATCH on the new — ENGINE fixed it in between, and the old instrument could not have said so because that row was CONTROL-NOT-QUIET. Every other movement below is the instrument.

	before (old instrument, a4c7f898ad0e)	instrument only (a4c7f898ad0e)	published (bfefdb697454)
FIRED-AND-BOARDS-DIFFER	0	2	2
DID-NOT-FIRE	0	6	5
FIRED-AND-BOARDS-MATCH	84	76	77
CONTROL-NOT-QUIET	15	18	18
COULD-NOT-STAGE	217	214	214
clause	PASS	FAIL	FAIL

1. A CONTROL THAT IS ITSELF A LIVE ABILITY IS NOW VARIED, NOT CAPTIONED (#121).

The 15 rows were controlled by a second real mechanic because their carrier species has no quiet alternative — and this format has only 8 quiet abilities, none of which shares a species with any of the 15, so "pick a quieter one" cannot reach them and never will. A quiet ability cannot be lent from another species either: buildPair clamps an ability to its species' own list and falls back to slot 0, so an illegal pairing silently becomes the subject arm again.

What CAN reach them is a SECOND CONTROL: where the carrier has a third ability, the identical scenario is played a third time against it and the two deltas are compared leaf for leaf in both engines. A leaf that survives both does not depend on which ability was removed, so it is not the control's. 98 rows were varied this way. It is the noise-floor discipline (§9 of LESSONS) applied to a control arm: vary the knob that is supposed not to matter, and believe the effect only if it survives.

IT IMMEDIATELY CAUGHT EIGHT VACUOUS GREENS. Anger Point, Justified, Rivalry, Keen Eye, Shell Armor, Stalwart, Sticky Hold and Slush Rush were FIRED-AND-BOARDS-MATCH and the agreement was about the CONTROL's work — Intimidate's -1 Attack on three of them, Weak Armor's drop, Gooley's drop, Stamina's boost, Supersweet Syrup's evasion drop. Anger Point needs a crit and the pin lands none; Rivalry needs a gender and buildPair sets every body to N by construction. Those rows could never have fired.

AND IT CAUGHT THE INSTRUMENT'S OWN FIRST ANSWER, WHICH IS WORTH RECORDING. The first version called a row with nothing surviving both controls inert. Two rows come out with identical arithmetic — 28 leaves against the first control, 0 against the second — and mean opposite things. ANGER POINT really is inert. SLUSH RUSH is live: against Swift Swim (inert in snow) Beartic's kill lands before the foe's stat drop; against SNOW CLOAK the drop misses, because evasion in snow turns a

100-accuracy move into a guaranteed miss under this pin — the same board by a different mechanism. **Two arms cannot separate "the subject did nothing" from "the subject and this control did the same thing"**, and Beartic has exactly three abilities so there is no third to ask with. Those rows are UNATTRIBUTABLE and say so; they are not inert and they are certainly not passes.

SIX ARE DECLARED UNTESTABLE, with the pool printed rather than asserted — Aroma Veil, Flower Veil, Fluffy, Imposter, Rain Dish, Solar Power. Every legal carrier of each has exactly one alternative ability and that alternative is live. The declaration is derived every run from the format, so a regulation change retires it without anybody remembering to.

A TIE-BREAK THAT WOULD HAVE BOUGHT TWO OF THOSE SIX WAS TRIED AND TAKEN BACK OUT. Ranking "carrier has a third ability" above bulk moves Rain Dish to Pelipper and Solar Power to Heliolisk — and it moved Water Absorb to Politoed, whose highest-ranked alternative is **DRIZZLE**, and Sand Rush and Sand Force to Excadrill, where the staging comes out inert in Showdown and the rows lose their coverage entirely. **Changing the fixture to suit the control is the wrong trade.** The second control is a measurement taken on whatever carrier the rule chose.

2. THE CLAUSE STATES ITS DENOMINATOR (#120). 84 of 316 is 26.6%; excluding the 115 rows whose ability has **no legal carrier in this format** — out of scope by regulation, not untested — it is 84 of 201 = **42%**. Neither number was on the line. The artifact now carries a `scope` block written at the refusal (`cannot(why, 'no-legal-carrier')`), tagged where the refusal happens, never matched out of prose afterwards) and `quarantine.js` reads it rather than re-deriving it. The clause reads `84 TESTED of 201 IN SCOPE`, of 316 total and **names the 18 unattributable rows in the text.**

An artifact predating the block says **DENOMINATOR NOT CARRIED** and names the re-run, rather than defaulting to zero — a missing count must not read as "none", which is the shape of the bug it replaces. `data/roster.items.json` and `data/roster.moves.json` are in that state now and are their owners' to re-run.

THE ONE JUDGEMENT CALL, stated so it can be overruled: an unattributable row is REPORTED and does not hold the gate shut. Six of the eighteen are untestable in this format by construction, and a clause that can never open is not a gate. They are named in the clause text on every run instead of sitting invisible inside a green. One `&&` in `rosterStage` changes that if Will wants it.

3. THE ABILITY STAGE CAN ASK FOR A SWITCH NOW, AND IT CLOSES ZERO ROWS (#122). No ability row had ever attempted one, so trapping abilities had never been asked the only question that distinguishes them from doing nothing. `ability/traps-and-somebody-tries-to-leave` reuses the moves stage's `switchVerdict` rather than building a second probe. **Asked of the format: Arena Trap and Magnet Pull have no legal carrier at all, and Shadow Tag's only carrier is Gengar-Mega — a forme whose ability the forme change WRITES, so it cannot be swapped and its only control is suppression, which `gastroWorks()` measures as dead in this simulator (6 leaves in Showdown, 0 here).** So the honest count of trapping rows closed is **zero**, and the rows say that instead of saying nothing.

The capability proves itself rather than being asserted, because a rule whose every member is COULD-NOT-STAGE never reaches `--reds` and a staging path that has never run is assumed broken: `abilitySwitchWorks()` plays the identical fixture with a NON-trapping carrier and requires both engines to complete the ask. It is green — the untrapped Kangaskhan leaves for Milotic in Showdown and in medicham2 — and it is a printed selftest line, so it can go red.

THE SEVEN NEW REDS ARE ENGINE DEFECTS AND ARE NOT FIXED HERE — filed to `@engine`, each with the second-control receipt that makes it attributable:

ability	verdict	what parted
Electromorphosis	DID-NOT-FIRE	Charge never applies — Showdown's Bellibolt hits 22 harder after being struck; ours does not move at all

ability	verdict	what parted
Ice Body	DID-NOT-FIRE	no hail/snow residual heal — Showdown 71 → 81 → 91 → 101, ours flat at 71
Curious Medicine	DID-NOT-FIRE	the ally's stat stages are not reset on entry (−1 Atk / −1 SpA stay)
Reckless	DID-NOT-FIRE	the recoil-move base-power boost is absent (Brave Bird 11 HP short, and the recoil with it)
Sweet Veil	DID-NOT-FIRE	the ally is not protected from sleep — Spore lands in ours and is refused in Showdown
Mirror Armor	FIRE-AND-BOARDS-DIFFER	the drop is not reflected: Showdown puts Scrafty at −1/−2 Atk and −1 SpA, ours leaves 0
Supersweet Syrup	FIRE-AND-BOARDS-DIFFER	evasion drops twice here and once in Showdown — an entry effect firing on both foes and again, or not once-per-battle

One process note. `data/roster.abilities.json` was rewritten (its bytes are at `.prev.json`); `data/roster.json` — the labelled convenience copy of whichever stage ran last — was also rewritten and previously mirrored the moves stage. Nothing is lost: `data/roster.moves.json` is that stage's artifact and `quarantine.js` reads the per-stage file first. `tests/roster.js --keep-shared` now exists so a division running beside another can leave that file alone.

oo. THE QUARANTINE IS A MECHANISM NOW, NOT A PARAGRAPH — 2026-08-08

`engine/quarantine.js`. Will's standing call: *"all engines that take medicham's output should be regarded as out of date and we should stop referencing them until medicham is up to date and we can rerun them."* CLAUDE.md states the rule; this file executes it, and `engine/status.js` no longer prints the figures it covers.

THE BUG IT CLOSES IS THIS DIVISION'S OWN. `status.js` has printed `PRE-CHANGE — measured` against a different build of: ... beside the leaf-calibration number for days, and the number went on being quoted anyway — by the sessions that printed the caption. That is the identical failure to a red gate reported for two days as "one of the two known failures". **A caption is not a quarantine. The figure is WITHHELD.**

The gate, today: 4 of 4 clauses FAIL.

clause	state
game differential	1 of 150 comparisons disagree with Showdown — <code>chesnaught woodhammer -> mimikyu</code> , showdown 0-0, medicham 120-130
deliberate roster / items	NO ARTIFACT — a missing stage is a FAILING clause
deliberate roster / abilities	2 FIRE-AND-BOARDS-DIFFER, 4 DID-NOT-FIRE
deliberate roster / moves	NO ARTIFACT

A MISSING STAGE FAILS. `tests/roster.js --write` writes `data/roster.json` whatever stage it ran, so the file holds only the newest one; reading it three times and calling that three stages would be "a capability was absent and everything reported success" inside the guard written to stop it. A stage counts only when an artifact's own `stage` field names it — `data/roster.<stage>.json`, `data/roster.all.json`, or `data/roster.json` when it matches.

Membership is derived from one root, and it is not a list of filenames. The PLAY LAYER is the transitive closure of *requires the simulator*, seeded with `engine/medicham2-browser.js` alone: 63 modules, `board.js` among them through `damageEngine()`. An artifact is quarantined if its generator is in that closure, if it reads a dump one of our own runs wrote, or if it reads a quarantined artifact. **34 of 114** artifacts are held. The transitive arm is what holds `mag.js`, `scoreboard.js`, `ladder.json` and `weight-multiplicity.json` behind `policy-weights.json`.

The strict direction is the dangerous one, and nothing that measures MEDICHAM is withheld. The census, the interaction matrix, the differential, the roster, the release ladder and everything OPS counts off the ingested store all print. Most fall out for free — they are written by `tests/`, or they drive the authority through a subprocess. Two do not and are **declared with a reason**, the `RAW-STORE-OK` convention: `game_differential.js` (MEDICHAM is its subject; it is the gate's own first clause) and `derive_protocol_events.js` (it loads the simulator only to read the event list it claims to emit, and quarantining it would have withheld the differential downstream of it). Both declarations are **checked** — an exemption naming a module no longer in the play layer fails the gate.

Shown RED before being trusted. The leaf-calibration withhold was removed by hand; `--check` exited 1 naming `data/winrate-backtest.json` and the leaked verdict sentence. And the negative was verified too: driving the real `withholder` with an open gate releases **34 of 34** and withholds none, so this can lift.

~~**What the graph still cannot see, stated rather than papered over.** `provenance.js` finds a writer only in `engine/` and `build/`, so ~50 artifacts written by `tests/` or through an unfollowed path variable have no row and are neither cleared nor withheld.~~ **CLOSED 2026-08-09 — see §00a.** The graph is 115 → 160 artifacts and the unknown set is 61 → 16. `data/rollout-r1-explore1.json` classifies on its own now and is HELD by its own derived reason, so the both-files workaround at `status.js:665` is retired-able. The instruments stayed clear and the consumers went behind the gate: **34 of 114** → **40 of 160** withheld.

Where a withheld number is still cited — measured, not remembered, and ratcheted in `data/quarantine-stamp.json` (may shrink, never grow): `docs/ENGINE.md`, `docs/MEASURE.md`, `docs/SEARCH.md` and `web/status-data.js`. Not edited from here; `web/` is WEB's and a gate its owner cannot satisfy becomes a known failure.

00a. ROADMAP #105 — FIFTY ARTIFACTS HAD NO ROW IN THE GRAPH, AND ONE OF THEM WAS THE ARM MILTANK RUNS — 2026-08-09

`engine/provenance.js` discovered an artifact's writer by looking for its literal name beside a write call in `engine/` or `build/`. Anything written by `tests/`, or through a path the scan did not follow, had **no row at all** — not `ok`, not UNSAFE, absent. Nothing could compare it to its source, which is the condition `CLAUDE.md`'s derived-artifact rule exists one level above.

	before	after
artifacts with a writer	115	160
artifacts with NO writer	61	16
withheld by <code>quarantine.js</code>	34 of 114	40 of 160
UNSAFE	13	20

Four arms, ranked by how directly each proves a write, and `via` **is now part of the graph** so a consumer can tell a write call from a sentence: `write` line (83), `path` variable (52), `path` template (12), near a `write` (1), `DECLARED BY THE ARTIFACT` (10). The scan reads `tests/` as well; a computed name such as `'roster.' + STAGE + '.json'`, ``exploitability-${TAG}.json` or `+'.meta.json'` becomes a regex with one wildcard per runtime value; and where no source can say — `fit_policy.js` takes its path from the `OUT_WEIGHTS` environment variable — the artifact's own `by`` is accepted, labelled as the weakest evidence, and only if the named script exists.

`data/rollout-r1-explore1.json` **CLASSIFIES ON ITS OWN and is HELD**, by the derived reason *"engine/rollout_r1_artifact.js reads rollout-r1-rows.jsonl — a dump of games MEDICHAM played"*. The both-files workaround at `status.js:665` — asking the quarantine about `rollout-r1.json` so the shipped arm could not slip through on a technicality — is a workaround for a hole that is now closed. **Not edited here; reported.**

Nothing was defaulted, and the strict direction held. The instruments print (`mechanics-census` , `engine-diff` , `interaction-matrix` , `roster.*` , `tag-walk` , `wire-ladder-census.pin` — all `ok`); the consumers went behind the gate (`exploitability-mag` , `exploitability-machamp` , `scoreboard` , `policy-weights-joint-presheet` , `ab-batch-effect` , `rollout-r1-explore1`). The **16 that remain UNKNOWN are printed every run, at zero as well as at sixteen**, because an empty list has to mean "every artifact has a writer" and never "we stopped looking". Seven are `policy-weights-*.json` variants written through `OUT_WEIGHTS` and declaring no `by`; `engine-release.json` is written through `writeJsonAtomic(S.pointer, ...)`, a helper no detector follows; `regulations.json` and `quality-filter.json` are CONFIG and correctly have no generator.

Four false attributions were found and closed on the way, three of them by this pass's own change. Each is written into the file beside the rule it produced.

- **A read** `open()` **on a line that later contains** `'w'` . `M=json.load(open('...'))`; ... `w=np.array(M['w'])` matched `open\s*(\.['"']wa)['"']` — so a test that only loads JOLTEON's weights outranked `engine/ditto.py`, which computes them, and flipped the artifact to not-store-derived. The mode string must now be an argument of the `open` call.
- **A write into a scratch tree is not a write into** `data/ . tests/test-miltank-release.js` writes `path.join(EMPTY, 'engine-release.json')` as a fixture and took ownership of the release pointer. A write now has to be rooted in `data/` — including through a helper whose own definition roots it, because requiring the literal word cost six correct rows on the first attempt.
- `writesNear` **never stripped comments, and this file's own new comment was its victim** — the example of what NOT to count credited `provenance.js` with generating the release pointer. The loose arm reads code now; the tight arms still read the source, because `build/build_browser_data.js` names its two targets in a trailing comment *deliberately* and stripping them took both artifacts back to "no generator".
- `data/regulations.json` **was credited to** `engine/analyze.js` , **which only reads it**. It has no generator: it is config, and `engine/conformance.js` already says so. One row lost, and the row was wrong.

A template match is CORROBORATED, not trusted. `exploitability-${TAG}.json` also reaches `data/exploitability-holdout.json`, which `exploit.js` did not write and cannot — it has no holdout mode. A template attribution must agree on top-level key shape with something the same generator writes by name, or it is revoked and the file stays UNKNOWN with that reason printed.

THE RATCHET GREW, AND THAT IS THE ONE JUDGEMENT CALL IN THIS PASS. `mtime_only` went **91 → 128**. None of the 38 regressed — 37 of them had no row for anything to ratchet, and a file that was never visible cannot have lost a stamp. The ratchet was measuring two different things: *"a generator shipped without recording what it read"* and *"this checker's coverage changed"*. They are opposite events and only the first is a fault. So the stamp now records `graph_files`, the diff splits REGRESSION from DISCOVERY, a regression still fails, and a discovery is printed in full and appended to `discoveries` in `data/provenance-stamp.json` with its date, reason and file list — the growth is permanent and auditable rather than laundered. The first run could not split (the old stamp carries no `graph_files`) and **says so instead of guessing**: an artifact-mtime heuristic was tried and accused five instruments other divisions had regenerated that afternoon. **Shown RED before being trusted** — dropping one file from the baseline while it stays in `graph_files` reproduces `RATCHET BROKEN`.

UNSAFE 13 → 20, and the eight movements are accounted for. Seven are newly visible and were always unsafe: `nature-arms.json` (older than `tags.json`, and `game_differential.js` moved under it) and six run sidecars pinned to superseded releases. The eighth was `quarantine-stamp.json`, which stamps `engine/provenance.js` by content — editing this file invalidated it by construction, and it returned to `ok` when `node engine/quarantine.js --check` re-ran. No artifact left the UNSAFE set.

Filed, not fixed:

- `engine/conformance.js` 's **S13 still hand-rolls the question**. It decides "no generator writes it" with `allSrc.includes(file)` over source text, so `data/roster.moves.prev.json` still trips it. The answer is derived now — `node engine/provenance.js --graph --json` carries by and via — and S13 should ask for it, the way `status.js` shells out rather than reimplementing staleness.
- `engine/rollout_r1_artifact.js` **writes the literal** `data/rollout-r1.json` **whatever dump it read**. The `explore=1` arm exists only because somebody renamed the output afterwards, which is why no pattern over that source can reach it and why the artifact's own `by` is the only witness. Derive `OUT` from `ROWS`, as `run_stamp.js` already derives its sidecar path.
- `readsNear` **has the same comment hole** `writesNear` **just had**. It decides `from`, and `from` decides staleness verdicts, so moving it changes what this tool SAYS rather than what it can SEE. Deliberately left; it belongs in a pass that re-derives the drift table.
- **ROADMAP #108 is easier to close but is not closed**. `status.js` printing a figure `provenance` calls UNSAFE is unchanged in kind — but every artifact `status.js` reads now HAS a row, so the two gates can finally be asked the same question about the same file. `status.js` is not edited here.

o. THE FORK IS DECIDED, AND THE ANSWER IS NO — 2026-08-07 (3.69.0)

A MORE CORRECT ENGINE DID NOT MAKE BETTER PREDICTIONS, AND NEITHER DEPTH METRIC PREDICTS LEAF ERROR. `engine/leaf_engine_contrast.js` → `data/leaf-engine-contrast.json`. Read every figure from the artifact; the rows are in `data/leaf-engine-contrast-rows.jsonl` so the curves can be re-cut without replaying 74 minutes of rollouts.

What was measured. MILTANK's live in-game leaf — `rolloutWinProb` at `n=200`, `explore=1.0`, `foePolicy` uniform, horizon 60, read from `miltank.js` `DEFAULTS` rather than retyped — on **8,883 positions**, the whole clean scorable corpus at the photograph, scored through **two frozen releases**:

	release	medicham2-browser.js	cut
BASELINE	<code>cf6a68fa412c</code>	<code>795be0c58cd7</code>	2026-08-07T02:34:46Z
TOP	<code>dc3c43336539</code>	<code>d10db6714fc9</code>	2026-08-07T17:27:09Z

The run **refuses to start** unless the two manifests differ in `engine/medicham2-browser.js` and in nothing else, and they do — same weights, same board, same damage table, same tag file, same Showdown commit `20ad99f`. Identical positions, identical per-position seeds. So the contrast is the simulator and nothing else.

1. THE TWO ENGINES ARE INDISTINGUISHABLE AT THE LEAF, AND THE NULL IS POWERED.

	TOP – BASELINE	95% CI	noise floor	detectable at 80% power
Brier	0.0000	[−0.0007, +0.0007]	0.000642	0.001013
log-loss	+0.0013	[−0.0007, +0.0033]	—	—

The confidence interval is **narrower than the smallest effect this n can detect**, so this is a tight null and not an underpowered one. McNemar on the 7,994 positions where both engines made a decisive call: **37 discordant for TOP, 36 for BASELINE**, $z = 0.117$, $p = 0.91$. The two leaves correlate $r = 0.9881$, mean $|\Delta p| = 0.0254$, max 0.175 — they are not the same function, they just score the same.

2. BOTH LEAVES ARE STILL WORSE THAN A COIN, ON A SAMPLE 6.4× THE PRIOR ONE.

Brier vs coin, paired: **+0.0325 [0.0281, 0.0372]** (TOP) and **+0.0325 [0.0281, 0.0371]** (BASELINE). Positive is worse. The 2026-08-04 held-out reading of **+0.0502 [0.0371, 0.0628]** is reproduced in direction and sign at `n = 1,778` held-out: **+0.0382 [0.0279, 0.0485]**.

3. DISCRIMINATION IS REAL ONLY IN-SAMPLE, AND IT SITS ON ITS OWN NOISE FLOOR.

On the full corpus the leaf names the winner on **52.48% of 8,320 decisive calls [51.40, 53.55]**, $p < 1e-4$ — and the split-half accuracy spread for that same arm is **2.49 points** against an effect of **2.48 points**. By

this division's own rule (LESSONS §9) that is not an effect. On the **held-out newest fifth it is 50.48%, p = 0.70** — no ranking at all, for both engines.

4. CALIBRATION IS THE FAILURE, AND IT IS A COMPRESSION. ECE **0.1514**, MCE 0.405. The reliability curve is monotone and almost flat: 88 points of predicted range map onto **13 points of observed range**.

leaf says	0.06	0.16	0.25	0.35	0.45	0.55	0.65	0.75	0.84	0.94
it wins	.466	.443	.494	.494	.498	.513	.528	.564	.544	.594

When it says 94% it wins 59%. A maximiser lives in that column.

5. AND THE JOINT ANSWER — NEITHER LINES NOR TURNS PREDICTS LEAF ERROR. Per-position leaf Brier against that position's first-divergence depth against the official simulator, on the same bodies the leaf rolls out. Spearman, bootstrapped over positions; **negative rho is the hypothesis**.

predictor	BASELINE engine	TOP engine
divergence depth in LINES	+0.0313 [0.0118, 0.0541] p=0.003	+0.0010 [-0.019, 0.022] p=0.92
divergence depth in TURNS	+0.0287 [0.0072, 0.0514] p=0.007	-0.0000 [-0.021, 0.023] p=1.00

MDE 0.0298 at this n. Under the engine that ships, **both are zero**. Under the baseline both are significant, **both have the WRONG SIGN** (more correct simulation → *larger* leaf error), and both sit essentially *at* the detection threshold. The sharpest form — **Δdepth against Δerror**, where every position-level confound constant across the two engines cancels — is **rho -0.0115 [-0.0307, +0.0082]** on 8,601 positions that parted in both.

6. THE TWO INSTRUMENTS THE PROJECT THOUGHT WERE DISAGREEING BEHAVE IDENTICALLY. 0.0313 against 0.0287; 0.0010 against -0.0000. **The turn metric is not degenerate on this sample** — 13 distinct values, 0 through 12, modal share 0.68 — so "turns cannot predict because it has no spread" is measured and rejected. Lines and turns are two readings of one thing, and neither reads on the leaf.

7. THE BINS ARE FLAT, AND THE BEST-FIDELITY BIN IS THE WORST. TOP engine, by line depth:

depth	0-4	5-9	10-14	15-19	20-29	30-49	50+	NEVER PARTED
n	130	1,421	2,165	1,505	1,474	1,126	788	246
mean Brier	.263	.280	.281	.286	.284	.290	.261	.321

The 246 positions where MEDICHAM matched the authority for the whole game have the **worst** leaf error and the only sub-chance accuracy (46.9%). At n=228 decisive that interval spans 0.5 and the claim is *no trend*, not *inverted*. The largest fidelity improvement available — **241 positions that used to part and now never part** — moved the leaf error by **+0.00213**, one standard error, in the wrong direction.

8. THE FIDELITY GAIN ITSELF IS REAL, AND IT REPLICATES THE LADDER ON AN INDEPENDENT SAMPLE. These are corpus positions, not the swarm's team pool, and the engine still improved: games that never part **13** → **246**, median first-divergence line **12** → **16**, games diverging 8,842/8,855 → 8,609/8,855. **Median completed turns: 1** → **1**. So the night's work was real. It simply does not reach the leaf.

9. AN INCIDENTAL, CONTROLLED SPEED NUMBER — the first this division has. §0a says three readings of engine throughput disagree by an order of magnitude and none is reproducible. This one is like-for-like by construction: identical positions, identical seeds, identical rollout budget, identical 6-shard layout, same machine, back to back. **One MILTANK in-game leaf call costs 1,478 s / 8,883 on the pre-**

WIRE-1 engine and 2,591 s / 8,883 on WIRE 10 — the shipping engine is 1.75× SLOWER per leaf call. It is not a battles/sec or turns/sec figure and must not be quoted as one; it is the cost of the thing the search actually spends its budget on. Filed to §0a rather than published as the missing artifact.

WHAT THIS MEANS, PLAINLY. Ten WIRE rungs made the simulator measurably more correct and bought **nothing** at the leaf, on the largest sample this division has ever run, with the null tighter than the smallest detectable effect. **Engine correctness is not what limits the leaf.** The leaf's failure is calibration — an 88-point predicted range compressed onto 13 observed points — and grinding the differential further cannot touch that. The remaining candidates are the ones §1 already lists and this measurement does not settle: the leaf is scored at **turn 0**, where a game-outcome label exists and where the position genuinely carries little information; and the **playout policy** is uniformly random at `explore=1.0`, which is a policy question rather than a mechanics one.

WHAT WOULD FALSIFY THIS. A depth metric that is not first-divergence — the differential stops at the first parting, so a position scoring 16 lines has an unmeasured remainder. If somebody builds a *cumulative* divergence measure and it predicts leaf error where these two do not, this section is wrong and should say so.

Four things about the run itself, because the run nearly went wrong twice.

- `engine/leaf_scoring.js` **is new, and it is not a second implementation.** It holds the Brier, log-loss, interval, reliability, ECE and noise-floor definitions that lived as private functions inside `backtest_winrate.js`. `node engine/leaf_scoring.js --verify` replays `data/winrate-backtest-rows.jsonl` and reproduces **749 of 749 scalars** of the published `data/winrate-backtest.json`, exactly. The generator refuses to run if that fails. `backtest_winrate.js` **was deliberately NOT edited to import it** — it is the generator of a published artifact, it has no `--out`, and it cannot be smoke-tested without overwriting the file everybody quotes. Filed below.
- **The first full run was killed by the harness at 65 minutes** with the baseline arm finished and on disk. Resume support was added and the arm recovered. The guard written while doing so **caught a real hole**: a reuse check on COUNTS passes when a re-derived sample has the same size and different members, and the store grew 8,887 → 9,003 during the run. It checks the id set now, and a resumed run reads the photograph rather than today's store.
- **A reused arm is REPRODUCED, not trusted.** 24 positions of each reused leaf arm are re-run by the current code and must be bit-identical; they were, for both engines.
- **The depth arm has no per-position reproduction, and that is a finding rather than an exemption.** The first version of that check refused the depth arm, 16 of 24 positions disagreeing. `game_differential`'s driver is coverage-seeking and its `CLICKS / COV_HITS` carry across games on purpose, so a divergence depth is a function of the position **and of every game played before it**. A 24-position slice starts from an empty click history and plays different games by construction. The substitute is the **reversed-order control**, which is why it was built: same release, same positions, driver history deliberately changed, **rho 0.836 [0.825, 0.846]** on 8,855 positions. That is the ceiling on every correlation in §5, and it is high — the nulls above are the world, not the ruler.

Filed from this pass, not fixed:

- `backtest_winrate.js` **should import** `engine/leaf_scoring.js`. One line, in the pass that next re-runs it. Two copies of a scoring rule is how `data/guru.js` came to say 0 where its source said 6.
- `provenance.js` **classes this artifact as OPEN-SHEET and it is a LADDER artifact.** The corpus detection reads one require hop deep, and `engine/game_differential.js` *mentions* `data/games.bo3.jsonl` in a comment — so a comment one file away picks the denominator. Drift is reported as 10.7% against the open-sheet ceiling where the ladder ceiling gives ~2%. Consequence is nil today because the POWER line correctly says the missing games move a proportion by at most 0.33 points, below the 0.43-point floor —

but the mechanism is the same one §5 records for `winrate-backtest.json`, recurring through comments rather than code.

- `docs/ABRA-whitepaper.md` 's **3.68.0 block quotes** `wire-ladder` **figures that the artifact does not carry** — "1,995 games per arm" and "mean 15.0 → 24.0" against `data/wire-ladder.json` 's 1,997 and 14.78 → 33.98. Another division was mid-pass on that file while this ran. Flagged, not edited.

ob. THE HEADLINE METRIC IS NOW EXPLOITABILITY, AND THIS DIVISION DOES NOT HAVE ONE — 2026-08-06 (3.59.0)

ADR-003 is accepted and published across the docs in this pass. **Exploitability replaces win rate as the project's headline metric, and the published comparator is VGC-Bench's approximately 100%.** That lands on this division, because the instrument is `engine/exploit.js` → `data/exploitability.json`, and that artifact is the one thing in the repository that is *declared void*.

Nothing here is a new measurement. The reframe rests on a number somebody else published and on a speed benchmark run earlier tonight. This entry records what the reframe costs MEASURE.

1. The headline metric has no value. `data/exploitability.json` is `void: true` and `provenance.js --strict` exits non-zero on it — §5g of this file has the full account. The 2026-07-26 figure was fitted on 17 features against the 58 shipped, on an engine 25 wire-fixes old, before the quality filter existed. The 2026-08-04 re-run is void because `data/policy-weights.json` — the defender — was refitted at 22:15:24 UTC while it was running. **So the comparison this project now leads with is a comparison it has set up and not made.** Say that, in those words, until it is made.

2. It raises the bar on the frozen-release discipline, and the evidence is our own. An exploitability run needs a best response trained against a *frozen* agent over thousands of games. The 2026-08-04 void **was** an exploitability run. That is a demonstrated failure mode on this exact measurement, not a hypothetical, and it is why the release boundary is a precondition rather than a nicety. `engine/exploit.js` stamps no engine digest and no digest of the target vector, which is why nothing caught it at the time.

3. The comparability argument is sound and should be stated whenever the number is. Their checkpoints are Reg M-A, ours Reg M-B, and their own paper shows policies do not transfer across team sets — a head-to-head is impossible. **Exploitability is intrinsic:** it is defined against a best response trained against *you*, in *your* format. So the two numbers sit on one scale although the two agents can never meet. This is the rare case where a comparator is legitimate without a shared population, and the reason has to travel with the figure or somebody will eventually read it as a head-to-head.

4. A noise floor for exploitability does not exist and is not obvious. §6 of this file says the floor belongs to the measurement rather than to a global constant. For a hill-climb the natural split-half does not apply — the arms are not exchangeable. The 2026-08-04 run gives the shape of the problem rather than a floor: it accepted **1 of 24** steps and its step scale decayed to 0.0168, so from round ~10 it was perturbing a near-copy of MAG. **An attack that dies in 58 dimensions returns an uninformative null on a still tree too**, which means a low exploitability figure from this tool is not yet distinguishable from the tool failing. That is the first thing to fix, and it is upstream of producing any number at all.

5. What this division must NOT do with the reframe. It must not report an exploitability number computed against a moving target, and it must not report an interim one. The SPRT rule applies with full force here: 66.7% became 44% and 57.7% became 50% in this project because somebody looked early. An adversarial search that is watched while it climbs is the same failure with a different name.

0a. ENGINE SPEED IS UNMEASURED BY THIS DIVISION'S STANDARDS, AND IT DECIDED AN ARCHITECTURE

Three readings of MEDICHAM's throughput are on record and they disagree by an order of magnitude: 3,401 battles/sec (ADR-001, July), 1,606 battles/sec (ROADMAP #61) and 13,041 turns/sec (the 3.59.0 re-measurement). The published ratio against `champions_sim` moves with them: **117x, corrected to 24.9x**. All three are now stated in ADR-001, ADR-002, `docs/MODELS.md`, the white paper, the deck, `docs/SUMMARY.md` and the technical docs, with the July figures kept and annotated rather than rewritten.

Everything wrong with this is a MEASURE problem, and none of it is an ENGINE problem.

- **No artifact holds any of the three.** There is no `data/*.json` for engine speed. A figure that decided the project's largest architectural decision has never been through the machinery every win rate in this repository goes through.
- **No generator exists.** There is no script in the repository that runs the comparison, so none of the three can be reproduced, and the two that disagree cannot be adjudicated.
- **No ratchet.** Nothing fails when engine speed regresses. ROADMAP #61 already recorded that a 2x regression went unseen for that reason; the same hole let a 4.7x error in a published ratio survive two weeks.
- **The unit was doing damage.** `turns/sec` compares the two engines and `battles/sec` does not, because MEDICHAM was driven to its 60-turn cap and Showdown with `choose('default')` to a natural end — so a "battle" is not the same quantity of work on the two sides. The 7.7x battles/sec ratio measured tonight is **not** a like-for-like number and must not be quoted as one. Two of the three historical readings are in the wrong unit for the comparison they were used to justify.

The correct fix is a stamped artifact with a declared method, not a fourth reading, and it is the same fix this division applied to the four rollout gates: a generator, a sidecar recording the machine and the configuration, and a floor. Filed here rather than done — it is a new measurement, and this pass was a publication pass.

One location is left uncorrected and is flagged rather than edited: `engine/champions_sim.js` lines 10 and 26 still state the 117x in the file header. It is ENGINE's file and ENGINE is working in the tree tonight.

1. LEAF CALIBRATION — MEASURED 2026-08-04. The leaf is not calibrated, and the claim is now powered.

`data/winrate-backtest.json` was re-derived against the current engine, on the whole clean corpus instead of a 350-game subsample, and it publishes a reliability curve. The finding is worse than the verdict string it replaces, and worse in a specific and actionable way.

The old number scored a leaf no live decision calls. It measured `winProb2` — `battle()` at MEDICHAM's default 20-turn horizon with entry effects re-fired. MILTANK calls neither: its team-preview leaf is a greedy payout at `maxTurns=60` with `seeded:true`, and its in-game leaf is `rollout_leaf.rolloutWinProb` at `explore=1.0 / foePolicy=uniform / maxTurns=60`. All three are now scored on identical positions, so the difference between them is about the leaf.

Confidence carries no information. The curve is close to a horizontal line:

leaf	says 0-10%	says 90-100%	discrimination	Brier vs coin (paired)
in-game, 200 rollouts, held-out n=1,378	wins 53.8% (n=52)	wins 53.6% (n=56)	50.99% [48.3, 53.7]	+0.0502 [0.0371, 0.0628]
preview, 40 rollouts, full clean n=6,886	wins 45.7% (n=831)	wins 55.3% (n=933)	53.22% [52.0, 54.4]	+0.0740 [0.0668, 0.0813]

Positive is worse. Both leaves are decisively **worse than a coin** on Brier and on log-loss, and worse than player-Elo, paired on the games where both have an opinion. The preview leaf puts **25.6% of all its predictions into the two extreme buckets**, where it is wrong by ~40 points.

Discrimination and calibration are separate failures needing separate answers. The preview leaf does rank — 53.22% on 6,700 decisive calls, $p < 1e-4$ — real, but only ~1.9 points above the split-half noise floor. The in-game leaf does not rank at all: 50.99%, $p = 0.47$. Randomising the payout bought variance and spent the signal.

What this is not. Nothing here says the engine is broken. The legacy `winProb2` leaf reproduces the 2026-08-02 number closely on the current engine — held-out log-loss **1.0243** against the **1.0748** published, discrimination **51.94%** against **52.63%** — so the twenty-two engine commits in between did not move the headline. Do not spend this finding on a mechanics hunt.

Also fixed here: the split was cutting on **store append order**, not date, and the store carries 4,775 date inversions. A side-symmetry witness scores 400 boards from both sides and reports $\text{mean}(p_1 + p_2 - 1) = -0.0099$, so no side advantage inside the engine is contaminating the result.

Still open, in order:

- the leaf is scored at **turn 0**, because that is where a game-outcome label exists. `rollout_r1.js` scores mid-game positions and is the other half of this; the two have never been read together.
- the **horizon** is the first suspect. `battleResult` falls back to bodies-then-HP whenever the payout does not finish, so a confident number can be a material count wearing a probability's clothes. That is a SEARCH change, not a MEASURE one — file it, do not fix it here.
- `data/winrate-backtest-rows.jsonl` holds the per-game predictions, so the curve can be re-cut without re-running the ~15 minutes of rollouts.

2. R4 has an artifact — CLOSED 2026-08-04

`engine/rollout_r4.js` writes `data/rollout-r4.json`, and `status.js` now prints the verdict out of that file instead of `NO ARTIFACT`. It does not re-pair anything: it shells out to `sprt.js --verify` and `paired_h2h.js` and refuses to write if they disagree, the same way `status.js` shells out to `provenance.js`.

The remembered 55.5% held. **ACCEPT H1 — MILTANK takes 55.5% of 535 DECISIVE PAIRS, decided after 522 of them, LLR 3.00 against a 2.94 bound.** The corpus is 5,248 lines, which is 2,624 games, which is 1,312 seed pairs — the store writes a log-only companion record under the same id, so a line count double-counts every game and the handoff's "5,248 games" was exactly twice the truth. The artifact records all four numbers and asserts the invariant that makes them relate.

Two things it is not. The point estimate is **stopped at a boundary**, so it is biased high and the 95% CI beside it is a fixed-n formula quoted for context, not the inference — read the verdict. And `status.js` classes the corpus **PRE-CHANGE**: `engine/medicham2-browser.js` moved 04:47, the games were played 04:41. Both arms shared the pre-fix rollout model, so the contrast is fair and the run stands as a measurement *of that build*; that the edge survives into HEAD is an assumption. It gets re-run at the next frozen engine release.

No A/A run exists for this comparison, so the noise floor is **not established**. The substitute in the artifact is three independent split-half cuts of this run: spreads of 0.2, 3.9 and 1.3 points against an effect of 5.5. One cut alone would have been useless — the spread of a single split-half is itself a draw with sd about 4.3 points at this sample size.

3. R1 has an artifact, and it does not say what the docs said — CLOSED 2026-08-04

`engine/rollout_r1_artifact.js` writes `data/rollout-r1.json` from the committed row dump, and `status.js` prints its verdict. The gate previously read a file of the same name that `engine/rollout_r1_join.py` wrote for the **withdrawn** cross-language join — nothing was hidden, the join prints its own withdrawal, but the gate

read it because it owned the filename. The join is now `data/rollout-r1-withdrawn-join.json` with `withdrawn: true`, and `status.js` refuses to print any artifact carrying that field.

The recomputation does not reproduce the published PASS. `docs/ROLLOUT-design.md` claimed 68.18% against material's 65.26%, +2.91 [1.79, 4.04]. From `data/rollout-r1-rows.jsonl` the same formulas give **65.72% against 65.26%, +0.46, 95% CI [-0.72, +1.63] — UNDECIDED** on 9,201 positions.

The material column matches the published figure to the digit, so it is the same sample. The rollout column reproduces §4.2.1's **greedy** calibration table bin-for-bin, so the surviving dump is the `explore=0` incumbent and the `explore=1` run that produced 68.18% left no file. That is the lesson, not the arithmetic: **the dump stamped no N, no explore and no build digest, so two runs four accuracy points apart were byte-indistinguishable.** `rollout_r1.js` now writes `data/rollout-r1-rows.meta.json` beside every dump. R2 and R3 still have the same hole.

The split-half spread of this run ranges 0.43 to 2.01 points against an effect of 0.46 — the effect is inside its own noise floor, which is an independent route to the same UNDECIDED.

Open consequence, filed to SEARCH: `--rollout-explore` defaults to 1.0 and `engine/rollout_leaf.js:147`, `engine/mag_bot.js:145` and `docs/MILTANK.md` all cite 68.18% as the reason. Re-running `EXPLORE_LIST=1 DUMP=rollout-r1-rows.jsonl node engine/rollout_r1.js` and then `node engine/rollout_r1_artifact.js` settles it, and R2 says the leaf is cheap.

SEARCH RAN IT, 2026-08-04, AND EVERY FIGURE IT PRODUCED IS QUARANTINED — withheld, not annotated. `data/rollout-r1-explore-sweep.json` and `data/rollout-r1-explore1.json`

are downstream of MEDICHAM: they are built from dumps of games MEDICHAM played, which is R1 leaf accuracy, named in CLAUDE.md's quarantine list. MEDICHAM is not correct — `node`

`engine/status.js` names the failing clauses. So the sample size, the judged accuracy, the lift over the material baseline, the paired interval against the greedy arm and the monotone-in-explore ladder are all absent here, and whether the published figure reproduces is a question this ledger currently may not answer either way. It becomes quotable again when the gate opens AND this is re-run: `node`

`engine/rollout_explore_sweep.js`. Three things this division should act on: 1. **The command in this section would have destroyed the evidence.** `DUMP` resolves under `data/`, so `DUMP=rollout-r1-`

`rows.jsonl` overwrites the committed greedy dump — the only record of the incumbent arm, committed for exactly that reason. SEARCH used a new filename. Worse, the sidecar path in `rollout_r1.js:338` is

the hardcoded literal `data/rollout-r1-rows.meta.json` whatever `DUMP` is set to, so it lands beside the wrong dump. `rollout_r1_artifact.js` rejects it on the name-and-row-count check, which is the check

working — but the fix is to derive the sidecar path from `DUMP`. 2. `status.js:229` **still prints the greedy arm as "R1 leaf accuracy"**. SEARCH did not overwrite `data/rollout-r1.json`, deliberately: it is this division's artifact, written hours earlier, and it is the only record of the incumbent. But the line now reads

UNDECIDED for a configuration the bot does not run. One line — point it at `rollout-r1-explore-sweep.json`, or print both arms. 3. `engine/mew.js` **exposes no** `--miltank-explore`. So the question this

settles — which playout JUDGES better — cannot be escalated to the one that matters, which playout WINS more. R4 was itself run at `explore=1.0` and cannot arbitrate its own setting. Two parsed flags on `mew.js` and the A/B becomes runnable. One hypothesis this division filed to SEARCH is **measured and rejected**: `battleResult` scoring bodies-then-HP on unfinished playouts is real but is not the mechanism.

Over 1.1M playouts, 99.5–99.8% end by an actual wipeout at every explore setting and at horizons 20 and 60; cap-hits are 0.2–0.5%. Exploration makes playouts longer (4.4 → 6.1 mean turns), not truncated. Filed to ENGINE as a latent hazard. The flat reliability curve in `data/winrate-backtest.json` needs another explanation — and note that on human corpus positions the `explore=1.0` leaf is **not** flat: its ECE is 0.104 with a monotone curve running 0.166 → 0.842, against 0.196 for greedy.

4. R2 and R3 stamp their configuration now — and R3's published number has no control

`engine/run_stamp.js` is one implementation of the sidecar `rollout_r1.js` hand-rolled inline, so the next gate cannot grow a second format. It writes `<artifact>.meta.json` — the same convention that makes `data/rollout-r1-rows.meta.json` describe `data/rollout-r1-rows.jsonl` — carrying N, explore, every knob including the ones left at a default, sha256 content digests of every source the gate reaches, the commit, and **whether the tree was dirty**. A clean commit id over a dirty tree is a lie of exactly the kind this exists to stop.

Both published numbers reproduce as arithmetic, and neither reproduction is worth much. That is the finding, and it is a different finding from R1's.

gate	published	recomputed from committed evidence	reproduces
R2	477 boards over 200 games; 5.83 ms median at n=10	the affordability table (K=3 → 0.47 s median, 1.75 s worst; K=4 → 1.49 s / 5.53 s) reproduces to the digit from <code>leafCostMs</code>	derived layer yes, base layer NOT CHECKABLE
R3	72.9% over 70 decisions (19 agreed, 20 skipped)	$100 \times (70 - 19) / 70 = 72.857142857142854$, bit-identical to the stored float	yes, and it is a tautology

R3's divergence is a pure function of two fields in the same file. There are no per-decision rows, so "it reproduces" means the artifact is internally consistent — nothing more. R2 dumps no per-leaf timing at all, and a duration cannot be recomputed by anyone in principle: it is a fact about a machine under a load, and nothing records the CPU, the node version or what else was running. **R2 is the one rung that is re-run or it is nothing.**

THE R3 RESULT IS NOT INTERPRETABLE AS PUBLISHED, and this outranks the sidecar work. `rollout_r3.js` computes the only control that makes a divergence rate mean anything — the same search on a different seed disagreeing with **itself**, where the truth is 0.00 by construction — and it `console.log`s it and does not write it. Its own verdict branches on that number: `rate <= floor` prints NOT A RESULT. So `data/rollout-r3.json` cannot say which branch its own run took.

`docs/ROLLOUT-design.md` §5 does publish floors — 71.7 / 50.0 / 45.5 / 43.8% — but **for four earlier runs, none of them this one**. At N=20 that floor measured *higher* than the divergence. The committed artifact is a fifth run at N=600 on 70 decisions, and its floor was printed to a terminal and lost. `engine/status.js` and `docs/MILTANK.md` both quote its 72.9%, and `MILTANK.md` spends it on a decision: "so it does diverge, and the equilibrium version is worth building."

Read plainly: **the divergence is probably real** — the doc's floor fell from 71.7% to 43.8% as N rose from 20 to 200, and this run used N=600, so its floor should be lower still. But *probably* is an inference from a different run, and the Wilson interval on 51/70 is **[61.5%, 81.9%]**, which is wide enough that a 44%-class floor is the only thing separating a result from an artefact of the argmax. The next run writes the floor; until one does, the 72.9% is a headline with its control missing.

A second defect, found on the way: `data/rollout-r3.json` 's own caveat is false about the run it describes. It reads "Switch candidates are excluded and counted". Commit `b4ec80b` deleted the `if (ca.switchTo || cb.switchTo) continue;` line — switches went **on** the menu, which is what that commit was *for* — and left the string alone. It has shipped that way since 2026-08-03, and the `withSwitch` / `choseSwitch` counters that commit added were printed and never written, so its own headline ("4 of 12 when one is on the menu") lives in a commit message.

R2 timed a leaf the bot does not run. `rollout_r2.js` called `RL.rolloutWinProb` without `explore` or `maxTurns`, inheriting `engine/rollout_leaf.js:197` 's `explore = 0` and `engine/medicham2-browser.js:1079` 's `maxTurns = 20`. MILTANK's in-game leaf is **explore=1.0 at maxTurns=60** — a randomised payout at

three times the horizon. That is R1's hole in cost form: two library defaults, written down nowhere, deciding the number. Both are now explicit, overridable and stamped, with defaults that preserve the old behaviour exactly so nothing re-dates the committed artifact by accident.

Also corrected in the generators, all of them visible in `status.js` :

- `games` was the `GAMES` environment **cap**, not a count. `status.js` printed "477 boards over 200 games", so an environment variable was being read as a measurement. It is now the distinct games actually traversed, with the cap beside it as `games_requested`.
- `leafCostMs` quantiles per N were computed over possibly-different board sets — a leaf returning null at one N and not another silently misaligns the columns — and only the n=10 count was recorded. `samples_per_n` now records all of them.
- R3's disagreement-gap median was computed twice, once to print and once to store. One variable now.
- `docs/ROLLOUT-design.md` §5's "roughly 200x the simulated turns per millisecond" is **155x** by the arithmetic of the two artifacts it cites (10×20 turns / 5.83 ms against 1 / 4.52 ms), and 155x is itself a ceiling because it assumes no playout ends early. `rollout_r2.js` now prints the division instead of a remembered figure. **The doc still says 200x** — see filed, below.

Two retrospective sidecars were written by `node engine/run_stamp.js --reconstruct`, which infers the build from the commit that carried the artifact and marks every field `reconstructed: true`. Both score HIGH: `data/rollout-cost.json` was written 25 s before `05248f2`, `data/rollout-r3.json` 159 s before `b4ec80b`. That is evidence about a commit, not a record of a run, and it says so on every line — a stamp that hashed today's sources would describe the file rather than the run, which is `data/rollout-r1.json`'s own stated reason for recording null.

Filed, not fixed:

- `data/rollout-cost.json` **should be** `data/rollout-r2.json`. It is the only rung whose file does not carry its gate's name. Four readers: `engine/status.js:230`, `web/build-status.js:200` and `:265`, and the generated `web/status-data.js`. Three of the four are under `web/`, which MEASURE does not own. A rename that misses a reader prints NOT DERIVED and reads as "nobody ran this", which is worse than the inconsistency. Needs WEB in the same pass.
- `~~ n / n_unit` **on R1 and R4.~~ DONE 2026-08-04** — see §7 below.
- `~~ engine/rollout_r1_join.py` **writes a naked** `isoformat()` **.~~ DONE 2026-08-04, and it was five files, not one** — see §8 below.
- `docs/ROLLOUT-design.md` **§5's 200x, and §R3's PASS**. Both are SEARCH's document and a SEARCH explore sweep is live. §5 should read 155x-at-most, and the R3 PASS should name which run it is quoting, because the floors in its table belong to runs that are not the committed artifact.
- `docs/MILTANK.md:70` **spends R3's 72.9% on a build decision** without its control. Same owner, same reason.
- `engine/rollout_r1.js` **should call** `engine/run_stamp.js` instead of its inline copy. SEARCH holds that file for the explore sweep. The shapes are identical today; two copies is how they stop being identical.

5. The possibly-stale artifacts, and the one class the checker could not see

`node engine/provenance.js` lists them; `node engine/provenance.js --graph` now prints the derived artifact graph itself, which is the part of that tool that could be silently wrong. Most entries are ordering artefacts inside a single run and are already annotated as such. The ones to actually chase are those older than `policy-weights.json`, those recording no game count at all, and — new — those carrying a **CORPUS DRIFT** note.

THE CANONICAL READER WAS HIDING ARTIFACTS FROM THE CHECKER. `provenance.js` derived an artifact's inputs by looking for a filename beside a read verb. A generator that loads the store the *recommended* way — `loadGames()` / `load_games()`, which resolve the path inside `engine/quality.js` and `engine/store.py` — never names `games.ladder.jsonl`, so it recorded **no dependency on the store at all** and was reported `ok` forever. Doing the right thing was the thing that made you invisible. Store derivation is now detected by the `LOADER CALL` (or an import of the reader), which is what a generator actually does.

Three more attribution faults surfaced with it, each of which had the same effect of exempting real artifacts from every corpus check:

- **A read** `open()` **looked like a write.** `pokemon-roles.json`, `role-matchups.json` and `roles-eval.json` were credited to `engine/build_roles_js.py`, which READS them to build a browser bundle, instead of `engine/roles.py`, which computes them from the store. `build_roles_js.py` touches no games, so all three were classed not-store-derived. A write test at line scope now requires a mode string.
- **The Python** `OUT = os.path.join(...)` **idiom hid a writer entirely.** `data/guru-matchups.json` — the source file at the centre of the `guru.js` divergence — had **no detected generator and was absent from the audit**. One level of variable indirection is now resolved.
- **Following into** `engine/quality.js` **classified everything as open-sheet.** `quality.js` names every store by construction, in its comments and in the error message that tells a caller how to pick one. `data/winrate-backtest.json`'s **ladder** games were being judged against the open-sheet ceiling. It is a named exception now, with that reason. The two corpus sizes that made the point are withheld with their artifacts: `data/winrate-backtest.json` is downstream of `MEDICHAM` (`engine/backtest_winrate.js` reaches `engine/medicham2-browser.js` through `require`), and it becomes quotable again when the gate opens AND this is re-run: `node engine/backtest_winrate.js`.

The graph went from 76 artifacts (49 store-derived) to **84 artifacts (57 store-derived)**. One of the eight newly visible files, `data/counters.json`, was **older than the quality filter** — the `UNSAFE` condition this tool exists to catch, invisible for nine days. Regenerated (15 s); `provenance.js --strict` is green.

5g. AND IT WAS HIDING SEVEN MORE — the write detection had the same hole in the other language

84 artifacts → 91, 57 store-derived → 60, and two of the seven were UNSAFE. Found 2026-08-04 while writing the missing `docs/MODELS.md` entries: the roster guard reported `data/move-priors.json` as generated by `engine/state_encoder.py`, which only **reads** it. Three defects, each the same shape as §5's and each found by chasing the previous one.

- **const broke the path-indirection arm.** §5 taught this file the Python idiom `OUT = os.path.join(...)` ... `json.dump(..., open(OUT, "w"))`. The JavaScript spelling is `const OUT = process.argv[3] || path.join(...)`, and the capture took the KEYWORD and then failed on the `=`. Every generator using it scored zero. The cost is the §5 cost exactly: `engine/policy.js` loads the store through `quality.js`, `state_encoder.py` opens no game file, so **the behaviour clone that nine files read was classed not-store-derived and exempt from every corpus check here**. It is 2026-07-31 vintage — the 5,269-game era — and nothing could say so.
- **A READ assignment is not a writer, and accepting** `const` **proved it immediately.** `const r = JSON.parse(fs.readFileSync(...'regulations.json'...))` followed later by an unrelated `fs.writeFileSync(file, r.body)` credited `engine/fetch_smogon_stats.js` with generating the format registry — a one-letter identifier matching `\br\b` inside any later write. An assignment whose own right-hand side is a read verb now never establishes a writer.
- **named() was a substring test, and the comment claiming that was safe was FALSE for the most-read file in the repository.** `ladder.json` is a substring of `games.ladder.jsonl`, so every

generator that opens the game store was recorded as naming `data/ladder.json` — which is how `engine/refresh-site-data.NOARCH.py` was credited with generating **MACHAMP's hill-climb artifact**, whose real writer is `engine/ladder.js` (the on-disk keys are `ladder.js` 's to the letter). The same fault hung a **phantom** `ladder.json` **input** on every store reader, and `roles.js` inside `pokemon-roles.json` did it again across eight more artifacts. Those show up as "older than its input" notes about dependencies that do not exist. An occurrence now only counts when the name is not the prefix of a longer one.

Corrected attributions, each verified against the generator rather than trusted: `move-priors.json` → `engine/policy.js`, `ladder.json` → `engine/ladder.js`, `dynamics.json` → `engine/dynamics.js`, `rollout-r4.json` → `engine/rollout_r4.js`, `smogon-priors.json` → `engine/smogon_priors.js`. Newly visible: `bring-bias.json`, `bring-priors.json`, `brood.json`, `core-matchups.json`, `exploitability.json`, `playstyle-matchups.json`, `smogon-priors-bo3.json`.

Two of the seven were UNSAFE, and the split between them is the point.

- `data/bring-priors.json` **was genuinely UNSAFE** — five minutes older than the quality filter, so computed under a different definition of which games count. It reads the store through `quality.js`. Regenerated (30 s), and it moved a figure a long way: `n_sides` **5,368** → **14,456**, and the format's **mega rate had been measured on 62 sides and is now measured on 12,442**, `p_side_megas` **0.9355** → **0.8785**, `p_mega_is_lead` **0.5345** → **0.5159**. `CLAUDE.md` sets a domain RATE floor there — "a game without a mega should be rare" — and the floor was being checked against 62 sides.
- `data/exploitability.json` **is a FALSE POSITIVE of the filter rule and a TRUE negative anyway, and it is left RED**. `engine/exploit.js` reads no game store at all — it plays self-play games from `policy-weights.json` — so the quality filter has no bearing on it, and the honest fix is a `not_store_derived` declaration, which only a re-run can write. **I did not add one, deliberately**. Stamping it would make a **genuinely unquotable** artifact look clean: it is PRIORITIES #18, WOBBUFFET's headline share — **withheld here, not annotated** — fitted on **17 features against the 53 we ship**, and it is rendered on `web/stadium.html` and `app/stadium.html` today. Re-running it is a ~4,000-game adversarial search against a mid-flight MAG, which the engine release boundary forbids.

So `node engine/provenance.js --strict` **exits 1 on one artifact, and** `tests/run-all.js` **gates on it. That is stated, not filed.** The artifact has been invalid since 2026-07-26; the only thing that changed today is that something can see it. It needs Will's call between re-running WOBBUFFET after the release boundary and pulling the number off the two stadium pages.

5a. CORPUS DRIFT — and the answer to "two definitions of clean games"

There are not two definitions. There is one, and the other number is four days old.

`data/live.js` and `data/winrate-backtest.json` said 6,943; `data/meta-usage.json`, `data/roles-eval.json` and `data/guru-matchups.json` said ~5,269. Measured rather than argued:

figure	written	what it is
6,943	2026-08-04 03:09	<code>load_games(clean=True)</code> over the store as it stood then
6,890	2026-08-04 02:52	the same predicate, 17 minutes earlier — <code>backtest_winrate.js</code> began its run then and the collector appended exactly 53 clean games while it ran
6,886	—	6,890 minus 4 games whose <code>winner</code> matches neither player's name. A genuinely narrower question, and it is already NAMED: <code>scorable</code> / <code>dropped_no_label</code>
5,269	2026-07-31 16:42	the same predicate over a store holding 29,117 collected instead of 38,587
5,265	2026-07-31 16:43	5,269 minus the same 4 unlabelable games

Collected grew $\times 1.325$ and clean grew $\times 1.318$ over those four days. A changed predicate does not scale with the corpus; a snapshot does. And `tests/test-quality.js` run tonight has the JS and Python readers selecting the **identical** 6,943 ids, sha `60aab8e1978e7554` on both sides. So renaming anything would have been wrong: **the defect was a date, not a word.**

Why nothing caught it: mtime cannot. The store is append-only and its mtime moves every hour, so an mtime rule would mark every store-derived artifact stale within an hour of being rebuilt — a gate that cries wolf.

`provenance.js` now compares the **declared count** against the clean corpus and warns past 10%, which the measured growth rate ($\sim 7\%$ /day) makes "roughly a day and a half behind". Thirteen artifacts are flagged, including all three named above at 24.1–24.2%.

Two supporting fixes, both of which were the reason the headline artifact escaped:

- `declaredGames` now reads an explicit corpus claim first — `provenance.funnel.clean`, `provenance.usable`, `corpus.clean_games`. `data/meta-usage.json` states its population more carefully than any other file in the repository and had **no key the checker looked at**, so the file that started this question was the one it could not see a count for.
- The drift check ignores a bare `games` key, and that is deliberate: `rollout_r2.js` published `games` as the GAMES environment **cap**. Until every writer says whether `games` is a corpus or a sample, a drift figure computed on it is a guess, and the fix belongs in the generator.

Do not touch `data/quality-filter.json` **to record the new funnel**. Its mtime is `FILTER_MT`; bumping it marks every older artifact UNSAFE and turns `--strict` red across the repository.

5b. `data/guru.js` **said 0 where** `data/guru-matchups.json` **said 6 — and both were misleading** `build/build_guru.js` read `g.decisive`. `engine/guru.py` writes the list as `decisive_matchups`. A missing key gave `[]`, the generator then recomputed `n_decisive` from **its own empty fallback**, and shipped a provenance note asserting "ZERO statistically-decisive matchups on this population" as though it were a finding. The 144-cell matrix was byte-identical throughout, which is why nobody noticed. `venusaurmega` / `venusaur-mega`, in a new pair of files.

The true value is 6 directed = 3 distinct matchups, and the generator carries the source's count now instead of recomputing it. Three things stop it recurring, all derived: every source key must be projected or named in `DELIBERATELY_UNUSED` with a reason; the source must agree with itself (`decisive_matchups.length === min(n_decisive, 20)`); and `build_guru.js --check` rebuilds the bundle in memory and diffs it, run by `tests/test-guru-derived.js` on every suite run.

And the measurement underneath it: 3 of 66 pairs is what chance produces. Each cell is its own 95% test. Over 66 unordered pairs the expected number clearing that bar with no real effect is **3.3**, and **3** clear it. The smallest exact two-sided binomial p-value in the matrix is **$6.1e-3$** against a Benjamini-Hochberg threshold of **$7.6e-4$** , so **zero survive FDR at $q=0.05$** and zero survive Bonferroni. The bundle publishes both counts (`n_decisive`, `n_decisive_corrected`) plus the arithmetic. The old file's "ZERO decisive" string was accidentally right and arrived there by a bug — which is worse than being wrong, because it cannot be checked.

`web/index.html:1845` gates a panel headed "*These are the matchups we can actually trust*" on `GURU.decisive.length`, so it will now render three matchups that do not survive multiplicity. It should read `decisive_corrected`. WEB's file; flagged, not edited. Note the same issue already affects `isSig()` in the matrix and the "statistically significant loop" claim, independently of this fix.

`data/guru-matchups.json` is itself 24.2% behind the corpus. Regenerating it is a separate, deliberate refresh — it moves every number the GURU booth renders — and is not done here.

5c. The thirteen drifting artifacts — TRIAGED, and NONE of them is a silent refresh

`node engine/provenance.js` flags thirteen artifacts 10.6–47.2% behind the clean corpus. The question that decides what to do with each is *does regenerating it move a published figure*, and it was measured rather than guessed: every scalar in each artifact was matched against the living docs and the site pages, at headline depth, with the universal constants (0.693, 0.25, 50%) excluded because they appear for reasons that have nothing to do with the artifact.

The answer is that there are zero safe silent refreshes in the set. Nine carry a verdict string or an interval-based claim that regenerating could flip; the other four have headline figures typed into `MODELS.md`, the white paper or `SUMMARY.md`, which `engine/sanity_check.py` §5 cross-checks. By this project's own living-docs rule, regenerating any of them is a docs pass, not a refresh.

artifact	behind	what regenerating moves	act
<code>war.json</code>	47.2%	verdict <i>"WORSE THAN A COIN AT EVERY REGULARISATION STRENGTH TESTED"</i> ; <code>held_out.log_loss</code> 0.694 in <code>MODELS</code> + white paper	STOP — a null on a corpus that has since doubled is the most interesting one here
<code>policy-eval.json</code>	43.8%	verdict <i>"phase-conditioning did not help; species-only prior retained"</i>	STOP
<code>pory-eval.json</code>	33.4%	<code>log_loss.pory</code> 0.6298 in white paper + <code>SUMMARY</code> , gated by <code>sanity_check</code>	STOP — restamped instead, see §5d
<code>pory-nn.json</code>	29.4%	Blast radius OVERSTATED in this row and corrected 2026-08-04. <code>val_logloss</code> and <code>auc</code> are not keys in the file — it holds an <code>arms</code> array with per-arm <code>logloss / acc / auc</code> . And the 71.6% in <code>MODELS.md</code> and the white paper is the policy clone's top-3 accuracy , a different measurement that happens to match. No living doc cites PORY-NN , so regenerating moves zero published figures. Regenerated	done
<code>xatu-belief.json</code>	29.3%	<code>n_games</code> 4,910 and <code>top1_accuracy.belief</code> 31.2% in <code>MODELS</code> ; an improvement CI clear of zero	STOP
<code>guru-matchups.json</code>	24.2%	every number the GURU booth renders; <code>log_loss_matchup_prior</code> 0.712 in the white paper	STOP — and explicitly not in this pass , WEB is in that booth
<code>roles-eval.json</code>	24.1%	headline <i>"0.6935 vs a coin 0.6931 and rating 0.6967"</i> — a knife-edge that regeneration can flip either way; six figures in <code>MODELS</code>	STOP
<code>pokemon-roles.json</code> , <code>role-matchups.json</code>	24.1%	same generator as <code>roles-eval</code> (<code>engine/roles.py</code>); all three move together or not at all	STOP
<code>vocab-usage.json</code>	24.1%	<code>role_coverage_of_battle_usage</code> 97.2% in <code>MODELS</code>	STOP (one-line docs pass)
<code>xatu-context.json</code>	24.1%	improvement CI [0.022, 0.042] rendered on the site	STOP
<code>meta-usage.json</code>	24.1%	nothing typed — the closest thing to a clean refresh, and PRIORITIES #16 names <code>node engine/analyze.js data/games.ladder.jsonl</code> as its closing command	ASK — <code>engine/mag_bot.js</code> and <code>engine/mew.js</code> read it, so it is the live bot's meta prior, and moving that is not MEASURE's call with an engine release boundary pending. It is not a refit trigger: <code>engine/feature_fixture.js</code>

artifact	behind	what regenerating moves	act
			excludes it by name and <code>board.js</code> never reads it
<code>counterplay.json</code>	10.6%	<code>result.mean_coverage_gap</code> 0.0321, CI [0.0086, 0.0563] — an interval that currently excludes zero	STOP

A false positive in the drift check itself, measured not argued. `pory-eval.json` is reported 33.4% behind, and it cannot be less than ~21% behind however often it is regenerated. Its population is not *clean ladder games*; it is *clean ladder games whose raw log is present and names a winner*, a strict subset. Running the generator over the whole current corpus reaches **5,456 games, not 6,943**, so its true drift is 15.3%. Every artifact reading `games.ladder.raw-logs.jsonl` has this. `provenance.js` 's existing escape hatches (`gate`, `games_requested`, `sampled`) do not cover it, because this is neither a gate nor a deliberate sample — the artifact needs to declare the ceiling its population can reach, in the same style. Not fixed here: `provenance.js` was built tonight and a second hand on its drift arithmetic is how two files come to disagree about one fact.

5d. PORY — the artifact restamped, and the coefficients were wrong for ten days

The verdict was not stale. The generator was answering the wrong question, correctly, every run. `data/pory-eval.json` still read "*a real, calibrated value net*" ten days after PORY was retracted, and `engine/pory.py` 's gate was `hi < coin` and `hi < material_heuristic` — which is TRUE on this sample (`hi` 0.6456 against 0.6931 and 0.6550). Restamping the file alone would have been undone by the next run. `material_heuristic` is a crude 0.75/0.25/0.5 **sign** rule; beating it is arithmetic.

The tie is now measured, not inferred. Against a logistic on [`alive_diff`, `hp_diff`] alone — same gradient descent, same standardisation, same temporal split — PORY scores **0.629799 to 0.629778**: paired difference **+0.000021 (POR Y worse)**, **95% CI [−0.000013, +0.000056]** clustered by game over 925 held-out games. On the current corpus (5,456 games) it is **−0.000001**, **CI [−0.000031, +0.000030]**. The retraction is robust to the corpus growth.

The reduction is structural, so no amount of data changes it. Every state is emitted from both perspectives with the label flipped, so the gradient on any column identical across the two rows cancels exactly: intercept and `turn/10` are pinned to `0.000000000`, not shrunk to it. `my_alive` and `foe_alive` swap and come back exactly antisymmetric (sum `0.000000000`). Five features, two degrees of freedom.

`engine/pory.py` **reproduces its own artifact bit-for-bit** — replayed on the identical first-4,623 clean-game sample it returned this file's weights, `feat_std` and log-loss exactly. So the fault was never the arithmetic. The gate now reads the paired difference, the withdrawn string travels under `withdrawn_verdict`, and `reduced_form` is derived from the file's own weights.

REGENERATED 2026-08-05 on the current corpus, deliberately (the dispatch Will approved). `data/pory-eval.json` now describes **5,883 games / 97,732 board-states** (was 4,623 / a 925-game test split) and **declares** `population_ceiling: 5883` — the §5f hatch, written by the generator on its first deliberate run since the hatch existed. Everything below survives the growth: paired difference vs the two-feature logistic **+0.000001**, **95% CI [−0.000026, +0.000029]** clustered over 1,177 held-out games; the verdict string is unchanged. The numbers a document would quote moved: log-loss **0.6298** → **0.6236** [0.607, 0.6387], sign-rule heuristic **0.655** → **0.6428**, two-feature baseline **0.629778** → **0.623623**, reduced form **0.9943 / 1.4080** → **0.9809 / 1.4093**, accuracy 0.6264 → 0.630, ECE 0.017 → 0.0138. Doc locations quoting the old figures are listed in §13b; propagation is the router's pass, not this file's.

The documented coefficients had no artifact behind them — the P1 class. 1.256 / 1.544 in docs/MODELS.md and web/stadium.html:342 is commit 44e0fb0 (2026-07-24, n_games 7,381), the last run fitted on the **unfiltered store with bot games in it.** 7f74236 put every model behind the clean filter on 2026-07-26 and the coefficients moved to 1.0259 / 1.4347, then 0.9946, 0.9962, **0.9943 / 1.4080**. The retraction has been citing bot-contaminated coefficients as its evidence ever since. MODELS.md is corrected with the history; web/stadium.html:342 **still says 1.256 — WEB's file, flagged not edited.**

5e. tests/test-site-data-fresh.js — two rules in it were wrong

It kept a second definition of stale, and it was the one provenance.js **had already rejected.** The verdict-input check compared each artifact's mtime against the newest games.*.json1 and failed past a day. The store is append-only and the collector runs hourly, so that clock cannot be beaten: five artifacts were red and **four of them are clean** by the canonical rule. It delegates to provenance.js now, the same way status.js does. The founding case survives the change — chomp-ev.json four days behind is ~28% drift, well past the 10% threshold.

What delegation loses is stated rather than dropped: drift can only see an artifact that declares a corpus, and chomp-ev.json, eval-report.json, policy-weights.json, policy-weights-joint.json and damage-validation.json declare none. They are **listed every run without failing**, so the pressure is on the generator to record a count — the shape tests/test-timestamps.js already uses.

--fix **would have refitted two models to make a freshness check go green.** The guard that stops it detected a publisher by the filename suffix -eval.json, which is not a property of anything. It caught engine/pory.py. It did **not** catch engine/nmf_roles.py (writes nmf-roles.json) or engine/xatu.py (writes xatu.json) — both fitted models quoted in MODELS.md, both on the auto-run list. The rule now is that a bundle writes only browser files and a generator that also writes a data/*.json is a publisher; checked against all ten generators this test names.

Seven of the ten stale bundles were regenerated. **Four were byte-identical apart from a date stamp** (mew.js, move-effects.js, mega-formes.js, status.js) and two entirely so (abra-meta.js, roles.js) — pure mtime, the check crying wolf. **Two had really rotted:** mag.js was serving standard errors from before the last weight change (0.02452 against policy-weights.json's 0.02363) and scoreboard.js was rendering superseded weights (1.1887 against 1.0884). That is the class this check exists for and it was real.

Also found: build_mag_data.js and build_scoreboard.js **crash without** SHOWDOWN_PATH, so --fix fails on them in any shell that has not exported it, and the test does not say so. Same shape as Po #40 — two ratchets that crashed rather than failed for the same reason.

Still red, not filed: data/pory-nn.json at **29.4% corpus drift.** The command is python engine/pory_nn.py; it is a neural-net train and it republishes val_logloss 0.612 and auc 71.6%, which MODELS.md, the white paper and SUMMARY.md all quote. That is a stop-and-ask, not a refresh.

5f. IS THE DRIFT THRESHOLD A TREADMILL? Yes, and the unit is wrong — DECIDED 2026-08-04

data/pory-nn.json was regenerated on the current corpus and tests/test-site-data-fresh.js immediately reported **CORPUS DRIFT 15.7% — declares 6,008, 7,123 clean now.** The store grew during the retrain. That is not a bug in either tool; it is what a percentage of an unbounded append-only corpus does.

The verdict: a percentage is the wrong unit, and it is not a fraction at all — it is an age. The collector runs hourly and clean games grew 5,269 → 7,123 in four days. For an artifact of age Δt , drift is $1 - n(t_0)/n(t)$, which depends only on elapsed time. A 10% threshold is therefore "about a day and a half old", stated in a unit that hides the fact that it is a clock — which is exactly what §5e removed from test-site-data-

fresh.js and put back through the front door. A freshly-regenerated artifact failing its own freshness check on the day it was made is the shape of a check that gets filed as *known*, and CLAUDE.md names normalisation, not invisibility, as how the docs-currency guard rotted.

The unit that answers the real question is absolute power, and provenance.js now prints it. Every drift note carries a POWER line beside it:

- ci_gain — how many percentage points narrower the 95% interval would be. Precision goes as $1/\sqrt{n}$, so $1.96 \times 0.5 \times (1/\sqrt{n_{dec}} - 1/\sqrt{n_{now}})$.
- max_shift — how far the pooled point estimate could move if every missing game arrived. The pooled mean shifts by $(m/n_{now})(\bar{x}_{new} - \bar{x}_{old})$ and $se(\bar{x}_{new}) = sd/\sqrt{m}$, so a 2sd bound is $2 \times 0.5 \times \sqrt{m} / n_{now}$. Worst case, not expected case.

Measured across all thirteen drifting artifacts, the percentages span 4× and the power spans 2×:

artifact	drift	missing	CI gain	max shift (2sd)
war.json	48.6%	3,460	0.46 pts	0.83 pts
policy-eval.json	45.2%	3,220	0.41	0.80
pory-eval.json	35.1%	2,500	0.28	0.70
xatu-belief.json	31.1%	2,213	0.24	0.66
guru-matchups.json	26.1%	1,858	0.19	0.61
roles-eval.json and family	26.0%	1,854	0.19	0.60
pory-nn.json	15.7%	1,115	0.10	0.47
counterplay.json	12.8%	915	0.08	0.42

No artifact in this repository has enough missing data to move a proportion by one percentage point. war.json is missing *half its corpus* and can move 0.83 points. The smallest split-half noise floor this division has published is **0.43 points** (R1's cuts run 0.43–2.01; R4's three run 0.2 / 1.3 / 3.9), so counterplay.json is already **below the noise floor** and the rest sit inside a factor of two of it. "24% behind" and "the games it lacks cannot move it past its own noise floor" are different statements and only the second one is actionable.

And it self-extinguishes, which is the property the percentage lacks. max_shift = $\sqrt{f}/\sqrt{n}$, so the same 15.7% drift that moves 0.47 points at n=7,123 moves 0.33 at n=14,000 and 0.24 at n=28,000. The treadmill stops on its own as the corpus grows instead of being switched off by hand.

What was NOT changed, deliberately: the 10% trigger. Two reasons, and the second is the honest limit of this work.

1. Lowering the bar changes thirteen artifacts' status and that is not a call to make inside a measurement pass.
2. max_shift **still cannot see the thing that decided every row of §5c's hand triage** — the DISTANCE from an artifact's headline estimate to its decision boundary. roles-eval.json publishes 0.6935 against a coin's 0.6931; that 0.0004 margin is flippable by any new data at all, while war.json's null is not flippable by 0.83 points. That margin is not computable from n , and the artifact is the only thing that knows it. max_shift is also stated for a **proportion** at sd = 0.5; a log-loss lives on another scale and the number is not directly comparable there. The next rung is a declared decision_margin, in the same convention as below.

A grace period measured in regenerations is rejected. It is a clock with extra steps, it cannot tell chomp-ev.json from pory-nn.json, and a fixed window is a licence for a genuinely flippable artifact to sit quiet inside it.

The `pory-eval.json` **false positive is fixed on the reader side and still needs its generator.**

`provenance.js` now honours a declared `population_ceiling` (`j.population_ceiling`, `provenance.population_ceiling` or `corpus.population_ceiling`) and measures drift against it — the same declaration convention as `not_store_derived`, `raw_store_ok`, `gate` and `games_requested`. `pory-eval.json` is a strict subset (clean ladder games whose raw log exists AND names a winner: 5,456, not 7,123), so it can never get below ~21% against the wrong denominator. **This is a separate defect from the unit question and the answer above does not fix it by itself:** the hatch exists, and `engine/pory.py` must write the key on its next deliberate run. **DONE 2026-08-05** — the deliberate run happened (§5d addendum) and the generator now writes `population_ceiling` with a note naming its predicate. Every artifact reading `games.ladder.raw-logs.jsonl` has the same shape.

WHAT THE NEW RULE WOULD AND WOULD NOT HAVE CAUGHT, plainly.

- `data/counters.json`, **older than the quality filter, UNSAFE for nine days — CAUGHT, and it was never a drift case.** That is the `FILTER_MT` check: the PREDICATE changed, so the artifact answers a different question, and no amount of power makes it valid. It is untouched, it is still `bad` rather than `warn`, and it still fails `--strict`. Confusing the two checks is how a volume rule gets credit for a correctness rule's catch.
- `data/chomp-ev.json` **four days behind, publishing "does not beat a coin" about a model with a directional edge — CAUGHT, and better than before.** That verdict sits *at* its boundary, so its decision margin is ≈ 0 and any `max_shift` exceeds it. The percentage caught it at 28% > 10%; the power rule catches it for the right reason.
- **An artifact recording a corpus it did not use — NOT CAUGHT, by either rule, and this file already says so on every run.** Only re-running the generator can.
- `data/slowking-playstyle.js`, **a GURU run written under the playstyle name — NOT CAUGHT by anything, and this is why the crude mtime rule in** `test-site-data-fresh.js` **was left alone.** See §9.

6. The noise floor is not a standing artifact

Split one arm in half and measure the spread. An effect smaller than that is not an effect. This gets re-derived by hand every time somebody needs it, which means it usually is not derived at all.

Two consumers now emit their own and neither is general: `rollout-r4.json` carries three split-half cuts of the H2H, and every block of `winrate-backtest.json` carries a `noise_floor` on Brier and on accuracy. That is the right shape — the floor belongs to the measurement, not to a global constant — but there is still no A/A run for the H2H, and a floor computed inside the arm being judged cannot see between-run variance.

7. All four rungs carry `n_measured` / `n_unit` — CLOSED 2026-08-04

`engine/rollout_r1_artifact.js` and `engine/rollout_r4.js` now write the pair R2 and R3 already carried, and both artifacts were regenerated from committed evidence (no rollouts): `data/rollout-r1.json` carries scored positions and `data/rollout-r4.json` carries decisive pairs; the counts themselves are

QUARANTINED — withheld, not annotated, because both artifacts are built from dumps of games MEDICHAM played and MEDICHAM is not correct (`node engine/status.js` names the failing clauses). What the slot is FOR does not depend on the numbers: choosing which of R4's four quantities goes in the common slot is the whole point of having one — the SPRT is computed on decisive pairs and nothing else, and a handoff that quoted a line count of a store that writes two lines per game is exactly the confusion the slot prevents. They become quotable again when the gate opens AND these are re-run: `node engine/rollout_r1_artifact.js` and `node engine/rollout_r4.js`.

Still **not** called `n`: `data/rollout-r3.json` has published `n` as the rollout BUDGET since 2026-08-03, and one key meaning a sample size in one rung and a budget in the next is worse than no common key.

`tests/test-rollout-gates.js` derives the rung list from the filenames `engine/rollout_r*.js` write, asserts every generator emits both keys, and then permits exactly one artifact state beyond "carries them": *its generator does, awaiting a re-run*. `data/rollout-cost.json` is in that state and cannot leave it here — it is a set of TIMINGS, and R2 is re-run or it is nothing. What the test forbids is the state that actually goes wrong: missing in the artifact **and** in the generator, which is nobody having done it.

8. Naive timestamps — CLOSED 2026-08-04, and it was FIVE writers, not one

`engine/rollout_r1_join.py` was the reported case. The real answer to "is one occurrence a typo or a pattern" is that `datetime.now().isoformat(timespec="seconds")` appeared in **five** generators — `rollout_r1_join.py`, `lookahead_bound.py`, `lookahead_clock_control.py`, `nmf_rank.py`, `polygon2.py` — which makes it the house style rather than a slip. Eight committed artifacts carry one, and all eight come from exactly those five.

Correct the diagnosis, not just the bug. JavaScript does not misparse it. ECMA-262 gives the two ISO forms opposite defaults — date-TIME with no offset is read as LOCAL, date-ONLY is read as UTC:

```
new Date('2026-08-03T04:14:10') -> 2026-08-03T08:14:10.000Z (local, this box is UTC-4)
new Date('2026-08-03')          -> 2026-08-03T00:00:00.000Z (UTC)
```

So the four-hour figure is the RENDERED string, not the parse, and on this machine the value round-trips. The defect is that the stamp means something different to every reader, that the two forms this project already uses side by side follow opposite rules, and that it is wrong by the reader's UTC offset the moment it is compared against a `Z` stamp — which is what every JavaScript writer here emits and what `status.js` and `provenance.js` exist to do.

`engine/isotime.py` is the single home (`utc_now()`, `utc_today()`); all five call it. `data/rollout-r1-withdrawn-join.json` is deliberately NOT regenerated — it is a withdrawn result kept so the withdrawal can be checked. `tests/test-timestamps.js` gates the WRITERS, asserts the two ISO forms really do disagree on the running machine rather than quoting a comment about it, and lists the artifacts still carrying a naive stamp without failing on them, because an artifact is fixed by re-running its generator and that pressure is how "KNOWN FAILURE" gets typed.

9. The two "stale bundles" — one was a no-op and the other was never the file it claims to be

Both were regenerated with the verify-before-trusting step first. That step is the entire finding.

`data/engine-data.js` — **BYTE-IDENTICAL. Nothing was landed.** `SHOWDOWN_PATH=... node build/rebuild_sets_from_sheets.js` reports 318 species, 195 rebuilt from real sheets, 123 left alone under 10 sheets, **materially changed 0, illegal abilities fixed 0**. Run with `--write` and diffed against a preserved copy: **identical to the byte**. The generator reproduces its own artifact and the 0.9-day staleness was `mtime` and nothing else.

The original `mtime` was then RESTORED, and that is the point of the entry. Writing the identical file moved `engine-data.js` forward and immediately turned `counterplay.json`, `scoreboard.js` and `winrate-backtest.json` — this division's own leaf-calibration artifact — into "older than its input `engine-data.js`". Three false staleness flags manufactured by a regeneration that changed nothing. A restamp with negative information content is still a restamp; `status.js` 's refit edge is hash-based (`feature_fixture --check`) and was never at risk, but `provenance.js` 's input-ordering rule is `mtime`-based and was.

REPAIRED AND LANDED 2026-08-04. The section below is kept as the diagnosis; this is the result. Run with **both** variables set — `TAG=playstyle MATRIX_FILE=data/playstyle-matchups.json` — every figure predicted below reproduces to the digit: `n_games` 5,265 → **2,860**, `archetypes` 12 → **8**, `mixture` → **Rain 0.8079 / Setup 0.1657 / FakeOutBalance 0.0255**, `greedy-Nash` 0.0409 [−0.0001,

0.1735] → **0.026 [-0.0001, 0.1498]**, uniform 0.0761 → **0.0338**, triples 1,320 → **336**, cycle → **TailwindOffense** → **Sand** → **TrickRoom** (legs on 40, 5 and 140 games, still supported: false), and **the verdict flips** to "no material exploitability gap ... close to transitive at this granularity." It reproduces itself on a second run, byte-identical, and is no longer byte-identical to data/slowking.js. **The GURU arm was re-run first and reproduces its own artifact bit-for-bit** — every shared key unchanged, which is what licensed trusting the playstyle run from the same code. **THE FIX IS THE DEFAULT, NOT THE FILE.** engine/slowking_preview.py now REFUSES to write a TAG-named artifact from the default matrix and prints the two-variable command. The rule is narrow on purpose: a TAG names a NON-default run, so TAG-set-with-default-matrix is the one combination that cannot mean anything; the ordinary GURU run and the correct playstyle run are both untouched. A relative MATRIX_FILE now resolves against the repo rather than the shell's cwd — the documented command only worked from the repository root, and a path that works from one directory and not another is how the wrong matrix gets reached for. **A second half of the same bug, found on the way:** source_matrix was the hardcoded literal "data/guru-matchups.json". So even a CORRECT playstyle run would have stamped the GURU matrix as its source — the one field that could have exposed the clobber was pinned to agree with it. It is derived now, and tag is recorded beside it. **What moves on the page** (app/index.html:907-923 / web/index.html , WEB's files, not edited): games 5,265 → 2,860; the cycle legend from three species pairs to TailwindOffense → Sand → TrickRoom; leg edge 10% → 5%; the mixture chips from Gengar-Incineroar 66% / Charizard-Garchomp 22% / Pelipper-Archaludon 12% to **Rain 81% / Setup 17% / FakeOutBalance 3%**; greedy exploitability 4% → 3%. **And two TYPED literals in that paragraph are now wrong on both pages** — "these matchups rest on 49, 37 and 15 games" (really 40, 5 and 140) and "the strongest of 1,320 candidate triples" (really 336). **Worse than the numbers: the room's whole thesis is now contradicted by its own artifact.** The panel is headed "The meta looks like rock-paper-scissors" and argues "picking one single playstyle is exploitable while a mixture isn't — the reason to mix", while the artifact it renders now says mixing buys little here. That is a WEB pass, and it is a rewrite rather than a number swap. **Three consumers checked.** engine/sanity_check.py passes, and its \$5 check "site mixture top == report top" now reads **Rain == Rain**; before the repair it compared two copies of the same wrong file and passed for that reason. tests/test-docs-current.js \$1b likewise now reads the real playstyle artifact (0.026, CI[-0.0001, 0.1498]) where it had been reading GURU's numbers under the playstyle name — a guard built to track this artifact was tracking the other one. And engine/build-status.js:18 reads slowking-playstyle-eval.json into a variable ex that **nothing in the file ever uses**; it is a consumer in name only.

data/slowking-playstyle.js — **STOP. It is not stale; it is the wrong file, and has been since 2026-08-03 15:15.**

engine/slowking_preview.py takes its OUTPUT NAME from TAG and its MATRIX from MATRIX_FILE, which defaults to data/guru-matchups.json. Run with TAG=playstyle and MATRIX_FILE unset it writes a GURU result under the playstyle name. Measured:

- data/slowking-playstyle.js has a payload **byte-identical** to data/slowking.js;
- data/slowking-playstyle-eval.json is a **byte-identical file** to data/slowking-eval.json;
- both read 5,265 games / 12 species-pair archetypes / 1,320 candidate triples — GURU's shape.
data/playstyle-matchups.json holds **2,860 games over 8 playstyles.**

Regenerating it correctly moves published figures, so it was **restored and not landed**: n_games 5,265 → **2,860**, archetypes 12 → **8**, mixture Gengar-Incineroar 0.66 / Charizard-Garchomp 0.22 / Pelipper-Archaludon 0.12 → **Rain 0.81 / Setup 0.17 / FakeOutBalance 0.03**, greedy-Nash 0.0409 [-0.0001, 0.1735] → **0.026 [-0.0001, 0.1498]**, uniform 0.0761 → **0.0338**, cycle Charizard-Garchomp→Kingambit-Garchomp→Incineroar-Whimsicott → **TailwindOffense**→**Sand**→**TrickRoom**, triples searched 1,320 →

336, and the verdict string flips from *"substantially less exploitable... the meta is non-transitive here (rock-paper-scissors)"* to *"no material exploitability gap between Nash and greedy — this meta is close to transitive at this granularity."*

The corrected numbers are the ones `docs/MODELS.md` **already publishes** in the MACHAMP entry — 336 triples, a leg on 5 games, 0.026, `[-0.0001, 0.1498]` — to the digit. So the docs are right and the artifact is wrong, which is the rare direction, and the 2026-08-02 withdrawal of the SLOWKING cycle still rests on measured evidence. `engine/build-status.js:18` and `engine/sanity_check.py:32` both read the clobbered file, and `app/index.html:907-923` renders its mixture and cycle legs. The repair is one command with **both** variables set; it is a WEB pass, not a refresh. The generator should refuse to write a `TAG`-named artifact from the default matrix.

Two things this costs the checkers, and both are recorded rather than fixed here.

- `provenance.js` reports `slowking-playstyle.js` as `ok`, correctly by its own rules — it is co-generated with `slowking.js`, so the ordering carries no information. Provenance sees ordering and declared counts; it cannot see that a file's CONTENT came from the wrong input. This is why the crude mtime bundle rule in `tests/test-site-data-fresh.js` was **left alone** despite §5f: delegating it to drift today would have marked this file clean.
- `tests/test-site-data-fresh.js` printed the repair as `node engine/slowking_preview.py` — the wrong interpreter, in the STALE table only; the `--list` path already derived it from the extension. Fixed. The command it names is still incomplete, which the test's own comments already admit, and running it with `TAG` set and `MATRIX_FILE` unset is a plausible route to the clobber that is on disk.

10. `train_value.py` was discarding a fifth of the corpus — PRIORITIES #13, FIXED 2026-08-04

`idn()` normalises punctuation and nothing else, so the event stream's `charizardmegay` never matched the bring list's `charizard`, `side_of` returned None, and the event was thrown away without a word. **Measured on 4,000 clean games before the fix: 21.7% of faints, 22.7% of damaging events, 20.8% of all damage, at least one discard in 96.5% of games, and 97.6% of discarded targets are megas.** The visible symptom is the one to remember: **88.9% of clean games ENDED with both sides still holding bodies.** The value net was learning from trajectories in which almost nobody ever loses their team.

The fix is a verb on the one resolver, not a fourth copy of it. `engine/mc_key.js` gained `mcKey.base` (which body is this) and `mcKey.bases` (the whole map, for a caller that cannot call into JavaScript per name), reading the `base` field the generator already wrote from the dex into `MC.mons`. Three properties were deliberate:

- **It is not a string strip.** `re.sub(r'mega[xy]?$', '', s)` is the obvious three lines and it is wrong — `mc_key.js` already records that the identical JavaScript strip answered Victreebel for Victreebel-Mega. The table carries the answer; this reads it.
- **It returns a flat BODY id, not a table key.** `MC.mons` holds `floette-mega` with `base: "floette"` and holds no `floette` row, so a version resolving the base back through the table returned null for exactly the 1,613 events this exists to rescue, one layer down. Being in our damage table is a fact about our table; the body is a fact about the game.
- **It touches no dex, so it cannot need `SHOWDOWN_PATH`** — the crash-instead-of-fail mode PRIORITIES #40 records for two other ratchets. `train_value.py` shells it once per run and **fails loudly** if it cannot, rather than reverting to the behaviour above.

After: 22.7% → 1.7% of damaging events dropped, 21.7% → 1.5% of faints, and games ending with both sides intact 88.9% → 26.3%.

What it moved, and the honest size of it. Paired on identical held-out states — 1,445 games, 10,120 states, both arms fitted on the same split:

	before	after	paired difference
log-loss	0.6634	0.6520	-0.0114, 95% CI [-0.0183, -0.0041] (bootstrap clustered by game)
accuracy	59.72%	61.47%	+1.75 pts, 95% CI [0.50, 2.94]

The mechanism is legible in the weights: `hpDiff` moved **0.169** → **0.377**, because a fifth of all damage had never been applied and the feature was attenuated toward zero. The shipped artifact (`ladder + self-play`, as `main()` defaults) moved `test_logloss` **0.6638** → **0.6536**.

Both intervals clear zero and the effect is still inside the noise floor for an unpaired comparison. Twenty split-half cuts of the fixed arm alone spread by a **median 1.87 accuracy points** (range 0.30–5.37) against a 1.75-point effect. The pairing is what buys the resolution; two runs on different samples could not tell these value nets apart. And the ceiling is unchanged — 61.4% sits below the **66.92%** in-sample ceiling for this feature class and below the live leaf's **67.97%**. **This is a correctness fix, not a capability change, and it was worth making on the first ground alone.**

Residual, measured rather than assumed: 1.7% still drops, and 1,613 of the 1,625 are one species. `MC.mons` carries `floette-mega` → `floette`, the store's bring lists hold `floetteeternal`, and no `floette` row exists — the chain does not close. The rest are in-battle formes the mega table does not cover (`mimikyubusted` 274, `morpekohangry` 48, `castformsnowy / castformrainy` 11). **Filed to ENGINE, not patched here:** closing them means reaching for the Showdown dex, which would make `train_value.py` produce different numbers depending on whether `SHOWDOWN_PATH` is set. That is *fitting environment and playing environment must match* in a new place, and it is not worth 1.5% of events.

11. THE THIRD COPY OF THE WEATHER MAP IS IN `board.js`, IT IS WRONG, AND THE FIXTURE CANNOT SEE IT

LANDED 2026-08-04, with the two fixture boards it needed, and refitted. The diagnosis below is kept because it is the reason the fixture grew; the result is here. `engine/board.js` no longer keeps a weather map. Both reads — `dmgFractions` and the `punishExposure` call in `featuresFor` — go through a `weatherKind(board, D)` helper that calls the damage engine's exported `weatherId`, the same consolidation ENGINE made in `tag_dex.js`. A damage engine that cannot answer is counted in `dmgFailures.weatherUntranslated` rather than defaulted to clear skies; measured over 234,873 candidate vectors it is **0**, on both the node and the browser export path (`tests/test-board-browser.js` : 58 of 58 features agree to 6 dp). **THE PRE-LANDING MEASUREMENT REPRODUCES ON THE CURRENT ENGINE TO THE ROW.** Re-run before relying on it, because the engine had moved underneath it: `fit_policy.decisionsFor` over the first 1,200 open-sheet corpus games, **32,054 decisions / 234,873 candidate vectors**, one process holding both builds of `board.js` : | | measured 2026-08-04 (pre) | re-measured after the ENGINE band | |---|---|---| | candidate vectors that move | 1,768 (0.75%) | **1,768 (0.75%)** | | decisions that move | 892 of 32,054 (2.78%) | **892 (2.78%)** | | columns that move | 14 of 58 | **14 of 58** | | games containing a moved vector | not measured | **238 of 1,200 (19.83%)** | The 19.83% is the new number and it is the one that matches the census: 18.5% of games carry sand or snow at some turn. **THE FIXTURE NOW SEES IT — 0 columns before, 10 after.** `sand-is-up` and `snow-is-up` join `SCENARIOS`. `Tyranitar` takes special hits under sand (the Rock special-defence 1.5x) and `Ninetales-Alola` and `Weavile` take physical hits under snow (the Ice defence 1.5x); `Hippowdon` and `Ninetales-Alola` both carry Weather Ball, whose TYPE resolves off the same field, so a translation that mapped one weather and not the other cannot pass both. A Rock body and an Ice body sit on each bench, because the switch family prices a body that is not on the field and would otherwise be untouched — that one choice took the detection from 6 columns to 10. Scored against the pre-fix map:

koTarget , dmgFrac , killIsRoll , killsThreat , koFirst , switchSurvives1 , switchKOFast , switchDiesFirst , benchRisk , and the joint partnerCoversMe . **State the limit rather than the win: the fixture catches 10 of the 14 columns the corpus moves.** protectThreatened , diesBeforeMoving , screenValue and switchKOSlow move on corpus boards and not on these ten. A fixture is evidence for the restamp rule, never proof of it, and this is the measured size of the gap rather than a caveat in prose. Coverage did not regress: 40 slots, 324 candidates, 1,309 pairs, and **o features that never fire**. Two properties of the landing worth keeping. weatherId still does not know desolateland / primordialsea ; on 339,483 corpus turn-boards that costs nothing and adding them is ENGINE's call on SD2WEATHER , not a fourth table here. And adding scenarios re-stamps every hash, so it went in the same pass as the refit — a fixture change and a restamp cannot be separated.

This is a refit trigger. It is measured, it is NOT landed, and the gate for landing it is the Po/P1 band, which is not met.

engine/board.js:1190 carries WEATHER_KIND , a third private copy of the Showdown-weather → engine-weather translation that medicham2-browser.js owns as SD2WEATHER / weatherId . It is read at two sites — dmgFractions (:1247 , every damage-derived MAG feature) and the punishExposure call in featuresFor (:2937 , clickCost). What it holds:

```
{ sunnyday: 'sun', desolateland: 'sun', raindance: 'rain', primordialsea: 'rain' }
```

It maps the two weathers this format cannot produce and misses the two it does. desolateland and primordialsea are primal weather: **0 occurrences in 339,483 corpus turn-boards.** sandstorm and snowscape are not in the table at all, so WEATHER_KIND[board.weather] is undefined , || '' makes it clear skies, and **every damage feature under sand or snow has been computed in no weather**. The engine reads field.weather === 'sand' for the Rock special-defence 1.5×, === 'snow' for the Ice defence 1.5×, and Weather Ball's type comes off the same field.

Exposure, re-measured rather than inherited — a census of board.weather at every turn-board across games.ladder + games.bo3 + games.ots (52,441 games):

weather at a turn-board	share	WEATHER_KIND gives
clear	64.15%	'' correct
sunnyday	14.90%	'sun' correct
raindance	10.23%	'rain' correct
snowscape	5.43%	'' — WRONG
sandstorm	5.29%	'' — WRONG

10.72% of turn-boards, and 18.5% of games contain at least one.

THE FIXTURE PASSES ANYWAY, AND THAT IS THE FINDING. The patch was applied, measured and reverted. engine/feature_fixture.js --check returns feature semantics OK on both policy-weights.json and policy-weights-joint.json **before and after** — because the fixture's only two weather scenarios are RainDance and SunnyDay , the two WEATHER_KIND already gets right. All 58 columns are hash-identical while the feature function has moved.

What actually moves, measured on the fit's own rows — fit_policy.decisionsFor over the first 1,200 open-sheet corpus games, **32,054 decisions / 234,873 candidate vectors**:

candidate vectors that move	1,768 (0.75%)
decisions that move	892 of 32,054 (2.78%)
columns that move	14 of 58

`dmgFrac (1,182)`, `koTarget (238, max |Δ| 0.93)`, `killIsRoll`, `killsThreat`, `diesBeforeMoving`, `benchRisk`, `koFirst`, `protectThreatened`, `switchKOFast`, `switchKOSlow`, `screenValue`, `switchDiesFirst`, `switchSurvives1`, `switchSurvives2` — several flipping a full 0→1.

So `feature_fixture --check` is necessary and not sufficient, and `status.js`'s `refit edge: CLEAN` inherits that limit. The hash covers the boards the fixture builds; the feature FUNCTION lives over every board the corpus contains. This division's own rule — *a restamp is only valid if the feature FUNCTION is unchanged* — is the binding one, and a green fixture is evidence for it, not proof of it. **The fixture needs a sand board and a snow board.** Adding them is itself a fixture change that re-stamps every hash, so it belongs in the same pass as the refit, not before it.

The patch is small and is recorded here rather than left on disk: replace both reads with a `weatherKind(board, D)` helper that calls the damage engine's exported `weatherId`. Do not write a fourth map. Note that `weatherId` does not know `desolateland / primordialsea`; on the measured corpus that costs nothing, and adding them is ENGINE's call on `SD2WEATHER`, not a second table here.

11a. `position_features.js` — the same defect at the second boundary. LANDED, no refit owed.

`engine/position_features.js:292` built its field object as `B.norm(board.weather || '')` — the board's *move name* — and handed it to `M.dmgRange` and `M.effSpeed`, which compare against the engine's words. Truthy and meaningless, exactly as at the leaf boundary. Now `M.weatherId(...)`.

No refit is owed and that was checked, not assumed: nothing in the repository fits, reads or renders these columns. The only callers of `positionFeatures` are four tests, and no artifact contains any of its feature names.

Exposure re-measured on the boards this module is actually asked about — the `joint_rows.build` walk, 400 open-sheet games, **3,202 mid-game boards / 6,404 scored positions. 49.94% carry a weather**, higher than the 35.85% turn-board figure because weather accumulates as a game runs.

- **1,962 of 6,404 positions moved (30.64%):** sun 73.6%, rain 66.7%, sand 49.5%, snow 29.5%.
- 7 of 16 columns: `raceEdge` (29.67%, max |Δ| 0.42), `killFirstEdge`, `iKillNext`, `theyKillNext`, `benchAnswersDiff`, `speedEdge`, `pinnedDiff`.
- **0 of 3,206 clear-weather positions moved** — the control that says this is the weather.

`:296`'s `terrain` now goes through `M.terrainId` as well. That one is a **confirmed no-op**: 16 of the 3,202 boards carry a terrain, 11 of them under clear weather, and all 22 positions scored on those are bit-identical. Every downstream reader already calls `terrainId` itself. It is translated here so the field object leaves in one vocabulary rather than two, which is the condition that let the weather half go unnoticed. The four probed keys are a list of BOARD KEYS, not a translation table — the same justification `rollout_leaf.terrainOnBoard` records.

ENGINE's "*o of 400 boards*" for this file was terrain-only and is reproduced (0.50% here). It was being read as though it covered the weather too, and the weather number is 49.94%.

11b. `git checkout -- <file>` INVALIDATES EVERY CONTENT STAMP ON THIS MACHINE

Found by doing it. `core.autocrlf=true`, the committed blobs are LF, and the worktree files are LF — so a checkout REWRITES them as CRLF. The bytes change, nothing in git notices (`git diff` is empty), and `sha256(worktree)` moves. `winrate-backtest.json`'s `measured_against` and `run_stamp.js`'s `source_digests` both hash worktree bytes, so `engine/board.js` immediately began printing `PRE-CHANGE` — measured against a different build of: ... `engine/board.js` after a revert that changed no code. Converted back to LF; the digest returns to `bcf2dab9dc6f`, which is the stamp exactly, and `board.js` drops out of the `PRE-CHANGE` list.

This is the mirror of the warning already in *Reading a stamp — never compare source digests to git.blobs*, they differ by line-ending translation. The new half is that an ordinary git operation can move one of them. **After any git checkout -- or git stash pop on this box, check node engine/status.js before believing a PRE-CHANGE line.**

12. tests/test-no-silent-failure.js — the 32 MEASURE entries, and what two of them were hiding

Worked through without --update, which would have laundered the SEARCH, OPS and WEB entries in the same command. **NEW since the baseline: 52 → 20.** Every MEASURE-owned entry is cleared; the 20 that remain are 13 SEARCH (miltank.js, rollout_leaf.js, rollout_r1.js), 3 OPS (mag_bot.js) and 4 in tests/test-{site-data-fresh, stadium-roster, web-parses}.js.

Two were hiding something real.

engine/run_stamp.js:92 recorded "the tree was clean" whenever git refused to answer.

```
const porcelain = git(['status', '--porcelain', '--'].concat(watched)) || '';
```

git() returns null when the command throws — an index lock, an interrupted rebase (CLAUDE.md documents this repository reaching one 43 commits into a 45-commit replay), git not on PATH. An EMPTY porcelain is git's way of saying CLEAN. || '' collapsed the two, so a stamp written while git was unavailable published dirty: false beside a commit id that described nothing on disk — and this module's own header says "a clean commit id over a dirty tree is a lie of exactly the kind this module exists to stop", while docs/MEASURE.md's *Reading a stamp* tells every reader to trust the commit when dirty is false. rev-parse HEAD was already guarded; status --porcelain was not. Now a third state: dirty: null with git_errors, and status.js renders it as **DIRTINESS UNKNOWN** rather than as the clean case.

engine/backtest_winrate.js:71 + engine/status.js:176 composed into a false clean bill. stampOf returns {mtime: null, error} with no sha256_12 when a source cannot be read. status.js then did if (!st || !st.sha256_12) continue; and, finding nothing in moved, printed "CURRENT — every engine source the leaf reads still hashes to what it was measured against". With every stamp failed, that sentence was printed over **zero comparisons**. Two silent catches, neither wrong on its own, producing a clean provenance line on this division's headline number. The count is now stated (all N engine sources), unstamped sources are named, NOT DERIVED is printed when N is zero, and a source that has been DELETED is reported as gone rather than as "a different build of".

The rest were latent rather than active, and the honest answer is that they were hiding nothing today — which is a measurement, not an absence of one.

- engine/rollout_r4.js:279 — the split-half scan that produces the NOISE FLOOR discarded a torn line in silence, while countLine, reading the same file for the header counts, keeps torn and publishes it. A row lost here shrinks an arm and moves the spread, and the spread is the entire output. Counted; the split now refuses to report if anything was lost. **Measured on games.r4-decided.jsonl: 0 torn, 0 bad-seed, across all three cuts.** Before the counter existed, 0 and 500 looked identical from there. A second hole found while adding it: a non-numeric seed put every such record on side B, because NaN % 2 !== 0 — now rejected rather than piled up. Incidental: the seed hash parity cut splits **1,382 / 1,242**, an 11% imbalance, and it is the cut that produced the largest of the three spreads (3.9 pts) quoted as this run's noise-floor range.
- tests/test-timestamps.js:49/54 — a directory that would not list and a file that would not read were both skipped with continue, so "no Python generator writes a naive datetime" was also the answer when **zero generators had been looked at**. That is CLAUDE.md's *a capability that cannot prove it ran* inside a

guard written for a different failure. Now asserts a floor on files scanned (39 in `engine/`, 1 in `build/`; `tools/` has no `.py` and `scripts/` does not exist) and fails on anything skipped unread.

- `tests/test-timestamps.js:92` — an artifact that would not parse was silently excluded from the published "*N artifacts still carry a naive stamp*" list, so the files most likely to be broken were the ones the survey could not see. Now counted and named: **0 of 108** `data/*.json` **fail to parse**, so the list of 8 is complete — a statement that could not previously be made at all.
- `tests/test-web-status.js:181` — `catch { return false }` on the freshness filter means "not newer than the board", the same answer a perfectly fresh artifact gets. A source that had been **deleted or renamed** read as up to date, in the test whose job is that every rendered figure traces to an artifact. Missing is now its own failure. None are missing today.
- `tests/test-rollout-gates.js:81` and `engine/rollout_r1_artifact.js:228` — both collapsed "no such file" into "will not parse". The first then granted a CORRUPT gate artifact the one tolerated state ("*awaiting a re-run*"); the second made a broken sidecar indistinguishable from a run nobody ever stamped, which is the exact distinction §4 and §7 exist to preserve.
- `engine/status.js` (9), `engine/rollout_explore_sweep.js` (3), `engine/rollout_r1_artifact.js` (3 more), `engine/run_stamp.js:60`, `tests/test-web-status.js:58/112`, `tests/test-guru-derived.js:56` — each conflated *absent* with *unreadable*. `status.js` now carries a `DIAGNOSTICS` block, printed on screen and deliberately **outside** the section bodies so `--write` never stamps a transient into a ledger.

Two defects in the ratchet itself, filed not fixed:

- `--update` **is all-or-nothing, so the tool's own guidance cannot be followed.** It says "*if a silent fallback is genuinely right here, say why in the code and re-baseline with --update so the exception is deliberate and visible*" — but `--update` re-baselines every silent catch in the repo, including other divisions'. There is no way to bless ONE. It needs a per-entry allow with a reason string, in the shape `build_guru_js.js` 's `DELIBERATELY_UNUSED` already uses.
- `isSilent` **cannot see a recorder it does not recognise by name.** `error: e.message` inside a returned object is a colon, not an `=`, so an artifact that carries its own reason still reads as silent; and a named helper that pushes onto a list looks like nothing from inside the catch body. Four surviving entries are this. Widening the regex would launder real ones, so the code was moved to the documented convention instead — `status.js` 's recorder is named `logUnreadable`, and two locals are named `errWhy` / `errBundle`.

`--all` was added to the ratchet: the 25-line cap is right for a gate, but "*... and 27 more*" is how the tail of a list stops being anybody's job.

13. THE REFIT RAN — and it moved nothing measurable. That is the result, not a preamble to one.

`node --max-old-space-size=4096 engine/fit_policy.js` then `engine/fit_joint.js`, on the weather landing in §11. **8,759 clean open-sheet games, 229,339 usable decisions** (up from 8,414 / 220,613), 183,679 train / 45,660 held out, lambda selected on held-out likelihood at 0. Both weight files carry a fresh `featureHashes` over the 10-scenario fixture, and `feature_fixture --check` exits 0 on both.

The before/after in the artifacts is not the comparison to read, because the two fits have different corpora and different held-out sets — 44,033 decisions against 45,660. Quoting `heldOut.boardAware` 32.269% against 32.271% would be comparing two samples, which is the confound this division keeps finding in other people's work. The comparison that means something scores the SAME held-out decisions three ways, with the split reproduced exactly (`hash(game) % 5 === 0`):

arm	what it is	logL/decision	top-1
A	old weights + old features	-1.732548	32.204%

arm	what it is	logL/decision	top-1
B	old weights + NEW features	-1.732200	32.252%
C	NEW weights + NEW features — what ships	-1.732276	32.178%

1,772 held-out games, 46,162 decisions, paired per decision, bootstrapped over 10,000 resamples of GAMES (decisions inside one game share a team and a board):

paired difference	logL/decision	top-1 points
B – A the weather fix alone, weights frozen	+0.000348 [0.000075, 0.000623]	+0.048 [0.009, 0.093]
C – B the refit, given the fixed features	-0.000076 [-0.000172, +0.000021]	-0.074 [-0.155, +0.004]
C – A everything, against what shipped	+0.000273 [-0.000010, +0.000556]	-0.026 [-0.117, +0.064]

Read plainly, three statements:

- 1. The correctness fix is detectable and it is a quarter of the noise floor.** B – A clears zero on both metrics, and twenty split-half cuts of arm C alone spread by a **median 0.192 top-1 points** (range 0.005–0.770) against an effect of 0.048. The pairing is the whole reason it resolves at all; two runs on different samples could not tell these builds apart. Same shape as §10's value net, one order of magnitude smaller.
- 2. Refitting the weights on the corrected features bought nothing.** C – B contains zero on both metrics and its point estimate is NEGATIVE. Only **1 of 58 weights moved more than 2 SE** (`dmgFrac` +0.0592, 2.45 SE — the column the fix touches most), 6 moved more than 1 SE, and the L2 norm of the whole weight change is **0.216**. The joint file moved less: largest term `terrainSetupHelpsPartner` +0.102, L2 **0.128**.
- 3. The combined change is indistinguishable from zero on held-out human-click prediction.** C – A contains zero on both. The fix was worth making because it is a fact about the game that the feature function was getting wrong on 10.72% of turn-boards — that is the whole justification and it does not need a metric to support it.

What this does NOT say. Top-1 agreement with a human click is not a win rate. Nothing here measures whether MILTANK plays better; that is an H2H and it belongs to SEARCH. And the leaf is a separate model from MAG — §1's finding that the leaf is worse than a coin is untouched by any of this.

13a. THE FIT IS NOT RUN IN THE ENVIRONMENT THE BOT PLAYS IN, and it is 20x the weather defect

This is the headline of the refit, not a caveat on it. CLAUDE.md's rule is *fitting environment and playing environment must match*, recorded because MAG's weights were once fitted with the sheet visible while the bot played without it. **The mismatch is back, pointing the other way, and nothing was watching for that direction.**

```
engine/fit_policy.js:376   board.setSheet(side, m.species, { nature, item })
engine/magnemite.js:522   this.board.setSheet(m[1], sp, { nature, item, ability, moves })
```

`Board.switchIn` copies all four onto the active mon, and `dmgMon`, `effAbility` and `movePriority` read them. So the LIVE player sees a sharper board than the fit ever did: the fit prices every opponent on the dataset's representative moveset and on Smogon's per-species ability odds, while an open-sheet game hands the player the declared four moves and the declared ability.

The sheets carry it. 14,400 sheet entries over 1,200 corpus games: 100.0% declare an ability, 100.0% declare four moves. This is not information the fit lacks — it is information the fit is handed and drops.

Measured the same way §11 was, one process holding both builds of `fit_policy` , identical games and identical decisions:

	weather defect (§11)	the sheet-channel gap
candidate vectors that move	1,768 (0.75%)	37,460 (15.95%)
decisions that move	892 (2.78%)	16,177 of 32,054 (50.47%)
columns that move	14 of 58	20 of 58
games containing a moved vector	238 (19.83%)	1,197 of 1,200 (99.75%)

`switchDiesFirst` (10,013), `diesBeforeMoving` (9,466), `switchSurvives1` (6,632), `dmgFrac` (4,188), `killsThreat` , `switchKOSlow` , `switchSurvives2` , `switchKOFast` , `protectThreatened` , `priority` (1,958 — the declared ability reaching `movePriority`) , `screenValue` , `movesFirst` , `koTarget` , `benchRisk` , `killIsRoll` , `koFirst` , `clickCost` , `passTurnAccrues` , `switchFaster` , `deadNoLastMove` . The choice set is unchanged — the row counts match game for game — so this is purely what the board KNOWS, not what it offers.

It is NOT landed and no second refit was started. Landing it is a one-line change to `fit_policy.js:376` plus a full refit of both files, and it would invalidate the refit reported above on the day it was published. It also needs a question answered first that this measurement does not answer: the fit's decisions come from games where the sheet was public, but MAG must also play the ~half of ladder games where the opponent declines OTS, and a model fitted on four channels degrades differently from one fitted on two when a channel goes missing. That is the Focus Sash lesson — *replacing a hedge with a certainty is only an improvement if you also track what invalidates the certainty* — and it is a decision, not a refresh.

Filed with its size stated, which is the part that was missing: **half of every decision the fit trains on is priced against a board the player does not see.**

13b. THE JOINT LAYER IS REFITTED ON FOUR CHANNELS, AND THE CHANNELS ARE WORTH A LIKELIHOOD GAIN, NOT AN ACCURACY GAIN — 2026-08-05

The other half of §13a's debt, run under Will's go. Three results, each with its instrument named.

The joint refit (`engine/fit_joint.js` , four-channel `joint_rows.js`): 8,856 clean open-sheet games, 101,459 joint turns → 95,886 usable, 77,975 train / 17,911 held out by game, `lambda 0` on held-out. The artifact now carries a `fitEnvironment` block and it says `matches_player: true` by measurement: the declared ability and moves reached the board on **202,343 of 202,918 scored slots (99.7%)** and **395,130 of 396,288 live foe actives (99.7%)**. Held out, predicting the pair: separate decisions `logL` -3.3294 / top-1 10.3%, refit with joint terms zeroed -3.3199 / 9.8%, with the joint terms **-3.2308 / 12.2%**. The chosen pair fell outside the top-6 menu on 11.1% of kept turns. `feature_fixture --check` passes on the new artifact. No before/after against the presheet joint vector is quoted because none was measured — the presheet run published no held-out table to its artifact, and comparing two logs would be comparing two samples.

What the two extra channels are worth at the marginal layer (`engine/sheet_channel_value.js` , arm A = release `d3d04b669e18` 's two-channel incumbent, 44,982 paired held-out decisions over 1,789 games, 10,000 game-bootstrap resamples). The first run **VOIDED itself** — ENGINE saved `engine/medicham2-browser.js` mid-run and the instrument recorded `void: true` — and the second run is clean, with every deterministic figure identical between the two:

paired difference	logL/decision	top-1 points
B – A the information alone, weights frozen	+0.002853 [0.001611, 0.004072]	+0.009 [-0.140, +0.157]
C – B the refit, given the information	+0.002234 [0.001638, 0.002831]	+0.165 [0.029, 0.299]
C – A everything vs what shipped	+0.005087 [0.003854, 0.006331]	+0.173 [-0.011, +0.360]

Split-half noise floor of the shipping arm, 20 cuts: **median 0.331 top-1 points** (range 0.012–1.385; the earlier refit's floor was 0.192 on a smaller paired set). Read plainly: **the sheet channels buy a real likelihood gain — every logL interval clears zero — and no demonstrable top-1 gain.** The one top-1 interval that clears zero ($C - B$, +0.165) is half its own noise floor and resolves only because the comparison is paired; the total effect against what shipped contains zero. Same shape as §13: correctness and information first, metric second, and the honest metric statement is "better calibrated per decision, not measurably more often right on the argmax."

The degradation budget did not move, and cannot move by this lever. `fit_joint.turnsDropped` is **5.4929% (5,573 of 101,459) against a 5.49% ceiling — still red.** The dropped turns are unmatched clicks (5,555) and ambiguous mirrors (18); the chosen pair is kept regardless of its rank, so the four-channel w_1 changes which ALTERNATIVES are on the menu, never which turns are kept. The rate crept from 5.4811% when the ceiling was ratcheted (86,242 turns) because the newly ingested games unmatch at 5.56%. The ceiling is untouched; the call on it is Will's.

PORY family regenerated, and `tests/test-site-data-fresh.js` is GREEN (7/7) — §5d addendum has the pory-eval numbers; `data/nmf-roles.json` moved 13,258 → 14,808 team-docs, 258 → 263 moves, recon-err 0.8346 → 0.8356, rank 10 unchanged, and both site bundles (`data/pory.js` , `data/nmf.js`) were rewritten by their own generators in the same runs. `data/pory-nn.json` retrained at 6,289 games / 106,782 states (was 6,008 / 102,296): every arm ordering holds — N6 0.6201, NR 0.6132 and LR 0.6064 all still beat the two-feature bar at 0.6229, nonlinearity is still worth ~0.003 and representation ~0.016 — and no living doc quotes these figures (§5c). The first retrain immediately re-red the drift check at 15.1%, the §5f false-denominator class to the letter: its population is the raw-logs subset. Both `engine/pory.py` and `engine/pory_nn.py` now declare `population_ceiling` (the artifact's own generator wrote it; the retrain reproduced every arm to the digit under its seeds), which is what turned the check green rather than a threshold being moved.

Doc and site locations quoting superseded PORY figures, for the propagation pass (grep-verified, historical HANDOFF files excluded as history): `docs/ABRA-whitepaper.md:113` (0.6298 [0.6125, 0.6456]), `docs/SUMMARY.md:77` (same + 0.655), `docs/MODELS.md:358/360/364` (0.9943/1.4080, 0.629799/0.629778, +0.000021 [-0.000013, +0.000056], 925, 4,623), and WEB's `web/stadium.html:506,:728` + `app/stadium.html:506,:728` (the `kadabra` data object and its prose), which render every one of those numbers and are flagged, not edited.

14. THE OUTPLAYED TURNS — 1,336 recorded actions were not clicks, and the model was learning from every one of them. LANDED 2026-08-05.

`docs/CLICK-CENSORING-FIX.md` is the spec, ordered by Will: *"i def dont like just tossing turns because they got outplayed with a move liek encore or follow me, these are the basis of vgc man."* Four artifacts: `data/click-censoring-census.json` , `data/partial-label-em.json` , `data/censoring-value.json` , and the refitted `data/policy-weights{,-joint}.json` .

LEAD WITH THE RESULT, INCLUDING THE HALF THAT DID NOT WORK. Two headline classes were measured and only one moved:

held-out class	what changed, after – before	verdict
COERCED (n=284) — Encore replaced the click, or the mon was dragged in	P(model picks the action no human chose) -0.002614 , 95% CI [-0.003663, -0.001637]	the poison is unlearned, and it is the only headline that moved
PARTIAL (n=643) — a redirector soaked the attack	mass on the true candidate set +0.000109 [-0.000286, +0.000491] ; logL on the set -0.002646 [-0.004037, -0.001377]	no improvement. The likelihood is very slightly WORSE.
CONTROL, CLEAN (n=46,268)	logL +0.000447 [0.000142, 0.000743] ; top-1 +0.002 [-0.094, 0.098]	as the spec predicted: no top-1 change

47,195 paired held-out decisions over 1,809 games, 10,000 bootstrap resamples **clustered by game**, `engine/censoring_value.js` . The spec disclaims a corpus-wide top-1 improvement in advance and none is claimed here; the CLEAN row is a control.

***RE-MEASURED 2026-08-05 on the current engine and a corpus grown to 9,230 games — every figure in this section reproduces inside its interval.** The table above is the 3.42.0 run and is kept as published; the artifact on disk now holds **n=48,274 over 1,851 held-out games**, **COERCED -0.002613 [-0.003650, -0.001672]**, **PARTIAL mass +0.000122 [-0.000261, +0.000514]**, **CLEAN logL +0.000485 [0.000189, 0.000777]**. §17 has the full comparison and the reason the engine move could not have touched it.*

Say the negative result plainly: Stage C bought nothing measurable, and the reason was predicted by Stage C's own validation before the refit ran. The EM harness recovers **97.4%** of a planted censoring bias when the censoring is heavy, and at the rate the corpus actually censors, the bias in weight space is **-0.0030 against a 0.2600 noise floor** — unmeasurable. The redirection correction is right in principle, and the class is 1.35% of actions with a candidate set of exactly two, so there was almost nothing to recover. Both instruments agree, which is the only reason to believe either.

Stage A — the census. 241,927 recorded human actions over 8,942 clean open-sheet games (the FIT corpus; the census artifact has since been re-run twice with the store, at 9,022 and then **9,230** games, and the shares are stable to a hundredth of a point — see §17):

class	n	share	mechanism	
CLEAN	229,555	94.886%	—	
PARTIAL	3,260	1.3475%	Follow Me / Rage Powder 3,231; Lightning Rod 29. Every candidate set is size 2	
COERCED	1,336	0.5522%	Encore's application turn 1,116; a ` \	drag\ ` 220 (Roar 184, Dragon Tail 33, Whirlwind 3)
dropped, not a censoring class	7,776	3.214%	unmatched 6,937, trivial 809, ambiguous 30	

The mechanism list is read from the running format, never typed — moves with `condition.onOverrideAction` , moves with `forceSwitch` , abilities with `onFoeTryMove` , items assigning `switchFlag / forceSwitchFlag` (**empty here:** Eject Button, Eject Pack and Red Card are all `isNonstandard: 'Past'`), plus `data/tags.json` 's `redirects / redirectsType` . Every set refuses to be empty, and a zero on either counter is fatal in both fitters.

THE CLASSIFIER WAS SCORED AGAINST THE PROTOCOL, NOT ASSERTED. The census has a second arm that reads `data/games.*.raw-logs.jsonl` and compares per (game, turn, slot):

class	protocol says	classifier flagged	both	recall	precision
Encore application	619	642	617	99.68%	96.11%
drag	86	86	83	96.51%	96.51%

The 25 Encore false positives are 0.01% of all actions and the asymmetry is the right way round: a false positive deletes one real click, a false negative keeps a poisoned one, and there are two of those. Most likely cause is an Encore blocked by Protect, which the extractor records no failure flag for. Stated, not chased — the classifier was frozen while the refit that depends on it ran.

Two corrections to the spec, both measured. (1) §1's first row is wrong and `engine/redirect_audit.js` said so on 2026-08-02: redirection does **not** drop the turn. The redirector is a legal candidate target, so the matcher matches it and the click enters the fit with a CONFIDENT WRONG TARGET. It is label noise, not censoring — which makes Stage C a poison fix as much as a recovery. (2) A `\ | drag \ |` is a third coerced class

the spec does not list. `engine/durable-ingest.js:67` parses `\|switch\|`, `\|drag\|` and `\|replace\|` with one regex, and `fit_policy`'s `forcedSlot` guard only knows about faints, so every phased arrival was fitted as a voluntary switch decision.

WILL'S FARIGIRAF CASE IS ANSWERED: PARTIAL, NOT ERASED.

```
|cant|p1a: Farigiraf|ability: Armor Tail|Aqua Jet|[of] p2b: Basculegion
```

The blocker is named first, the attempted **move** is named, and `[of]` names the **attacker. 284 of 284 priority-block lines carry the attacker slot (100.0%)**, so the user and the move are exact and only the target is ambiguous — between the blocker and its ally, and nowhere else, because the ability blocks nothing aimed elsewhere.

It is counted and NOT recovered, and that is a judgement with a reason. Showdown emits no `\|move\|` line for a blocked attempt, so the class leaves no event and lives only in the raw logs — which cover **66.17% of the fit corpus (5,917 of 8,942 games)**, and the gap is one SOURCE, `data/games.ots.jsonl`, an external archive with no log file. Recovering these 284 clicks, and the 126 more that `\|cant\|` states outright (Taunt 59, Disable 58, Heal Block 5, Imprison 4), would add outplayed turns from two stores and none from the third. That is a corpus reweighting wearing a bug fix's clothes. Closing it means re-ingesting the ots archive with its logs — OPS work, filed.

A FOURTH THING, FOUND ON THE WAY, AND IT IS A WRONG DENOMINATOR RATHER THAN A WRONG LABEL. `engine/board.js`'s `candidates()` narrows the choice set for a **Choice item**, derived from the dex's `isChoice`, with its own comment saying why: *"that is not a scoring error, it is a WRONG DENOMINATOR. A conditional logit divides by the sum over the choice set."* It does nothing about the other family that shrinks a menu — the `onDisableMove` set. **2,280 of 139,769 logged actions (1.6313%) were taken with a menu-sealing volatile up:** Encore 1,276, Throat Chop 375, Taunt 329, Disable 239, Heal Block 94. A human left one legal move by Encore is priced as having chosen it over nine. **NOT FIXED HERE** — narrowing the menu moves every feature row and owes its own refit, and it is a different defect from the one this dispatch was for. Counted so the decision has a size.

Stage C — the estimator, shown failing on known-bad input before it was believed.

`engine/em_validation.js`, 31,940 real corpus feature rows over 1,200 games with SYNTHETIC labels drawn from a known planted vector, 3 seeds, the real censoring process applied to the planted labels:

regime	rows censored	oracle	naive	EM	noise floor	verdict
amplified	20.961%	0.9978	1.8913	1.0208	0.2600	bias 0.8935 clears the floor; EM recovers 97.4%
observed	0.439%	0.9978	0.9948	1.0021	0.2600	bias -0.0030 — inside its own noise floor

Distances are $\|\hat{w} - w^*\|_2$. The noise floor is the spread of the ORACLE arm across the three seeds, so it carries no information about the contrast. The **first** amplified regime censored EVERY eligible row and EM recovered only 45% — correctly, because with every same-move row collapsed there is nothing left to identify the target features from. That is Cour et al.'s identifiability condition failing, not the estimator; the eligibility is now exogenous and the collapse label-dependent, which is what the corpus does. `engine/em_validation.js --check` re-verifies the recorded verdict AND re-hashes every source, so editing `engine/click_class.js` turns the gate red instead of leaving a stale PASS; it is registered in `tests/run-all.js`.

Stage D — what the refit moved, and the confound stated rather than buried. `data/policy-weights.json` records the corpus, the train/held-out split, lambda 0 on held-out and the reweighted vector that ships. **The corpus sizes are QUARANTINED — withheld, not annotated:** the artifact is downstream of MEDICHAM (`engine/fit_policy.js` reaches `engine/medicham2-browser.js` through

require), and it becomes quotable again when the gate opens AND this is re-run: `node --max-old-space-size=4096 engine/fit_policy.js` . What the refit MOVED is legible without them: $\|new - old\|_2 = 0.8030$ and **9 of 58 weights moved more than 2 SE**, and the mechanism is legible in which ones:

feature	before → after	
deadStall	-1.3114 → -1.4763	5.44 SE
stallIntoEncore — <i>"I am about to Protect and something across from me can Encore me for it"</i>	-1.0502 → -1.6281	3.10 SE, the largest single movement
deadSide	-2.7606 → -3.1414	4.01 SE

That is the predicted direction. The poisoned rows were victims "choosing" their last move under an active Encore; deleting them makes clicking into an Encore threat look worse, and the Encore/stall family is exactly where the vector moved. Same shape as §10's `hpDiff` $0.169 \rightarrow 0.377$.

THE CONFOUND, NAMED: the two vectors differ in four ways, not one. The incumbent was fitted on 8,856 games and the new one on 8,942 — the collector never stops — so the Stage D contrast carries the coerced removal, the partial-label EM, 86 extra games and the refit itself. The weight-movement pattern above is evidence for attribution and is not proof of it. `CENSORING=off` now exists in `engine/fit_policy.js` for exactly this: it fits the OLD way on the NEW corpus, and it records `censoring: "off (CONTROL ARM – not shippable)"` in its own artifact so a control can never be mistaken for a ship. **That arm has not been run** — it is a second full refit and free RAM was 1.3 GB with the joint fit in flight. It is the next thing this section owes.

AND EVERY EFFECT HERE IS SMALLER THAN ITS OWN CLASS'S NOISE FLOOR. The COERCED contrast is 0.002614 against a split-half floor of 0.007635; the CLEAN logL gain is 0.000447 against 0.007855. They resolve only because the comparison is **paired per decision** — two runs on different samples could not tell these builds apart. This is the same statement §13 and §13b make, and it must travel with the numbers.

Stage B — the budgets are RE-DERIVED, not renumbered, and `turnsDropped` is retired.

`fit_joint.turnsDropped` was $(turnsSeen - kept)/turnsSeen$ and sat at 5.4929% against a 5.49% ceiling. Stages B–C change what "dropped" MEANS: coerced turns used to be inside `kept` , carrying a wrong label, and now leave the labelled set — so the old total would have gone UP while the artifact got strictly better, and a ceiling that may only tighten would have gone red for an improvement. **Raising or lowering the number would have been the wrong move in either direction.** Three counters now, each with its granularity stated in `data/degradation-budgets.json` and its ceiling ratcheted from a measured run:

counter	what it counts	denominator
<code>fit_policy.decisionsUnreadable</code> / <code>fit_joint.turnsUnreadable</code>	the click existed and could not be recovered. A LOSS. Successor to the old totals, and directly comparable because it is the same quantity minus a term that was misfiled as <code>kept</code>	human actions seen by <code>fit_policy</code> / joint turns seen by <code>fit_joint</code>
<code>fit_policy.coercedActions</code> / <code>fit_joint.coercedTurns</code>	the recorded action was not a click and was removed. A CORRECTION, not a loss — it should track the metagame's use of Encore and phazing and nothing else	same
<code>fit_policy.decisionsDropped</code> / <code>fit_joint.turnsDropped</code>	RETIRED. Carried in a new <code>superseded</code> block with its old ceiling intact, so the history is not deleted	—

`measured_at` also used to read *"over 120 corpus games"* on every row, which is true of the three `board.js` counters and **false of every fitter rate** — those come out of an artifact written over the whole corpus. A ceiling whose denominator is misdescribed cannot be re-derived by anyone.

What this does not say. Top-1 agreement with a human click is not a win rate; whether MILTANK plays better is an H2H and belongs to SEARCH. The COERCED class has no ground-truth label by construction, so its contrast measures a change in the MODEL, not an improvement in accuracy — it cannot be otherwise, and inventing an agreement number for it would have been the dishonest option.

15. THE TWO LEAF ARTIFACTS DO NOT CONTRADICT EACH OTHER. RESOLVED 2026-08-05, and the answer is a decomposition, not a winner.

`data/winrate-backtest.json` and `data/rollout-r1-explore-sweep.json` rank the same leaf at very different accuracies — **both figures are QUARANTINED here, withheld and not annotated**, because both artifacts are downstream of MEDICHAM and the gate is closed. The RECONCILIATION is what this section is for and it survives without them. The sweep flagged the conflict itself in `reading_against_the_leaf_calibration` and refused to treat "explore=1.0 spent the signal" as established while a second measurement disagreed. **It was right to refuse, and both artifacts are correct.** They score the same function on positions of very different difficulty, and the gap decomposes cleanly.

The sweep named two differences — rollout budget and horizon — and **those are the two that do not matter.** There are six, and the three that carry the gap are position, corpus and sheet, in that order.

`engine/leaf_position_contrast.js` holds five of the six fixed at a time: one leaf (`rolloutWinProb` , `explore=1.0`, `n=40`, `horizon 20`), one frozen release (**6b0e4117d964**), the same seeds, and both accuracy definitions on every arm. `data/leaf-position-contrast.json` , with `data/leaf-position-contrast-rows-6b0e4117d964.jsonl` beside it so any cut can be re-derived without re-running the rollouts.

arm	corpus	position	sheet	n	maj. class	accuracy	Brier vs coin (paired)	ECE	MCE	curve slope
D = the sweep	open-sheet bo3	mid-game	yes	9,201 pos / 2,500 g	52.5%	69.83% [68.6, 71.1]	-0.0440 [-0.0513, -0.0360]	0.0925	0.162	0.703
C	open-sheet bo3	mid-game	no	9,201	52.5%	68.73% [67.5, 69.9]	-0.0401 [-0.0470, -0.0329]	0.0939	0.146	0.693
B	open-sheet bo3	turn o	yes	2,500	52.4%	58.20% [56.4, 60.2]	+0.0101 [0.0020, 0.0182]	0.1120	0.332	0.402
A	open-sheet bo3	turn o	no	2,500	52.4%	55.92% [54.0, 57.8]	+0.0166 [0.0088, 0.0243]	0.1284	0.351	0.331
E = the backtest	closed ladder	turn o	no	1,499	52.2%	51.17% [48.6, 53.6]	+0.0456 [0.0344, 0.0567]	0.1793	0.458	0.068

Intervals are game-clustered bootstraps, because 9,201 mid-game positions come from 2,500 games and an unclustered interval on them is too narrow by about $\sqrt{3.7}$.

The decomposition telescopes exactly. 69.83 – 51.17 = 18.66 points:

term	contrast	points	how measured
POSITION	A → C	+12.81	mid-game vs turn o, sheet off, same 2,500 games
CORPUS	E → A	+4.75	closed ladder vs open-sheet bo3, turn o, sheet off, same config
SHEET	C → D	+1.10 [0.31, 1.88]	paired McNemar, same 9,201 boards from two walks

C and D come from two passes of `joint_rows.build` over the same games, the second with `Board.prototype.setSheet` disabled — so suppressing the sheet changes what the leaf KNOWS and must not change which boards are scored. **The original run asserted that and it PASSED:** all 9,201 positions

agree across the two walks on gid, turn, label, `aliveDiff` and the continuous HP witness, and the run aborts rather than report a pairing it did not check. The artifact on disk is a re-cut of that run's rows, so its `pairing_check` says the result is carried rather than re-performed — a process that did not do the check does not get to say PASSED.

The sheet at turn 0 is worth **+2.28** [0.57, 3.99] (B – A, paired), so taking the other path through the square gives position +11.63 instead of +12.81. Either way position is two-thirds of it and the sheet is the smallest of the three.

Three independent things say the config is not the explanation. Arm E re-runs the backtest's condition at the SWEEP's budget and horizon (n=40, h=20) and lands on 51.17% / Brier +0.0456 / ECE 0.1793 against the published 51.66% / +0.0466 / 0.1827 — inside E's own split-half floor of 1.54 points. The sweep re-ran itself at h=60 and got 69.86% against 69.84%. And §1 already recorded that 40 and 200 rollouts give the same turn-0 answer. **The horizon and the budget are settled: they move nothing.**

Two independent routes reach the same turn-0 number, which is why I believe the decomposition. Cutting the sweep's own committed dump down to `turn ≤ 1 AND aliveDiff == 0 AND |hpDiff| < 0.02` — its nearest thing to a preview board, sheets on — gives **55.70%** on n=237. Arm B measures a real turn-0 board on the same corpus with sheets on and gives **58.20%** on n=2,500. The subset's split-half spread runs 0.47 to 21.47 points across ten random by-game cuts (median ≈ 4.2), so those two agree.

THE HEADLINE 50.99% IS THE UNDERPOWERED READ, and the better number is not better news. It is the held-out fifth at n=200. Re-cut from `data/winrate-backtest-rows.jsonl`, the same leaf at n=40 over the **full** 6,886-game clean corpus ranks at **51.66%** (51.80% on 6,570 decisive calls) — real by p, and its majority class is 51.25%, so its edge over *always say p1* is **0.41 points against a median split-half floor of 0.75**. LESSONS §9: an effect smaller than the noise floor is not an effect. **On the closed-sheet ladder at turn 0 the leaf does not beat the majority class.** "Cannot rank at all" was reported off the wrong n and happens to survive the correction.

Now the answer to the three options, plainly.

- **(a) "fine mid-game, broken at turn 0" — the largest term, and "fine" is too kind.** Mid-game the leaf is genuinely not broken: it beats a coin on Brier by 0.044 [0.036, 0.051], its curve is monotone with slope 0.703, and on boards where the material baseline has *collapsed to the majority class* (`aliveDiff` 0, `|hpDiff|` < 0.02, n=411) it scores 62.29% against material's 51.09% — **+11.19 [5.05, 17.34] over counting**. It is reading real non-material structure. But it still puts **31.4%** of positions in the two extreme bins, and its top bin predicts 97% and wins 86%. A slope of 0.70 is not calibration; it is a leaf that ranks well and lies about how sure it is.
- **(b) "broken everywhere, the sweep measures something easier" — right that it is easier, wrong that it is only easier.** The +11.19 over a collapsed material baseline is not an artefact of easy positions. The sweep is not measuring material with extra steps.
- **(c) and there is a term nobody named: the CORPUS, +4.75 points — bigger than the sheet channel at either position.** The open-sheet corpus is `fit_policy.loadCorpus()`, which on its first 2,500 games is **99.9% our own** `gen9championsvgc2026regmbbo3` **scrape**, not the OTS archive as the generator's own comment implies. Its pool is lower-rated (median 1,174 against the ladder held-out fifth's 1,266) and it plays under **forced** open sheets, so both humans had full information and the outcome may be more determined by the matchup. Turn counts and forfeit rates are the same in both. **Neither mechanism is tested here** — the term is measured, its cause is not, and it is not a sampling artefact of the held-out slice, because the full-corpus backtest agrees with E.

DISCRIMINATION AND CALIBRATION FAIL SEPARATELY AND THE SPLIT WIDENS AS INFORMATION IS REMOVED. Arm A ranks at 55.92% against a 52.4% majority — its interval's lower bound is 54.0, clear of both the majority class and its 2.24-point split-half floor — and its Brier is still **worse than a coin**, with an MCE of 0.351. So a turn-0 leaf can carry real ranking signal and still be a liar about its confidence, which is the failure mode that matters to an argmax. By arm E even the ranking is gone and only the confidence is left.

A SIGN FLIP WORTH A RE-RUN, EXPLICITLY NOT ESTABLISHED. Exploration helps mid-game and may hurt at turn 0. Paired on the backtest's own 6,886 turn-0 ladder games at horizon 60, the **greedy** playout ranks at 53.09% against explore=1.0's 51.66% — **+1.44 [0.10, 2.77]** — while paired on the sweep's 9,201 mid-game boards explore=1.0 wins by **+3.20 [2.24, 4.15]**. The turn-0 lower bound is 0.10 against a median split-half floor of 0.75, so it is **inside its own noise floor and is not a result**; and greedy's Brier there is *worse* (0.3240 against 0.2966), so the two playouts differ in which failure they have rather than in quality. The two arms are also not the same code path (`battleInit` + `chooseAction` against `rolloutWinProb`). It needs `mew.js --miltank-explore`, which §3 already filed to SEARCH, and it needs the position held fixed.

WHAT THIS MEANS FOR PORYZ, since that is what the question was for. `docs/PORYZ-spec.md`'s representation is per-Pokémon HP fraction, status, every stat stage, revealed item and ability, and a threat matrix — and it is the leaf of $EV(a) = \sum P(\text{reply}) \times V(\text{board after})$. **Every one of those inputs is constant across games or absent at turn 0:** all eight bodies are at 1.0 HP, no status, no stages, and the closed-sheet ladder has revealed no items. The only feature that survives to turn 0 is the threat matrix over the brought four a side. So **PORYZ cannot move the turn-0 number, by construction of its own feature list** — if the target was "the leaf that reads 100% and loses", that leaf is the PREVIEW one and this spec is not aimed at it.

Aimed at what PORYZ-spec's engineering section actually says — making the mid-game EV sum affordable — it is well aimed, and this run hands it a bar measured on the same positions rather than quoted from another sample: **69.83% accuracy, Brier 0.2060, ECE 0.0925, slope 0.703 on 9,201 positions at release 6boe4117d964**, with the rows on disk. PORYZ's premise sentence, "the whole learned value function is worth 3.4 points over counting", is about PORY2. The rollout leaf is worth **+4.58 [3.47, 5.68]** over the same graded material baseline mid-game and **+11.19 [5.05, 17.34]** where material has nothing to say. That is the incumbent PORYZ has to beat, and it is a harder incumbent than the spec assumed. This is a measurement, not a build decision; the decision is SEARCH's.

Filed, not fixed.

- `engine/rollout_r1.js:436` **puts a prose note key inside** `source_digests`. `engine/provenance.js:648` calls `digestOf()` on every key in that map, so a key that is not a readable path marks the whole artifact `unverifiable` — which is why `data/rollout-r1-explore-sweep.json` cannot be digest-verified. This file made the same mistake and moved the prose to a sibling key; the artifact went from `stale?` to `ok`. SEARCH's file.
- `engine/rollout_r1.js:26-29` **'s corpus comment is misleading about what it samples.** `loadCorpus()` reads `bo3`, `OTS` and `ladder` in that order, and the head of that list — the whole `R1` and sweep sample — is `bo3` with a handful of `smogtours`. The exact counts are withheld with the artifact, which is downstream of `MEDICHAM`. Every `R1` number ever published is a **bo3 open-sheet** number, and §15 measures that this corpus is worth 4.75 accuracy points at turn 0. The published figures are not wrong; what they are *about* is narrower than the comment says.
- `data/censoring-value.json` **and** `data/click-censoring-census.json` **trip the provenance ratchet** — their generator ships without recording what content it read. Another division's files, written this session. Reported, not touched.

Re-cutting this artifact costs seconds, not half an hour:

```
RECUT=data/leaf-position-contrast-rows-6b0e4117d964.jsonl node engine/leaf_position_contrast.js
```

It opens the release named in the filename rather than the newest, refuses a dump that cannot name its engine, and never rewrites the dump it read. Verified: the re-cut reproduces every figure above bit-for-bit and leaves the row file byte-identical.

§16 — `censoring-value.json` is UNSAFE, and re-running it is not a repeat

ANSWERED IN §17, 2026-08-05 — and by none of the three options below. *The confound was measured instead of argued: all 58 feature columns are identical across the engine bundles on all 1,751,688 corpus rows, so the fitting environment and the playing environment are the same FUNCTION here. Both artifacts were re-run against the live tree and both are `ok`. The section below is kept as what was true before that was measured; do not read its three options as open.*

2026-08-05. `provenance.js` flags it: `medicham2-browser.js` was `e2bcff0db96f` when it was measured and is `80fe43fba1a9` now, because WIRES 114–116 landed underneath it. The flag is correct and the artifact should not be quoted.

Two things had to be fixed before it could even be re-run, and both are worth more than the number.

The comparison baseline lived in a session scratchpad. The run compares the pre-censoring incumbent against the post-censoring fit, and the incumbent existed only as a copy in a temp directory that gets cleaned. A published figure whose input is in `%TEMP%` is not reproducible by anyone, ourselves included, one cleanup later. It is now `data/policy-weights-pre-censoring.json`, sha12 `01bc43936324` — the digest the artifact itself records, verified to match.

The artifact was invisible to provenance despite recording more than most files that pass. It stamped `source_digests_before` and `source_digests_after`; `provenance.js` reads `source_digests` and nothing else, so it fell to "rests on mtime alone" while carrying better evidence than the files around it. Recording something correctly under a name the checker cannot see has the same outcome as not recording it. The generator now writes the canonical key too — and the moment it did, the artifact stopped being `ok` and became `UNSAFE`, which is the whole point.

THE RE-RUN IS BLOCKED ON A JUDGEMENT, NOT ON COMPUTE. Both weight vectors were fitted under the pre-WIRE-114 engine. Scoring them through the current one breaks the rule in `CLAUDE.md` that the fitting environment and the playing environment must match, and it would measure *the censoring change plus three wires* as one quantity. The options, none free:

- **Refit both vectors under the current engine, then re-measure.** Correct, and the expensive one.
- **Re-measure through a release frozen at `e2bcff0db96f`.** Reproduces the original honestly, but the artifact deliberately reads the live tree — `no_engine_release` says freezing it would measure the thing being tested — so this changes the design of the measurement.
- **Leave it UNSAFE until the next refit lands anyway,** and do not quote it. Cheapest, and the status quo, but only honest while nothing downstream depends on it.

Not chosen here. `engine/censoring_value.js` refuses to run without `WEIGHTS_OLD` and now points at the preserved baseline and at this section, so whoever picks it up is choosing rather than guessing.

§17 — THE CONFOUND WAS MEASURED AND IT IS EMPTY. Both artifacts re-run, both ok. 2026-08-05.

None of the three options in §16 was taken, and the reason is a measurement rather than an argument. The blocking question — "the vectors were fitted under one engine and would be scored through another" — is a claim about the FEATURE FUNCTION, and a feature function is a function from a board to a number. Two versions of it are the same function if they agree on every board. So they were run against each other on every board the fit actually uses.

Result: all 58 feature columns are hash-identical across the three engine bundles, over 1,751,688 candidate feature vectors from all 9,230 clean open-sheet games.

bundle	medicham2-browser.js	data/tags.json	what read it	58 column hashes
release 09acd3b404ef	e2bcff0db96f	c0bb781f47a8	censoring-value.json	identical
release 032b4a2979dd	80fe43fba1a9	c0bb781f47a8	click-censoring-census.json	identical
live	0cb911437fed	73c81e6421b8	the re-runs below	identical

The three bundles were loaded from the frozen releases and registered under the live module paths, so board.js, fit_policy.js and click_match.js are the same bytes in every arm and only the simulator and the tag dex move. engine/quality.js is deliberately NOT swapped — a snapshot copy of it resolves the store inside the release directory and the walk would have had no rows to disagree about.

A null result from an instrument that cannot see is worth nothing, so the instrument was shown seeing. Under a Psychic Terrain with a Levitate body, the two frozen engines return 0 — priority refused — and the live one returns Infinity. The harness reads the same call the feature code reads, so a difference of that kind would have moved a column.

Why the change is real in the simulator and invisible in the features: across the whole corpus board.js makes 173,478 guarded calls to priorityRefusedAbove, of which 424 are under a Psychic Terrain, and in 0 of them is every live defender airborne. WIRE 117 can only change an answer when no grounded body is left to hold the bar up.

Both artifacts were then re-run against the live tree, and both reproduce. The corpus had grown 8,942 → 9,022 → 9,230 clean open-sheet games in between, so this is a fresh measurement on a superset rather than a replay — which makes the agreement evidence rather than tautology:

held-out class	published 3.42.0 (n=47,195, 1,809 games)	re-run (n=48,274, 1,851 games)
COERCED P(the coerced action), lower is better	-0.002614 [-0.003663, -0.001637]	-0.002613 [-0.003650, -0.001672]
PARTIAL mass on the candidate set	+0.000109 [-0.000286, +0.000491]	+0.000122 [-0.000261, +0.000514]
PARTIAL log-likelihood of the set	-0.002646 [-0.004037, -0.001377]	-0.002662 [-0.004002, -0.001368]
CONTROL, CLEAN log-likelihood	+0.000447 [0.000142, 0.000743]	+0.000485 [0.000189, 0.000777]
CONTROL, CLEAN top-1	+0.002 [-0.094, 0.098]	-0.008 [-0.107, 0.085]

Every verdict in §14 stands, including the negative one: the redirection correction still buys nothing measurable, and every effect is still smaller than its own class's split-half floor (COERCED 0.002613 against 0.011909; CLEAN logL 0.000485 against 0.004820). They resolve because the comparison

is paired per decision, and that sentence must keep travelling with the numbers.

The census moved with the corpus and its shares did not: **249,404 actions over 9,230 games — CLEAN 94.9111%, PARTIAL 1.3344% (3,328), COERCED 0.5545% (1,383: Encore 1,152, |drag| 231)**, against 94.8916 / 1.3467 / 0.5509 at 9,022 games. The classifier still scores against the raw protocol at **encore recall 99.69% precision 96.31%, drag 96.74% / 96.74%** on the 67.23% of games that have a raw log.

`node engine/provenance.js --strict` **exited 0 at that point: 0 UNSAFE, 1 declared VOID (exploitability.json), 57 ok.** Both files carry `source_digests` over the tree they were computed on, so a next engine move flags them again by CONTENT rather than by mtime — **and one did, forty minutes later. See §17b, which is the more important half of this section.**

What this does NOT license. It says the four wires moved no feature on THIS corpus — it does not say the engine did not change, and it is not a general permit to score old weights through a new simulator. The next engine move gets the same treatment: run the columns, then decide.

The harness is `engine/feature_engine_contrast.js` **and it is in the repository, not in a session scratchpad — which is §16's own lesson applied to §17's evidence.** It writes `data/feature-engine-contrast.json` with `source_digests`, and it costs about four minutes per bundle over the whole corpus:

```
SHOWDOWN_PATH=... BUNDLES=live,09acd3b404ef,032b4a2979dd node engine/feature_engine_contrast.js
```

Each bundle runs in its own child process, because a module-cache swap cannot be undone in one. Two properties are worth more than the number it prints:

- **It refuses to report agreement unless its positive control disagreed.** `BUNDLES=live, live` returns *NOT A RESULT* — *the positive control did not separate the bundles*, verified before this was believed. A harness that silently loaded the same bytes twice would otherwise publish a confident "identical", which is the exact shape of every failure in this project's history.
- **It is not** `engine/feature_fixture.js` **and does not replace it.** The fixture hashes ~50 frozen boards so a weight file can *carry* the hashes, and its own header states the limit: a guard only guards what it exercises. This runs the same question over every board the fit actually uses, so a branch no fixture board stands on cannot hide in it. Both were green here, which is the first time they have been asked the same question on the same day.

§17b — AND THEN THE TREE MOVED AGAIN, AND THIS TIME THREE COLUMNS MOVED WITH IT. A REFIT IS OWED.

The instrument built in §17 found a real feature change forty minutes after it was written, and `engine/feature_fixture.js --check` **— the guard** `status.js` **prints the refit edge from — is BLIND to it.** That is the finding of this session, and it outranks everything above.

Between 15:40 and 15:44 on 2026-08-05, while this division was measuring, three files moved:

file	was	is	what it did
<code>engine/fit_policy.js</code>	45f545425420	caeeec21c560	<code>loadCorpus()</code> went 9,230 → 6,055 clean open-sheet games
<code>engine/medicham2-browser.js</code>	0cb911437fed	82bed8cdcf6b	—
<code>engine/board.js</code>	54e3d2ca9f85	5bdaa3923958	the feature file itself

Re-run with the sample pinned — **1,136,845 candidate vectors over the same 6,055 games, identical row_key_hash in all three arms** — the verdict is no longer IDENTICAL:

MOVED — `deadNoLastMove`, `movesFirst`, `diesBeforeMoving` **differ on identical rows. This is a REFIT, not a restamp.**

Both frozen bundles (`e2bcff0db96f`, `80fe43fba1a9`) agree with each other and disagree with the live tree in the same three columns, which is what a single new change looks like — ****and it is: CHANGELOG 3.49.0, "There were two implementations of who moves first. One is deleted, and the survivor is dynamic."**** Speed order is now re-sorted mid-turn, so `movesFirst` and everything downstream of it answers a different question than the weights were fitted against. The columns name the change without anyone having to guess, which is what a per-column hash is for.

**** node engine/feature_fixture.js --check data/policy-weights.json says "feature semantics OK — agrees with board.js on every fixture board" on that same tree. Both instruments are working; they are answering the question on different boards, and the ~50 frozen fixture boards do not stand on the branch that moved. The fixture's own header says a guard only guards what it exercises — this is the first time that limit has been shown with a number rather than stated. status.js prints refit edge: CLEAN from that check, so the refit edge is currently reported clean and is not.**** Two consequences, in order:

1. **A refit is owed on the 15:43–15:44 change** — three of the 58 columns changed meaning under weights fitted against the old ones. That is not WIRES 114–117; those were measured empty above.
2. **The refit edge needs both instruments.** The fixture is what a weight file can CARRY, and it should stay; the corpus contrast is what can DETECT. Wiring `feature-engine-contrast.json` into `status.js` beside `feature_fixture --check` is the obvious next move, and it is deliberately not done in this pass — a status line added at the end of a session that watched three files move is a line nobody has watched behave.

`data/click-censoring-census.json` and `data/censoring-value.json` are therefore UNSAFE again, now through `engine/fit_policy.js` rather than through the simulator, together with `data/partial-label-em.json`, which is the same cause and was not touched here. `node engine/provenance.js --strict` **exits 1 with 3 UNSAFE.** That is stated, not filed: the re-runs in §17 were valid photographs of the tree at 14:26–14:40 and they say so in their own digests; the tree they photographed no longer exists.

They were not re-run a third time, deliberately. The corpus definition changed by a third (9,230 → 6,055 open-sheet games) inside the same twenty minutes, so a third run would publish a different population under the same headline, attributable to neither the engine nor the censoring change. Re-run both against a still tree — the loader digest is in every artifact — and the numbers in §17 are the ones to compare against.

A measurement cannot be taken while the lens is being changed, and engine_release.js does not cover this case. A release freezes 23 files; it does not freeze `engine/fit_policy.js`, and it cannot freeze the store. That is why `feature_engine_contrast.js` pins its sample by game id and refuses when one goes missing: the first version of it reported **all 58 columns moved** purely because the corpus shrank between two children, which is a REFIT verdict manufactured out of somebody else's edit.

§17a — the board.js partial-body over-refusal is worth 0 rows, and here is the number

ENGINE filed it rather than fixing it: `engine/board.js:2565` and `engine/position_features.js:231` map their priority defenders to `{ability, fainted}`, so `isGrounded()` sees no type list and no item and a Flying-type foe is still over-refused **in the feature vector**. Widening that signature moves the feature vector, which is a refit, which is why it came here. Measured on the fit's own decisions over all 9,230 games, rebuilding every defender twice — once the way `board.js` does it, once with the types and item the board already holds:

	n	of
candidate feature vectors	1,751,688	—

	n	of
with a priority move	332,030	19.0% of candidates
aimed at a body, i.e. reaching <code>board.js:2560</code> 's guard	135,552	40.8% of those
under a Psychic Terrain	362	0.27% of guarded priority candidates
where a complete body changes the answer	0	—

The artifact on disk carries the same measurement over the post-15:40 corpus (6,055 games, 1,136,845 vectors, 220,932 with priority, 91,240 reaching the guard, **273** under a Psychic Terrain, **0** changed, upper bound 5). Two corpora a third apart give the same answer, which is the strongest thing that can be said about it without more Psychic Terrain in the metagame.

The only five rows in the entire corpus where types and item flip the bar are `protect` ×4 and `ragepowder` ×1 — **self-targeted moves, which** `board.js` **never routes through** `priorityRefusedAbove` **at all**, because the branch is guarded on `cand.targetMon`. Counting them as exposure would have overstated it by five rows out of 1.75 million; both counts are recorded here so the guard is visible rather than assumed.

So: NOT WORTH A REFIT, and the exposure is 0 rows in 1,751,688 (upper bound 5, of which 0 are reachable). Two things keep it from being closed. `fails.groundedBodyIncomplete` fires on **100% of 173,478** calls — every single feature-path call is made with a body that cannot answer — so the defect is total and only its consequence is nil; and the consequence is a property of THIS corpus, where 0.27% of guarded priority candidates stand on a Psychic Terrain. A metagame that pairs Psychic Surge with Flying bodies moves that number without anything in the code changing. The right time to widen the signature is the next refit, when the feature vector is moving anyway and the change is free. `engine/position_features.js`'s copy is a separate call site and is NOT measured here.

Reading a run

```
node engine/sprt.js <file>
```

Cat the shards together first. **Never read an interim SPRT** — 66.7% became 44%, 57.7% became 50%. The bound exists precisely so you do not have to look. SPRT is valid under continuous monitoring because its boundaries were derived for it; a Wilson interval read repeatedly is not the same thing and does not inherit that property.

The unit is the **decisive pair**, not the game. In a paired run a 1-1 split means the team decided it, not the policy.

Reading a stamp

```
node engine/run_stamp.js --show          data/rollout-r3.json
node engine/run_stamp.js --reconstruct data/rollout-cost.json
```

Every gate artifact has a `<name>.meta.json` beside it saying which configuration produced it. `status.js` prints the headline under the gate line, so the absence of a stamp is on the same screen as the number — R1's +2.91 was quoted for a day against a dump that could not say which of two runs four accuracy points apart it was, and nothing was hidden then either. The fact simply lived in a file nobody opened.

Three things to check before quoting any of it:

- `reconstructed: true` means **inferred from a commit, not observed**. Read `confidence`, which publishes the gap in seconds between the artifact's own timestamp and the commit that carried it.

- `git.dirty: true` means the commit id does not describe what ran. Trust `source_digests` .
- `source_digests` hashes **worktree bytes**; `git.blobs` names git objects. On Windows those differ by line-ending translation — `data/engine-data.js` does — so never compare one to the other.

`writeStamp()` is the only mode worth trusting, because only the run knows its own settings. `reconstruct()` exists for the artifacts that predate it and labels itself on every line.

18. THE PORYGON2 SEPARATION GATE — PRIORITIES #23. PASS, and the interesting number is the one that is not in the verdict. 2026-08-06

`engine/porygon2_separation_gate.py` → `data/porygon2-separation-gate.json` . The MILTANK leaf redesign (#24) is buildable. PORYGON2 does not collapse a subtree to one number.

39,843 same-game position pairs two turns apart, across 6,328 clean HUMAN ladder games, every interval bootstrapped with the GAME as the cluster. Thresholds were written to disk at **05:59:44Z**, the run wrote at **06:56:06Z**, and `--run` refuses to start unless the declaration on disk matches the block in the generator character for character.

	measured	declared bar			
T1 separation median \	Δ score\	over 2 turns	0.1628 [0.1600, 0.1653]	≥ 0.02	PASS
T2 locality same-game 0.1985 vs unrelated 0.2801; D	+0.0815 [0.0786, 0.0845]	CI lower > 0.0043	PASS		
T2 locality ratio R = same / unrelated	0.709 [0.700, 0.718]	≤ 0.75	PASS		
T3 direction agrees with the material sign	85.58% [85.16, 85.98]	CI lower > 50, point ≥ 60	PASS		
T3 secondary, moves toward the eventual winner	61.59% [61.12, 62.07]	reported, not gated			

All eight PORYGON2 arms pass — 17 and 19 features, plain and weighted, k=50 and k=200 — with R between 0.684 and 0.739. The verdict is read off **17f weighted k=50**, which is what `docs/MODELS.md` headlines.

THE NEGATIVE CONTROLS DID THEIR JOB, AND THE SECOND ONE IS THE ONE THAT MATTERS. A constant 0.5 leaf fails all three (median 0, R undefined, direction 0%). That was the required control and it is the weaker one. A **uniform-random** leaf **PASSES T1 with a median of 0.2924 — nearly twice PORYGON2's separation** — and fails T2 (R = 0.995 [0.985, 1.004], D CI [−0.0012, +0.0049] straddling zero) and T3 (50.28% [49.72, 50.87]). So separation alone cannot tell a value function from noise, which is precisely why T2 was written as the deciding test. A gate proved only against a constant would have been passed by static.

AND THE FINDING THE VERDICT DOES NOT CONTAIN. A bare material count — `0.5 + 0.15·alive_diff`, the same rule `porygon2.py` scores itself against — was run through the identical pipeline as a BASELINE rather than a control. At a two-turn gap **it passes the gate too**: R = 0.703 [0.692, 0.715], statistically indistinguishable from PORYGON2's 0.709. Read alone, that says the 17 features buy no locality at all.

It is not read alone, because the addendum below settles it. **At the ONE-turn gap the search actually operates at, the material count goes flat: its median $|\Delta|$ is 0.000 and it returns the identical number on 58% of adjacent positions**, while PORYGON2 moves on 99.3% of them with a median of 0.1154 and its locality gets *better*, R = 0.5464 [0.5392, 0.5534]. Every branch a material leaf cannot separate is a branch the argmax decides by tie-break. That is the case for #24, and it is a different case from the one the headline makes.

Two comparisons in that block that look like findings and are not, stated so nobody quotes them:

- the material baseline's *toward-the-eventual-winner* rate (64.77%) is **higher** than PORYGON2's (61.59%) — but its score moves on only 21,975 of 39,843 pairs, i.e. only where a Pokemon actually fainted. It is scoring the easy subset. The two rates are computed on different populations and are not comparable.
- the gate produces properly-intervalled accuracies for free: **17f weighted k=50 at 63.11% [62.32, 63.81]** against the material sign's 61.02% [60.06, 61.96] on the same 52,501 positions. These are **separate** game-clustered intervals, **not a paired test**. `docs/MODELS.md` 's 63.59% is still marked **NOT MEASURED** and this **supersedes nothing** — a paired difference with a split-half floor is what would close it, and nobody has run one.

WHAT WAS FROZEN, AND WHAT COULD NOT BE. The gate is stamped to engine release `4c73f9cafa4b` and that stamp is honest about its own limits: **none of PORYGON2's sources are in the frozen set** — not `engine/porygon2.py`, not `data/porygon2-species.json`, not either corpus. PORYGON2 is a Python model and `REL.require` is a JavaScript shim, so it cannot be loaded through a release at all. What the release *did* supply, through `REL.require`, is the thing that decides the population: the frozen `engine/quality.js` + `data/quality-filter.json`. For the rest the generator takes its own photograph — sources copied into a private tree and imported from the copy, live originals re-digested afterwards (none moved) — and the two append-only stores are pinned by the **clean id set** (7,992 ids, sha256 `4ccc0afc...`) rather than by a whole-file digest, because the collector appends hourly and a file digest would void any run longer than an hour.

A DEFECT FOUND ON THE WAY, AND IT IS THIS REPOSITORY'S SIGNATURE SHAPE. The first artifact carried `"R_same_over_unrelated": NaN` — Python's `json.dump` writes a bare `NaN`, which every Python reader accepts and which **is not valid JSON**. `JSON.parse` throws on it. The effect was that `engine/provenance.js` **could not read one field of the file and reported it** `ok`: a clean bill of health issued over a document it had never parsed, including the `void` flag that exists precisely so a generator can condemn its own run. Both dumps now pass `allow_nan=False`, so it raises instead of shipping. Worth a sweep: any Python generator here can emit this, and the artifact still looks fine from Python.

Three smaller things the gate needed and now does:

- it writes `corpus.clean_games` and `corpus.population_ceiling` **spelled the way** `provenance.js` **reads them**. Its prose `population` block was invisible to `declaredGamesFrom()`, which is the §5e state where an artifact "records no game count".
- the declaration timestamp survives every re-run. `--run` overwrites the file, so reading `generated` would report the last run as the moment the thresholds were fixed — drifting later than the numbers, every time.
- a `--run` re-run used to silently delete the `--addendum` block. It is carried forward now, each block keeping its own timestamp and digests.

Disclosed rather than omitted: a 150-game smoke run of this pipeline executed at 06:02Z, after the declaration and before the headline sample, to find bugs. Its numbers were seen first. No threshold changed — the equality check enforces that — but the smoke run's R landed at 0.735 against a 0.75 bar, close enough that saying nothing about it would be the omission this division exists to prevent.

What this gate does NOT establish, and #24 should not be read as having it:

- it says the leaf **separates**, not that swapping it in **wins**. The unit that answers that is the decisive pair, and it needs an SPRT against the incumbent playout.
- the pairs are consecutive positions from *real games*, not **sibling branches from one node**. Siblings differ by one action from an identical board and are more alike than anything measured here. The lag-1 addendum is the closest available proxy and it is a proxy.
- T3's ground truth is `alive_diff + hp_total_diff`, which are two of PORYGON2's own inputs (`alive_diff` carries a learned weight of 5.12 against a mean of 1.0). It asks whether the model respects its strongest

features. A k-NN guarantees no such thing, so it is not vacuous — but it is not independent, which is why the outcome-anchored secondary is reported beside it.

- **no split-half was run.** The noise floor here is built into the design instead: T2's unrelated-pair arm is the floor for the effect claimed, and every interval is game-clustered. The estimator is deterministic given the game set, so a split-half would re-measure what the bootstrap already reports. The one stochastic input — how the unrelated partner is drawn — was checked by a second mechanism: the any-turn control gives $R \approx 0.74$ against the turn-matched 0.709, so the conclusion does not depend on the draw.

19. THE STRONG-PLAYER BASELINE — the cutoff gradient exists and is free; §1.3's "real humans" column cannot be compared to it; and "flat in rating" is NOT MEASURED. 2026-08-06

`data/strong-player-baseline.json`, written by `build/strong_player_baseline.js` (one process, ~2 minutes). Built for task #46 out of `data/smogon-stats/` and this repository's own stores. **It reads no simulator, no leaf, no feature layer and no policy weights**, so it is not invalidated by an engine release or by a MAG refit and does not need re-running when either happens. It is invalidated by a new Smogon month or by a change to `data/quality-filter.json`.

The generator exists because of §19e. The claim this section retracts — *"measured move quality is close to flat in rating"* — has been quoted in a fitted model's own caveat block for weeks with no generator behind it. Shipping an artifact with the same defect would have been the same mistake in a new file. It lives in `build/` rather than `engine/` for a dated reason: it was written on 2026-08-06 while an ENGINE agent was rewriting the simulator, and this division does not add files to another agent's directory mid-flight. If that reason expires, `engine/` is the better home.

The question it was built for is Will's, 2026-08-06, reading `docs/ROADMAP.md` §1.3: *"outright failed' could be incompetence or a high level play and we dont know the difference."*

19a. What the 1630 weighting can and cannot support — stated before it is used

It CAN support what strong players BRING and RUN. Species usage at four skill weightings, and ability / item / spread / move frequencies within a species.

It CANNOT support what strong players CLICK. The files are team-composition aggregates. There is no turn, no board, no opponent and no click in them. All three of §1.3's metrics — *moves that outright failed*, *moves that hit an immune target*, *moves that were super effective* — are per-turn rates and **have no Smogon counterpart at any cutoff**. A 1630 column added to that table would look comparable and would not be, which is worse than a missing column.

"Cutoffs are weightings, not subsets" is now MEASURED rather than quoted. Across both months and both formats, every species' `Raw count` is identical at all four cutoffs — **3,117 of 3,117 species-cutoff pairs**, and 310 of 310 rows of the usage table's `Raw` column. Only `Avg. weight` moves. So no cutoff describes a *set of players*; each describes the same battles seen through a different lens.

And no rating number is mapped onto a cutoff number anywhere in this work. Nothing in this repository or in the Smogon files establishes that 1630 is the same ruler as the Showdown `|player|` rating field. The corpus is located on the cutoff axis by **composition**, which is scale-free.

19b. The gradient exists, and it was free

Four cutoffs (0 / 1500 / 1630 / 1760) × two months (2026-06, 2026-07) × two formats (Bo1, Bo3) are already on disk: 16 usage files and 16 moveset files. Nothing was collected.

Effective sample size is `Raw count × Avg. weight` per species per cutoff. Smogon weights lie in [0,1], so $\Sigma w^2 \leq \Sigma w$ and the true effective sample $(\Sigma w)^2 / \Sigma w^2$ is **at least** Σw — the intervals below are therefore too WIDE, not too narrow. The offsetting hazard is stated and not corrected for: the independent unit is a PLAYER, one strong player contributes many battles, and the files cannot measure that.

The 1760 column costs almost everything. Effective team slots, 2026-07 Bo1:

cutoff	avg weight/team	effective team slots	share of raw
0	1.000	3,529,372	100%
1500	0.512	1,807,038	51.2%
1630	0.068	239,997	6.8%
1760	0.002	7,059	0.2%

The noise floor is the same cutoff across two months, which is an *upper* bound because it contains real metagame drift as well as sampling. Total absolute species-usage difference, summed over 310 ranked species, 2026-07:

contrast	L1 (points)	vs the cutoff-0 month floor of 140.9
cutoff 0 vs 1500	69.4	inside it — 1500 is not distinguishable from the whole ladder
cutoff 0 vs 1630	151.6	above
cutoff 0 vs 1760	195.0	above
cutoff 1630 vs 1760	70.6	inside

83 of the 104 species at ≥1% base usage have a 0→1760 usage change whose 95% interval excludes zero, and the large ones are monotone across all four cutoffs. Charizard-Mega-Y 15.69% → 18.75% → 23.46% → **30.58%** ($\Delta +14.89$ [13.81, 15.96]); Kingambit 23.40 → 36.79 (+13.39 [12.27, 14.52]); Sableye 6.52 → 3.36 (−3.17 [−3.59, −2.74]). (*Population: Smogon 2026-07 `gen9championsvgc2026regmb` , 1,764,686 battles, reweighted.*)

Within a species, the gradient is real for **moves and spreads**, marginal for **items**, and absent for **abilities**. Mean total-variation distance over the top 20 species by raw count, cutoff 0 vs 1760, computed over the **intersection** of the two listed key sets:

section	cutoff gradient	month noise at cutoff 0
moves	17.89	11.52
spreads	9.17	3.72
items	6.98	4.99
abilities	1.63	5.12

The abilities row is the honest negative, and it needs one species named: the 5.12 month-noise is dominated by **Sneasler**, whose Unburden/Poison Touch split really did move 15.3 points between June and July. Excluding it, abilities barely move with skill either — the format's abilities are near-locked at every cutoff.

A correction made mid-run, recorded because a first pass shipped it. A key absent from a Smogon moveset list is **not 0%** — the file lists the top few plus `Other` , so an absent key is below that list's reporting floor. Treating it as zero manufactured an 18-point "gradient" on Venusaur/Energy Ball. Every distance here is over the intersection, with the unlisted mass reported separately.

19c. Where our corpora sit on that axis — and there are THREE of them, not one

They must never share a sentence with only one population named.

corpus	store	filter	rated slots	median	≥1400	≥1500
clean closed ladder	<code>games.ladder.jsonl</code>	clean, 8,047 games	11,852	1266	26.2%	14.4%

corpus	store	filter	rated slots	median	≥1400	≥1500
Bo3 open sheet	games.bo3.jsonl	none	14,539	1175	5.33%	0.72%
Bo1 open sheet	games.ots.jsonl	clean, 2,860 games	3,414	1087	0.23%	0.09%
unfiltered ladder	games.ladder.jsonl	none, 45,006 lines	83,668	1130	6.21%	3.09%

Two things fall out of that table.

The dispatch's figures are confirmed and they belong to the Bo3 store. 14,465 rated slots at p10 1043 / median 1175 / p90 1355 / max 1707, ≥1400 5.3%, ≥1500 0.7% is `data/games.bo3.jsonl`, which now holds 14,539 slots at exactly those quantiles. Same file, 74 slots of growth.

The clean filter moves the population a long way, and that is the population §1.3 benchmarks against. Unfiltered ladder median 1130 with 6.21% ≥1400; clean ladder median **1266** with **26.2%** ≥1400. The bot and behavioural-bot rules remove a large low-and-flat block. So §1.3's "real humans" are not the median-1175 population — that is the corpus MAG was *fitted* on. Any sentence of the form "the ladder is median X" has to say which of the two it means.

On the cutoff axis, by composition (L1 between the corpus's team-preview species vector and each cutoff, mega formes collapsed to base on both sides):

corpus	vs cutoff 0	1500	1630	1760	own split-half floor
clean closed ladder, vs 2026-07	138.6	174.9	230.9	274.6	36.9 – 75.8 (median 53.4)
Bo1 open sheet, vs 2026-06	211.2	184.0	159.2	157.1	45.9 – 91.8 (median 67.3)

The closed ladder is nearest cutoff 0 and moves monotonically away from every higher one. Its 138.6 is about one month of drift at cutoff 0 (140.9) and its games run 2026-07-22 to 2026-08-06 — later than the newest published month — so cutoff 0 is a fit and 1630/1760 are not. **The Bo1 open-sheet corpus cannot be placed at all:** its four distances span 157.1–211.2 against its own split-half floor of 45.9–91.8, so the cutoff ordering is inside its noise. NOT MEASURED for that store, and the apparent preference for 1760 is not a finding.

19d. The Fake Out / Armor Tail case, answered as far as it can be

Farigiraf runs Armor Tail on **97.80%** of sets at cutoff 0 [97.77, 97.83] and **99.11%** at 1760 [98.59, 99.45] — a real +1.31-point gradient [0.89, 1.73] against a month noise of 0.11, and **completely useless for the question**, because the ability was already near-universal everywhere. Incineroar carries Fake Out on 98.89% of sets at cutoff 0 and 99.82% at 1760. Expected Fake Out carriers per team of six: 0.728 at cutoff 0 → 0.751 at 1760.

So the collision is **at least as available** at the top as at the bottom. The aggregate can say the two sides are both brought slightly more by strong players; it cannot say who clicked, because there is no turn in the file.

What this implies for task #44 part 1, which is owned elsewhere and NOT attempted here. The denominator of a failed-move rate is composition-confounded, and the split has to condition on it. Protection is the largest single source of a `|-fail|` line — a repeated Protect fails by rule — and it moves with the cutoff: expected protection carriers per team of six run **4.00** → **4.14** → **4.34** → **4.39** across the four cutoffs, and Detect alone runs 0.103 → 0.144 (+39% relative). A raw failed-move rate therefore rises with how much protection the population runs, independently of anybody playing better or worse.

19e. `fit_policy.js:1264` 's "flat in rating" — NOT MEASURED, not false

The claim: "Open-sheet players also average ~185 rating points lower, though measured move quality is close to flat in rating." `docs/DEFENSE.md` §1 gives the numbers — failed moves 2.59% under 1100 against 2.30% at 1400–1600; blocked actions 4.66% to 3.43%.

Neither figure has a generator in this repository. `engine/realism_report.js` counts the same protocol lines but **pools both players of a game and never bands by rating**, and no `data/*.json` carries a rating-banded rate. The claim has sat in a fitted model's own caveat block with nothing behind it that can be re-run. That is the P1 class this division already named for PORY's coefficients.

Recomputed here from the protocol, attributed to the **acting** side, on the clean closed ladder — 8,047 clean games / 7,040 raw logs matched / 14,078 player-slots / **171,801 moves**. Intervals are a game-clustered bootstrap, 400 resamples.

rating band	moves	failed % [95%]	immune % [95%]	super-effective % [95%]
<1100	21,098	2.218 [1.93, 2.57]	2.318 [2.01, 2.63]	20.84 [20.02, 21.59]
1100–1199	27,390	2.472 [2.08, 3.03]	2.234 [1.98, 2.47]	20.62 [19.97, 21.33]
1200–1299	22,699	2.291 [1.99, 2.59]	2.071 [1.83, 2.31]	20.62 [19.90, 21.32]
1300–1399	22,577	2.197 [1.95, 2.48]	2.210 [1.98, 2.46]	20.84 [20.16, 21.56]
1400–1599	24,669	2.165 [1.83, 2.53]	2.027 [1.81, 2.25]	21.24 [20.55, 21.91]
1600+	7,486	2.658 [1.93, 3.52]	2.151 [1.72, 2.57]	19.16 [17.95, 20.36]

The direction reproduces. <1100 against 1400–1599 on failed moves: **−0.054 points, 95% CI [−0.539, +0.403]**. Nothing here contradicts "close to flat".

But the design was powered for a 0.674-point difference at 80% power on a 2.2% base rate — a 31% RELATIVE change. Anything smaller was never detectable, so "*flat*" and "*an effect up to 30% of the base rate*" are the same observation in this corpus. Immunity, same two bands: −0.291 [−0.679, +0.057], MDE 0.517 on a 2.32% base. Super effective: +0.396 [−0.606, +1.461], MDE 1.527 on a 20.8% base. All three contrasts contain zero.

And the whole between-band spread is inside the within-band noise floor. Failed-move rate across all six bands spans 2.165% to 2.658% — 0.49 points. Cutting a *single* band eight ways produces spreads from **−0.489 to +0.806** points. The observed effect is the noise floor.

The binding constraint is not the rating range, which is what the dispatch expected. The clean closed ladder holds 32,155 moves at ≥1400, 17,551 at ≥1500 and 7,486 at ≥1600. The binding constraint is that the metric is a ~2% event: 25,000 moves is only ~530 failures, and separating two bands by a fifth of a point on that needs far more.

So the verdict is NOT MEASURED. The claim is not shown false and it should not be repeated as though it were established. `fit_policy.js:1264` and `docs/DEFENSE.md §1` should say *not measured at this power* — filed, not edited, because `engine/` is being rewritten in parallel.

What would settle it is not more games at this metric. Either a metric with a higher event rate (super-effective is 21%, so its 1.53-point MDE is **7.3% relative** against failed-move's 31% — four times better), or a paired design that holds the board fixed, which is what a click-level model gives and a rate does not.

19f. What §1.3 should say instead

The gap §1.3 reports between MAG and humans is **3.87 points** (6.34% against 2.47%). The entire measurable rating effect inside the human population is at most **0.67 points** and its point estimate is **0.05**. The gap is ~5.7× anything skill does to this metric within the corpus. **Closing it makes MAG resemble a human; it cannot make MAG resemble a strong human, because on this metric strong and weak humans are not separated at all.**

Five changes, in order:

1. **Name the population on the same line as every figure.** The column is *clean closed-sheet ladder games, both players pooled, no rating condition*. It was 1,905 clean games when written; the same

predicate now selects 8,047 clean games / 7,040 raw logs / 171,801 moves at median rating 1266 with 26.2% of rated slots ≥ 1400 . It is **not** the open-sheet corpus MAG was fitted on.

2. **Give an interval.** A bare percentage invites a comparison it cannot support.
3. **Say the human column is a REALISM target, not a skill target** — and give the reason rather than the assertion: within this corpus the same rate does not separate a sub-1100 player from a 1400–1599 player at a detectable size (-0.054 , 95% CI $[-0.539, +0.403]$, MDE 0.674).
4. **Do not add a Smogon 1630 column to that table.** It has no per-turn rate at any cutoff.
5. **If a strong-player column is wanted, it belongs in a different table about bring and build** — where 1630 and 1760 genuinely answer the question. That table is what this artifact provides.

For reference, the human column recomputed on the current corpus with attribution to the acting side (*clean closed ladder, 171,801 moves over 7,040 logs, 2026-08-06*): failed **2.275%**, immune **2.074%**, super-effective **21.005%**, against §1.3's 2.47 / 1.91 / 21.37 on 1,905 games. Close, and stated so the population and the count travel with the numbers — not as a substitute for re-running `realism_report.js`, which pools rather than attributes.

19g. Filed, not fixed — a real defect found on the way

`data/smogon-priors.json` 's `teammates` array is polluted on 275 of its 284 species. Kingambit's holds 32 entries where the source file lists 10, and the extras are `intimidate`, `blaze`, `sitrusberry`, `passhoberry` — the next species' Abilities and Items rows. The cause is `engine/smogon_priors.js:160: grab('Teammates', 'Checks and Counters')` terminates on a section these moveset files **do not contain**, so the regex falls through to `$` and swallows the rest of a 14-chunk window that already spans into the following species block. The same file's `abilities`, `items`, `spreads` and `moves` are unaffected — each has a section that really does follow it.

Blast radius today: zero. Nothing in the repository reads `teammates` out of that file; the only occurrence is the writer. It is latent, not live. Not fixed here because `engine/` is being rewritten in parallel and this division does not patch another agent's file mid-flight.

`provenance.js` **classes this artifact's corpus as open-sheet, and it is closed.** `provenance.js:268` is `/games\.(ots|bo3)\.jsonl/.test(withDeps) ? 'opensheet' : 'ladder'`, and the generator names both open-sheet stores because it reads their rating summaries. Its **primary** corpus is the clean CLOSED ladder, so drift is judged against the wrong ceiling and the artifact prints *CORPUS DRIFT — declares 8,047 games; 9,177 are clean open-sheet now, 12.3%*. That is the same class as the named exception §5 already records for `winrate-backtest.json`, and it needs the same treatment in `provenance.js`, which is `engine/` and was in flight.

It is left flagged on purpose. `provenance.js` honours a declared `population_ceiling`, and declaring one here would set the ceiling equal to the count and silence this artifact's drift check permanently — which is §5e's `--fix` failure exactly: refitting the world so a check goes green. A false-positive warning that says why is better than a real check switched off.

19h. What this artifact does and does not stamp

`source_digests` is at the **top level**, because `provenance.js:685` reads `j.source_digests` and not `j.provenance.source_digests` — a first run put it in the wrong place and broke the content-digest ratchet in `data/provenance-stamp.json`, which may fall and may never rise. It is closed now; the artifact is one of the **2 of 93** verified by content rather than by mtime.

Only the stable inputs are stamped, deliberately. The generator, `engine/quality.js`, `data/quality-filter.json` and all 32 Smogon monthly dump files — every input that is supposed to be frozen, so a change in one is a real event. The three game stores are **not** stamped and the artifact lists them under `provenance.unstamped_inputs` with the reason: they are append-only, the collector runs hourly, and their

digest changes every hour by construction. Stamping them would hang a permanent mismatch on the artifact, which is §5a's "mtime cries wolf" wearing a hash. The instrument for an append-only corpus is the declared count, and that is what `corpus.clean_games` is for.

One more small trap, recorded because it is generic: `run_stamp.sourceDigests()` adds a prose `note` key to the map it returns, and `provenance.js` iterates every key of `source_digests` as a path to re-hash. Left in, it prints "stamped input note cannot be read to verify" on every run. The generator deletes it and carries the prose in `source_digests_note` instead.

FIXED IN THE READER, 2026-08-07 (§20). Working around it in each generator is a fix that has to be remembered once per generator, and the next one to follow `provenance.js` 's own printed advice pays the same tax — `data/wire-ladder.json` did, on its first run. `run_stamp.js` now `EXPORTS STAMP_NOTE_KEY` and `provenance.js` imports it and skips exactly that key. One place knows which key is prose, which is the `FACTS-ARE-GLOBAL` rule applied to a two-line string. Same shape as the frozen-release false positive already handled twenty lines below it in that loop: a checker that penalises the workflow it recommends gets ignored.

Living-docs obligations this pass did NOT discharge, because the dispatch scoped the write to two files and forbade `status.js --write`: the `CHANGELOG` entry and version bump, and whether `SUMMARY.md` or the white paper should carry the §1.3 correction. `docs/ROADMAP.md` §1.3 itself is unedited — the rewrite is specified in 19f and is the router's call to place.

20. THE RELEASE LADDER — what one night of WIRE fixes bought, controlled. MEASURED 2026-08-07.

`engine/wire_ladder.js` → `data/wire-ladder.json`. Nine frozen releases plus a repeat of the baseline, **1,995 games each**, one pinned census (`data/wire-ladder-census.pin.json` , `f63179105d3c`) and one team pool (`3d0112fce455`). `arms_comparable.compare()` cleared **all nine** arms against the baseline; the eleven watched inputs were byte-identical before and after; the planted-divergence proof and all seven equivalence rules passed on every arm. This replaces the pairwise before/after retracted in `CHANGELOG 3.62.1`, and it replaces WIRE 6's own artifact, which pinned its census to a path inside an agent scratchpad.

The instrument is deterministic and that is now demonstrated, not assumed. The pre-WIRE-1 baseline ran first and last with eight arms between: every measured field identical and the per-game divergence depth identical game for game. The whole ladder was then run three times end to end and reproduced. So every difference in the table is the engine change.

The headline is a negative and it should be read first. The median game still parts after ONE completed turn, at every rung. Six wires did not move it. Whole-game agreement went **2 → 22 games of 1,995** — 98.9% of games still diverge. The divergence rate is saturated and says almost nothing; what moves is the `DEPTH`, and the useful unit there is the protocol line, not the turn:

	baseline	after six wires
median first-divergence line	13	14
mean	15.01	23.97
p90	30	57
games that never diverge	2	22
median completed turns	1	1

Paired over the same 1,995 games, the top rung parts **later on 742, earlier on 141, at the same line on 1,112**. More than half the sample is untouched by the entire night, and the median delta is zero.

Per rung, the one number each is worth (paired against the rung before it; "net" is games parting later minus games parting earlier):

rung	net later	its own class
mega resolution order (unpublished intermediate)	+73	ordering 247 → 170
Knock Off base-power truncation (intermediate)	0	4 games reclassified, nothing else
WIRE 1 Protect/crash	+65	-miss field 3 18 → 1
WIRE 2 stall counter	+156	unrelated event mismatch 700 → 562
WIRE 3 refused stat drops	+99	event missing 673 → 636
WIRE 4 fixed-point damage	+16	-damage field 3 216 → 179
WIRE 4 recoil/drain rounding	+48	-damage field 3 179 → 141
WIRE 6 priority brackets	+287	turn order 85 → 3

Three things in that table are worth more than the ladder itself:

- **An unpublished intermediate outranks a named wire.** The mega-resolution-order cut (28e66a7c9ab8 , 02:36) is worth **more than WIRE 1 on every measure here** — net +73 against +65, and ordering 247 → 170 (−77 games) against -miss field 3 18 → 1 (−17). It sits between the baseline release and WIRE 1, so a pairwise baseline→WIRE-1 comparison credits WIRE 1 with all of it and reports WIRE 1 as more than twice its true size. (The largest rung of the night is WIRE 6 at +287; the point is the misattribution, not a new champion.)
- **An unambiguously correct arithmetic fix can be worth zero at the whole-game level.** The Knock Off base-power truncation moved the divergence position in **0 of 1,995 games** and reclassified four. It is right, it is not measurable here, and those are different statements.
- **A class count can fall because a game parts EARLIER on something else.** -damage field 3 RISES 170 → 216 over WIRES 1-3 and then falls to 141 — the earlier wires push games deeper, which exposes damage divergences that had been masked. So a per-class delta is only readable beside the depth column, and event missing from medicham2 (604 → 627) growing is not a regression.

141 games part EARLIER than the baseline after six correct fixes, most of it appearing at the mega and WIRE 1 rungs. That is not a contradiction: changing a trajectory surfaces a different pre-existing bug sooner. It is 7.1% of the sample and it is the reason "net later" is reported rather than "later".

Coverage, controlled: distinct moves connected 224 → 261 (+37). The wires' own reports claimed 173 → 197 on ~346-game uncontrolled arms. The absolute levels are not comparable — a different census steers a different sample — but the controlled delta is **larger** than the claimed one, which is WIRE 4's pattern again: the findings were real and the numbers were wrong.

What remains at the top rung, in cause order, because a ladder should end in something actionable: [from] hospitality heals medicham2 emits and Showdown does not (127 games across two classes, the single largest cause in the file), Illusion (zoroarkhisui on every switch: a different body), -prepare for two-turn moves, -activate|feint , and -end|throatchop . All of it is data/wire-ladder.json → what_remains_at_the_top_rung , and all of it is ENGINE's.

Carried forward, not fixed here: every arm reports trace_body_off_field 54-69 — a ?? identifier reaching the medicham2 stream, which tests/test-protocol-trace.js PART 6 says must read 0. It is present in all nine releases including WIRE 6, so it is not something the night introduced. ENGINE's.

What the ladder still cannot see, restated rather than implied: an uncommitted edit inside SHOWDOWN_PATH . The other two blind spots arms_comparable.js declares — the driver itself and data/protocol-events.json — are digested before and after every arm and recorded in the artifact.

21. THE VIEWER WAS AN INSTRUMENT AND NOBODY HAD AUDITED IT. 2026-08-13

engine/divergence_cards.js renders diverging games for a human to read. It computes nothing by design, which is exactly why it was never checked — and it carried three claims that were false.

what the page said	what was true	why
"Sixty diverging games out of 209"	80 of 655	typed into the template weeks ago
"209 / 815 diverged"	top 309 / 1,539, bottom 346 / 1,539	same
"80 of 160" (after the first fix)	80 of 655	of_diverged tallied the POOL, which is capped at 2× the dump

The third is the interesting one, because it survived a fix. Deriving a number is not enough if the field you derive it from is itself a cap wearing a population's name. It is the same error as the coverage credit before ROADMAP #91: a figure that is SMALLER than the truth still misleads, because it is read as the truth.

And the sample itself was one corner. The dump drew from the primary arm alone, so the page answered "where do the engines part when every sub-100 move misses" while presenting itself as "where do the engines part". Both arms now feed it, interleaved and labelled.

Rule this leaves behind: a page that renders a measurement is part of the measurement. It gets the same treatment as an artifact — every number derived, every narrowing declared, and the population distinguished from the sample by name.

Running the release ladder

```
SHOWDOWN_PATH=C:/Users/willj/Projects/Pokemon/pokemon-showdown node engine/wire_ladder.js --write
```

About 7 minutes, one process, ten arms. `--keep-arms <dir>` keeps each arm's full differential artifact; without it they go to a temp directory and the ladder file carries the numbers. It exits non-zero if any arm is refused as incomparable or if the two baseline runs disagree. **It does not run** `tests/test-mechanics.js` **and neither should anything measuring beside it** — that regenerates the census, which is the steering input the ladder pins.

Running the backtest

```
SHOWDOWN_PATH=C:/Users/willj/Projects/Pokemon/pokemon-showdown node engine/backtest_winrate.js
```

About 13-18 minutes, one process, on the clean-game corpus; the corpus size and the per-run wall clock are withheld with `data/winrate-backtest.json`, which is downstream of MEDICHAM. `MAXG=200` thins it for a smoke run, and the artifact records the n it actually scored, so a thinned run cannot be mistaken for a published one. It writes `data/winrate-backtest.json` and the per-game rows beside it. Every seeded configuration reproduces bit-identically across runs; the unseeded legacy `winProb2` arm moved 0.26 accuracy points between two full runs, which is its run-to-run floor.

It stamps the sha256 of every source the leaf reads. `status.js` re-hashes those and prints `CURRENT` or `PRE-CHANGE`, which is a comparison rather than an mtime inference — a checkout moves an mtime without moving code, and the 2026-08-02 artifact was quoted for two days against an engine that had gained "one mega per side" in between.

Done looks like

- `status.js` prints a leaf calibration line that is fresh, adequately powered, and states the reliability curve rather than a verdict string. **Done 2026-08-04** — and the answer is that the leaf is worse than a coin.
- `provenance.js --strict` exits zero.

- Every gate R1–R4 has an artifact. **Done 2026-08-04** — and R1's turned out to disagree with the prose it replaced. An artifact per gate is the floor, not the goal.
- Every gate artifact says which configuration produced it. **Done 2026-08-04 for R2 and R3** via `engine/run_stamp.js` ; R1's dump has its own inline copy of the same shape and should call the module. An artifact that records its build still is not enough on its own: R3 records its build and **not its control**, and a divergence rate without the self-disagreement floor beside it is a headline, not a result.
- `REFIT OWED` is either clear or has a dated reason next to it.