

MILTANK — the search player

Named for the classic Rollout user. `engine/rollout_leaf.js` is the mechanism; MILTANK is the player that uses it.

1. What it is, and why it is not MAG

ABRA has two players and until 2026-08-03 only one of them had a name.

MAG (`engine/magnemite.js`) is a fitted linear policy over `board.js` features, trained to predict **what a human clicks**. It is an imitator. Given a position it scores every legal option against a weight vector fitted on 60,000 corpus decisions and takes the best — one step, no lookahead, no notion of what happens next.

MILTANK decides by **playing the position out**. It forces a candidate pair of clicks, lets the game run to a conclusion, and takes the line that wins most. It has no fitted weights of its own and no opinion about what a person would do. It can therefore prefer a move that appears nowhere in the corpus, and it routinely does.

They are not rivals and MILTANK is not a replacement. The shipped arrangement is:

*MILTANK searches. When the search **cannot separate its options** — the finalists fall inside one standard error, or the position is already decided — it hands the decision back to MAG.*

That fallback is the point rather than a hedge. An argmax over estimates that overlap is a coin flip wearing a search's clothes, and a fitted human prior is a better tiebreak than dice. **Imitation is the floor; search is the ceiling.**

2. Where MILTANK decides

Decision	Before	Now
Which 4 to bring, which 2 to lead	<code>return 'default'</code> — packed order	90 brings enumerated and played out
Both clicks each turn	MAG's argmax	successive halving over every legal pair
Post-KO replacement	one-step heuristic	each candidate played out

The preview one matters most. Showdown's `RandomPlayerAI.chooseTeamPreview` returns the literal string `default`, `magnemite` never overrode it, so **every game ever played against this bot had its lead decided by where a Pokémon happened to sit in a team string** — on what `prior_player.js` itself calls *"the single largest branch in the game"*.

All three now call the same leaf. Until 2026-08-04 the preview did not: it held its own playout loop, calling `battleInit` / `battleTurn` directly with deterministic greedy on both sides, and never reached `rolloutWinProb` at all. MILTANK therefore shipped **two players**, and only the in-game one was ever swept — so the measured 53.22%-preview against 50.99%-in-game contrast was partly a contrast between two *implementations* rather than between two settings of one. `rollout_leaf.runPlayout` is the single implementation now and the differences are parameters. Greedy is still reachable; it is `previewExplore: 0`. See [SEARCH.md](#) open item 0 for exactly what moved.

The preview also seeded a game that had not started. It passed `battleInit({seeded: true})`, which exists to stop a *mid-battle* leaf re-firing entry effects that already happened. At preview nobody has entered yet, so it deleted the whole switch-in class from the one decision they matter most for: no Intimidate, no

Drought, no Drizzle, no Sand Stream, no Snow Warning, no terrain setter on turn one. Deciding a lead is largely deciding who eats an Intimidate and the search could not see one. `seeded` is a parameter of `rolloutWinProb` now, default unchanged, and preview passes `false`.

3. The academic position, and where MILTANK sits in it

3.1 This is a simultaneous-move game, and MILTANK does not solve it

Pokémon is not turn-based. Both players commit without seeing the other, which puts it in the **simultaneous-move** class where a deterministic best response is exploitable by construction — the matrix game at each node can have no pure equilibrium.

The literature offers several MCTS variants for this setting: **Decoupled UCT (DUCT)**, sequential UCT, Exp3, and regret matching (Lanctot et al., CIG 2014; Tak, Lanctot & Winands). DUCT lets each player pick from their own statistics without modelling the other. It is known **not to converge to an optimal strategy even in a single-state game** — and it still performs best empirically across their test set.

MILTANK is weaker than all of these and the docs should say so. It computes a *best response to a fixed opponent policy*: the opponent plays `chooseAction` during the stepped turn, identically for every candidate. That makes the comparison across candidates like-for-like, which is what the measurement needs, but it is not an equilibrium and it is exploitable by anyone who models it.

It is the right first cut for one stated reason: if a best response does not diverge from MAG, a mixture over the same cells will not either. `data/rollout-r3.json` records 72.9% divergence on 70 decisions, so it does diverge, and the equilibrium version is worth building.

3.2 The playout is fully random, and that is deliberate

The single most counter-intuitive result in this literature is that **a stronger playout policy can make the search worse**. Policies that maximise MCTS performance as playouts are *weak* at playing the game (Are Strong Policies Also Good Playout Policies?, AIIDE; An Analysis of Monte Carlo Tree Search). What matters is how a policy *ranks* actions relative to each other, and **low-variance policies are dangerous** — a deterministic greedy playout replays the same line from one position, so N samples carry barely more information than one.

ABRA measured exactly this pathology before reading the literature. `chooseAction` is deterministic greedy; 53% of positions landed in the 0–10% or 90–100% bin and those bins were wrong by 22–29 points, and accuracy was **flat in N**. R1 then found a fully random playout judges a position at **68.18%** against **64.42%** for the greedy one.

That comparison was withdrawn on 2026-08-04 as UNCHECKABLE. The 68.18% was never written to an artifact — `engine/rollout_r1.js` printed it — and the one committed row dump turns out to hold the *greedy* column, so the figure could not be recomputed from anything in the repository. `data/rollout-r1.json` records what survived, the greedy arm, and **that figure is QUARANTINED — withheld, not annotated**: the artifact is built from a dump of games MEDICHAM played, MEDICHAM is not correct, and this is R1 leaf accuracy, which CLAUDE.md names in the quarantine list. `node engine/status.js` names the failing clauses.

It was re-run the same day, and the re-run is quarantined with it. `data/rollout-r1-explore-sweep.json` scored a three-arm sweep over one identical set of positions. The table that stood here carried accuracy, Brier, log-loss, ECE and a saturation share for greedy, explore 0.5, explore 1.0 and the material baseline, plus the paired lift and its interval; all of it is absent rather than captioned, and the reader may not

infer from the absence whether the retracted comparison reproduced. Both become quotable again when the gate opens AND these are re-run: `node engine/rollout_r1_artifact.js` and `node engine/rollout_explore_sweep.js`.

What survives is the DESIGN and the retraction's REASON: the original figure was never written to an artifact, the one committed row dump held the greedy column, and a figure that cannot be recomputed from anything in the repository is uncheckable whatever its value.

So the retraction was right about the provenance. Whether the CLAIM survives it is not stated here while the gate is closed: `--rollout-explore = 1.0` **ships and its defence rests on an artifact that is currently unquotable**. One correction that came with it does not need a figure: the greedy accuracy quoted in the comments at `rollout_leaf.js:147` and `mag_bot.js:145` does *not* reproduce on either the committed dump or a fresh run, so those comments overstate the gap and should cite the sweep artifact instead of a typed number. And the saturation story is confirmed rather than merely argued: the greedy playout puts **half** its positions in the two extreme bins and `explore=1.0` puts under a third, halving the calibration error. A deterministic playout replays one line, exactly as §3.2 claims.

Caveat, measured 2026-08-03: that result is about *judging a position*. Judging and *choosing an action* are different jobs, and a random opponent never punishes a wasted turn. A direct test on 40 corpus positions found Protect scoring -4.3pt at `explore 1.0` and -5.2pt at `0.0` — i.e. the randomness is **not** obviously overpricing tempo. That is one measurement on one mechanic and it should not be read as settling the question.

3.3 The picker is a fixed-budget best-arm identification problem

A turn offers up to ~ 90 legal pairs. Estimating each and taking the max is **not** the same as finding the best one: the max over K noisy estimates is biased upward by roughly the spread of the noise, so the winner is disproportionately the arm that got lucky. At $n=120$ a win rate near 0.7 carries a standard error around 4 points, and the max over 63 such arms is inflated by ~ 2 SE of pure dice — larger than most real differences between two reasonable clicks.

The symptom was visible in play: **the bot clicked Recover at full HP**, a move MEDICHAM correctly heals 0 with. It was not choosing Recover; it was choosing whichever arm rolled well.

This is textbook **fixed-budget best-arm identification**, and the standard remedy is **successive halving**: spend the budget where it discriminates, not uniformly ([sequential halving](#) / [fixed-budget BAI](#)). MILTANK screens every pair at $N/3$, keeps the top 8, and re-tests those at $2N$ **with fresh seeds** — so an arm that advanced on lucky dice has to roll them twice. It is also cheaper than the flat version it replaced: $63 \times 40 + 8 \times 240$ beats 63×120 .

3.4 On-line policy improvement

The whole shape — take a base policy, improve it at decision time by rolling out — is Tesauro and Galperin's on-line policy improvement ([arXiv:2501.05407](#)). MAG is the base policy; MILTANK is the improvement operator. That framing predicts the fallback behaviour: where the rollout has no signal, the improved policy should return the base policy, which is what the noise-floor check does.

3.5 What the field says about the domain

[VGC-Bench](#) puts the team-configuration space at $\sim 10^{139}$, larger than Chess, Go, Poker, StarCraft or Dota, and ships 700,000 human battle logs with heuristic, LLM, behaviour-cloning and multi-agent-RL baselines. Two findings bear directly on ABRA:

1. In **single-team mirror matches** their methods beat a professional VGC competitor. Strength in a narrow team distribution is achievable; that is roughly MILTANK's situation, playing a fixed pool of 230 open-sheet teams.

2. The agent that is best on **one** team is **more exploitable and generalises worse** across teams. This is a warning aimed squarely at ABRA: a bot tuned against the 230-team pool should not be assumed to hold up outside it, and any strength claim has to name the team distribution it was measured on.

3.6 The seeded leaf assumed the mega and forgot the weather — fixed 2026-08-04

`dmgMon` builds a stone-holder with the **mega's** ability: a Charizard carrying Charizardite Y arrives in the playout with `drought`, not `blaze`. `battleInit({seeded:true})` then suppresses entry effects, correctly — the actives are already standing there and re-firing Intimidate would drop the same Attack twice — so the sun that ability implies never fired. `applyMegaWeather` existed to close exactly that gap and **its write was discarded by the very next line** for as long as it existed, so Mega Charizard Y stood in clear weather in every mid-battle rollout this project ever ran.

Fixed by applying the caller's field first and the mega's weather second, so the function's own guard arbitrates. **It is now mega-gated, and that gate is the load-bearing half:** the discarded version had no mega check, and un-discarding it without one would have invented a Torkoal's sun on a board the tracker says is clear. The board is authoritative for a body that is what the board says it is; a mega is the one case where the rollout's body is something else.

Measured, in one process so the engine cannot differ between arms: **15 of 250 corpus boards moved, every one of them a mega weather setter standing in no weather, and nothing else moved at all.** Charizard-Mega-Y's sun is worth **+11.0 pt to the side that owns it and -12.5 pt to the side facing it.** Full parity table, the direct playout counter, and the much larger weather defect this uncovered are in [SEARCH.md](#) items 5 and 5b.

Every leaf number in this document predates it. R1, R2, R3, R4 and the calibration were all computed with the mega's weather missing.

3.7 The search spoke a third dialect of terrain — fixed 2026-08-04

Same shape as 3.6 and worth stating beside it, because it is the second time a field the leaf was *given* never reached the engine. ENGINE routed every terrain read in `medicham2-browser.js` through `terrainId()` — `weatherId()`'s sibling — and then measured that **0 of 863 terrain-carrying boards reached the leaf at all.** The cause was upstream of the engine entirely: `miltank.js` and `rollout_r1.js` built their field object with

```
terrain: ['electric', 'grassy', 'misty', 'psychic'].find(t => board.hasField(t)) || ''
```

Those are the **engine's** words. `board.startField` stores the dex's `move.terrain`, which is `electricterrain`, `grassyterrain`, `mistyterrain`, `psychicterrain`. So a third spelling was being probed against a Board that has never held it, it matched nothing on every board ever walked, and the field handed to every rollout carried `terrain: ''` whether or not a terrain was up. Terrain was correct in the engine and invisible to the search.

Fixed by one helper, `rollout_leaf.terrainOnBoard`, which probes the board's own keys and translates with `MEDI.terrainId`. **No fourth map was written**, and the post-KO replacement search — which had `terrain: ''` hardcoded — now reads the board too.

The honest size of it, measured rather than claimed: near zero, and that was the expected answer. On 3,256 corpus boards, 29 carry a terrain (0.9%; ENGINE's whole-corpus figure is 1.24%) and the old probe found 0 of them. Rolling out all 29 plus 176 terrain-free controls at n=40, **4 boards moved, all four Psychic Terrain, mean $|\Delta|$ 8.75 pt, and 0 of 176 controls moved.** Every Electric Terrain board was unchanged, because the engine's terrain readers are thin: the priority block, Grassy Glide's priority, Hadron Engine and the two `terrainScaled` moves. A rare condition times a thin reader is a small number, and reporting it as one is the point.

3.8 The seed dropped the dead, so every playout believed nobody had died — half fixed 2026-08-13

Same shape as 3.6 and 3.7 and the third instance of it: a fact the real position holds that never reached the engine. `battleInit` derives the field, the bench **and the roster** from one array, and `rollout_leaf.buildSide` dropped every fainted body before handing it over — correct for the first two, and it deleted the side's dead from `fallenCount` 's denominator. Every playout ran with a fallen count of **0**: Last Respects at 50 where the position says 150, on a move whose whole identity is that it grows as your team dies.

The corpses are reconstructed rather than the count being passed in, which is Will's statement of the design and not a style preference: *"miltanks rollout needs to just play the game out on medicham and have it match showdown perfectly thats the whole point. miltanks just chooses the actions."* A `fallen: 2` field on `battleInit` would have made the number right and left the seed handing MEDICHAM a side of four where the real side has six.

They come from `board.graveyard` , not from the fainted actives — a Pokémon that died and was replaced was never in `sideTeam` at all — and they are **appended after the living**, so the actives/bench split and every living body's index are unmoved.

Measured on 13,592 open-sheet games (`data/rollout-fallen-prevalence.json`): **8.75% of decision points** have a death already on the acting side *and* a fallen-count carrier brought, and among those the mean fallen count is 1.67 — Last Respects priced at 50 where it should have averaged 133.5.

It is fixed from turn TWO only. `battleInit` stamps `fainted: 0` and the recount is at turn end, so the turn `rolloutAfterActions` forces the candidate click on still prices the move at its floor. That is one line in `medicham2-browser.js` , ROADMAP #246, and it is ENGINE's — **closed the same day by ENGINE**, so the fallen count is now right from turn one. Four further seeding approximations were swept out beside it and were closed in one batch a few hours later; see §3.9.

3.9 The seed was wrong often enough to close four rows — 2026-08-13; the headline share is withheld

Fourth, fifth, sixth and seventh instances of the same shape as 3.6, 3.7 and 3.8, closed as **one batch** because they are one surface: ROADMAP #247, #248, #249, #250. Will's principle is what makes each of them a defect rather than a tradeoff — *"miltanks rollout needs to just play the game out on medicham and have it match showdown perfectly thats the whole point. miltanks just chooses the actions."* The seed is the only place a correct simulator can still produce a wrong game.

the seed said	the position says
a benched Pokémon is whole, unstateded, and carries the dataset's four moves	it is at 20% and burnt, and its sheet declares its own four
every body may Fake Out	a body that has already taken a move action may not
no hazards, no screens, no Gravity	Reflect is up and the rocks are down
Supreme Overlord entered with nobody dead	it walked in over two graves

Measured over 14,102 open-sheet games / 190,378 decision points (`data/rollout-seed-prevalence.json` , a store scan that plays no game): a first-turn-only move was offerable and illegal at **17.7%** of decision points, a benched body is hurt or stateded at **17.1%**, something is up on the field at **13.1%**, and a Supreme Overlord snapshot is wrong at **0.061%** — that last one correct and very nearly inert, because Kingambit is almost always brought with Defiant. Each is a ceiling on reach, not a count of flipped argmaxes.

THE MOVESET SHARE AND THE "AT LEAST ONE OF THE FIVE" SHARE ARE WITHHELD — ROADMAP #402, 2026-08-23. Both were in this paragraph and in this section's headline and both are removed rather than footnoted. The generator resolved a species through `norm()` , which strips the hyphen

MC.mons keys every forme with, so **every forme** was handed an empty dataset moveset and scored as *"the moves differ"* — inflation only, in one direction, by an amount nobody has published. The four figures above are computed independently and are unaffected. The pinned re-run that would restore them, and why an unpinned one would not, are in [SEARCH.md](#) R14.

Two details are worth carrying because getting them wrong looks like getting them right. **The board quantity Fake Out needs is not** `turnsActive` — that counts turns on the field, so a lead reads 1 on the turn it may still Fake Out — it is Showdown's `activeMoveActions`, counted where a move happens. And **whose side a hazard lands on is read, never re-derived**: ROADMAP #254 resolved it once at the write, so the seed takes the board's per-side record straight; a second flip would re-introduce that bug one layer up and cancel invisibly if only one seat were tested.

Five more were swept out and registered. **#271 closed 2026-08-14 — see §3.10**. Still open: #270 (the seeded field has no clock, so a weather with one turn left lasts sixty), #269 (every durable volatile), #268 (a permanent hazard is given a one-turn duration and layers are not counted), #267 (a status is seeded and its counter is not). Full account in [SEARCH.md](#) R14.

3.10 A knocked-off item was still on the board's body — fixed 2026-08-14

Eighth instance of the shape §3.6 to §3.9 record, and the first one that is **not** only about the seed:

`board.switchIn` **copied** the item off the team sheet onto the slot object, `board.noteItem` — the one thing `| -item| / |-enditem|` reaches — wrote only `itemNow`, and `sheetItem()` was the sole reader of `itemNow`. So the board held two answers to one question and `dmgMon` read the wrong one.

```
declare Life Orb -> noteItem('p1','garchomp','') ->
  sheetItem = '' (correct)      slot.item = dmgMon(...).item = 'lifeorb' (stale)
```

A Life Orb, a Choice Scarf, an eaten Citrus Berry and a spent Focus Sash therefore kept applying — **in the damage and speed features MAG scores with as well as in every seeded payout**. CLAUDE.md states the mechanism exactly and states the rule that was broken: *"the damage and speed calculations keep applying an Assault Vest, Choice Scarf or Life Orb that is gone"*, and **PREFER OBSERVED OVER DECLARED**.

The fix is ONE SOURCE and a write-through was refused. `mon.item` is now an accessor calling `board.sheetItem`, so every existing reader — `dmgMon`, `monSpeedMult`, the mega-stone check, the choice lock — asks the one function without being edited. Making `noteItem` patch the slot objects would have fixed the symptom and left two places answering "what is it holding", which is what created this. **Four more readers of the declared item were swept out first and all five landed together**; `miltank.js`'s team-preview builder also reads the sheet and is CORRECT, because at preview nothing has been knocked off yet.

Measured on 14,288 open-sheet games (`data/rollout-item-prevalence.json`, a store scan): **3.622% of decision points have a body priced holding an item it does not hold, 3.197% of them a body actually on the field**. That is a **FLOOR**: the store records no item consumption at all, so every spent Sash and eaten berry is missing from it by construction and the live path sees strictly more.

It does not invalidate the fit, and that was measured rather than assumed. `noteItem` has exactly one caller, `magnemite.js`, so the offline board never sees an item event and the accessor returns the same string the copy did — `engine/feature_fixture.js`'s 58 per-feature column hashes are byte-identical before and after. The remaining half is PRIORITIES 13e and is MEASURE's: the FIT sees no item event at all, so a declared item still stands for a whole stored game.

4. Open team sheets are what make the preview search possible

MILTANK's bring/lead search is only meaningful because `|showteam|` hands over **the opponent's entire team with full sets**. A fitted bring prior can say "*Whimsicott leads 53% of the time*" — a fact about the population, not about this game. The sheet is information no usage prior can encode.

The opponent's own bring is **marginalised, not assumed**: each playout samples one of their 15 possible fours. Fixing a guess would optimise against one opponent out of fifteen and call it a plan.

Known weakness: that sample is *uniform*. A real opponent is not equally likely to bring every four. Weighting it by `bring_priors` or by which four best answers MILTANK's own team is unbuilt.

The sample is now shared across brings. Each candidate bring used to be scored against its own independently seeded opponent draws, so the difference between two brings sat underneath that noise; one seed now serves the whole preview, so playout *i* of every bring faces the same four and the same dice. Common random numbers, the identical fix the post-KO replacement search needed for the identical reason — there, five of five candidates fell inside their own error bar and the search deferred to MAG on every replacement of a live game.

And the enumerator was wrong whenever a body would not build. It mixed positions in the buildable list with team indices; on a six-mon team with two unbuildable bodies it produced 19 brings where 6 exist, 18 of which named a Pokémon it had just declared unbuildable, and scored three-mon brings as four-mon ones. Full teams enumerated 90 correctly, which is why it survived. Fixed 2026-08-04.

5. What MILTANK is not

- **Not an equilibrium.** Best response to a fixed policy (§3.1). Exploitable by design.
- **Not a tree search.** One step of forced actions, then a playout. No selection, no backprop, no tree. Calling it MCTS would overclaim.
- **Not measured to be stronger than MAG.** R1 (leaf accuracy), R2 (cost), R3 (divergence) all pass. **R4 — the SPRT against shipped greedy — has never been run.** Every claim in this document is about mechanism, not about winning more games. As of 2026-08-03 the bot's record against its author is 0–8.

6. Files

File	Role
<code>engine/miltank.js</code>	the player — the three decision points, installed onto a magnemite instance
<code>engine/rollout_leaf.js</code>	<code>rolloutWinProb</code> , <code>rolloutAfterActions</code> — the leaf; <code>runPlayout</code> is the playout
<code>engine/mag_bot.js</code>	the live wiring and the CLI flags; <code>--miltank</code> (<code>--rollout</code> aliases it)
<code>engine/medicham2-browser.js</code>	the doubles engine the playouts run in
<code>engine/rollout_r1/r2/r3.js</code>	the gates. R5, the action-ranking backtest, is specified in SEARCH.md and unbuilt
<code>docs/ROLLOUT-design.md</code>	the original design note

Flags

Flag	Default	Meaning
<code>--miltank</code>	off	use the search player
<code>--rollout-n</code>	200	playouts per candidate at the final stage
<code>--rollout-explore</code>	1.0	playout randomness, now for the preview too . Re-earned 2026-08-04, <code>data/rollout-r1-explore-sweep.json</code> : +2.25 [1.31, 3.19] over greedy as a

Flag	Default	Meaning
		JUDGE. Not measured as a PLAYER — <code>mew.js</code> has no <code>--miltank-explore</code> (PRIORITIES #33)
<code>--rollout-turns</code>	60	horizon before a playout is scored
<code>--preview-n</code>	40	playouts per candidate bring
<code>--preview-ms</code>	15000	preview deadline; whatever was scored by then is reported
<code>--miltank-foe</code>	uniform	the playout opponent: a coin flip, or <code>prior</code> — what the species really clicks. Now reaches the preview ; before the unification the preview ignored it
<code>MILTANK_TIMING=<path></code>	off	write the per-decision wall-clock artifact. Reduce with <code>node engine/miltank.js -reduce <path></code>
<code>MILTANK_CLOCK=1</code>	off	adaptive spend against the bank. Env rather than flag because <code>mew.js</code> and <code>mag_bot.js</code> are other divisions' files
<code>MILTANK_EARLY_DEFER=1</code>	off	skip the finalist round when the screen says the position is heading for the tie band. Changes what is clicked — its own SPRT arm, not a timing claim

One option has **no flag yet**: `previewExplore` (`DEFAULTS.previewExplore` , null = follow `explore`) is the knob that runs the preview greedy again. Wiring it needs a one-line `arg()` in `mag_bot.js` and `mew.js` , neither of which is SEARCH's file — the same one-liner PRIORITIES #33 already owes `--miltank-explore` .

The budget — and the timer has never been on in a game anyone has watched

`budgetMs` is 20,000 and `previewMs` is 15,000. Our format's VGC Timer (`data/rulesets.ts:778`) is `Timer Starting = 420` , `Timer Grace = 90` , `Timer Add Per Turn = 0` , `Timer Max Per Turn = 55` , `Timer Max First Turn = 90` , `Timeout Auto Choose` , `DC Timer Bank` — one drawn bank, no refill, and the per-turn cap is `Math.min(bank, 55)` so the two are **one clock**.

The bank is 420 s, not 510, and this file said 510. `room-battle.ts:209` does start at `starting + grace = 510` , but `updateTurn ( :305-306 )` re-imposes `secondsLeft = Math.min(secondsLeft + 0, starting)` on **every new turn** and `starting` is 420. The grace is **use-it-or-lose-it on the first timed request** and is not bankable. So `budgetMs: 20000` buys **21 decisions**, not the 24 this file claimed, with preview free because it comes out of grace that is about to be clamped away regardless.

And the bank does not tick while the timer is off. `secondsLeft` falls only in `nextTick ( :353 )` , scheduled only by `nextRequest` , which returns at `:320` when nobody has requested the timer.

`Config.forcetimer` is off on the main server, so the timer is **per player and can be switched on mid-game** — Will, 2026-08-04: *"when i play the bot i have turned the timer off"* , which is why a constant has always looked fine. An opponent typing `/timer on` at turn 9 nevertheless finds a **full** bank, not one MILTANK has already spent. The mid-game switch-on is benign.

MILTANK can see the clock exactly and does not read it. `:332` privately sends the player `|inactive|Time left: X sec this turn | Y sec total | Z sec grace` on every request while the timer is on (the real bank is `Y + Z` ; `:330` subtracts the grace out of the total). `mag_bot.js` parses **zero** `|inactive|` lines. `bot.noteClock()` and the counter `bot.clockStats().notes` now exist here for OPS to wire; **that counter is o today**, and the adaptive rule is therefore designed to be correct without it.

Adaptive spend is implemented, behind `clock` , DEFAULT OFF — `(bank - reserve) / expected remaining requests` , clamped under the 55 s wall, only ever *lowering* the configured budget. The reserve is asymmetric because the failures are: an expired *turn* concedes one server-chosen move (`:451-453`) , an empty *bank* is `forfeitPlayer(..., 'lost due to inactivity.')` (`:455`). The unit is a **request**, not a turn — a post-KO replacement is its own request off the same bank (`:286`) , and that search had no deadline at all until 2026-08-04. Full derivation, the measured wall-clock distribution, the flags and the R6 validation spec are in [SEARCH.md](#).

Sources

- [Monte Carlo Tree Search Variants for Simultaneous Move Games — Lanctot et al., CIG 2014](#)
- [MCTS in Simultaneous Move Games — Tak, Lanctot & Winands](#)
- [Are Strong Policies Also Good Playout Policies? — AIIDE](#)
- [An Analysis of Monte Carlo Tree Search](#)
- [Fixed-Budget Best-Arm Identification in Structured Bandits](#)
- [On-line Policy Improvement using Monte-Carlo Search — Tesauro & Galperin](#)
- [VGC-Bench: Towards Mastering Diverse Team Strategies in Competitive Pokémon](#)