

A living dissemination doc. Maintained in the same pass as the work. It accumulates the publishable findings and keeps a platform-ready draft for each channel (academic, blog/Reddit, Twitter/X thread, YouTube). Numbers here must match the model JSON reports, the white paper, and the site. Last updated 2026-07-24.

You cannot predict who wins a Champions VGC game from the two team sheets — even a player-Elo model ties a coin. So we stopped predicting outcomes and built a *decision* stack; every claim ships a proper score, a confidence interval, and an honest baseline.

1. **Winner-from-sheets $\approx$ a coin.** Player-Elo held-out log-loss 0.687 vs coin 0.693. A cloned-policy rollout win% is *below* a coin — systematically inverted (backs fast offensive teams that lose).
2. **The damage engine is exact.** Within 5% of the Smogon damage calculator on 100% of 31 scenarios (median 0%).
3. **RETRACTED — the live board is predictable only in the sense that counting Pokémon predicts it.** PORY mid-game win% log-loss 0.567 vs coin 0.693, calibrated to ECE 1.6%, clustered CI [0.548, 0.583].
4. **Picking the team doesn't beat a coin (honest null).** CHOMP-EV on 1,205 held-out games: bring quality ties coin/Elo/usage; robust to dropping forfeits; a belief-weighted variant also ties.
5. ****The meta looks rock-paper-scissors, but it's a hint.**** Playstyle cycle (Trick Room $\rightarrow$ Hyper Offense $\rightarrow$ Sand) on only 13–18 games/leg, CIs cross 50% — suggestive, not settled.
6. **Intuitions flatten under data.** After fixing a classifier bug (Charizard = Sun/Mega-Y), Sun went 15 $\rightarrow$ 1,367 teams and Rain vs Sun is even (51/49, $n=236$); Tailwind vs no-Tailwind is 47% ($n=756$).
7. **Single-label archetypes were the bug; multi-label roles pool the data.** Tagging each team by the 26 functional roles it reveals (not one archetype) lifts the median matchup cell from $n \approx 15$ to **$n=20$** across 1,051 cells (2026-07-25, clean games) — still above the old single-label cells of 11–18, but only just. An earlier $n=7,971$ was retracted in 2.7.0 as an over-tagging artifact. But predicting the winner from preview roles still **ties a coin** (0.694 vs 0.693): roles describe and attribute, they don't predict.
8. ****WAR: which species you bring carries a small real signal. A ridge Adjusted-Plus-Minus (RAPM) model on preview species appeared to beat a coin (0.6875 vs 0.6931), a result withdrawn 2026-07-25**** — that figure was measured on the UNFILTERED store; on quality-filtered games WAR scores 0.7048 against a coin's 0.6931 (accuracy 0.502). The apparent signal was four bot accounts playing one team in 1,446 games — where roles and raw sheets do not. Leaders Basculegion / Kingambit / Sylveon; effect sizes small, ridge-shrunk.
9. **Roles can be discovered, not declared.** NMF of the team $\times$ role matrix recovers six interpretable archetypes (recon-err 0.53) — Intimidate+Fake-Out control, physical offense, special offense+sustain, bulky wall+screens+redirection, Tailwind+Encore, priority — with each team a *blend*, never one label. Move-level NMF is coarser (offensive cores dominate). Rank/weighting selection by topic coherence is next.
1. **The field really does tech for the metagame — measured.** Rare "tech" move slots carry **+0.0386** more meta-weighted type coverage than a species' standard kit (95% CI [0.0155, 0.0617], excludes zero; positive in 90/148 species). Against Kingambit (Dark/Steel, on 29% of teams) the top tech answers are Incineroar's Close Combat and Blastoise's Aura Sphere — both **4 $\times$** . Design note: the store spans 3 days, so the temporal test was impossible and a cross-sectional design was used instead; that limitation ships with the number.

2. **A silent data defect can hide an entire mechanic.** The extractor never parsed mega evolution, so 904 of 906 Charizard-Mega-Y sets had no ability, and the damage engine's dex contained ONE mega forme while computing damage for dozens. Raichu-Mega-Y is **No Guard** — which is the whole reason Zap Cannon (50% accuracy) appears on ~590 Raichu sets. A set that looks like a mistake is often a mechanic you are not modelling.
3. **Some facts are unobservable in principle, and should be labelled as such.** A mega's ability is never announced in a replay log, so it cannot be harvested at any sample size — it has to come from the game's own data files. Conversely, Zoroark's Illusion *can* be proven from logs alone: if the apparent species cannot learn the move and Zoroark can, the disguise is certain. **156 proven disguises** across 395 Zoroark team-sides.

Competitive Pokémon (VGC doubles) is a two-player, zero-sum, imperfect-information, simultaneous-move game with a non-transitive metagame and an astronomically large team-building space. We show empirically, on thousands of real ladder replays, that predicting a game's winner from the two team sheets is near the coin-flip ceiling — a strong player-rating model does not beat it. We therefore frame the problem as decision support rather than outcome prediction, and evaluate every component with a proper score, a clustered confidence interval, and an honest baseline. Our validated damage engine matches the community ground truth; a small mid-game value network beats a coin and is calibrated; and we report two honest negatives — team-preview bring selection ties a coin, and the "rock-paper-scissors" metagame is suggestive but underpowered. We outline a belief-state, depth-limited search capstone (ReBeL/Metamon-style) as future work.

"I collected 5,000+ real Champions games and tried to build an AI that predicts who wins from team sheets. It can't — it's a coin flip, and so is Elo. Here's what I built instead, and the two things I was sure about that the data proved wrong."

1/ You can't predict a competitive Pokémon game from the team sheets. Even a rating model ties a coin. 2/ So I stopped predicting winners and measured *decisions* — every number with error bars. 3/ Damage math: exact. Mid-game win%: beats a coin, calibrated. Team-picking: ties a coin (I report it). 4/ "Rain beats Sun"? Even (51/49). "Tailwind wins"? 47%. Intuitions flatten under data. 5/ The endgame: an in-battle coach (Stockfish-for-Pokémon). Repo + live site below.

- Cold open: "I tried to predict Pokémon games with AI. It failed — on purpose."
- The ceiling (why sheets can't call the winner) · the pivot (decisions, not outcomes)
- The town of models (MEDI, GURU, XATU, PORY, KING, CHOMP, ALAKAZAM) on the live site
- The honest nulls (why reporting them is the point) · the road to ALAKAZAM · call to action.

Permanent winner-prediction ceiling; meta-structure results are small-sample; rollout policy is the residual GIGO; the capstone's strongest version needs cloud-GPU training.

Metamon (offline RL + transformers, UT-Austin-RPL); VGC-Bench; PokéAgent Challenge (NeurIPS 2025); Aaron Traylor (competitive-Pokémon framing); ReBeL/CFR/DeepStack/Libratus; Nash-averaging (Balduzzi); blade-chest (Chen & Joachims); HodgeRank (Jiang–Lim–Yao–Ye); sports xG/EPV.

- [] arXiv / workshop write-up · [] r/VGC + r/MachineLearning · [] X/Twitter thread
- [] YouTube explainer · [] Smogon forums / PokéAgent · [] portfolio + live site link